
Docker and Kubernetes for Java EE Developers

v1.0, Jun 16, 2015

Table of Contents

1. Preface	2
2. Setup Environments	3
2.1. Attendees	3
3. Docker Basics	6
3.1. Docker Machine	9
3.2. Docker Client	9
3.3. Verify Docker Configuration	10
4. Run Container	11
4.1. Pull Image	11
4.2. Run Container	13
4.3. Run Container with Default Port	15
4.4. Run Container with Specified Port	16
4.5. Enabling WildFly Administration	17
4.6. Stop and Remove Container	20
5. Deploy Java EE 7 Application (Pre-Built WAR)	21
6. Deploy Java EE 7 Application (Container Linking)	22
7. Build and Deploy Java EE 7 Application	24
7.1. Build Application	24
7.2. Start Application Server	24
7.3. Configure JBoss Developer Studio	25
7.4. Deploy Application Using Shared Volumes	29
7.5. Deploy Application Using CLI	31
7.6. Deploy Application Using Web Console	32
7.7. Deploy Application Using Management API	35
8. Docker Maven Plugin	39
8.1. Run Java EE Application	39
8.2. Understand Plugin Configuration	40
9. Docker Tools in Eclipse	42
9.1. Install Docker Tools Plugins	42
9.2. Docker Explorer	44

9.3. Docker Images	48
9.4. Docker Containers	50
9.5. Details on Images and Containers	51
10. Test Java EE Applications on Docker	53
11. Multiple Containers Using Docker Compose	55
11.1. Configuration File	56
11.2. Start Services	57
11.3. Verify Application	59
11.4. Stop Services	60
11.5. Scale Services	61
12. Java EE Application on Mod Cluster	61
13. Java EE Application on Docker Swarm Cluster	65
14. Java EE Application on Kubernetes Cluster	71
14.1. Key Concepts	72
14.2. Start Kubernetes Cluster	73
14.3. Deploy Java EE Application	78
14.4. Access Java EE Application	83
14.5. Self-healing Pods	83
14.6. Application Logs	83
14.7. Delete Kubernetes Resources	84
14.8. Stop Kubernetes Cluster	84
14.9. View Logs	84
14.10. Debug Kubernetes Master (OPTIONAL)	85
15. Common Docker Commands	86
16. References	87

1. Preface

Containers are enabling developers to package their applications (and underlying dependencies) in new ways that are portable and work consistently everywhere? On your machine, in production, in your data center, and in the cloud. And Docker has become the de facto standard for those portable containers in the cloud.

Docker is the developer-friendly Linux container technology that enables creation of your stack: OS, JVM, app server, app, and all your custom configuration. So with all it offers, how comfortable are you and your team taking Docker from development to production? Are you hearing developers say, “But it works on my machine!” when code breaks in production?

This lab offers developers an intro-level, hands-on session with Docker, from installation, to exploring Docker Hub, to crafting their own images, to adding Java apps and running custom containers. It will also explain how to use Swarm to orchestrate these containers together. This is a BYOL (bring your own laptop) session, so bring your Windows, OSX, or Linux laptop and be ready to dig into a tool that promises to be at the forefront of our industry for some time to come.



Latest content of this lab is always at <https://github.com/javaee-samples/docker-java>

2. Setup Environments

This section describes the relevant steps for both attendees and instructors to setup the environments. Please follow the parts, that are appropriate for you.

2.1. Attendees

This lab is designed for a BYOL (Brying Your Own Laptop) style hands-on-lab.

Hardware

1. Operating System: Mac OS X (10.8 or later), Windows 7 (SP1), Fedora (21 or later)
2. Memory: At least 4 GB+, preferred 8 GB

Software

1. Java: [Oracle JDK 8u45¹](http://www.oracle.com/technetwork/java/javase/downloads/jdk8-downloads-2133151.html)
2. Webbrowser
 - a. [Chrome²](https://www.google.com/chrome/browser/desktop/)
 - b. [Firefox³](http://www.getfirefox.com)

A Word About Licenses

This tutorial only uses software which is open source or at least free to use in development and/or education. Please refer to the individual products/tools used in this tutorial.

¹ <http://www.oracle.com/technetwork/java/javase/downloads/jdk8-downloads-2133151.html>

² <https://www.google.com/chrome/browser/desktop/>

³ <http://www.getfirefox.com>

Git Client

Install Git Client as explained at: <https://git-scm.com/book/en/v2/Getting-Started-Installing-Git>

Maven

1. Download Apache Maven from <https://maven.apache.org/download.cgi>
2. Unzip to a directory of your choice and add it to the `PATH`.

VirtualBox

Docker currently runs natively on Linux. It can be configured to run in a virtual machine on Mac or Windows. This is why Virtualbox is a requirement for Mac or Windows.

Downloads are available from <https://www.virtualbox.org/>.



Linux Users

1. Have your kernel updated
2. Users should have the GNU compiler, build and header files for your current Linux kernel
3. Create a `/usr/src/linux` link to the current kernel source

Vagrant

Download Vagrant from <https://www.vagrantup.com/downloads.html> and install.

Docker Machine

Docker Machine makes it really easy to create Docker hosts on your computer, on cloud providers and inside your own data center. It creates servers, installs Docker on them, then configures the Docker client to talk to them.

Install as explained at: <https://docs.docker.com/machine/#installation>

On Windows copy the file in a convenience directory that is included in the `PATH` variable.

Create Lab Docker Host

1. Create Docker Host to be used in the lab:

```
docker-machine create --driver=virtualbox lab
eval "$(docker-machine env lab)"
```

2. To make it easier to start/stop the containers, an entry is added into the host mapping table of your operating system. Find out the IP address of your machine:

```
docker-machine ip lab
```

This will provide the IP address associated with the Docker Machine created earlier.

3. Edit `/etc/hosts` (Mac OS or Linux) or `C:\Windows\System32\drivers\etc\hosts` (Windows) and add:

```
<IP ADDRESS> dockerhost
```

Docker Client

Docker Client is used to communicate with Docker Host.

```
# Mac
curl -L https://get.docker.com/builds/Darwin/x86_64/docker-latest > /usr/local/bin/docker
chmod +x /usr/local/bin/docker

# Linux
curl -L https://get.docker.com/builds/Linux/x86_64/docker-latest > docker-latest-linux
chmod +x /usr/local/bin/docker

# Windows
curl -L http://test.docker.com.s3.amazonaws.com/builds/Windows/x86_64/docker-1.7.0.exe -o docker.exe
```

Kubernetes

1. Download Kubernetes (0.21.0) from <https://github.com/GoogleCloudPlatform/kubernetes/releases/tag/v0.21.0>
2. Install it by extracting the archive.

WildFly

1. Download WildFly 9.0 from <http://download.jboss.org/wildfly/8.2.0.Final/wildfly-9.0.0.Final.zip>
2. Install it by extracting the archive.

JBoss Developer Studio 9 - Beta 1

To install JBoss Developer Studio stand-alone, complete the following steps:

1. Download [9.0.0 Beta 1⁴](#) [Latest Nightly⁵](#) (used with the lab)
2. Start the installer as:

```
java -jar <JAR FILE NAME>
```

Follow the on-screen instructions to complete the installation process.

3. Docker Basics

PURPOSE: This chapter introduces the basic terminology of Docker.

Docker is a platform for developers and sysadmins to develop, ship, and run applications. Docker lets you quickly assemble applications from components and eliminates the friction that can come when shipping code. Docker lets you get your code tested and deployed into production as fast as possible.

— docs.docker.com/

Docker simplifies software delivery by making it easy to build and share images that contain your application's entire environment, or *application operating system*.

What does it mean by an application operating system ?

Your application typically require a specific version of operating system, application server, JDK, database server, may require to tune the configuration files, and similarly multiple other dependencies. The application may need binding to specific ports and certain amount of memory. The components and configuration together required to run your application is what is referred to as application operating system.

⁴ <http://tools.jboss.org/downloads/devstudio/mars/9.0.0.Beta1.html>

⁵ https://devstudio.redhat.com/9.0/snapshots/builds/devstudio.product_master/latest/installer/

You can certainly provide an installation script that will download and install these components. Docker simplifies this process by allowing to create an image that contains your application and infrastructure together, managed as one component. These images are then used to create Docker containers which run on the container virtualization platform, provided by Docker.

Main Components of Docker

Docker has three main components:

1. *Images* are **build component** of Docker and a read-only template of application operating system.
2. *Containers* are **run component** of Docker, and created from, images. Containers can be run, started, stopped, moved, and deleted.
3. Images are stored, shared, and managed in a *registry*, the **distribution component** of Docker. The publically available registry is known as Docker Hub.

In order for these three components to work together, there is **Docker Daemon** that runs on a host machine and does the heavy lifting of building, running, and distributing Docker containers. In addition, there is **Client** that is a Docker binary which accepts commands from the user and communicates back and forth with the daemon.




Figure 1. Docker architecture

Client communicates with Daemon, either co-located on the same host, or on a different host. It requests the Daemon to pull an image from the repository using `pull` command. The Daemon then downloads the image from Docker Hub, or whatever registry is configured. Multiple images can be downloaded from the registry and installed on Daemon host. Images are run using `run` command to create containers on demand.

How does a Docker Image work?

We've already seen that Docker images are read-only templates from which Docker containers are launched. Each image consists of a series of layers. Docker makes use of union file systems to combine these layers into a single image. Union file systems allow files and directories of separate file systems, known as branches, to be transparently overlaid, forming a single coherent file system.

One of the reasons Docker is so lightweight is because of these layers. When you change a Docker image—for example, update an application to a new version—a new layer gets built. Thus, rather than replacing the whole image or entirely rebuilding, as you may do with a virtual machine, only that layer is added or updated. Now you don't need to distribute a whole new image, just the update, making distributing Docker images faster and simpler.

Every image starts from a base image, for example `ubuntu`, a base Ubuntu image, or `fedora`, a base Fedora image. You can also use images of your own as the basis for a new image, for example if you have a base Apache image you could use this as the base of all your web application images.



By default, Docker obtains these base images from Docker Hub.

Docker images are then built from these base images using a simple, descriptive set of steps we call instructions. Each instruction creates a new layer in our image. Instructions include actions like:

1. Run a command.
2. Add a file or directory.
3. Create an environment variable.
4. What process to run when launching a container from this image.

These instructions are stored in a file called a Dockerfile. Docker reads this Dockerfile when you request a build of an image, executes the instructions, and returns a final image.

How does a Container work?

A container consists of an operating system, user-added files, and meta-data. As we've seen, each container is built from an image. That image tells Docker what the container

holds, what process to run when the container is launched, and a variety of other configuration data. The Docker image is read-only. When Docker runs a container from an image, it adds a read-write layer on top of the image (using a union file system as we saw earlier) in which your application can then run.

3.1. Docker Machine

Machine makes it really easy to create Docker hosts on your computer, on cloud providers and inside your own data center. It creates servers, installs Docker on them, then configures the Docker client to talk to them.

Once your Docker host has been created, it then has a number of commands for managing them:

1. Starting, stopping, restarting
2. Upgrading Docker
3. Configuring the Docker client to talk to your host

You used Docker Machine already during the attendee setup. We won't need it too much further on. But if you need to create hosts, it's a very handy tool to know about. From now on we're mostly going to use the docker client.

Find out more about the details at the [Official Docker Machine Website⁶](#)

Check if docker machine is working using the following command:

```
.....  
docker-machine -v  
.....
```

It shows the output similar to the one shown below:

```
.....  
docker-machine version 0.3.0 (0a251fe)  
.....
```

3.2. Docker Client

The client communicates with the demon process on your host and let's you work with images and containers.

Check if your client is working using the following command:

⁶ <https://docs.docker.com/machine/>

```
docker -v
```

It shows the output similar to the following:

```
Docker version 1.7.0, build 0baf609
```

The most important options you'll be using frequently are:

1. `run` - runs a container
2. `ps` - lists containers
3. `stop` - stops a container
4. `rm` - Removes a container

Get a full list of available commands with

```
docker
```

A more comprehensive list of commands is also available in [Section 15, “Common Docker Commands”](#).

3.3. Verify Docker Configuration

Check if your Docker Host is running:

```
docker-machine ls
```

You should see the output similar to:

NAME	ACTIVE	DRIVER	STATE	URL
SWARM				
lab		virtualbox	Running	tcp://192.168.99.101:2376

This machine is shown in “Running” state. If the machine state is stopped, start it with:

```
docker-machine start lab
```

After it is started you can find out IP address of your Docker Host with:

```
docker-machine ip lab
```

We already did this during the setup document, remember? So, this is a good chance to check, if you already added this IP to your hosts file.

Type:

```
ping dockerhost
```

and see if this resolves to the IP address that the docker-machine command printed out. You should see an output as:

```
> ping dockerhost
PING dockerhost (192.168.99.101): 56 data bytes
64 bytes from 192.168.99.101: icmp_seq=0 ttl=64 time=0.394 ms
64 bytes from 192.168.99.101: icmp_seq=1 ttl=64 time=0.387 ms
```

If it does, you're ready to start over with the lab.

4. Run Container

The first step in running any application on Docker is to run a container from an image. There are plenty of images available from the official Docker registry (aka [Docker Hub](#)⁷). To run any of them, you just have to ask the Docker Client to run it. The client will check if the image already exists on Docker Host. If it exists then it'll run it, otherwise the host will download the image and then run it.

4.1. Pull Image

Let's first check, if any images are available:

```
docker images
```

At first, this list is empty. Now, let's get a vanilla `jboss/wildfly` image:

```
docker pull jboss/wildfly
```

⁷ <https://hub.docker.com>

By default, docker images are retrieved from [Docker Hub⁸](#).

You can see, that Docker is downloading the image with it's different layers.



In a traditional Linux boot, the Kernel first mounts the root File System as read-only, checks its integrity, and then switches the whole rootfs volume to read-write mode. When Docker mounts the rootfs, it starts read-only, as in a traditional Linux boot, but then, instead of changing the file system to read-write mode, it takes advantage of a union mount to add a read-write file system over the read-only file system. In fact there may be multiple read-only file systems stacked on top of each other. Consider each one of these file systems as a layer.

At first, the top read-write layer has nothing in it, but any time a process creates a file, this happens in the top layer. And if something needs to update an existing file in a lower layer, then the file gets copied to the upper layer and changes go into the copy. The version of the file on the lower layer cannot be seen by the applications anymore, but it is there, unchanged.

We call the union of the read-write layer and all the read-only layers a *union file system*.




Figure 2. Docker Layers

⁸ <https://hub.docker.com/>

In our particular case, the [jboss/wildfly⁹](#) image extends the [jboss/base-jdk:7¹⁰](#) image which adds the OpenJDK distribution on top of the [jboss/base¹¹](#) image. The base image is used for all JBoss community images. It provides a base layer that includes:

1. A jboss user (uid/gid 1000) with home directory set to `/opt/jboss`
2. A few tools that may be useful when extending the image or installing software, like unzip.

The 'jboss/base-jdk:7' image adds:

1. Latest OpenJDK distribution
2. Adds a JAVA_HOME environment variable

When the download is done, you can list the images again and will see the following:

docker images			
REPOSITORY	TAG	IMAGE ID	CREATED
jboss/wildfly	latest	e908c8c95a8b	5 days ago
	581.5 MB		

4.2. Run Container

Interactive Container

Run WildFly container in an interactive mode.

```
docker run -it jboss/wildfly
```

This will show the output as:

```
=====
JBoss Bootstrap Environment
```

⁹ <https://github.com/jboss-dockerfiles/wildfly/blob/master/Dockerfile>

¹⁰ <https://github.com/jboss-dockerfiles/base/blob/master/Dockerfile>

¹¹ <https://github.com/jboss-dockerfiles/base/blob/master/Dockerfile>

Docker and Kubernetes for Java EE Developers

```
JBOSS_HOME: /opt/jboss/wildfly
```

```
JAVA: /usr/lib/jvm/java/bin/java
```

```
JAVA_OPTS: -server -XX:+UseCompressedOops -  
server -XX:+UseCompressedOops -Xms64m -Xmx512m -  
XX:MaxPermSize=256m -Djava.net.preferIPv4Stack=true -  
Djboss.modules.system.pkgs=org.jboss.byteman -Djava.awt.headless=true
```

```
=====
```

```
OpenJDK 64-Bit Server VM warning: ignoring option MaxPermSize=256m;  
support was removed in 8.0  
00:44:43,895 INFO [org.jboss.modules] (main) JBoss Modules version  
1.4.3.Final  
00:44:44,184 INFO [org.jboss.msc] (main) JBoss MSC version 1.2.6.Final  
00:44:44,267 INFO [org.jboss.as] (MSC service thread 1-2) WFLYSRV0049:  
WildFly Full 9.0.0.Final (WildFly Core 1.0.0.Final) starting  
  
.  
.  
.  
00:46:54,241 INFO [org.jboss.as] (Controller Boot Thread) WFLYSRV0060:  
Http management interface listening on http://127.0.0.1:9990/management  
00:46:54,243 INFO [org.jboss.as] (Controller Boot Thread) WFLYSRV0051:  
Admin console listening on http://127.0.0.1:9990  
00:46:54,250 INFO [org.jboss.as] (Controller Boot Thread) WFLYSRV0025:  
WildFly Full 9.0.0.Final (WildFly Core 1.0.0.Final) started in 4256ms -  
Started 203 of 379 services (210 services are lazy, passive or on-demand)
```

This shows that the server started correctly, congratulations!

By default, Docker runs in the foreground. `-i` allows to interact with the STDIN and `-t` attach a TTY to the process. Switches can be combined together and used as `-it`.

Hit Ctrl+C to stop the container.

Detached Container

Restart the container in detached mode:

```
docker run -d jboss/wildfly  
972f51cc8422eec0a7ea9a804a55a2827b5537c00a6bfd45f8646cb764bc002a
```

`-d` runs the container in detached mode.

The output is the unique id assigned to the container. Check the logs as:

```
> docker logs  
972f51cc8422eec0a7ea9a804a55a2827b5537c00a6bfd45f8646cb764bc002a  
=====  
  
JBoss Bootstrap Environment  
  
JBOSS_HOME: /opt/jboss/wildfly  
  
.  
.
```

We can check it by issuing the `docker ps` command which retrieves the images process which are running and the ports engaged by the process:

```
> docker ps  
CONTAINER ID        IMAGE               COMMAND  
CREATED             STATUS              PORTS  
NAMES              jboss/wildfly       "/opt/jboss/wildfly/"  3 seconds  
922abbb9c63a        Up  2 seconds      8080/tcp           desperate_lovelace  
.  
.
```

Also try `docker ps -a` to see all the containers on this machine.

4.3. Run Container with Default Port

Startup log of the server shows that the server is located in the `/opt/jboss/wildfly`. It also shows that the public interfaces are bound to the `0.0.0.0` address while the admin interfaces are bound just to `localhost`. This information will be useful to learn how to customize the server.

`docker-machine ip <machine-name>` gives us the Docker Host IP address and this was already added to the hosts file. So, we can give it another try by accessing: <http://dockerhost:8080>. However, this will not work either.

If you want containers to accept incoming connections, you will need to provide special options when invoking `docker run`. The container, we just started, can't be accessed by our browser. We need to stop it again and restart with different options.

```
docker stop `docker ps | grep wildfly | awk '{print $1}'`  
.
```

Restart the container as:

```
docker run -d -P jboss/wildfly
```

`-P` map any exposed ports inside the image to a random port on Docker host. This can be verified as:

```
> docker ps
CONTAINER ID        IMAGE               COMMAND
CREATED             STATUS              PORTS
NAMES
63a69bff9c69      jboss/wildfly      "/opt/jboss/wildfly/   14 seconds
ago                Up 13 seconds       0.0.0.0:32768->8080/tcp   kickass_bohr
```

The port mapping is shown in the `PORTS` column. Access the WildFly server at <http://dockerhost:32768>. Make sure to use the correct port number as shown in your case.

4.4. Run Container with Specified Port

Lets stop the previously running container as:

```
docker stop `docker ps | grep wildfly | awk '{print $1}'`
```

Restart the container as:

```
docker run -it -p 8080:8080 jboss/wildfly
```

The format is `-p hostPort:containerPort`. This option maps container ports to host ports and allows other containers on our host to access them.



Docker Port Mapping

Port exposure and mapping are the keys to successful work with Docker. See more about networking on the Docker website [Advanced Networking¹²](#)

Now we're ready to test <http://dockerhost:8080> again. This works with the exposed port, as expected.

¹² <https://docs.docker.com/articles/networking/>



Welcome to WildFly 9

Your WildFly 9 is running.

[Documentation](#) | [Quickstarts](#) | [Administration Console](#)

[WildFly Project](#) | [User Forum](#) | [Report an issue](#)

[JBoss Community](#)

Figure 3. Welcome WildFly

4.5. Enabling WildFly Administration

Default WildFly image exposes only port 8080 and thus is not available for administration using either the CLI or Admin Console.

Default Port Mapping

The following command will override the default command in Docker file, explicitly starting WildFly, and binding application and management port to all network interfaces.

```
.....  
docker run -P -d jboss/wildfly /opt/jboss/wildfly/bin/standalone.sh -b  
0.0.0.0 -bmanagement 0.0.0.0  
.....
```

Accessing WildFly Administration Console require a user in administration realm. A pre-created image, with appropriate username/password credentials, is used to start WildFly as:

```
.....  
docker run -P -d arungupta/wildfly-management
```

-P map any exposed ports inside the image to a random port on Docker host.

Look at the exposed ports as:

```
.....  
docker ps  
CONTAINER ID        IMAGE               COMMAND            CREATED             STATUS              PORTS  
                NAMES  
af7d6914a1f9        arungupta/wildfly-management   "/opt/jboss/wildfly/  
2 seconds ago      Up 1 seconds          0.0.0.0:32770->8080/tcp,  
0.0.0.0:32769->9990/tcp   happy_bardeen
```

Look for the host port that is mapped in the container, `32769` in this case. Access the admin console at <http://dockerhost:32769>.

The username/password credentials are:

Field	Value
Username	admin
Password	docker#admin

This shows the admin console as:

The screenshot shows the WildFly 9.0.0.Final dashboard. At the top, there's a navigation bar with tabs: Home, Deployments, Configuration, Runtime, Access Control, and Patching. The Home tab is selected. Below the navigation bar, the page title is "WildFly". Under "View and Manage Settings", there are sections for Deployments (Add and manage deployments), Runtime (Monitor server status, retrieve diagnostic information and perform other operational tasks), Configuration (Configure your server profiles, including the subsystems and resources available to your server), and Access Control (Perform administrative tasks for your server, including security and control). In the "Common Tasks" section, there are four expandable cards: "Deploy an application" (To deploy an application. 1. Add a deployment. 2. Enable the new deployment. > Create Deployment), "Create a datasource" (Add a new datasource and follow the steps in the wizard. > Datasources), "Apply a patch" (Update to the latest patch level. The patch file must be downloaded to your local machine prior to starting this process), and "Assign user roles" (Assign roles to individual users or user groups and resources. 1. Add user or group).

Figure 4. Welcome WildFly

Additional Ways To Find Port Mapping

The exact mapped port can also be found as:

1. Using `docker port`:

```
docker port 6f610b310a46
```

to see the output as:

```
0.0.0.0:32769->8080/tcp
0.0.0.0:32770->9990/tcp
```

2. Using `docker inspect`:

```
docker inspect --format='{{(index (index .NetworkSettings.Ports "9990/tcp") 0).HostPort}}' <CONTAINER ID>
```

Fixed Port Mapping

This management image can also be started with a pre-defined port mapping as:

```
docker run -p 8080:8080 -p 9990:9990 -d arungupta/wildfly-management
```

In this case, Docker port mapping will be shown as:

```
8080/tcp -> 0.0.0.0:8080  
9990/tcp -> 0.0.0.0:9990
```

4.6. Stop and Remove Container

Stop Container

1. Stop a specific container:

```
docker stop <CONTAINER ID>
```

2. Stop all the running containers

```
docker rm $(docker stop $(docker ps -q))
```

3. Stop only the exited containers

```
docker ps -a -f "exited=-1"
```

Remove Container

1. Remove a specific container:

```
docker rm 0bc123a8ece0
```

2. Containers meeting a regular expression

```
docker ps -a | grep wildfly | awk '{print $1}' | xargs docker rm
```

3. All running containers, without any criteria

```
docker rm $(docker ps -aq)
```

5. Deploy Java EE 7 Application (Pre-Built WAR)

Java EE 7 Movieplex¹³ is a standard multi-tier enterprise application that shows design patterns and anti-patterns for a typical Java EE 7 application.

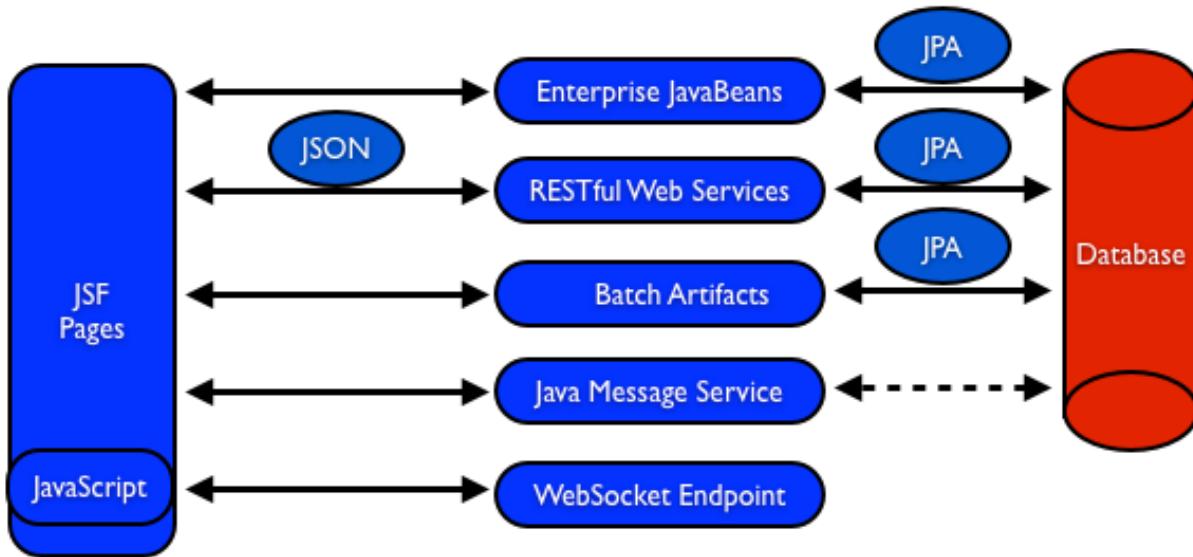


Figure 5. Java EE 7 Application Architecture

Pull the Docker image that contains WildFly and pre-built Java EE 7 application WAR file as shown:

```
docker pull arugupta/javaee7-hol
```

The [javaee7-hol Dockerfile](#)¹⁴ is based on [jboss/wildfly](#) and adds the movieplex7 application as war file.

Run it as:

```
docker run -it -p 8080:8080 arungupta/javaee7-hol
```

See the application in action at <http://dockerhost:8080/movieplex7/>. The output is shown as:

¹³ <https://github.com/javaee-samples/javaee7-hol>

¹⁴ <https://github.com/arun-gupta/docker-images/blob/master/javaee7-hol/Dockerfile>

Figure 6. Java EE 7 Application Output

This uses an in-memory database with WildFly application server as shown in the image:

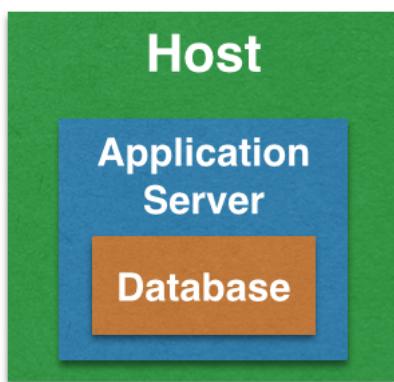


Figure 7. In-memory Database

Only two changes are required to the standard `jboss/wildfly` image:

1. By default, WildFly starts in Web platform. This Java EE 7 application uses some capabilities from the Full Platform and so WildFly is started in that mode instead as:

```
CMD ["/opt/jboss/wildfly/bin/standalone.sh", "-c", "standalone-full.xml", "-b", "0.0.0.0"]
```

2. WAR file is copied to the `standalone/deployments` directory as:

```
RUN curl -L https://github.com/javaee-samples/javaee7-hol/raw/master/solution/movieplex7-1.0-SNAPSHOT.war -o /opt/jboss/wildfly/standalone/deployments/movieplex7-1.0-SNAPSHOT.war
```

6. Deploy Java EE 7 Application (Container Linking)

Section 5, “Deploy Java EE 7 Application (Pre-Built WAR)” explained how to use an in-memory database with the application server. This gets you started rather quickly but becomes a bottleneck soon as the database is only in-memory. This means that any changes made to your schema and data are lost when the application server shuts down. In this case, you need to use a database server that resides outside the application server. For example, MySQL as the database server and WildFly as the application server.

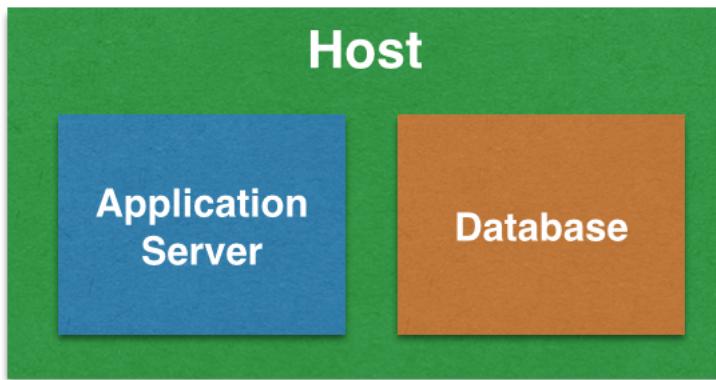


Figure 8. Two Containers On Same Docker Host

This section will show how [Docker Container Linking](#)¹⁵ can be used to connect to a service running inside a Docker container via a network port.

1. Start MySQL server as:

```
docker run --name mysqldb -e MYSQL_USER=mysql -e MYSQL_PASSWORD=mysql -  
e MYSQL_DATABASE=sample -e MYSQL_ROOT_PASSWORD=supersecret -p 3306:3306  
-d mysql
```

`-e` define environment variables that are read by the database at startup and allow us to access the database with this user and password.

2. Start WildFly and deploy Java EE 7 application as:

```
docker run -it --name mywildfly --link mysqldb:db -p 8080:8080  
arungupta/wildfly-mysql-javaee7
```

`--link` takes two parameters - first is name of the container we're linking to and second is the alias for the link name.



Container Linking

Creating a link between two containers creates a conduit between a source container and a target container and securely transfer information about source container to target container.

In our case, target container (WildFly) can see information about source container (MySQL). When containers are linked, information about a source container can be sent to a recipient container. This allows the recipient to see

¹⁵ <https://docs.docker.com/userguide/dockerlinks/>

selected data describing aspects of the source container. For example, IP address of MySQL server is exposed at \$DB_PORT_3306_TCP_ADDR and port of MySQL server is exposed at \$DB_PORT_3306_TCP_PORT. These are then used to create the JDBC resource.

See more about container communication on the Docker website
[Linking Containers Together¹⁶](#)

3. See the output as:

```
> curl http://dockerhost:8080/employees/resources/employees
<?xml version="1.0" encoding="UTF-8" standalone="yes"?
<><collection><employee><id>1</id><name>Penny</name></
employee><employee><id>2</id><name>Sheldon</name></
employee><employee><id>3</id><name>Amy</name></
employee><employee><id>4</id><name>Leonard</name></
employee><employee><id>5</id><name>Bernadette</
name></employee><employee><id>6</id><name>Raj</name></
employee><employee><id>7</id><name>Howard</name></
employee><employee><id>8</id><name>Priya</name></employee></collection>
```

7. Build and Deploy Java EE 7 Application

[Java EE 7 Simple Sample¹⁷](#) is a trivial Java EE 7 sample application.

7.1. Build Application

1. Clone the repo:

```
git clone https://github.com/javaee-samples/javaee7-simple-sample.git
```

1. Build the application:

```
mvn clean package
```

7.2. Start Application Server

Start WildFly server as:

¹⁶ <https://docs.docker.com/userguide/dockerlinks/>

¹⁷ <https://github.com/javaee-samples/javaee7-simple-sample>

```
docker run --name wildfly -d -p 8080:8080 -v /Users/youruser/tmp/  
deployments:/opt/jboss/wildfly/standalone/deployments/:rw jboss/wildfly
```

Make sure to replace `/Users/youruser/tmp/deployments` to a directory on your local machine. Also, make sure this directory already exists. For example, on my machine this directory is `/Users/arungupta/tmp/deployments`.

This command starts a container named ``wildfly''.

The `-v` flag maps a directory from the host into the container. This will be the directory to put the deployments. `rw` ensures that the Docker container can write to it.



Windows users, please make sure to use `-v /c/Users/` notation for drive letters.

Check logs to verify if the server has started.

```
docker logs -f wildfly
```

Access <http://dockerhost:8080> in your browser to make sure the instance is up and running.

Now you're ready to deploy the application for the first time.

7.3. Configure JBoss Developer Studio

Start JBoss Developer Studio, if not already started.

1. Select 'Servers' tab, create a new server adapter

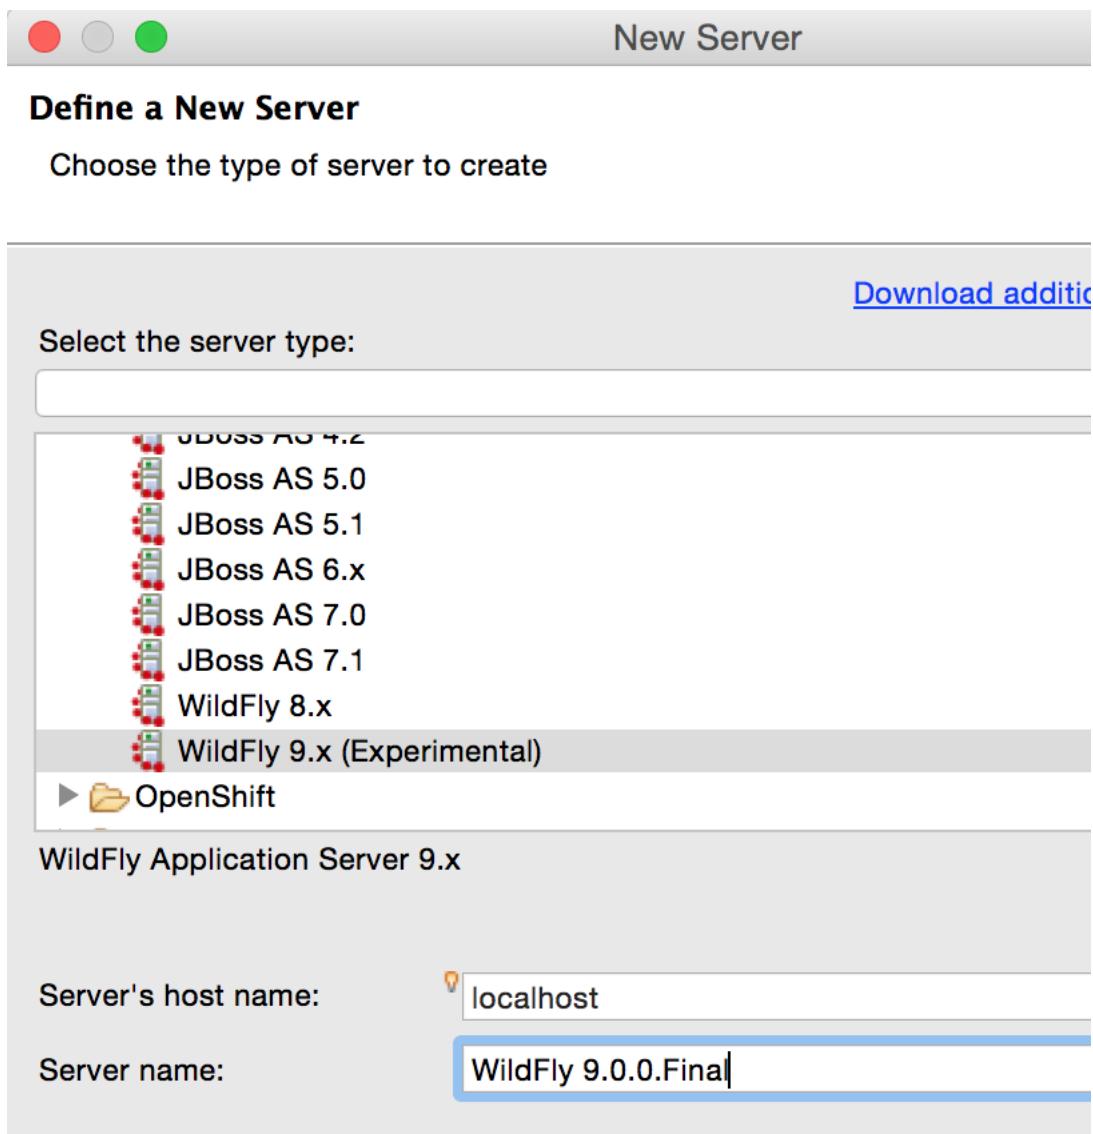


Figure 9. Server adapter

2. Assign an existing or create a new WildFly 9.0.0 runtime (changed properties are highlighted.)




Figure 10. WildFly Runtime Properties

3. If a new runtime needs to be created, pick the director for WildFly 9.0.0:




Figure 11. WildFly 9.0.0.Final Runtime

Click on 'Finish'.

- Double-click on the newly selected server to configure server properties as shown in the following image.

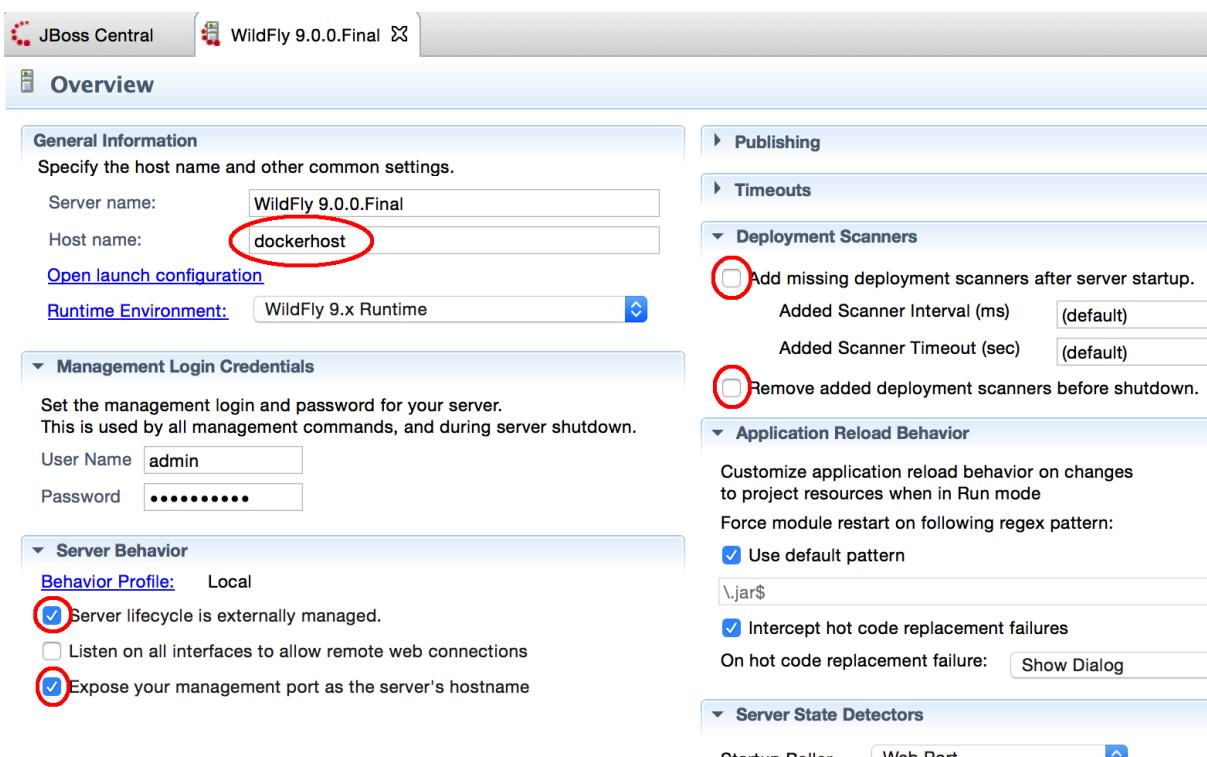


Figure 12. Server properties

The host name is specified to 'dockerhost'. Two properties on the left are automatically propagated from the previous dialog. Additional two properties on the right side are required to disable to keep deployment scanners in sync with the server.

5. Specify a custom deployment folder on Deployment tab of Server Editor



Figure 13. Custom deployment folder

6. Right-click on the newly created server adapter and click ``Start''.

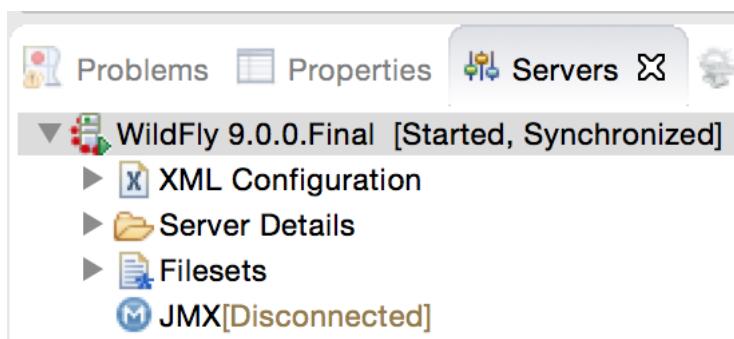


Figure 14. Started server

7.4. Deploy Application Using Shared Volumes

1. Import javaee7-simple-sample application source code using Import → Existing Maven Projects.
2. Right-click on the project, select ``Run on Server'' and chose the previously created server.

The project runs and displays the start page of the application.



GET all the [persons](#).

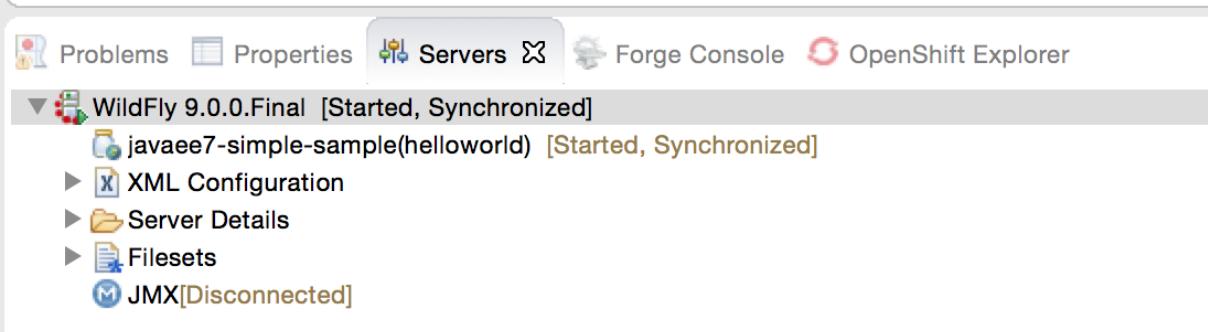


Figure 15. Start Server

Congratulations! You've just deployed your first application to WildFly running in a Docker container from JBoss Developer Studio.

Stop WildFly container when you're done.

```
docker stop wildfly
```

7.5. Deploy Application Using CLI

The Command Line Interface (CLI) is a tool for connecting to WildFly instances to manage all tasks from command line environment. Some of the tasks that you can do using the CLI are:

1. Deploy/Undeploy web application in standalone/Domain Mode.
2. View all information about the deployed application on runtime.
3. Start/Stop/Restart Nodes in respective mode i.e. Standalone/Domain.
4. Adding/Deleting resource or subsystems to servers.

Lets use the CLI to deploy `javaee7-simple-sample` to WildFly running in the container.

1. CLI needs to be locally installed and comes as part of WildFly. Download WildFly 9.0 from ifdef::public[<http://download.jboss.org/wildfly/9.0.0.Final/wildfly-9.0.0.Final.zip>]. Unzip into a folder of your choice (e.g. `/Users/arungupta/tools/`). This will create `wildfly-9.0.0.Final` directory here. This folder is referred to `$WILDFY_HOME` from here on. Make sure to add the `/Users/arungupta/tools/wildfly-9.0.0.Final/bin` to your `$PATH`.
2. Run the ``wildfly-management'' image with fixed port mapping as explained in the section called “Fixed Port Mapping”.
3. Run the `jboss-cli` command and connect to the WildFly instance.

```
jboss-cli.sh --controller=dockerhost:9990 -u=admin -p=docker#admin -c
```

This will show the output as:

```
[standalone@dockerhost:9990 /]
```

4. Deploy the application as:

```
deploy <javaee7-simple-sample PATH>target/javaee7-simple-sample-1.10.war --force
```

Now you've sucessfully used the CLI to remote deploy the Java EE 7 sample application to WildFly running as docker container.

7.6. Deploy Application Using Web Console

WildFly comes with a web-based administration console. It also relies on the same management APIs that are used by JBoss Developer Tools and the CLI. It provides a simple and easy to use web-based console to manage WildFly instance. For a Docker image, it needs to be explicitly enabled as explained in [Section 4.5, “Enabling WildFly Administration”](#). Once enabled, it can be accessed at <http://dockerhost:9990>.

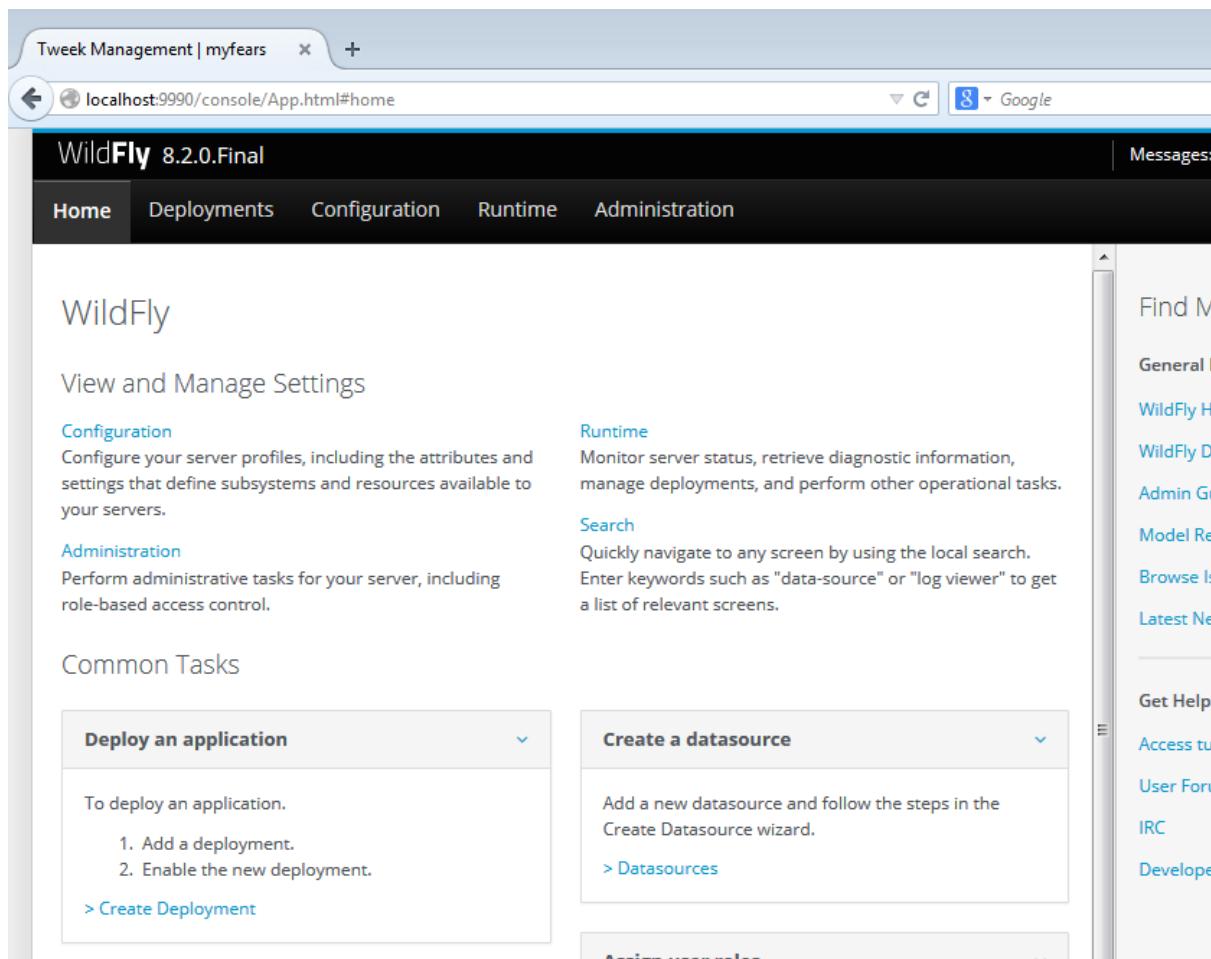


Figure 16. WildFly Web Console

Username and password credentials are shown in ???.



You may like to stop and remove the Docker container running WildFly. This can be done as `docker ps -a | grep wildfly | awk '{print $1}' | xargs docker rm -f`. Start a new container as `docker run -d --name wildfly -p 8080:8080 -p 9990:9990 arungupta/wildfly-management`.

Deploy the application using the console with the following steps:

1. Go to 'Deployments' tab.

The screenshot shows the WildFly 9.0.0.Final Web Console interface. At the top, there is a navigation bar with tabs: Home, Deployments (which is currently selected), Configuration, Runtime, Access Control, and Patching. Below the navigation bar, the main content area has a title 'Deployment' and a sub-section 'Common Configuration Tasks'. A button labeled 'Add' is visible. The main content area displays the message 'No Items!'.

Figure 17. Deployments tab in WildFly Web Console

2. Click on 'Add' button.
3. On 'Add Deployment' screen, take the default of 'Upload a new deployment' and click 'Next>>'.
4. Click on `Choose File`, select the `<javaee7-simple-sample PATH>/javaee7-simple-sample.war` file on your computer. This would be ``javaee7-simple-sample/target/javaee7-simple-sample.war` from [Section 7.1, “Build Application”](#).
5. Click on 'Next>>'.
6. Select 'Enable' checkbox.

Add Deployment ✖

Verify Upload

[Need Help?](#)

Name:

Runtime Name:

Enable:

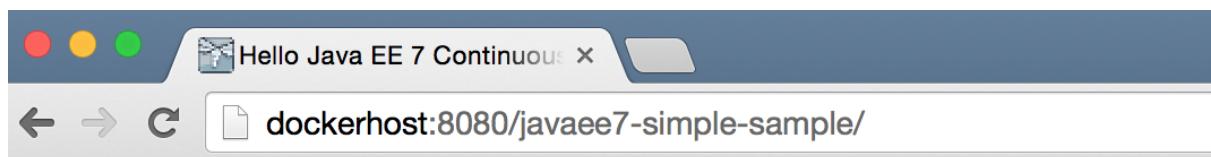
Figure 18. Enable a deployment

7. Click 'Finish'.

The screenshot shows the WildFly 9.0.0.Final Web Console. The top navigation bar includes Home, Deployments (selected), Configuration, Runtime, Access Control, and Patching. The main content area is titled 'Deployment' and shows a single entry: 'javaee7-simple-sam...'. An 'Add' button is located at the top right of this section. To the right, under the heading 'Deployment', it states: 'The deployment javaee7-simple-sample.war is enabled and has the status OK.' Below this, the 'Details' section lists: 'Last enabled at 2015-07-10 01:20:06,552 UTC', 'The deployment was never disabled', and 'Runtime name: javaee7-simple-sample.war'. Under 'Subsystems', it lists 'jaxrs' and 'undertow'.

Figure 19. Java EE 7 Simple Sample Deployed

This will complete the deployment of the Java EE 7 application using Web Console. The output can be seen out <http://dockerhost:8080/javaee7-simple-sample> and looks like:



Hello Java EE 7 Continuous Delivery!

GET all the [persons](#).

Figure 20. Java EE 7 Simple Sample Output

7.7. Deploy Application Using Management API


A standalone WildFly process, process can be configured to listen for remote management requests using its 'native management interface'. The CLI tool that comes with the application server uses this interface, and user can develop custom clients that use it as well. By default, WildFly management interface listens on 127.0.0.1. When running inside a Docker container, the network interface should be bound to all publicly assigned addresses. This can be easily changed by biding to 0.0.0.0 instead of 127.0.0.1.

1. Start another WildFly instance again:

```
docker run -d --name wildfly -p 8080:8080 -p 9990:9990 arungupta/  
wildfly-management
```

In addition to application port 8080, the administration port 9990 is exposed as well. The WildFly image that is used has tweaked the start script such that the management interface is bound to 0.0.0.0.

2. Create a new server adapter in JBoss Developer Studio and name it 'WildFly 9.0.0-Management'. Specify the host name as 'dockerhost'.



3. Click on 'Next>' and change the values as shown.

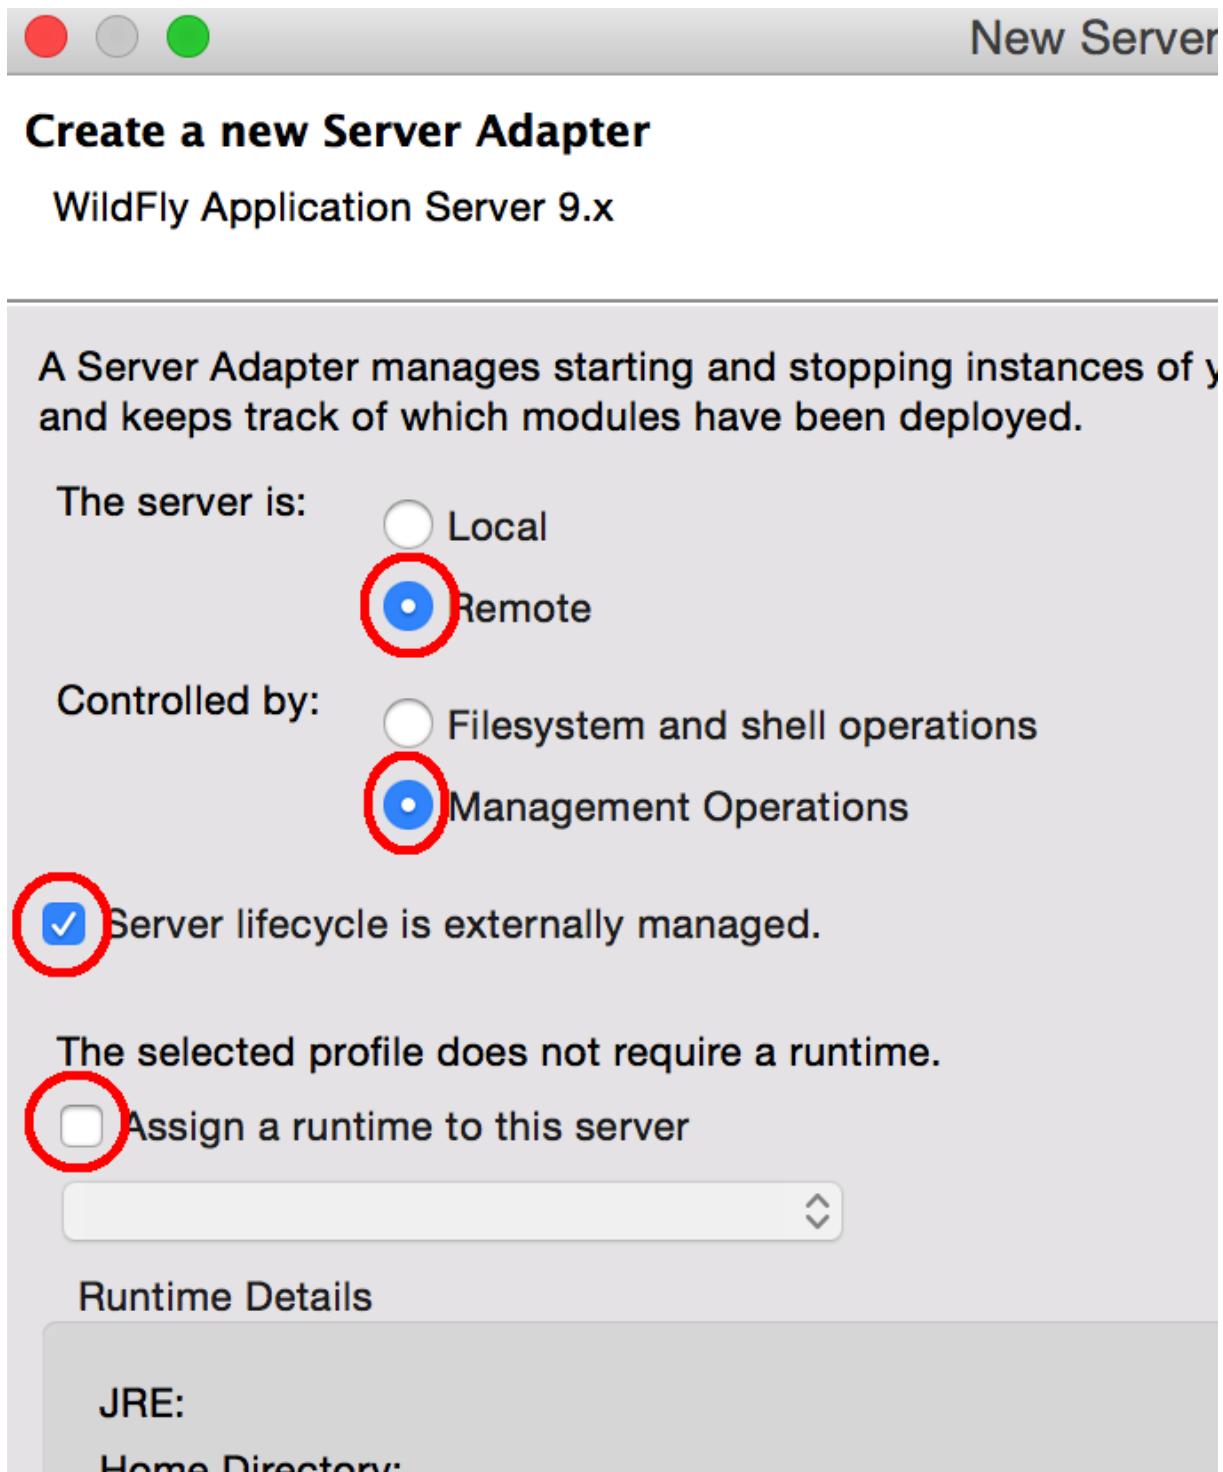


Figure 21. Create New Server Adapter

4. Take the default values in 'Remote System Integration' and click on 'Finish'.
5. Change server properties by double clicking on the newly created server adapter. Specify admin credentials (username: docker, password: docker#admin). Note, you need to delete the existing password and use this instead:

The screenshot shows the JBoss Central interface for managing a WildFly 9.0.0-Management server. The top navigation bar includes links for JBoss Central and WildFly 9.0.0-Management. The main content area is titled 'Overview' and contains a 'General Information' section where the host name is set to 'dockerhost'. Below this is a link to 'Open launch configuration' and a 'Runtime Environment' section indicating '(No Runtime)'. A red box highlights the 'Management Login Credentials' section, which contains fields for 'User Name' (set to 'docker') and 'Password' (represented by a series of dots). The 'Server Behavior' section includes a 'Behavior Profile' set to 'Remote (Use Management Operations)' and three checked checkboxes: 'Server lifecycle is externally managed.', 'Listen on all interfaces to allow remote web connections', and 'Expose your management port as the server's hostname'. At the bottom, there is a 'Host' dropdown menu with a 'New Host..' button.

JBoss Central *WildFly 9.0.0-Management X

Overview

General Information

Specify the host name and other common settings.

Server name: WildFly 9.0.0-Management

Host name: dockerhost

[Open launch configuration](#)

[Runtime Environment:](#) (No Runtime)

Management Login Credentials

Set the management login and password for your server.
This is used by all management commands, and during server shutdown.

User Name: docker

Password: [REDACTED]

Server Behavior

Behavior Profile: Remote (Use Management Operations)

Server lifecycle is externally managed.

Listen on all interfaces to allow remote web connections

Expose your management port as the server's hostname

Host

Figure 22. Management Login Credentials

6. Right-click on the newly created server adapter and click `Start''`. Status quickly changes to `Started` as shown.

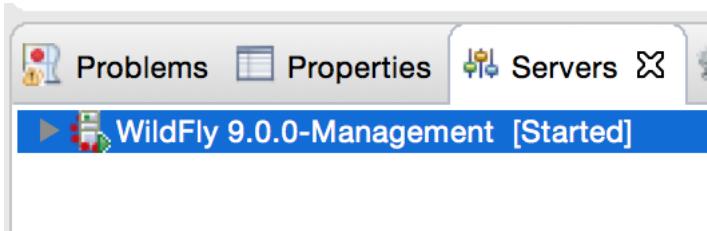


Figure 23. Synchronized WildFly Server

7. Right-click on the javaee7-simple-sample project, select ``Run on Server'' and choose this server. The project runs and displays the start page of the application.
8. Stop WildFly when you're done.

```
docker stop wildfly
```

8. Docker Maven Plugin

Maven plugin allows you to manage Docker images and containers from `pom.xml`. It comes with predefined goals:

Goal	Description
<code>docker:start</code>	Create and start containers
<code>docker:stop</code>	Stop and destroy containers
<code>docker:build</code>	Build images
<code>docker:push</code>	Push images to a registry
<code>docker:remove</code>	Remove images from local docker host
<code>docker:logs</code>	Show container logs

8.1. Run Java EE Application

1. Clone the workspace as:

```
git clone https://github.com/javaee-samples/javaee7-docker-maven.git
```

2. Build the image as:

```
mvn package -Pdocker
```

3. Verify the image as:

```
docker images
REPOSITORY           TAG      IMAGE ID
CREATED             VIRTUAL SIZE
arungupta/javaee7-docker-maven    latest   2e51b3fca40f
4 seconds ago       581.5 MB
```

4. Run the container as:

```
mvn install -Pdocker
```

5. Access your application at <http://dockerhost:8080/javaee7-docker-maven/resources/persons>. It shows the output as:

```
curl http://dockerhost:8080/javaee7-docker-maven/resources/persons
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>
<collection><person><name>Penny</name></person><person><name>Leonard</name></person><person><name>Sheldon</name></person><person><name>Amy</name></person><person><name>Howard</name></person><person><name>Bernadette</name></person><person><name>Raj</name></person><person><name>Priya</name></person></collection>
```

8.2. Understand Plugin Configuration

`pom.xml` is updated to include docker-maven-plugin as:

```
<plugin>
  <groupId>org.jolokia</groupId>
  <artifactId>docker-maven-plugin</artifactId>
  <version>0.11.5</version>
  <configuration>
    <images>
      <image>
        <alias>user</alias>
        <name>arungupta/javaee7-docker-maven</name>
        <build>
          <from>arungupta/wildfly:8.2</from>
          <assembly>
            <descriptor>assembly.xml</descriptor>
            <basedir>/</basedir>
          </assembly>
        <ports>
          <port>8080</port>
        </ports>
      </image>
    </images>
  </configuration>
</plugin>
```

```
</build>
<run>
  <ports>
    <port>8080:8080</port>
  </ports>
</run>
</image>
</images>
</configuration>
<executions>
  <execution>
    <id>docker:build</id>
    <phase>package</phase>
    <goals>
      <goal>build</goal>
    </goals>
  </execution>
  <execution>
    <id>docker:start</id>
    <phase>install</phase>
    <goals>
      <goal>start</goal>
    </goals>
  </execution>
</executions>
</plugin>
```

Each image configuration has three parts:

1. Image name and alias
2. `<build>` that defines how the image is created. Base image, build artifacts and their dependencies, ports to be exposed, etc to be included in the image are specified here. [Assembly descriptor format¹⁸](#) is used to specify the artifacts to be included and is defined in the `src/main/docker` directory.

`assembly.xml` in our case looks like:

```
<assembly . . .>
  <id>javaee7-docker-maven</id>
  <dependencySets>
    <dependencySet>
      <includes>
        <include>org.javaee7.sample:javaee7-docker-maven</include>
```

¹⁸ <http://maven.apache.org/plugins/maven-assembly-plugin/assembly.html>

```
</includes>
<outputDirectory>/opt/jboss/wildfly/standalone/deployments/</
outputDirectory>
<outputFileNameMapping>javaee7-docker-maven.war</
outputFileNameMapping>
</dependencySet>
</dependencySets>
</assembly>
```

3. `<run>` that defines how the container is run. Ports that need to be exposed are specified here.

In addition, `package` phase is tied to `docker:build` goal and `install` phase is tied to `docker:start` goal.

9. Docker Tools in Eclipse

The Docker tooling is aimed at providing at minimum the same basic level features as the command-line interface, but also provide some advantages by having access to a full fledged UI.

9.1. Install Docker Tools Plugins

As this is still in early access stage, you will have to install it first:

1. Download and Install [JBoss Developer Studio 9.0 Beta 2¹⁹](#), take defaults through out the installation.
Alternatively, download [Eclipse Mars latest build²⁰](#) and configure JBoss Tools plugin from the update site <http://download.jboss.org/jbosstools/updates/nightly/mars/>.
2. Open JBoss Developer Studio 9.0 Nightly
3. Add a new site using the menu items: `Help'' > Install New Software...'' > Add...''`. Specify the `Name:"` as `Docker Nightly''` and `Location:"` as <http://download.eclipse.org/linuxtools/updates-docker-nightly/>.

¹⁹ <http://classroom.example.com:8082/downloads/jboss-devstudio-9.0.0.Beta2-v20150609-1026-B3346-installer-standalone.jar>

²⁰ <http://www.eclipse.org/downloads/index-developer-default.php>

Docker and Kubernetes for Java EE Developers




Figure 24. Add Docker Tooling To JBoss Developer Studio

4. Expand Linux Tools, select 'Docker Client' and 'Docker Tooling'.

Available Software

Check the items that you wish to install.

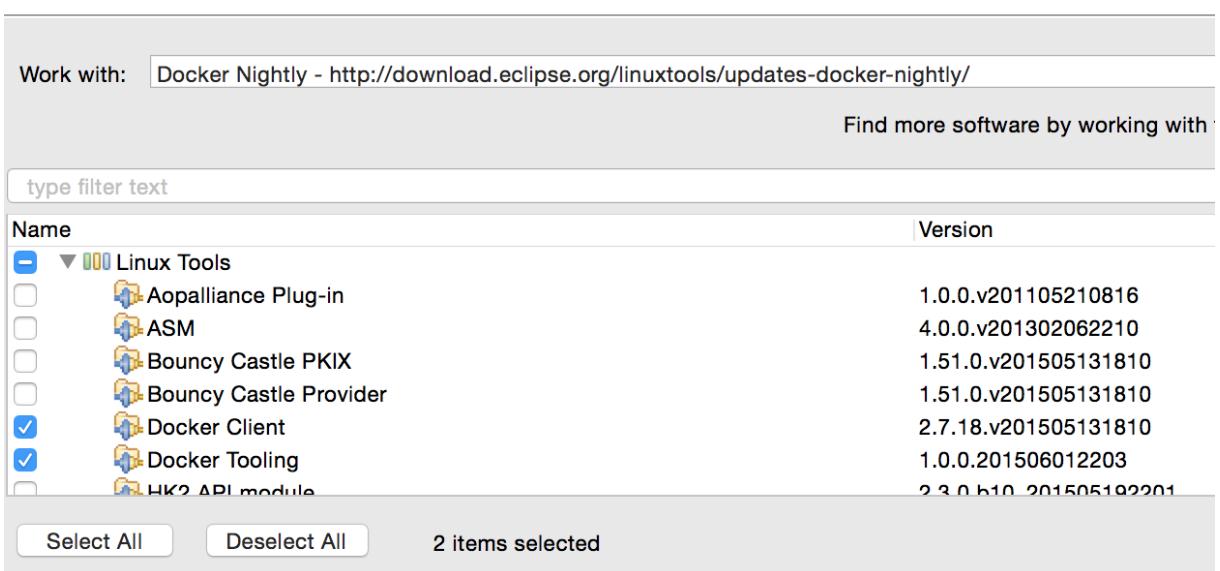


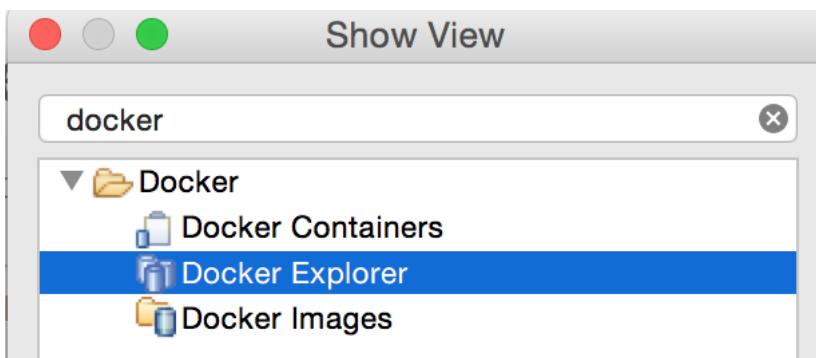
Figure 25. Add Docker Tooling

5. Click on 'Next >'' and 'Finish'. This will complete the installation of plugins. Restart the IDE for the changes to take effect.

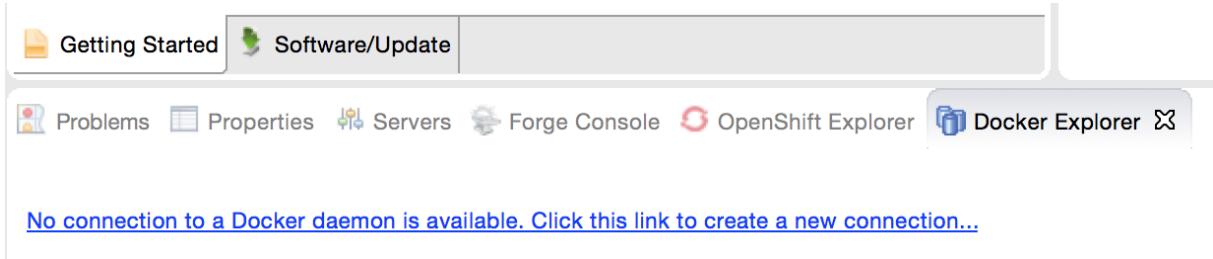
9.2. Docker Explorer

The Docker Explorer provides a wizard to establish a new connection to a Docker daemon. This wizard can detect default settings if the user's machine runs Docker natively (such as in Linux) or in a VM using Boot2Docker (such as in Mac or Windows). Both Unix sockets on Linux machines and the REST API on other OSes are detected and supported. The wizard also allows remote connections using custom settings.

1. Use the menu `Window`'', `Show View`'', `Other...`''. Type `docker`" to see the output as:



2. Select ``Docker Explorer" to open Docker Explorer.



3. Click on the link in this window to create a connection to Docker Host. Specify the settings as shown:




Figure 26. Docker Explorer

Make sure to get IP address of the Docker Host as:

```
docker-machine ip lab
```

Also, make sure to specify the correct directory for `.docker` on your machine.

4. Click on ``Test Connection" to check the connection. This should show the output as:

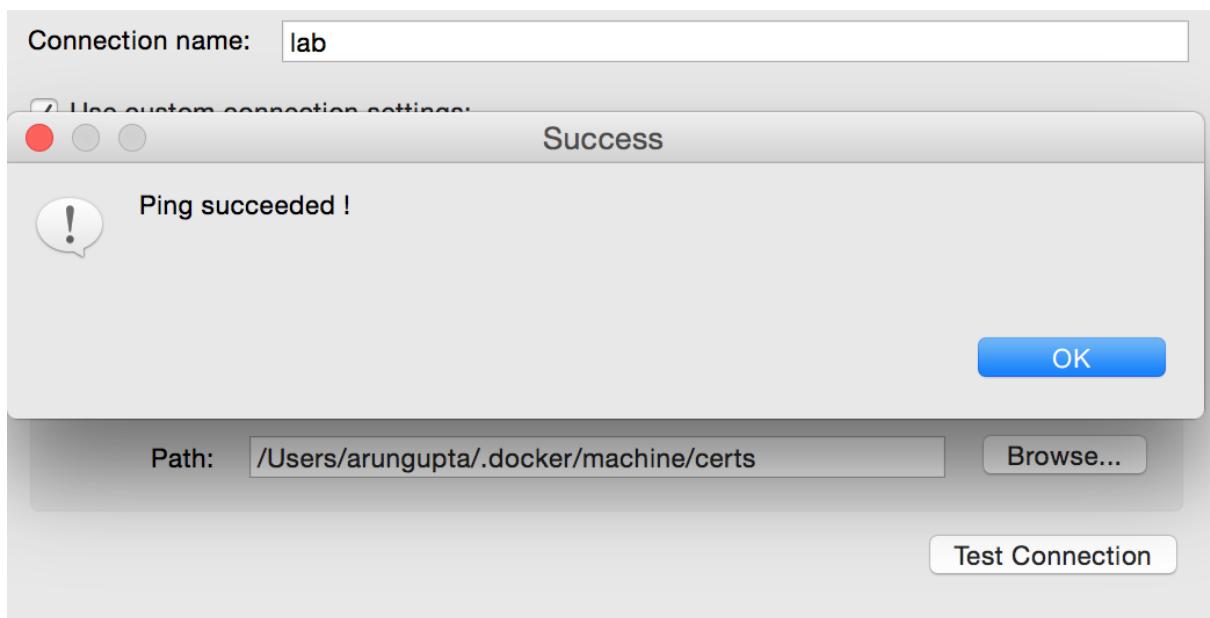


Figure 27. Docker Explorer

Click on 'OK' and 'Finish' to exit out of the wizard.

5. Docker Explorer itself is a tree view that handles multiple connections and provides users with quick overview of the existing images and containers.

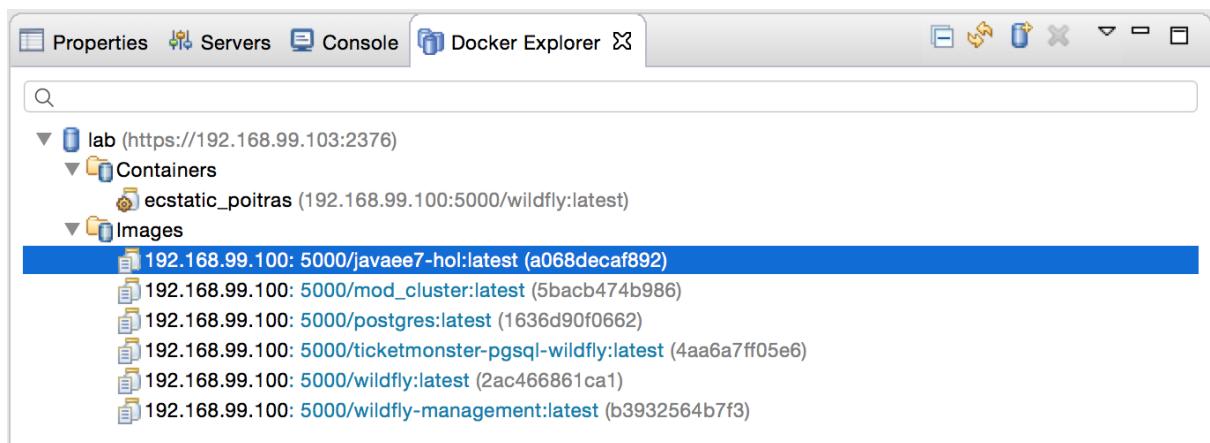


Figure 28. Docker Explorer Tree View

6. Customize the view by clicking on the arrow in toolbar:




Figure 29. Docker Explorer Customize View

Built-in filters can show/hide intermediate and 'dangling' images, as well as stopped containers.




Figure 30. Docker Explorer Customize View Wizard

9.3. Docker Images

The Docker Images view lists all images in the Docker host selected in the Docker Explorer view. This view allows user to manage images, including:

1. Pull/push images from/to the Docker Hub Registry (other registries will be supported as well, #469306²¹)
2. Build images from a Dockerfile
3. Create a container from an image

Lets take a look at it.

1. Use the menu `Window ''`, `Show View`", `Other...''`, select `Docker Images`". It shows the list of images on Docker Host:

Id	Repo Tags	Created	Virtual Size
4aa6a7ff05e6fd90d9d2a3c9f8d871...	192.168.99.100:5000/ticketmonster-pgsql-wildfl...	2015-06-02	1.1 GB
b3932564b7f3b9753808d8533b82...	192.168.99.100:5000/wildfly-management:latest	2015-06-02	951.3 MB
1636d90f0662e567463a158b71581...	192.168.99.100:5000/postgres:latest	2015-05-26	214 MB
2ac466861ca121d4c5e17970f4939...	192.168.99.100:5000/wildfly:latest	2015-03-13	951.3 MB
5bacb474b986660cef08b9f4f335c7...	192.168.99.100:5000/mod_cluster:latest	2015-01-13	492.5 MB
a068decaf8928737340f8f08fbddf9...	192.168.99.100:5000/javaee7-hol:latest	2014-12-05	619.7 MB

Figure 31. Docker Images View

2. Right-click on the image ending with ``wildfly:latest" and click on the green arrow in the toolbar. This will show the following wizard:

²¹ https://bugs.eclipse.org/bugs/show_bug.cgi?id=469306

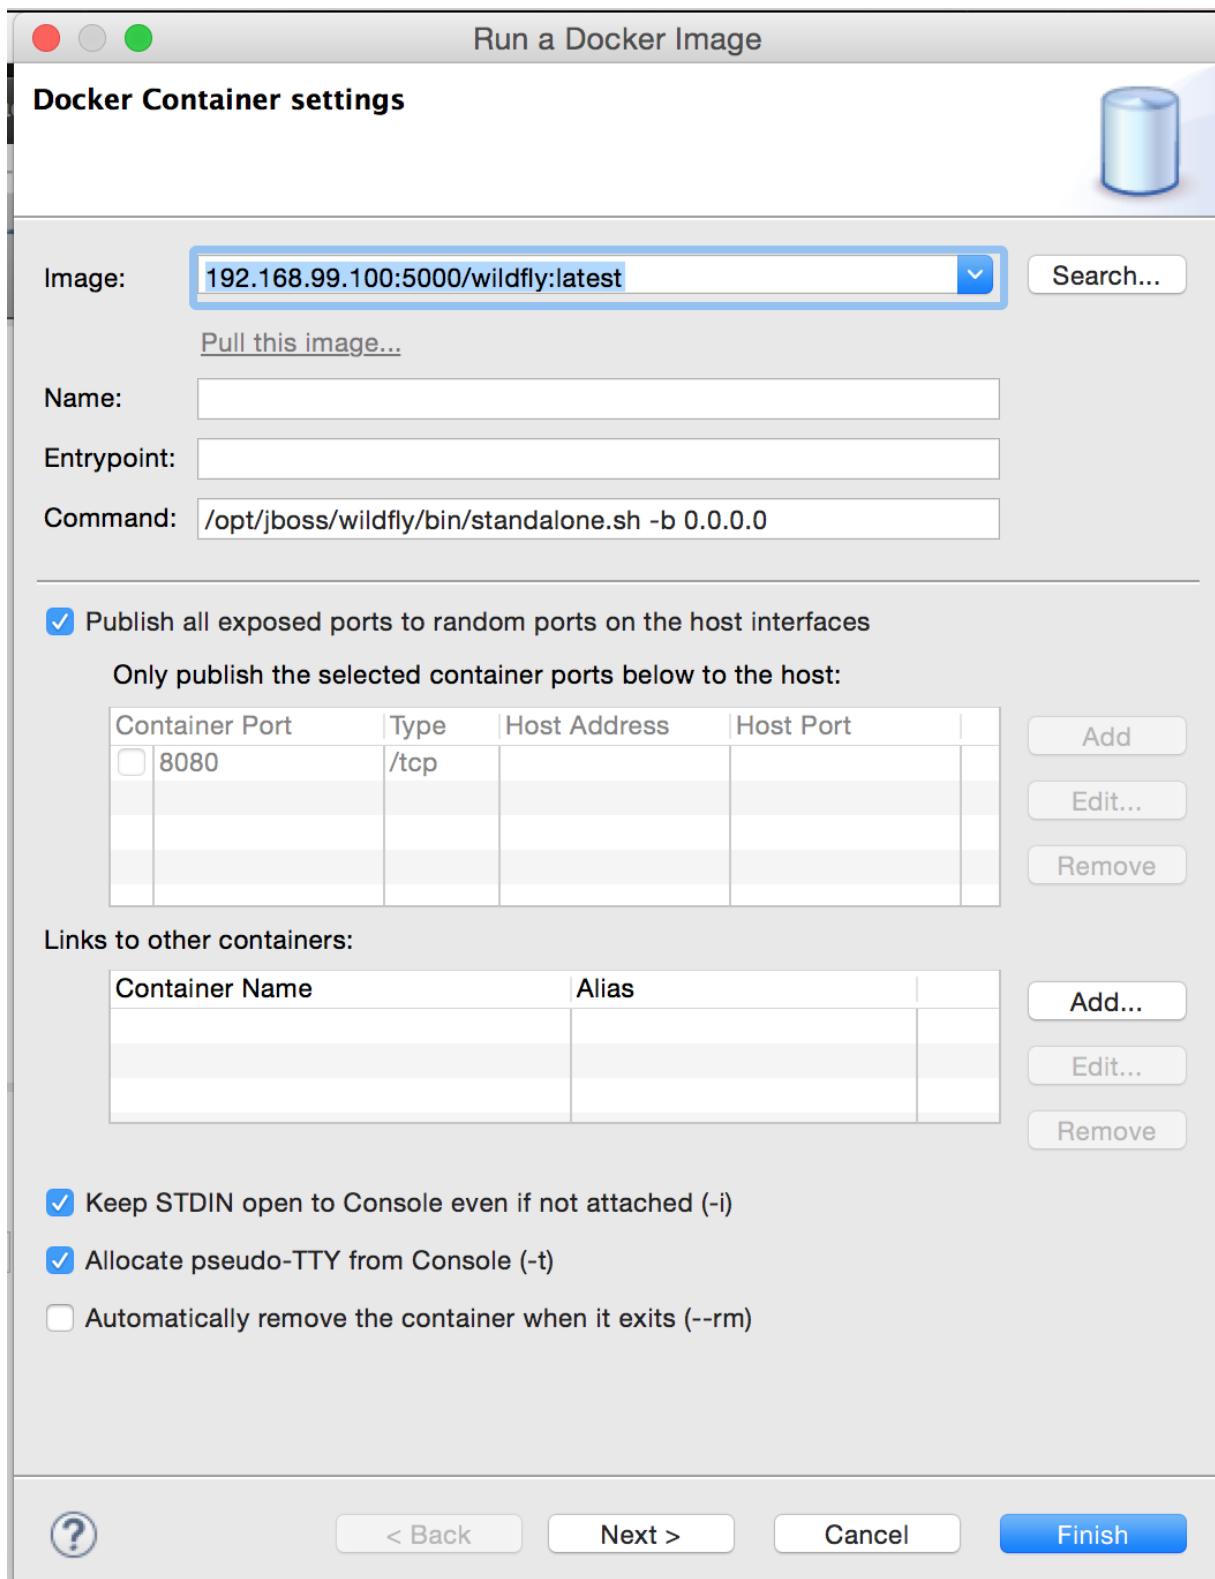
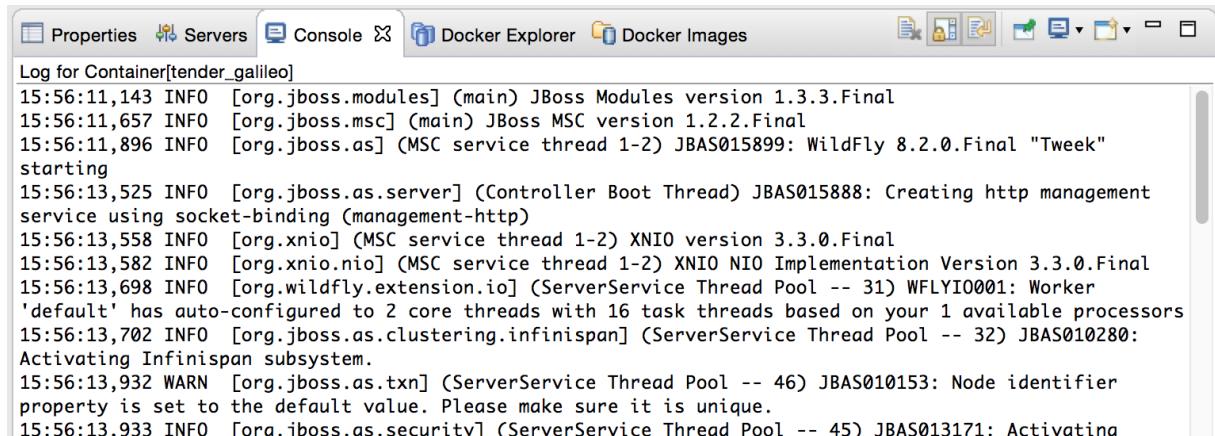


Figure 32. Docker Run Container Wizard

By default, all exports ports from the image are mapped to random ports on the host interface. This setting can be changed by unselecting the first checkbox and specify exact port mapping.

Click on ``Finish" to start the container.

- When the container is started, all logs are streamed into Eclipse Console:



The screenshot shows the Eclipse IDE interface with the Docker Explorer view selected. The title bar says "Log for Container[tender_galileo]". The log output is as follows:

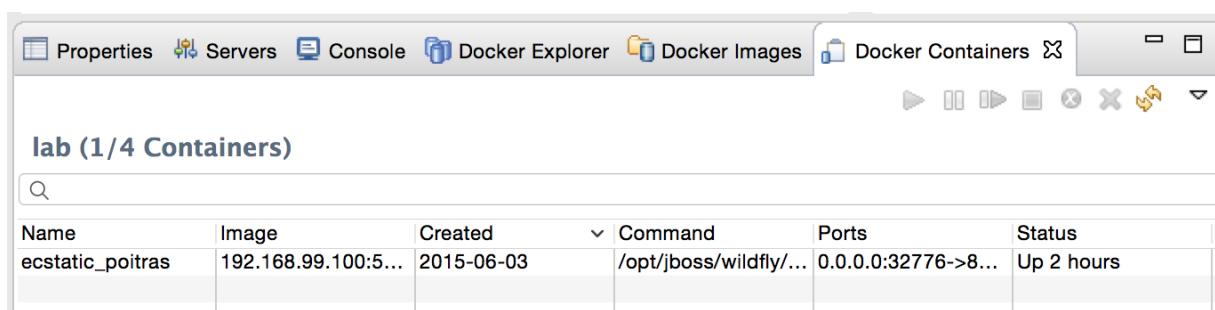
```
15:56:11,143 INFO [org.jboss.modules] (main) JBoss Modules version 1.3.3.Final
15:56:11,657 INFO [org.jboss.msc] (main) JBoss MSC version 1.2.2.Final
15:56:11,896 INFO [org.jboss.as] (MSC service thread 1-2) JBAS015899: WildFly 8.2.0.Final "Tweek"
Starting
15:56:13,525 INFO [org.jboss.as.server] (Controller Boot Thread) JBAS015888: Creating http management
service using socket-binding (management-http)
15:56:13,558 INFO [org.xnio] (MSC service thread 1-2) XNIO version 3.3.0.Final
15:56:13,582 INFO [org.xnio.nio] (MSC service thread 1-2) XNIO NIO Implementation Version 3.3.0.Final
15:56:13,698 INFO [org.wildfly.extension.io] (ServerService Thread Pool -- 31) WFLYIO0001: Worker
'default' has auto-configured to 2 core threads with 16 task threads based on your 1 available processors
15:56:13,702 INFO [org.jboss.as.clustering.infinispan] (ServerService Thread Pool -- 32) JBAS010280:
Activating Infinispan subsystem.
15:56:13,932 WARN [org.jboss.as.txn] (ServerService Thread Pool -- 46) JBAS010153: Node identifier
property is set to the default value. Please make sure it is unique.
15:56:13,933 INFO [org.jboss.as.security] (ServerService Thread Pool -- 45) JBAS013171: Activating
```

Figure 33. Docker Container Logs

9.4. Docker Containers

Docker Containers view lets the user manage the containers. The view toolbar provides commands to start, stop, pause, unpause, display the logs and kill containers.

- Use the menu "Window", "Show View", "Other...", select "Docker Containers". It shows the list of running containers on Docker Host:



The screenshot shows the Eclipse IDE interface with the Docker Containers view selected. The title bar says "Docker Containers". The table displays one container named "ecstatic_poitras" with the following details:

Name	Image	Created	Command	Ports	Status
ecstatic_poitras	192.168.99.100:5...	2015-06-03	/opt/jboss/wildfly/...	0.0.0.0:32776->8...	Up 2 hours

Figure 34. Docker Containers View

- Pause the container by clicking on the "pause" button in the toolbar (#469310²²). Show the complete list of containers by clicking on the "View Menu", "Show all containers".

²² https://bugs.eclipse.org/bugs/show_bug.cgi?id=469310

Name	Image	Created	Command	Ports	Status
ecstatic_poitras	192.168.99.100:500... 192.168.99.100:500...	2015-06-03	/opt/jboss/wildfly/bi...	0.0.0.0:32776->808...	Up 9 minutes (Paused)
high_fermi	jboss/wildfly:latest	2015-06-03	/opt/jboss/wildfly/bi...		Exited (130) 15 minu...
tender_galileo	192.168.99.100:500... 192.168.99.100:500...	2015-06-03	/opt/jboss/wildfly/bi...		Exited (137) 26 minu...
furious_goldstine	192.168.99.100:500...	2015-06-03	/opt/jboss/wildfly/bi...		Exited (137) 52 minu...

Figure 35. All Docker Containers

3. Select the paused container, and click on the green arrow in the toolbar to restart the container.
4. Right-click on any running container and select ``Display Log'' to view the log for this container.

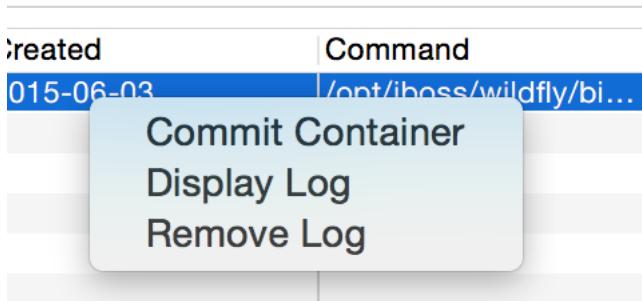


Figure 36. Eclipse Properties View

TODO: Users can also attach an Eclipse console to a running Docker container to follow the logs and use the STDIN to interact with it.

9.5. Details on Images and Containers

Eclipse Properties view is used to provide more information about the containers and images.

1. Just open the Properties View and click on a Connection, Container, or Image in any of the Docker Explorer View, Docker Containers View, or Docker Images View. This will fill in data in the Properties view.
Info view is shown as:

The screenshot shows the 'Properties' tab selected in the top navigation bar. Below it, the title bar reads 'ecstatic_poitras (192.168.99.100:5000/wildfly:latest)'. On the left, there are two tabs: 'Info' (selected) and 'Inspect'. The main area is a table with columns 'Property' and 'Value'. The properties listed are Id, Image, Command, Created, Status, Ports, and Names.

Property	Value
Id	24fc99c3ca92
Image	192.168.99.100:5000/wildfly:latest
Command	/opt/jboss/wildfly/bin/standalone.sh -b 0.0.0.0
Created	2015-06-03
Status	Up 23 minutes
Ports	0.0.0.0:32776->8080/tcp
Names	ecstatic_poitras

Figure 37. Docker Container Properties View Info

Inspect view is shown as:

The screenshot shows the 'Properties' tab selected in the top navigation bar. Below it, the title bar reads 'ecstatic_poitras (192.168.99.100:5000/wildfly:latest)'. On the left, there are two tabs: 'Info' (selected) and 'Inspect'. The main area is a table with columns 'Property' and 'Value'. The properties listed are Id, Name, Created, State (with sub-properties ExitCode, Finished at, Running, Paused, Pid), Args, Driver, ExecDriver, and Config (with sub-property AttachStderr).

Property	Value
Id	24fc99c3ca92
Name	/ecstatic_poitras
Created	2015-06-03
State	<ul style="list-style-type: none">ExitCode: 0Finished at: 2001-12-31T16:00:00.000PSTRunning: truePaused: falsePid: 14271
Args	<ul style="list-style-type: none">-b 0.0.0.0
Driver	aufs
ExecDriver	native-0.2
Config	<ul style="list-style-type: none">AttachStderr: false

Figure 38. Docker Container Properties View Inspect

10. Test Java EE Applications on Docker

Testing Java EE applications is a very important aspect. Especially when it comes to in-container tests, [JBoss Arquillian²³](#) is well known to make this very easy for Java EE application servers. Picking up where unit tests leave off, Arquillian handles all the plumbing of container management, deployment and framework initialization so you can focus on the task at hand, writing your tests.

With Arquillian, you can use [WildFly remote container adapter²⁴](#) and connect to any WildFly instance running in a Docker container. But this wouldn't help with the Docker container lifecycle management.

[Arquillian Cube²⁵](#), an extension of Arquillian, allows you to control the lifecycle of Docker images as part of the test lifecycle, either automatically or manually. This extension allows to start a Docker container with a server installed, deploy the required deployable file within it and execute Arquillian tests.

The key point here is that if Docker is used as deployable platform in production, your tests are executed in the same container as it will be in production, so your tests are even more real than before.

1. Check out the workspace:

```
git clone http://github.com/javaee-samples/javaee-arquillian-cube
```

2. Edit `src/test/resources/arquillian.xml` file and change the IP address specified in `serverUri` property value to point to your Docker host's IP. This can be found out as:

```
docker-machine ip lab
```

3. Run the tests as:

```
mvn test
```

This will create a container using the image defined in `src/test/resources/wildfly/Dockerfile`. The container qualifier in `arquillian.xml` defines the directory name in `src/test/resources` directory.

²³ <http://www.arquillian.org>

²⁴ <http://arquillian.org/modules/wildfly-arquillian-wildfly-remote-container-adapter/>

²⁵ <http://arquillian.org/modules/cube-extension/>



A pre-built image can be used by specifying:

```
wildfly:  
image: jboss/wildfly
```

instead of

```
wildfly:  
buildImage:  
dockerfileLocation: src/test/resources/wildfly
```

By default, the ``cube'' profile is activated and this includes all the required dependencies.

The result is shown as:

```
Running org.javaee7.sample.PersonDatabaseTest  
Jun 16, 2015 9:23:04 AM org.jboss.arquillian.container.impl.MapObject  
populate  
WARNING: Configuration contain properties not supported by the backing  
object  
org.jboss.as.arquillian.container.remote.RemoteContainerConfiguration  
Unused property entries: {target=wildfly:8.1.0.Final:remote}  
Supported property names: [managementAddress, password, managementPort,  
managementProtocol, username]  
Jun 16, 2015 9:23:13 AM org.xnio.Xnio <clinit>  
INFO: XNIO version 3.2.0.Beta4  
Jun 16, 2015 9:23:13 AM org.xnio.nio.NioXnio <clinit>  
INFO: XNIO NIO Implementation Version 3.2.0.Beta4  
Jun 16, 2015 9:23:13 AM org.jboss.remoting3.EndpointImpl <clinit>  
INFO: JBoss Remoting version (unknown)  
Tests run: 1, Failures: 0, Errors: 0, Skipped: 0, Time elapsed: 16.406  
sec - in org.javaee7.sample.PersonDatabaseTest
```

Results :

```
Tests run: 1, Failures: 0, Errors: 0, Skipped: 0
```

```
[INFO]
```

```
-----  
[INFO] BUILD SUCCESS
```

[INFO]

4. In `arquillian.xml`, add the following property:

```
<property name="connectionMode">STARTORCONNECT</property>
```

This bypasses the create/start Cube commands if a Docker Container with the same name is already running on the target system.

This allows you to prestart the containers manually during development and just connect to them to avoid the extra cost of starting the Docker Containers for each test run. This assumes you are not changing the actual definition of the Docker Container itself.

11. Multiple Containers Using Docker Compose

Docker Compose is a tool for defining and running complex applications with Docker. With Compose, you define a multi-container application in a single file, then spin your application up in a single command which does everything that needs to be done to get it running.

— github.com/docker/compose

An application using Docker containers will typically consist of multiple containers. With Docker Compose, there is no need to write shell scripts to start your containers. All the containers are defined in a configuration file using *services*, and then `docker-compose` script is used to start, stop, and restart the application and all the services in that application, and all the containers within that service. The complete list of commands is:

Command	Purpose
<code>build</code>	Build or rebuild services
<code>help</code>	Get help on a command
<code>kill</code>	Kill containers
<code>logs</code>	View output from containers
<code>port</code>	Print the public port for a port binding
<code>ps</code>	List containers

Command	Purpose
pull	Pulls service images
restart	Restart services
rm	Remove stopped containers
run	Run a one-off command
scale	Set number of containers for a service
start	Start services
stop	Stop services
up	Create and start containers

Docker Compose script is only available for OSX and Linux. <https://github.com/arun-gupta/docker-java/issues/3> is used for tracking Docker Compose on Windows.

11.1. Configuration File

1. Entry point to Compose is `docker-compose.yml`. Lets use the following file:

```
mysqldb:
  image: classroom.example.com:5000/mysql
  environment:
    MYSQL_DATABASE: sample
    MYSQL_USER: mysql
    MYSQL_PASSWORD: mysql
    MYSQL_ROOT_PASSWORD: supersecret
mywildfly:
  image: classroom.example.com:5000/wildfly-mysql-javaee7
  links:
    - mysqldb:db
  ports:
    - 8080:8080
```

This file is available in `../attendees/docker-compose.yml` and shows:

- Two services are defined by the name `mysqldb` and `mywildfly`
- Image name for each service is defined using `image`
- Environment variables for the MySQL container are defined in `environment`
- MySQL container is linked with WildFly container using `links`
- Port forwarding is achieved using `ports`

11.2. Start Services

1. All the services can be started, in detached mode, by giving the command:

```
.....  
docker-compose up -d
```

And this shows the output as:

```
.....  
Creating attendees_mysql_db_1...  
Creating attendees_mywildfly_1...
```

An alternate compose file name can be specified using `-f`.

An alternate directory where the compose file exists can be specified using `-p`.

2. Started services can be verified as:

```
> docker-compose ps  
.....  
      Name          Command           State  
      Ports  
-----  
attendees_mysql_db_1    /entrypoint.sh mysqld      Up      3306/  
tcp  
attendees_mywildfly_1   /opt/jboss/wildfly/customi ...      Up  
0.0.0.0:8080->8080/tcp, 9990/tcp
```

This provides a consolidated view of all the services started, and containers within them.

Alternatively, the containers in this application, and any additional containers running on this Docker host can be verified by using the usual `docker ps` command:

```
.....  
> docker ps  
CONTAINER ID        IMAGE               COMMAND  
      CREATED             STATUS              PORTS  
      NAMES  
3598e545bd2f        arungupta/wildfly-mysql-javae7:latest   "/opt/  
jboss/wildfly/      59 seconds ago       Up 58 seconds       0.0.0.0:8080-  
>8080/tcp, 9990/tcp   attendees_mywildfly_1
```

Docker and Kubernetes for Java EE Developers

```
b8cf6a3d518b      mysql:latest           "/  
entrypoint.sh mysq  2 minutes ago        Up 2 minutes      3306/tcp  
          attendees_mysqldb_1
```

3. Service logs can be seen as:

```
> docker-compose logs  
Attaching to attendees_mywildfly_1, attendees_mysqldb_1  
mywildfly_1 | => Starting WildFly server  
mywildfly_1 | => Waiting for the server to boot  
mywildfly_1 |  
=====  
mywildfly_1 |  
mywildfly_1 |     JBoss Bootstrap Environment  
mywildfly_1 |  
mywildfly_1 |     JBOSS_HOME: /opt/jboss/wildfly  
mywildfly_1 |  
mywildfly_1 |     JAVA: /usr/lib/jvm/java/bin/java  
mywildfly_1 |  
mywildfly_1 |     JAVA_OPTS: -server -Xms64m -Xmx512m -  
XX:MaxPermSize=256m -Djava.net.preferIPv4Stack=true -  
Djboss.modules.system.pkgs=org.jboss.byteman -Djava.awt.headless=true  
mywildfly_1 |  
  
...  
  
mywildfly_1 | 15:40:20,866 INFO [org.jboss.resteasy.spi.ResteasyDeployment] (MSC service  
thread 1-2) Deploying javax.ws.rs.core.Application: class  
org.javaee7.samples.employees.MyApplication  
mywildfly_1 | 15:40:20,914 INFO [org.wildfly.extension.undertow] (MSC  
service thread 1-2) JBAS017534: Registered web context: /employees  
mywildfly_1 | 15:40:21,032 INFO [org.jboss.as.server] (ServerService  
Thread Pool -- 28) JBAS018559: Deployed "employees.war" (runtime-  
name : "employees.war")  
mywildfly_1 | 15:40:21,077 INFO [org.jboss.as] (Controller  
Boot Thread) JBAS015961: Http management interface listening on  
http://127.0.0.1:9990/management  
mywildfly_1 | 15:40:21,077 INFO [org.jboss.as] (Controller Boot  
Thread) JBAS015951: Admin console listening on http://127.0.0.1:9990  
mywildfly_1 | 15:40:21,077 INFO [org.jboss.as] (Controller Boot  
Thread) JBAS015874: WildFly 8.2.0.Final "Tweek" started in 9572ms  
- Started 280 of 334 services (92 services are lazy, passive or on-  
demand)  
mysqldb_1 | Running mysql_install_db
```

```
mysqladb_1 | 2015-06-05 15:38:31 0 [Note] /usr/sbin/mysqld (mysqld  
5.6.25) starting as process 27 ...  
mysqladb_1 | 2015-06-05 15:38:31 27 [Note] InnoDB: Using atomics to  
ref count buffer pool pages  
.  
.  
.  
mysqladb_1 | 2015-06-05 15:38:40 1 [Note] Event Scheduler: Loaded 0  
events  
mysqladb_1 | 2015-06-05 15:38:40 1 [Note] mysqld: ready for  
connections.  
mysqladb_1 | Version: '5.6.25' socket: '/var/run/mysqld/mysqld.sock'  
port: 3306 MySQL Community Server (GPL)  
mysqladb_1 | 2015-06-05 15:40:18 1 [Warning] IP address '172.17.0.24'  
could not be resolved: Name or service not known
```

11.3. Verify Application

1. Access the application at <http://dockerhost:8080/employees/resources/employees/>. This is shown in the browser as:



This XML file does not appear to have any style information associated with it.

```
<collection>
  <employee>
    <id>1</id>
    <name>Penny</name>
  </employee>
  <employee>
    <id>2</id>
    <name>Sheldon</name>
  </employee>
  <employee>
    <id>3</id>
    <name>Amy</name>
  </employee>
  <employee>
    <id>4</id>
    <name>Leonard</name>
  </employee>
  <employee>
    <id>5</id>
    <name>Bernadette</name>
  </employee>
  <employee>
    <id>6</id>
    <name>Raj</name>
  </employee>
  <employee>
    <id>7</id>
    <name>Howard</name>
  </employee>
  <employee>
    <id>8</id>
    <name>Priya</name>
  </employee>
</collection>
```

Figure 39. Output From Servers Run Using Docker Compose

11.4. Stop Services

1. Stop the services as:

```
> docker-compose stop
Stopping attendees_mywildfly_1...
Stopping attendees_mysql_1...
```

11.5. Scale Services

<https://github.com/arun-gupta/docker-java/issues/51>

12. Java EE Application on Mod Cluster

A frequent requirement for Java EE based applications is running them on a cluster of application server. While setup and test can be complicated on developer machines, this is where Docker can play to its full potential. With the help of images and automatic port mapping, we're ready to test Ticket Monster on a couple of WildFly instances and add and remove them randomly.

The diagram below shows what will be achieved in this section:




Figure 40. Standalone Cluster with WildFly and mod_cluster

1. Start Apache HTTPD server
-

```
docker run -d --name modcluster -p 80:80 classroom.example.com:5000/
mod_cluster
```

2. Open http://dockerhost/mod_cluster_manager in your browser to see the empty console as:



Figure 41. Apache HTTPD runing mod_cluster_manager interface

3. Start Postgres database container, if not already running:

```
docker run --name db -d -p 5432:5432 -e POSTGRES_USER=ticketmonster -  
e POSTGRES_PASSWORD=ticketmonster-docker classroom.example.com:5000/  
postgres
```

4. Start the first WildFly instance:

```
docker run -d --name server1 --link db:db --link modcluster:modcluster  
classroom.example.com:5000/ticketmonster-pgsql-wildfly
```

Besides linking the database container using `--link db:db`, we also link the ``modcluster'' container. This should be done rather quickly and if you now revisit the [mod_cluster_manager](#)²⁶ in your browser, then you can see that the first server was registered with the loadbalancer:

²⁶ http://dockerhost/mod_cluster_manager/

The screenshot shows a web browser window with two tabs: 'Mod_cluster Status' and 'Ticket Monster'. The 'Mod_cluster Status' tab is active, displaying the URL `dockerhost/mod_cluster_manager?nonce=b726fe5d-7300-4445-bbfe-c203d7d7c042&refresh=10`. The page content includes:

- mod_cluster/1.2.6.Final**
- [Auto Refresh](#) [show DUMP output](#) [show INFO output](#)
- Node b30e2c6af434 (ajp://172.17.0.14:8009):**
- [Enable Contexts](#) [Disable Contexts](#)
- Balancer: mycluster, LBGroup: , Flushpackets: Off, Flushwait: 10000, Ping: 10000000, Smax: 2, Ttl: 60000000, Status: OK, Elected: 0, Read: 0, Transferred: 0, Connected: 0, Load: 96
- Virtual Host 1:**
- Contexts:**
`/ticket-monster, Status: ENABLED Request: 0 Disable`
- Aliases:**
`default-host
localhost`

Figure 42. First WildFly instance registered with Load Balancer

5. To make sure the Ticket Monster application is also running just visit `http://dockerhost/ticket-monster` and you will be presented with the Ticket Monster welcome screen.

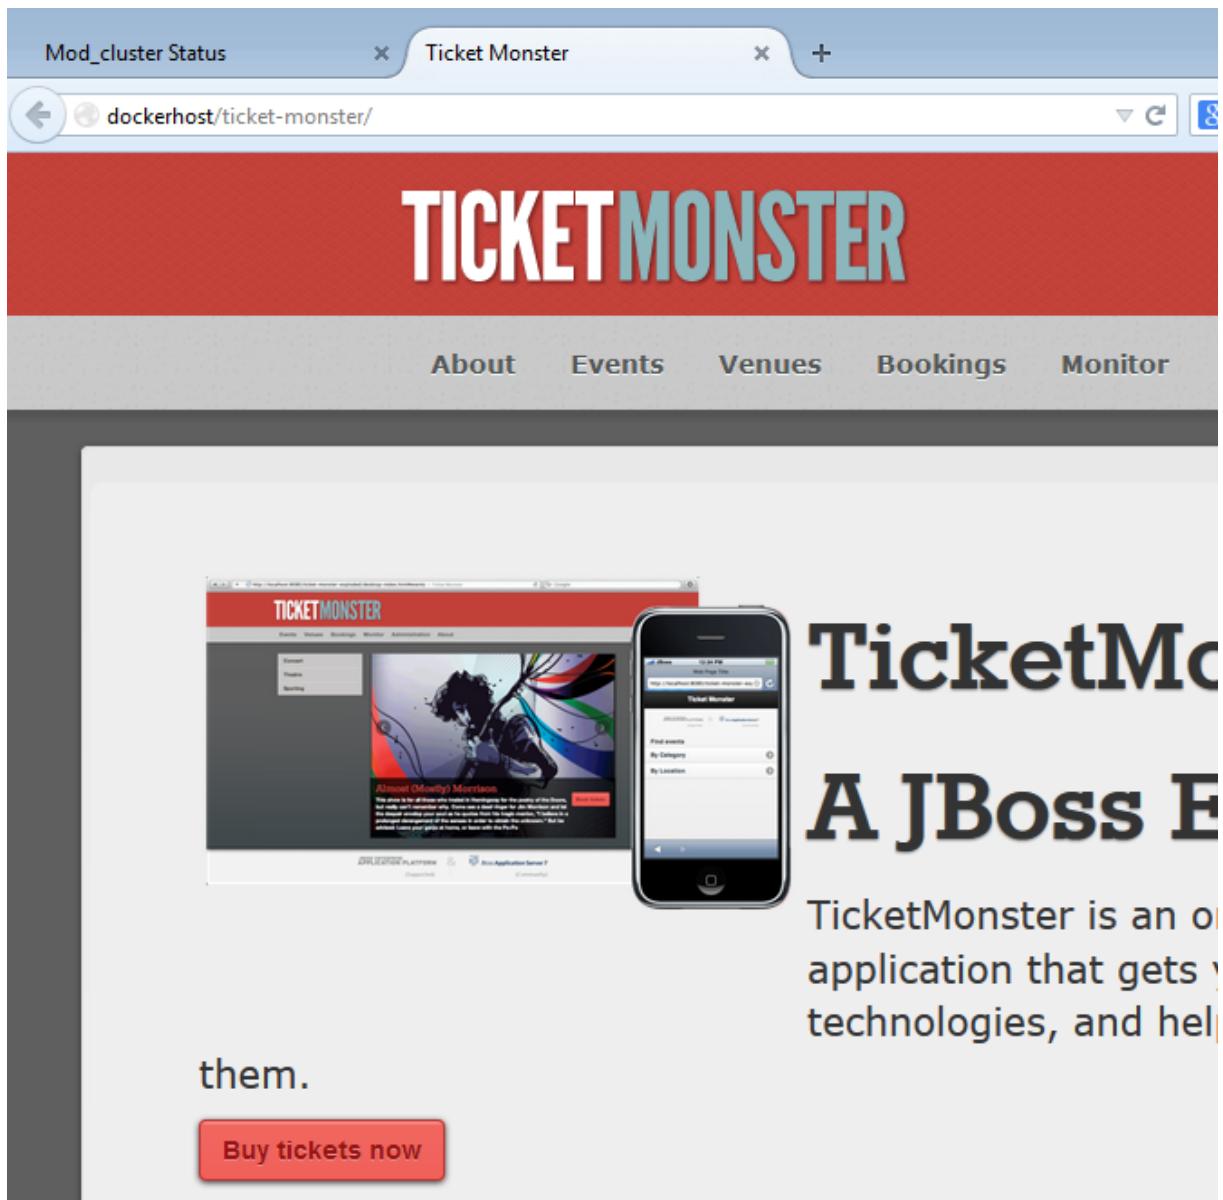


Figure 43. Clustered Ticket Monster Application

6. Start as many WildFly instances as you want (and your computer memory can handle):

```
docker run -d --name server2 --link db:db --link modcluster:modcluster
  classroom.example.com:5000/ticketmonster-pgsql-wildfly
docker run -d --name server3 --link db:db --link modcluster:modcluster
  classroom.example.com:5000/ticketmonster-pgsql-wildfly
docker run -d --name server4 --link db:db --link modcluster:modcluster
  classroom.example.com:5000/ticketmonster-pgsql-wildfly
```

7. Stop some servers:

```
docker stop server1
```

```
docker stop server3
```

Ensure that the application is still accessible at <http://dockerhost/ticket-monster>.

13. Java EE Application on Docker Swarm Cluster

Docker Swarm solves one of the fundamental limitations of Docker where the containers could only run on a single Docker host. Docker Swarm is native clustering for Docker. It turns a pool of Docker hosts into a single, virtual host.




Figure 44. Key Components of Docker Swarm

Swarm Manager: Docker Swarm has a Master or Manager, that is a pre-defined Docker Host, and is a single point for all administration. Currently only a single instance of manager is allowed in the cluster. This is a SPOF for high availability architectures and additional managers will be allowed in a future version of Swarm with #598. TODO: ADD LINK.

Swarm Nodes: The containers are deployed on Nodes that are additional Docker Hosts. Each Swarm Node must be accessible by the manager, each node must listen to the same network interface (TCP port). Each node runs a node agent that registers the referenced Docker daemon, monitors it, and updates the discovery backend with the node's status. The containers run on a node.

Scheduler Strategy: Different scheduler strategies (`binpack'', "spread"` (default), and ``random'') can be applied to pick the best node to run your container. The default strategy optimizes the node for least number of running containers. There are multiple kinds of filters, such as constraints and affinity. This should allow for a decent scheduling algorithm.

Node Discovery Service: By default, Swarm uses hosted discovery service, based on Docker Hub, using tokens to discover nodes that are part of a cluster. However etcd, consul, and zookeeper can be also be used for service discovery as well. This is particularly useful if there is no access to Internet, or you are running the setup in a closed network. A new discovery backend can be created as explained here. It would be useful to have the hosted Discovery Service inside the firewall and #660 will discuss this.

Standard Docker API: Docker Swarm serves the standard Docker API and thus any tool that talks to a single Docker host will seamlessly scale to multiple hosts now. That means if you were using shell scripts using Docker CLI to configure multiple Docker hosts, the same CLI would now talk to Swarm cluster and Docker Swarm will then act as proxy and run it on the cluster.

There are lots of other concepts but these are the main ones.

1. Create a Swarm cluster. The easiest way of using Swarm is, by using the official Docker image:

```
docker run swarm create
```

This command returns a <TOKEN> and is the unique cluster id. It will be used when creating master and nodes later. This cluster id is returned by the hosted discovery service on Docker Hub.



Make sure to note this cluster id now as there is no means to list it later.

2. Swarm is fully integrated with Docker Machine, and so is the easiest way to get started. Let's create a Swarm Master next:

```
docker-machine create -d virtualbox --swarm --swarm-master --swarm-discovery token://<TOKEN> swarm-master
```

Replace <TOKEN> with the cluster id obtained in the previous step.

--swarm configures the machine with Swarm, --swarm-master configures the created machine to be Swarm master. Swarm master creation talks to the hosted service on Docker Hub and informs that a master is created in the cluster.

3. Connect to this newly created master and find some more information about it:

```
eval "$(docker-machine env swarm-master)"  
docker info
```



If you're on Windows, use the docker-machine env swarm-master command only and copy the output into an editor to replace all appearances of EXPORT with SET and issue the three commands at your command prompt, remove the quotes and all duplicate appearances of "/".

This will show the output as:

```
> docker info  
Containers: 2  
Images: 7  
Storage Driver: aufs  
  Root Dir: /mnt/sda1/var/lib/docker/aufs  
  Backing Filesystem: extfs  
  Dirs: 11  
  Dirperm1 Supported: true  
Execution Driver: native-0.2  
Kernel Version: 4.0.3-boot2docker  
Operating System: Boot2Docker 1.6.2 (TCL 5.4); master : 4534e65 - Wed  
  May 13 21:24:28 UTC 2015  
CPUs: 1  
Total Memory: 998.1 MiB  
Name: swarm-master  
ID: USSA:35LS:WVRN:GKAZ:MLD4:XMQF:P7PL:VQ5D:7V4K:2QAH:5D2L:HC4K  
Debug mode (server): true  
Debug mode (client): false  
Fds: 24  
Goroutines: 37  
System Time: Wed Jun 10 03:40:00 UTC 2015  
EventsListeners: 1  
Init SHA1: 7f9c6798b022e64f04d2aff8c75cbf38a2779493  
Init Path: /usr/local/bin/docker  
Docker Root Dir: /mnt/sda1/var/lib/docker  
Username: arungupta  
Registry: [https://index.docker.io/v1/]
```

Labels:
provider=virtualbox

4. Create Swarm nodes.

```
docker-machine create -d virtualbox --swarm --swarm-discovery token://  
<TOKEN> swarm-node-01
```

Replace `<TOKEN>` with the cluster id obtained in the previous step.

Node creation talks to the hosted service at Docker Hub and joins the previously created cluster. This is specified by `--swarm-discovery token://...` and specifying the cluster id obtained earlier.

5. To make it a real cluster, let's create a second node:

```
docker-machine create -d virtualbox --swarm --swarm-discovery token://  
<TOKEN> swarm-node-02
```

Replace `<TOKEN>` with the cluster id obtained in the previous step.

6. List all the nodes / Docker machines, that has been created so far.

```
docker-machine ls
```

This shows the output as:

NAME	ACTIVE	DRIVER	STATE	URL
SWARM				
lab		virtualbox	Running	tcp://192.168.99.103:2376
swarm-master	*	virtualbox	Running	tcp://192.168.99.107:2376
swarm-node-01		virtualbox	Running	tcp://192.168.99.108:2376
swarm-node-02		virtualbox	Running	tcp://192.168.99.109:2376

The machines that are part of the cluster have the cluster's name in the SWARM column, blank otherwise. For example, `lab`' is a standalone machine where as all other machines are part of the `swarm-master`" cluster. The Swarm master is also identified by (master) in the SWARM column.

7. Connect to the Swarm cluster and find some information about it:

Docker and Kubernetes for Java EE Developers

```
eval "$(docker-machine env --swarm swarm-master)"  
docker info
```

This shows the output as:

```
> docker info  
Containers: 4  
Strategy: spread  
Filters: affinity, health, constraint, port, dependency  
Nodes: 3  
swarm-master: 192.168.99.107:2376  
  └ Containers: 2  
    └ Reserved CPUs: 0 / 1  
    └ Reserved Memory: 0 B / 1.023 GiB  
swarm-node-01: 192.168.99.108:2376  
  └ Containers: 1  
    └ Reserved CPUs: 0 / 1  
    └ Reserved Memory: 0 B / 1.023 GiB  
swarm-node-02: 192.168.99.109:2376  
  └ Containers: 1  
    └ Reserved CPUs: 0 / 1  
    └ Reserved Memory: 0 B / 1.023 GiB
```

There are 3 nodes – one Swarm master and 2 Swarm nodes. There is a total of 4 containers running in this cluster – one Swarm agent on master and each node, and there is an additional swarm-agent-master running on the master. This can be verified by connecting to the master and listing all the containers:

8. List nodes in the cluster with the following command:

```
docker run swarm list token://<TOKEN>
```

This shows the output as:

```
> docker run swarm list token://b9d9da9198c0facbeaae302242fb65a5  
192.168.99.109:2376  
192.168.99.108:2376  
192.168.99.107:2376
```

The complete cluster is in place now, and we need to deploy the Java EE application to it.

Swarm takes care for the distribution of the deployments across the nodes. The only thing, we need to do is to deploy the application as already explained in [Section 6, “Deploy Java EE 7 Application \(Container Linking\)”](#).

1. Start MySQL server as:

```
docker run --name mysqldb -e MYSQL_USER=mysql -e MYSQL_PASSWORD=mysql -  
e MYSQL_DATABASE=sample -e MYSQL_ROOT_PASSWORD=supersecret -p 3306:3306  
-d mysql
```

`-e` define environment variables that are read by the database at startup and allow us to access the database with this user and password.

2. Start WildFly and deploy Java EE 7 application as:

```
docker run -d --name mywildfly --link mysqldb:db -p 8080:8080  
arungupta/wildfly-mysql-javaee7
```

This is using the [Docker Container Linking](#)²⁷ explained earlier.

3. Check the state of the cluster as:

```
> docker info  
Containers: 7  
Strategy: spread  
Filters: affinity, health, constraint, port, dependency  
Nodes: 3  
swarm-master: 192.168.99.107:2376  
  ↳ Containers: 2  
  ↳ Reserved CPUs: 0 / 1  
  ↳ Reserved Memory: 0 B / 1.023 GiB  
swarm-node-01: 192.168.99.108:2376  
  ↳ Containers: 2  
  ↳ Reserved CPUs: 0 / 1  
  ↳ Reserved Memory: 0 B / 1.023 GiB  
swarm-node-02: 192.168.99.109:2376  
  ↳ Containers: 3  
  ↳ Reserved CPUs: 0 / 1  
  ↳ Reserved Memory: 0 B / 1.023 GiB
```

“swarm-node-02” is running three containers and so lets look at the list of containers running there:

²⁷ <https://docs.docker.com/userguide/dockerlinks/>

```
> eval "$(docker-machine env swarm-node-02)"  
> docker ps -a  


| CONTAINER ID | IMAGE                                  | COMMAND                                            |
|--------------|----------------------------------------|----------------------------------------------------|
| CREATED      | STATUS                                 | PORTS                                              |
| NAMES        |                                        |                                                    |
| f8022254703d | arungupta/wildfly-mysql-javaee7:latest | "/opt/jboss/wildfly/                               |
|              | About a minute ago                     | Up About a minute 0.0.0.0:8080->8080/tcp mywildfly |
| 7b6e9324c735 | mysql:latest                           | "/entrypoint.sh mysql                              |
|              | 7 minutes ago                          | Up 7 minutes 0.0.0.0:3306->3306/tcp mysql          |
| 6ed4c35c943b | swarm:latest                           | "/swarm                                            |
| join --addr  | 12 minutes ago                         | Up 12 minutes 2375/tcp                             |
|              | swarm-agent                            |                                                    |


```

4. Access the application as:

```
curl http://$(docker-machine ip swarm-node-02):8080/employees/  
resources/employees
```

to see the output as:

```
<?xml version="1.0" encoding="UTF-8" standalone="yes"?>  
<><collection><employee><id>1</id><name>Penny</name></employee><employee><id>2</id><name>Sheldon</name></employee><employee><id>3</id><name>Amy</name></employee><employee><id>4</id><name>Leonard</name></employee><employee><id>5</id><name>Bernadette</name></employee><employee><id>6</id><name>Raj</name></employee><employee><id>7</id><name>Howard</name></employee><employee><id>8</id><name>Priya</name></employee></collection>
```

TODO: <https://github.com/javaee-samples/docker-java/issues/55>

14. Java EE Application on Kubernetes Cluster

Kubernetes is an open source system for managing containerized applications across multiple hosts, providing basic mechanisms for deployment, maintenance, and scaling of applications.

— github.com/GoogleCloudPlatform/kubernetes/

Kubernetes, or ``k8s'' in short, allows the user to provide declarative primitives for the desired state, for example “need 5 WildFly servers and 1 MySQL server running”. Kubernetes self-healing mechanisms, such as auto-restarting, re-scheduling, and replicating containers then ensure this state is met. The user just define the state and Kubernetes ensures that the state is met at all times on the cluster.

How is it related to Docker?

Docker provides the lifecycle management of containers. A Docker image defines a build time representation of the runtime containers. There are commands to start, stop, restart, link, and perform other lifecycle methods on these containers. Kubernetes uses Docker to package, instantiate, and run containerized applications.

How does Kubernetes simplify containerized application deployment?

A typical application would have a cluster of containers across multiple hosts. For example, your web tier (for example Undertow) might run on a set of containers. Similarly, your application tier (for example, WildFly) would run on a different set of containers. The web tier would need to delegate the request to application tier. In some cases, or at least to begin with, you may have your web and application server packaged together in the same set of containers. The database tier would generally run on a separate tier anyway. These containers would need to talk to each other. Using any of the solutions mentioned above would require scripting to start the containers, and monitoring/bouncing if something goes down. Kubernetes does all of that for the user after the application state has been defined.

14.1. Key Concepts

At a very high level, there are three key concepts:

1. **Pods** are the smallest deployable units that can be created, scheduled, and managed. Its a logical collection of containers that belong to an application.
2. **Master** is the central control point that provides a unified view of the cluster. There is a single master node that control multiple minions.
3. **Minion** is a worker node that run tasks as delegated by the master. Minions can run one or more pods. It provides an application-specific “virtual host” in a containerized environment.

A picture is always worth a thousand words and so this is a high-level logical block diagram for Kubernetes:




Figure 45. Kubernetes Key Concepts

After the 50,000 feet view, lets fly a little lower at 30,000 feet and take a look at how Kubernetes make all of this happen. There are a few key components at Master and Minion that make this happen.

1. **Replication Controller** is a resource at Master that ensures that requested number of pods are running on minions at all times.
2. **Service** is an object on master that provides load balancing across a replicated group of pods. Label is an arbitrary key/value pair in a distributed watchable storage that the Replication Controller uses for service discovery.
3. **Kubelet** Each minion runs services to run containers and be managed from the master. In addition to Docker, Kubelet is another key service installed there. It reads container manifests as YAML files that describes a pod. Kubelet ensures that the containers defined in the pods are started and continue running.
4. Master serves **RESTful Kubernetes API** that validate and configure Pod, Service, and Replication Controller.

14.2. Start Kubernetes Cluster

1. In ``kubernetes'' directory, setup a cluster as:

```
cd kubernetes

export KUBERNETES_PROVIDER=vagrant
./cluster/kube-up.sh
```

The `KUBERNETES_PROVIDER` environment variable tells all of the various cluster management scripts which variant to use.



This will take a few minutes, so be patient! Vagrant will provision each machine in the cluster with all the necessary components to run Kubernetes.

It shows the output as:

```
Starting cluster using provider: vagrant
... calling verify-prereqs
... calling kube-up
Using credentials: vagrant:vagrant

.

Cluster validation succeeded
Done, listing cluster services:

Kubernetes master is running at https://10.245.1.2
KubeDNS is running at https://10.245.1.2/api/v1beta3/proxy/namespaces/
default/services/kube-dns
```

Note down the address for Kubernetes master, <https://10.245.1.2> in this case.

2. Verify the Kubernetes cluster as:

```
kubernetes> vagrant status
Current machine states:

master                  running (virtualbox)
minion-1                running (virtualbox)
```

This environment represents multiple VMs. The VMs are all listed above with their current state. For more information about a specific VM, run `vagrant status NAME`.


By default, the Vagrant setup will create a single kubernetes-master and 1 kubernetes-minion. Each VM will take 1 GB, so make sure you have at least 2GB to 4GB of free memory (plus appropriate free disk space).



By default, only one minion is created. This can be manipulated by setting an environment variable `NUM_MINIONS` variable to an integer before invoking `kube-up.sh` script.

By default, each VM in the cluster is running Fedora, Kubelet is installed into ```systemd`'', and all other Kubernetes services are running as containers on Master.

3. Access <https://10.245.1.2> (or whatever IP address is assigned to your kubernetes cluster start up log). This may present the warning as shown below:



Click on `Advanced` and then on `Proceed to 10.245.1.2` to see the output as:



```
{  
  "paths": [  
    "/api",  
    "/api/v1beta1",  
    "/api/v1beta2",  
    "/api/v1beta3",  
    "/healthz",  
    "/healthz/ping",  
    "/logs/",  
    "/metrics",  
    "/static/",  
    "/swagger-ui/",  
    "/swaggerapi/",  
    "/version"  
  ]  
}
```

Figure 46. Kubernetes Output from Master

Use `vagrant''` as the username and `vagrant"` as the password.

Check the list of nodes as:

```
> ./cluster/kubectl.sh get nodes  
NAME          LABELS           STATUS  
10.245.1.3   kubernetes.io/hostname=10.245.1.3   Ready
```

4. Check the list of pods:

```
kubernetes> ./cluster/kubectl.sh get po  
POD           IP             CONTAINER(S)      IMAGE (S)  
              HOST  
LABELS  
STATUS     CREATED      MESSAGE  
etcd-server-kubernetes-master  
  
kubernetes-master/  <none>  
                    Running   2 minutes
```

Docker and Kubernetes for Java EE Developers

```
          etcd-container gcr.io/
google_containers/etcdb:2.0.9

          Running  2 minutes
kube-apiserver-kubernetes-master

kubernetes-master/ <none>
          Running  2 minutes
          kube-apiserver gcr.io/
google_containers/kube-apiserver:465b93ab80b30057f9c2ef12f30450c3

          Running  2 minutes
kube-dns-v1-lxdof

10.245.1.3/           k8s-app=kube-dns, kubernetes.io/cluster-
service=true, version=v1   Pending  2 minutes
          etcd      gcr.io/
google_containers/etcdb:2.0.9
          kube2sky    gcr.io/
google_containers/kube2sky:1.7
          skydns     gcr.io/
google_containers/skydns:2015-03-11-001
kube-scheduler-kubernetes-master

kubernetes-master/ <none>
          Running  2 minutes
          kube-scheduler gcr.io/
google_containers/kube-scheduler:d1f640dfb379f64daf3ae44286014821

          Running  2 minutes
.....
```

5. Check the list of services running:

```
> ./cluster/kubectl.sh get se
NAME          LABELS
              SELECTOR          IP(S)        PORT(S)
kube-dns       k8s-app=kube-dns, kubernetes.io/cluster-
service=true, kubernetes.io/name=KubeDNS   k8s-app=kube-dns
10.247.0.10   53/UDP
                                         53/TCP
kubernetes     component=apiserver, provider=kubernetes
              <none>          10.247.0.2  443/TCP
kubernetes-ro  component=apiserver, provider=kubernetes
              <none>          10.247.0.1  80/TCP
.....
```

6. Check the list of replication controllers:

```
> ./cluster/kubectl.sh get rc
CONTROLLER      CONTAINER(S)      IMAGE(S)
          SELECTOR                      REPLICAS
kube-dns-v1     etcd            gcr.io/google_containers/etcd:2.0.9
                k8s-app=kube-dns,version=v1   1
                  kube2sky        gcr.io/google_containers/kube2sky:1.7
                    skydns         gcr.io/google_containers/
skydns:2015-03-11-001
```

14.3. Deploy Java EE Application

Pods, and the IP addresses assigned to them, are ephemeral. If a pod dies then Kubernetes will recreate that pod because of its self-healing features, but it might recreate it on a different host. Even if it is on the same host, a different IP address could be assigned to it. And so any application cannot rely upon the IP address of the pod.

Kubernetes services is an abstraction which defines a logical set of pods. A service is typically back-ended by one or more physical pods (associated using labels), and it has a permanent IP address that can be used by other pods/applications. For example, WildFly pod can not directly connect to a MySQL pod but can connect to MySQL service. In essence, Kubernetes service offers clients an IP and port pair which, when accessed, redirects to the appropriate backends.




Figure 47. Kubernetes Services



In this case, all the pods are running on a single minion. This is because, that is the default number for a Kubernetes cluster. The pod can very be on another minion if more minions exist in the cluster.

Any Service that a Pod wants to access must be created before the Pod itself, or else the environment variables will not be populated.

Start MySQL Service

1. Start MySQL service as:

```
./cluster/kubectl.sh create -f ../../attendees/kubernetes/mysql-  
service.yaml
```

It uses the following configuration file:

```
id: mysql  
kind: Service  
apiVersion: v1beta3  
metadata:  
  name: mysql  
  labels:  
    name: mysql  
    context: docker-k8s-lab  
spec:  
  ports:  
    - port: 3306  
  selector:  
    name: mysql  
  labels:  
    name: mysql
```

2. Check that the service is created:

```
> ./cluster/kubectl.sh get se -l context=docker-k8s-lab  
NAME      LABELS                      SELECTOR      IP(S)  
PORT(S)  
mysql     context=docker-k8s-lab, name=mysql   name=mysql  
10.247.141.208  3306/TCP
```

Note that the label used during the creation is used to query the service.

When a Pod is run on a node, the kubelet adds a set of environment variables for each active Service.

It supports both Docker links compatible variables and simpler `{SVCNAME}_SERVICE_HOST` and `{SVCNAME}_SERVICE_PORT` variables, where the Service name is upper-cased and dashes are converted to underscores.

Our service name is `'mysql'` and so `'MYSQL_SERVICE_HOST'` and `'MYSQL_SERVICE_PORT'` variables are available to other pods.

Send a Pull Request for #62²⁸.

Start MySQL Replication Controller

1. Start MySQL replication controller as:

```
> ./cluster/kubectl.sh --v=5 create -f ../../attendees/kubernetes/mysql.yaml
I0616 19:41:55.441461      8346 defaults.go:174] creating security
context for container mysql
replicationcontrollers/mysql
```

It uses the following configuration file:

```
kind: ReplicationController
apiVersion: v1beta3
metadata:
  name: mysql
  labels:
    name: mysql
    context: docker-k8s-lab
spec:
  replicas: 1
  selector:
    name: mysql
  template:
    metadata:
      labels:
        name: mysql
        context: docker-k8s-lab
    spec:
      containers:
        - name: mysql
          image: mysql:latest
```

²⁸ <https://github.com/javaee-samples/docker-java/issues/62>

```
env:  
  - name: MYSQL_USER  
    value: mysql  
  - name: MYSQL_PASSWORD  
    value: mysql  
  - name: MYSQL_DATABASE  
    value: sample  
  - name: MYSQL_ROOT_PASSWORD  
    value: supersecret  
ports:  
  - containerPort: 3306  
    hostPort: 3306
```

Once again, the ``docker-k8s-lab'' label is used. This simplifies querying the created pods later on.

2. Verify MySQL replication controller as:

```
> ./cluster/kubectl.sh get rc -l context=docker-k8s-lab  
CONTROLLER   CONTAINER(S)   IMAGE(S)      SELECTOR      REPLICAS  
mysql        mysql          mysql:latest   name=mysql   1
```

3. Check the status of MySQL pod as:

```
> ./cluster/kubectl.sh get po -l context=docker-k8s-lab  
POD           IP          CONTAINER(S)   IMAGE(S)      HOST  
LABELS  
mysql-7lq67               10.245.1.3/  
context=docker-k8s-lab, name=mysql     Pending     About a minute  
                                     mysql       mysql:latest
```

Start WildFly Replication Controller

1. Start WildFly replication controller as:

```
> ./cluster/kubectl.sh --v=5 create -f ../../attendees/kubernetes/  
wildfly.yaml  
I0616 18:59:00.563099    7849 defaults.go:174] creating security  
context for container wildfly  
replicationcontrollers/wildfly
```

It uses the following configuration file:

```
kind: ReplicationController
```

Docker and Kubernetes for Java EE Developers

```
apiVersion: v1beta3
metadata:
  name: wildfly
  labels:
    name: wildfly
    context: docker-k8s-lab
spec:
  replicas: 1
  selector:
    name: wildfly-server
  template:
    metadata:
      labels:
        name: wildfly-server
        context: docker-k8s-lab
    spec:
      containers:
        - name: wildfly
          image: arungupta/wildfly-mysql-javaee7:k8s
          ports:
            - containerPort: 8080
              hostPort: 8080
```

2. Verify WildFly replication controller using ``docker-k8s-lab" label as:

```
> ./cluster/kubectl.sh get rc -l context=docker-k8s-lab
CONTROLLER   CONTAINER(S)   IMAGE(S)
           SELECTOR           REPLICAS
mysql        mysql          mysql:latest
             name=mysql       1
wildfly      wildfly        arungupta/wildfly-mysql-javaee7:k8s
             name=wildfly-server 1
```

3. Check the status of WildFly pod as:

```
> ./cluster/kubectl.sh get pod -l context=docker-k8s-lab
POD           IP           CONTAINER(S)   IMAGE(S)
           HOST           LABELS
STATUS        CREATED      MESSAGE
mysql-7lq67   10.245.1.3/  context=docker-k8s-lab, name=mysql
             Pending      3 minutes
                           mysql          mysql:latest
wildfly-o0nw6  10.245.1.3/  context=docker-k8s-lab, name=wildfly-server
             Pending      45 seconds
```

```
wildfly      arungupta/wildfly-mysql-  
javaee7:k8s
```

Make sure the status of both WildFly and MySQL pod is changed to ``running''. It will look like:

```
> ./cluster/kubectl.sh get pod -l context=docker-k8s-lab  
POD          IP            CONTAINER(S)   IMAGE (S)  
HOST          LABELS  
STATUS        CREATED       MESSAGE  
mysql-7lq67    172.17.0.9  
              10.245.1.3/10.245.1.3  context=docker-k8s-lab,name=mysql  
Running     14 minutes  
                         mysql           mysql:latest  
  
Running     10 minutes  
wildfly-o0nw6  172.17.0.10  
              10.245.1.3/10.245.1.3  context=docker-k8s-lab,name=wildfly-  
server    Running     11 minutes  
                         wildfly         arungupta/wildfly-mysql-  
javaee7:k8s  
Running     26 seconds
```



Takes a while for all the pods to start. It took ~25 minutes on a 16 GB, i7 Mac OS X.

14.4. Access Java EE Application

http://<pod_ip>:8080/employees/resources/employees

14.5. Self-healing Pods

1. Delete the WildFly pod
2. Wait for k8s to restart the pod because of RC

14.6. Application Logs

1. Login to Minion-1 VM:

```
> vagrant ssh minion-1  
Last login: Fri Jun  5 23:01:36 2015 from 10.0.2.2  
[vagrant@kubernetes-minion-1 ~]$
```

2. Log in as root:

```
[vagrant@kubernetes-minion-1 ~]$ su -  
Password:  
[root@kubernetes-minion-1 ~]#
```

Default root password for VM images created by Vagrant is ``vagrant''.

3. See the list of Docker containers running on this VM:

```
docker ps
```

4. View WildFly log as:

```
docker logs $(docker ps | grep arungupta/wildfly | awk '{print $1}')
```

5. View MySQL log as:

```
docker logs <CID>
```

14.7. Delete Kubernetes Resources

Individual resources (service, replication controller, or pod) can be deleted by using `delete` command instead of `create` command. Alternatively, all services and replication controllers can be deleted using a label as:

```
kubectl delete -l se,po context=docker-k8s-lab
```

14.8. Stop Kubernetes Cluster

```
> ./cluster/kube-down.sh  
Bringing down cluster using provider: vagrant  
==> minion-1: Forcing shutdown of VM...  
==> minion-1: Destroying VM and associated drives...  
==> master: Forcing shutdown of VM...  
==> master: Destroying VM and associated drives...  
Done
```

14.9. View Logs

Viewing logs is causing <https://github.com/GoogleCloudPlatform/kubernetes/issues/9888>.

14.10. Debug Kubernetes Master (OPTIONAL)

1. Log in to the master as:

```
> vagrant ssh master
Last login: Thu Jun  4 19:30:04 2015 from 10.0.2.2
[vagrant@kubernetes-master ~]$
```

2. Log in as root:

```
[vagrant@kubernetes-master ~]$ su -
Password:
Last login: Thu Jun  4 19:25:41 UTC 2015
[root@kubernetes-master ~]
```

Default root password for VM images created by Vagrant is ``vagrant''.

3. Check the containers running on master:

CONTAINER ID	IMAGE	COMMAND	CREATED
STATUS	PORTS	NAMES	
2b92c80630d5	gcr.io/google_containers/etcd:2.0.9	"/usr/local/bin/etcd - 5 hours ago	
	k8s_etcd-container.ec4297e5_etcd-server-kubernetes-master_default_3595ac402f3a17c29dab95f3e0f64c76_56fa3dce64c375f8030b	kube-apiserver:465b93ab80b30057f9c2ef12f30450c3 - 5 hours ago	/bin/sh -c '/usr/lo 5 hours ago'
	gcr.io/google_containers/kube-apiserver.f4e485e1_kube-apiserver-kubernetes-master_default_c6b19d563bdbcfb0af80b57377ee905c_2f16c239d7d9d40bd479	k8s_kube-apiserver.f4e485e1_kube-apiserver-kubernetes-master_default_c6b19d563bdbcfb0af80b57377ee905c_2f16c239 - 5 hours ago	/bin/sh -c '/usr/lo 5 hours ago'
	gcr.io/google_containers/kube-controller-manager:572696d43ca87cd1fe0c774bac3a5f4b	kube-controller-manager:572696d43ca87cd1fe0c774bac3a5f4b - 5 hours ago	/bin/sh -c '/usr/lo 5 hours ago'
	gcr.io/google_containers/kube-scheduler:d1f640dfb379f64daf3ae44286014821	kube-scheduler:f53b6329_kube-scheduler-kubernetes-master_default_1f3b1657f7f1af67ce9f929d78c64695_de632a80b1809bdabd9c - 5 hours ago	/bin/sh -c '/usr/lo 5 hours ago'
	gcr.io/google_containers/pause:0.8.0	"/pause"	

```

k8s_POD.e4cc795_kube-apiserver-kubernetes-
master_default_c6b19d563bdbcfb0af80b57377ee905c_767dadbb1
280baf845b00      gcr.io/google_containers/pause:0.8.0
                    "/pause"
                    5 hours ago      Up 5 hours
k8s_POD.e4cc795_kube-scheduler-kubernetes-
master_default_1f3b1657f7f1af67ce9f929d78c64695_52a4ca74
615a314a35bf      gcr.io/google_containers/pause:0.8.0
                    "/pause"
                    5 hours ago      Up 5 hours
k8s_POD.e4cc795_kube-controller-manager-kubernetes-
master_default_8f8db766ebc90a00a99244c362284cf1_97cc1739
7a554eea05f3      gcr.io/google_containers/pause:0.8.0
                    "/pause"
                    5 hours ago      Up 5 hours
k8s_POD.e4cc795_etcd-server-kubernetes-
master_default_3595ac402f3a17c29dab95f3e0f64c76_593b9807

```

15. Common Docker Commands

Here is the list of commonly used Docker commands:

Purpose	Command
Build an image	<code>docker build --rm=true .</code>
Install an image	<code>docker pull \${IMAGE}</code>
List of installed images	<code>docker images</code>
List of installed images (detailed listing)	<code>docker images --no-trunc</code>
Remove an image	<code>docker rmi \${IMAGE_ID}</code>
Remove all untagged images	<code>docker rmi \$(docker images grep '^' awk '{print \$3}')</code>
Remove all images	<code>docker rm \$(docker ps -aq)</code>
Run a container	<code>docker run</code>
List containers	<code>docker ps</code>
Stop a container	<code>docker stop \${CID}</code>
Stop all running containers	<code>docker stop \$(docker ps -q)</code>
Find IP address of the container	<code>docker inspect --format '{{ .NetworkSettings.IPAddress }}' \${CID}</code>

Purpose	Command
Attach to a container	<code>docker attach \${CID}</code>
Remove a container	<code>docker rm \${CID}</code>
Remove all containers	<code>docker rm \$(docker ps -aq)</code>
Get container id for an image by a regular expression	<code>docker ps grep wildfly awk '{print \$1}'</code>
Remove container by a regular expression	<code>docker ps -a grep wildfly awk '{print \$1}' xargs docker rm -f</code>

16. References

1. JBoss and Docker: <http://www.jboss.org/docker/>