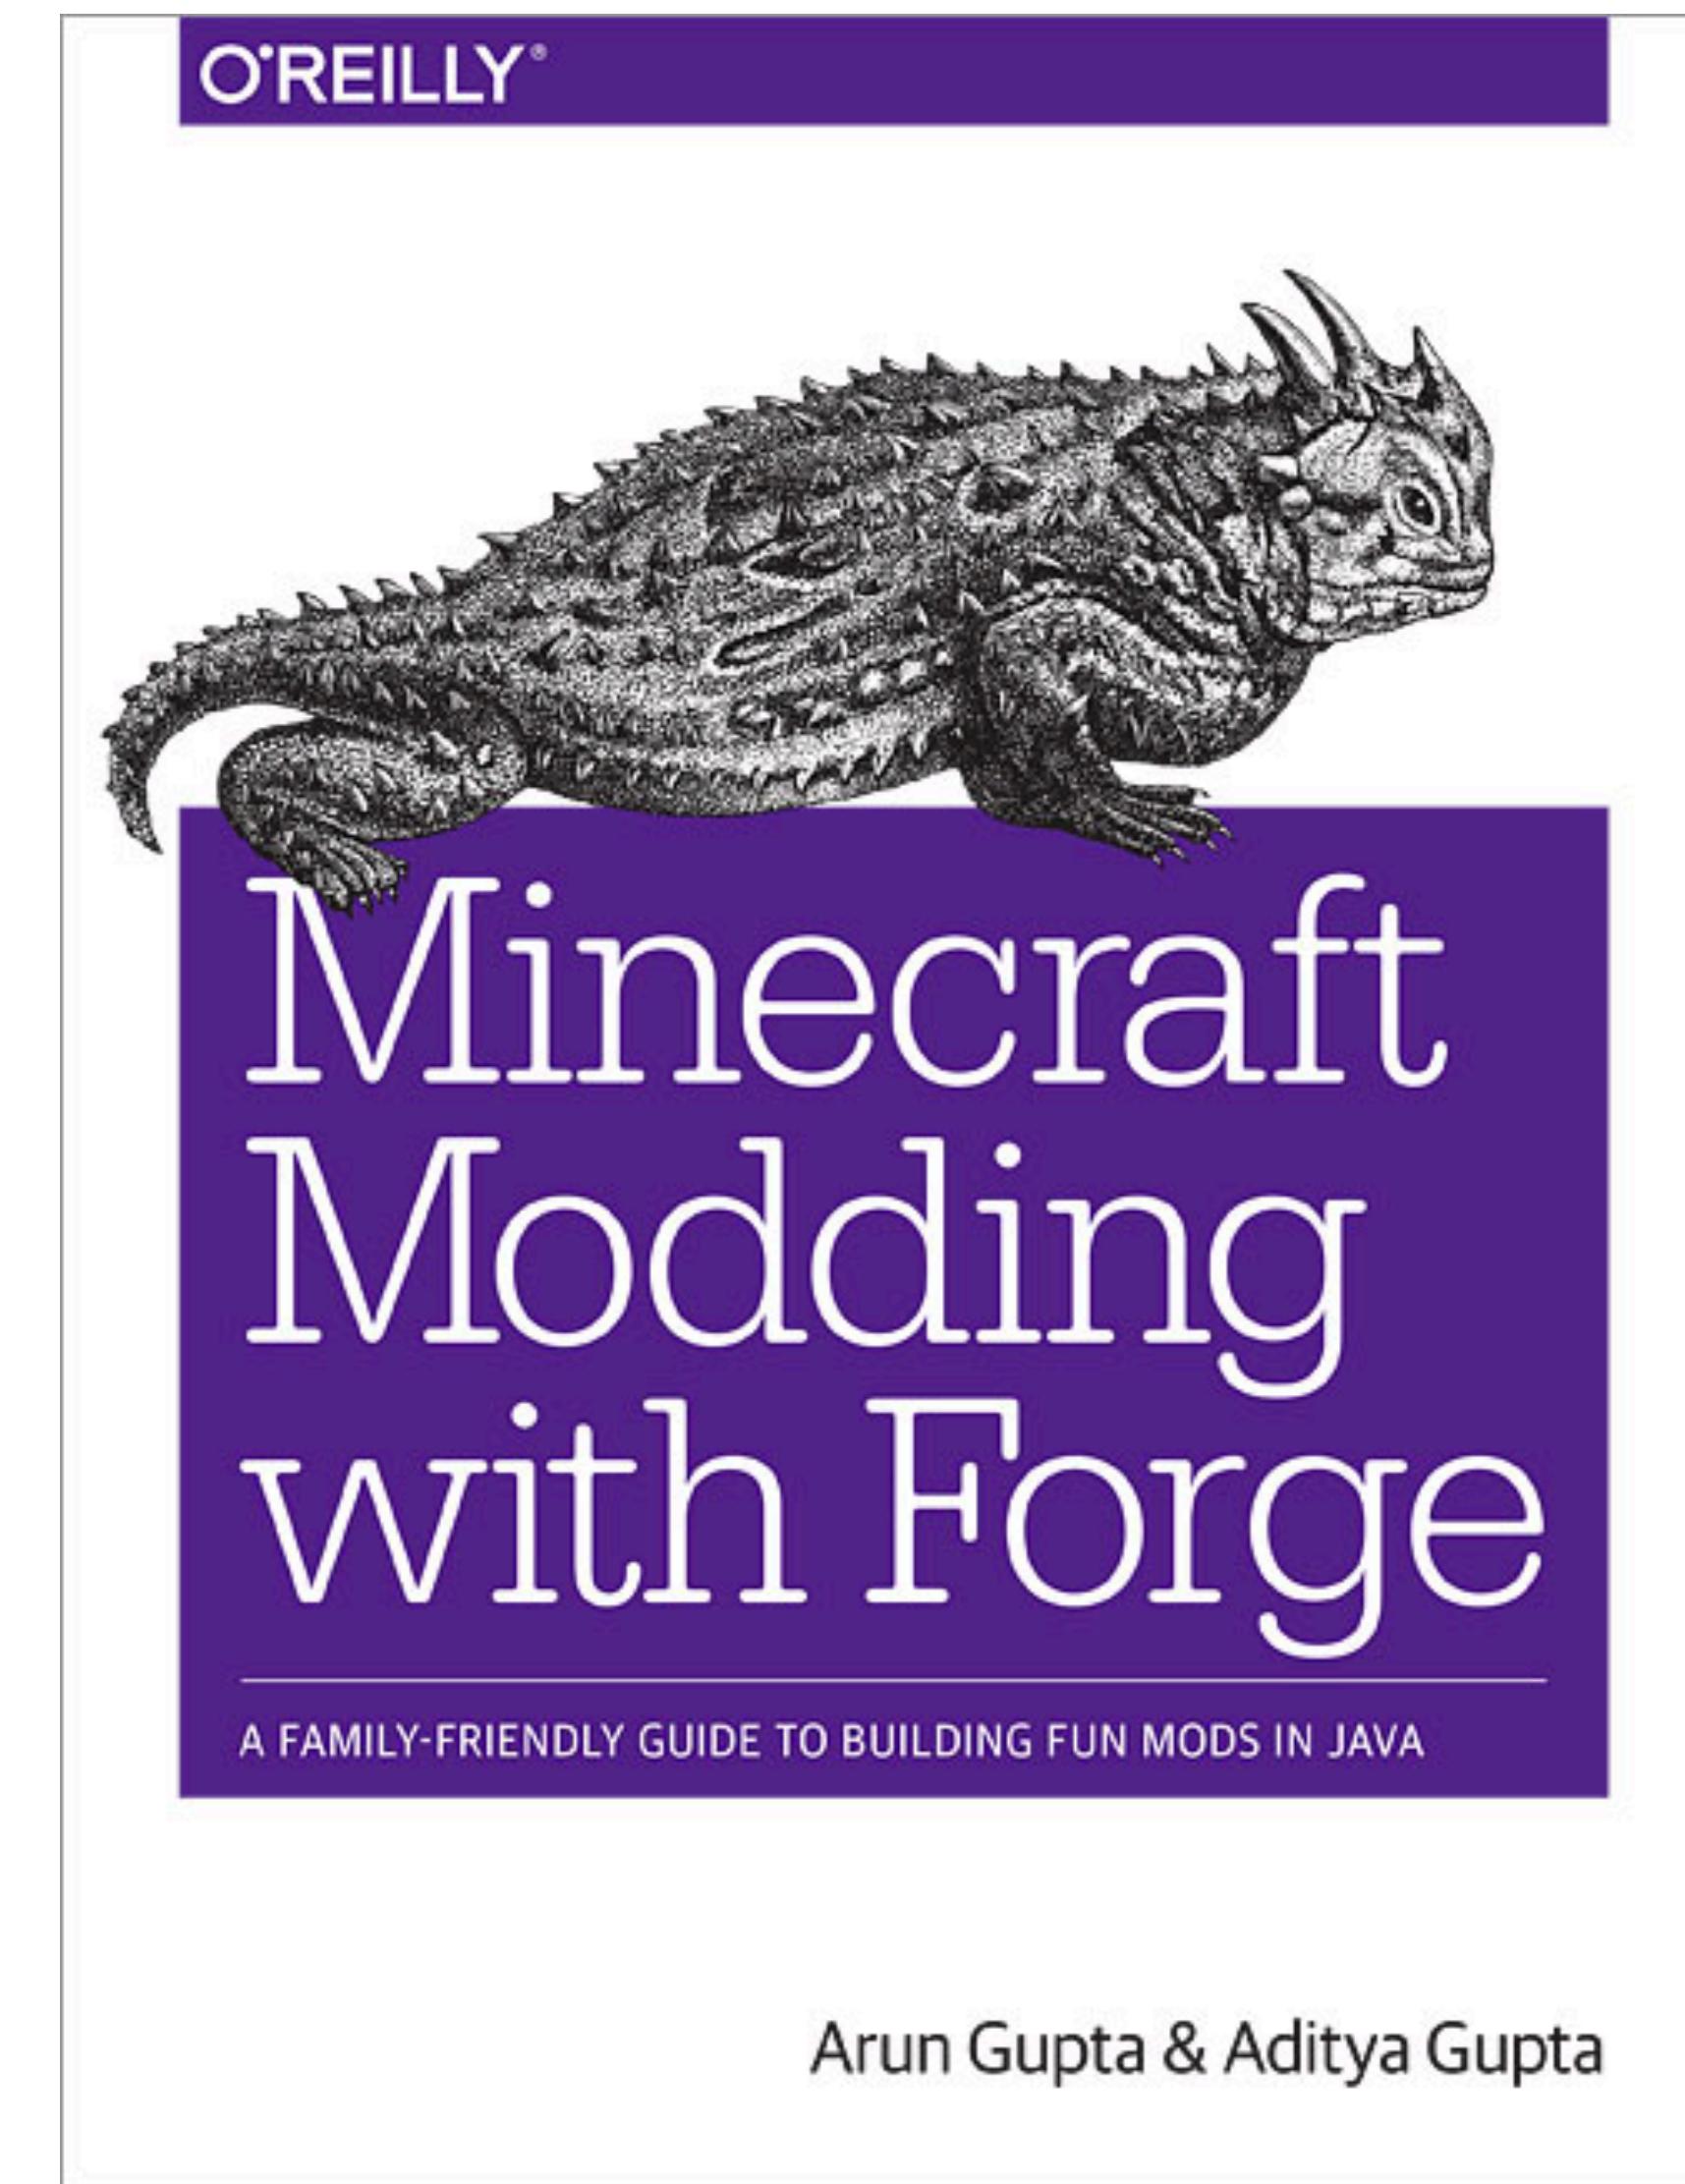




# Getting Started with Kubernetes

Arun Gupta, @arungupta  
VP Developer Advocacy, Couchbase



**CONTENTS**

- » What Is Kubernetes?
- » Kubernetes Architecture
- » Starting With Kubernetes
- » Run Your First Container
- » Scale Applications...and more!

Get More Refcardz! Visit [DZone.com/Refcardz](http://DZone.com/Refcardz)

**WHAT IS KUBERNETES?**

Kubernetes ([kubernetes.io](https://kubernetes.io)) is an open-source orchestration system for managing containerized applications across multiple hosts, providing basic mechanisms for the deployment, maintenance, and scaling of applications.

Kubernetes, or “k8s” or “kube” for short, allows the user to declaratively specify the desired state of a cluster using high-level primitives. For example, the user may specify that they want three instances of the Couchbase server container running. Kubernetes’ self-healing mechanisms, such as auto-restarting, re-scheduling, and replicating containers then converge the actual state towards the desired state.

Kubernetes supports Docker and Rocket containers. An abstraction around the containerization layer will allow for other container image formats and runtimes to be supported in the future.

**KEY CONCEPTS OF KUBERNETES****POD**

A Pod is the smallest deployable unit that can be created, scheduled, and managed. It’s a logical collection of containers that belong to an application.

Each resource in Kubernetes is defined using a configuration file. For example, a Couchbase pod can be defined with the following .yaml file:

A Replication Controller creating two instances of a Couchbase pod can be defined as:

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: couchbase-controller
spec:
  # Two replicas of the Pod to be created
  replicas: 2
  # Identifies the label key and value on the Pod that
  # this Replication Controller is responsible for
  # managing
  selector:
    app: couchbase-rc-pod
  # ‘cookie cutter’ used for creating new pods when
  # necessary
  template:
    metadata:
      labels:
        # label key and value on the pod.
        # These must match the selector above.
        app: couchbase-rc-pod
    spec:
      containers:
        - name: couchbase
          image: couchbase
          ports:
            - containerPort: 8091
```

**SERVICE**

Each Pod is assigned a unique IP address. If the Pod is inside a Replication Controller, then the pod is recreated but may be given

# Kubernetes

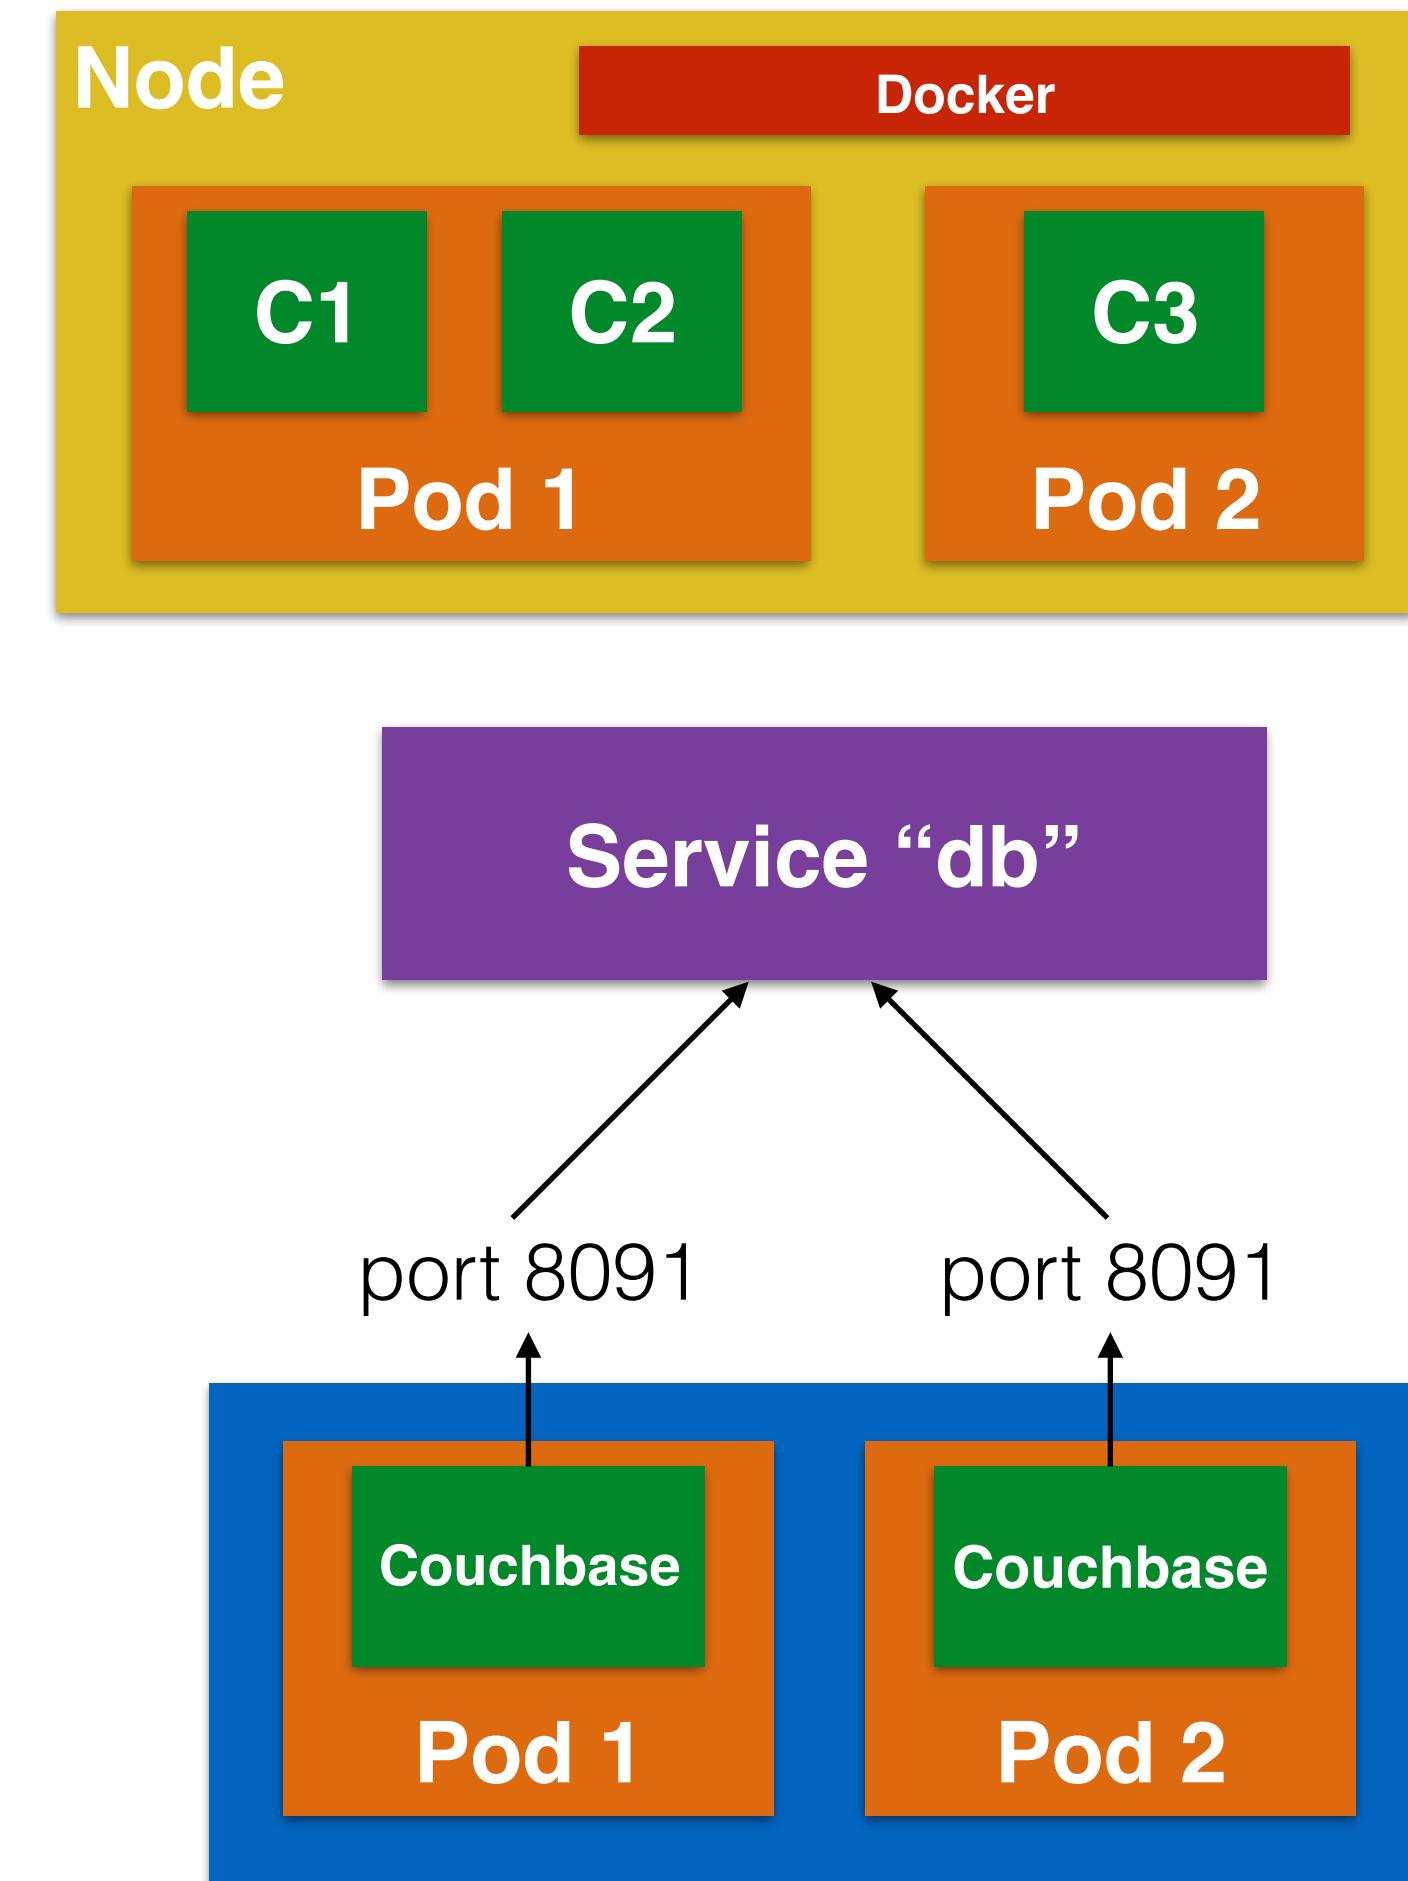
BY ARUN GUPTA

# Kubernetes

- Open source orchestration system for Docker containers
- Provide declarative primitives for the “desired state”
  - Self-healing
  - Auto-restarting
  - Schedule across hosts
  - Replicating

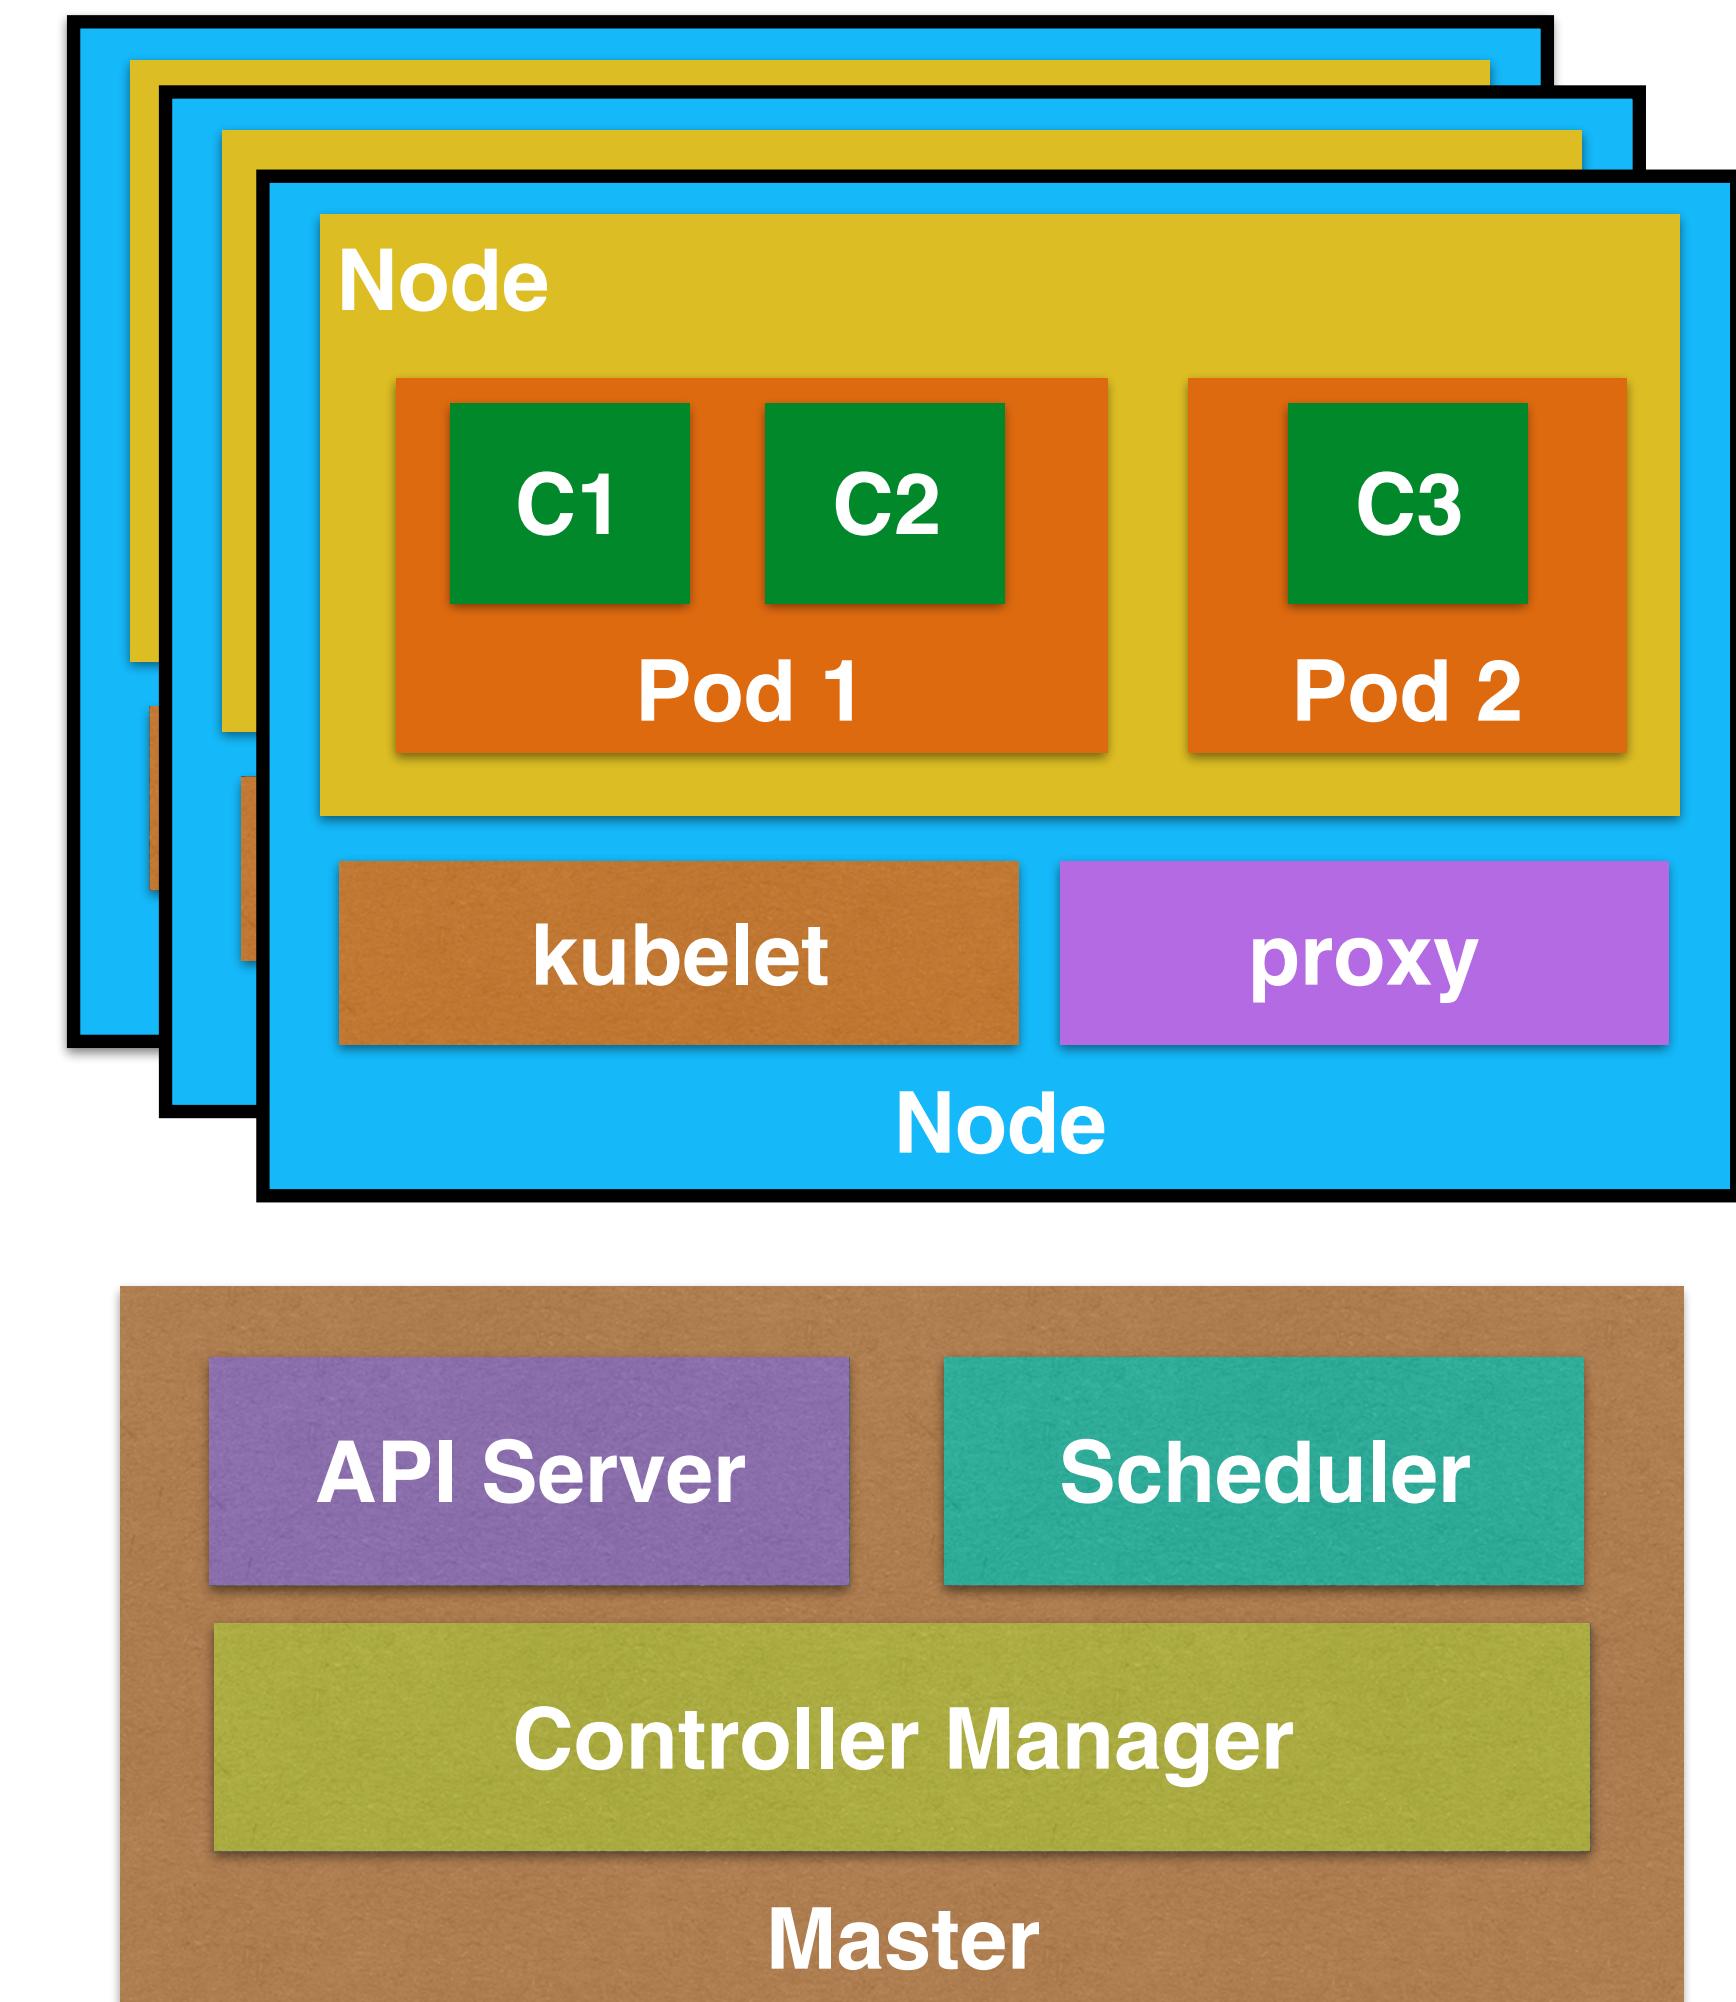
# Concepts

- **Pods**: collocated group of Docker containers that share an IP and storage volume
- **Service**: Single, stable name for a set of pods, also acts as LB
- **Label**: used to organize and select group of objects
- **Replication Controller**: manages the lifecycle of pods and ensures specified number are running

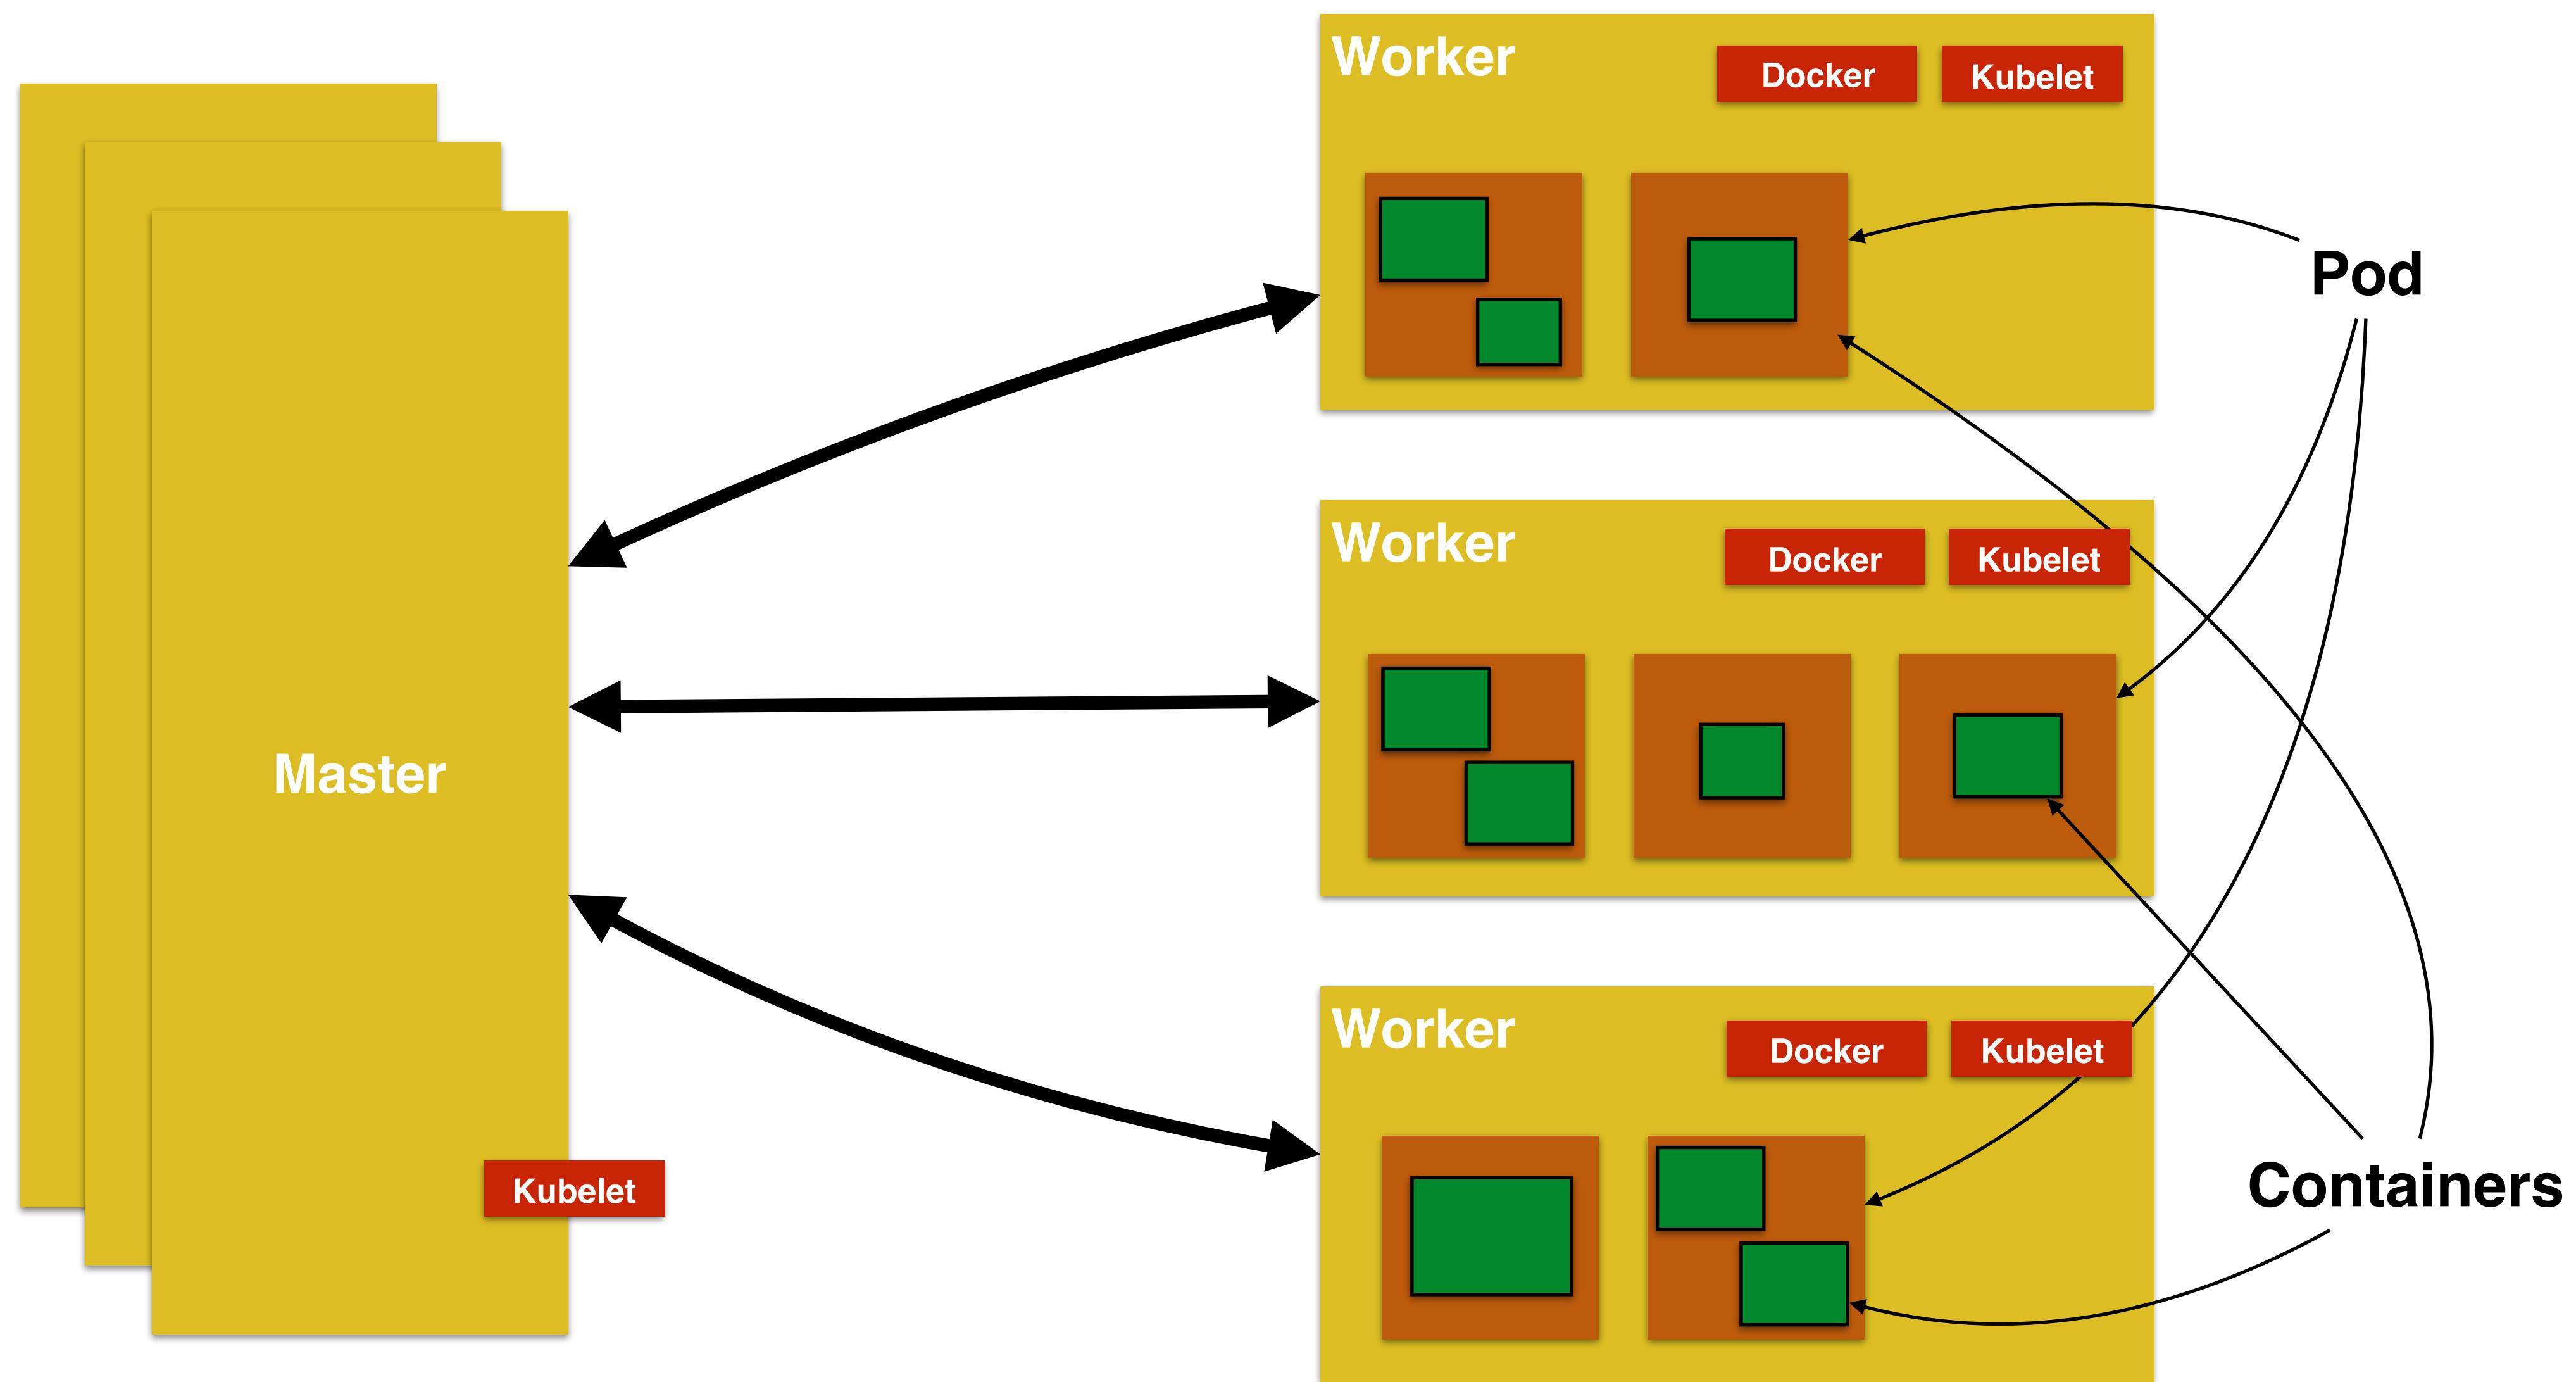


# Components

- **Node**: Docker host running *kubelet* (node agent) and *proxy* services
  - Monitored by *systemd* (CentOS) or *monit* (Debian)
- **Master**: hosts cluster-level control services, including the API server, scheduler, and controller manager
- **etcd**: distributed key-value store used to persist Kubernetes system state



# Architecture





# Master High Availability

- Hack by running a **podmaster** utility
- Proposal
  - Hot Standby
  - Warm Standby
  - Active-Active (Load Balanced)

# kubectl

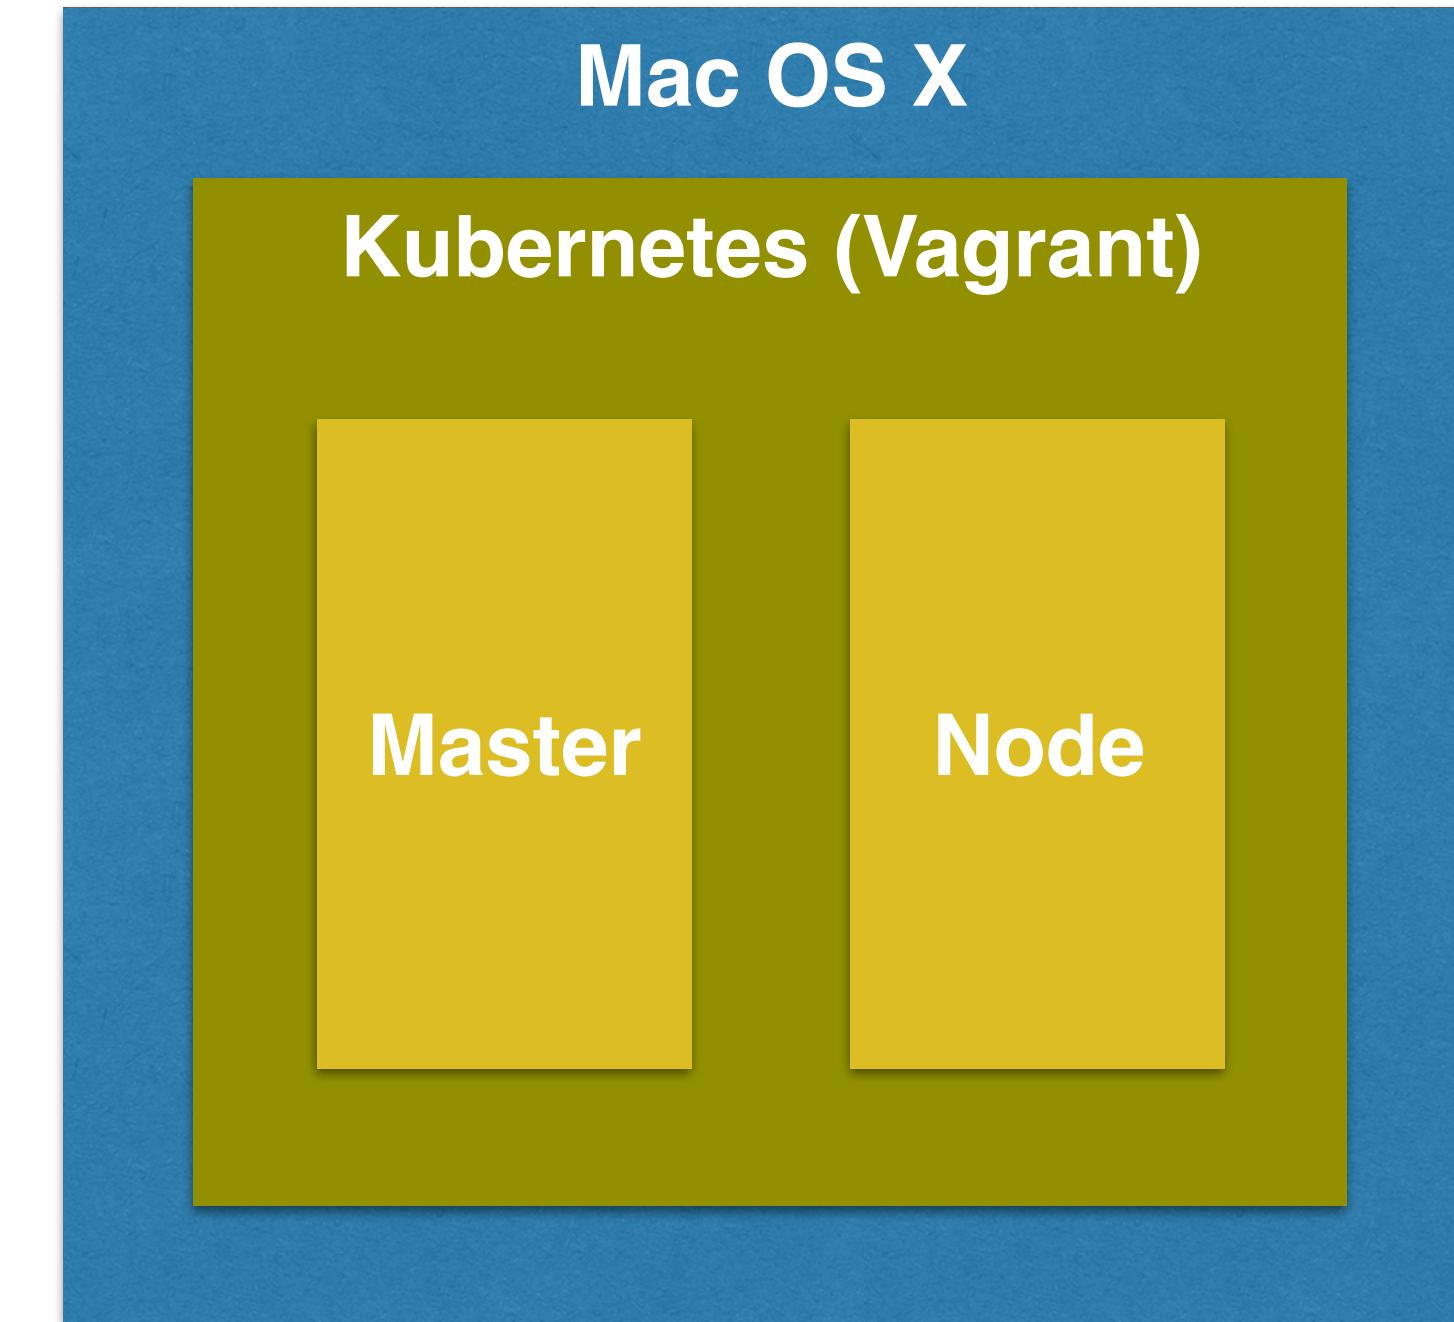
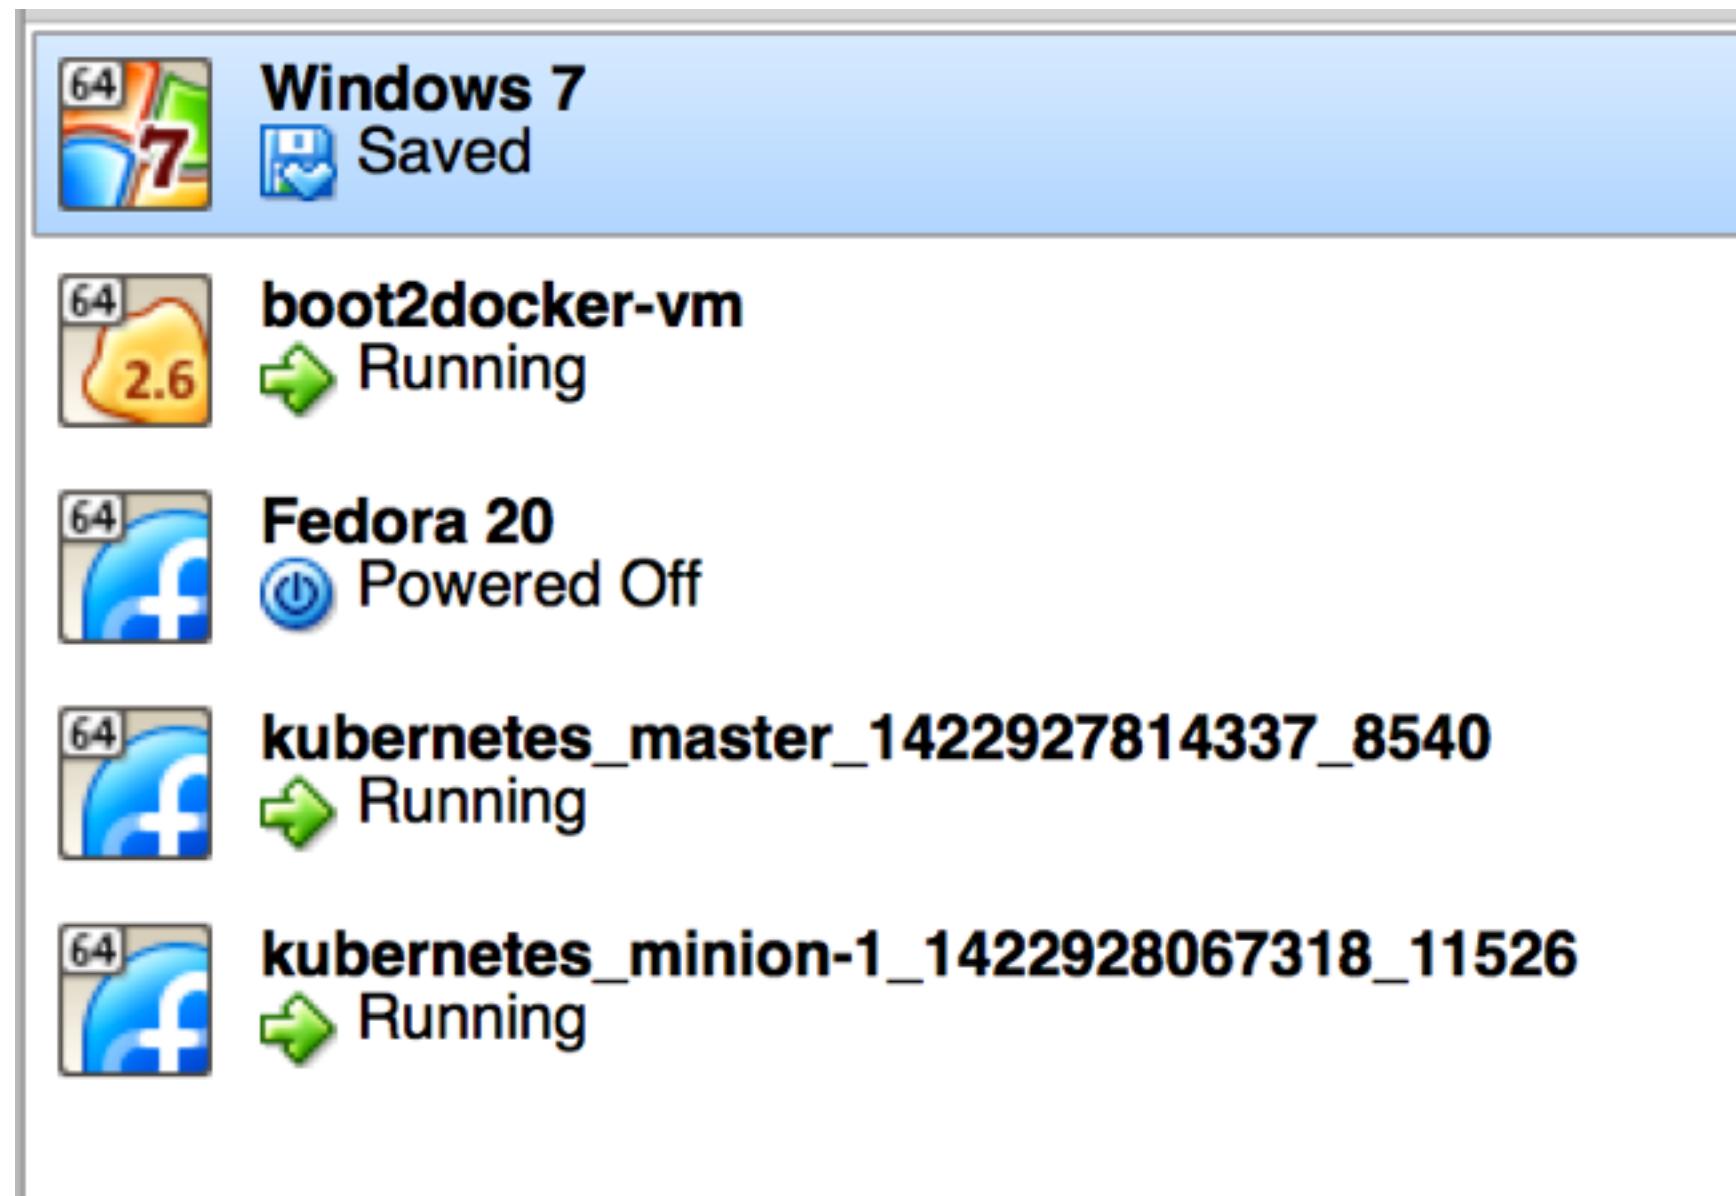
- Controls the Kubernetes cluster manager
- `kubectl get pods or minions`
- `kubectl create -f <filename>`
- `kubectl update or delete`
- `kubectl resize --replicas=3 replicationcontrollers <name>`

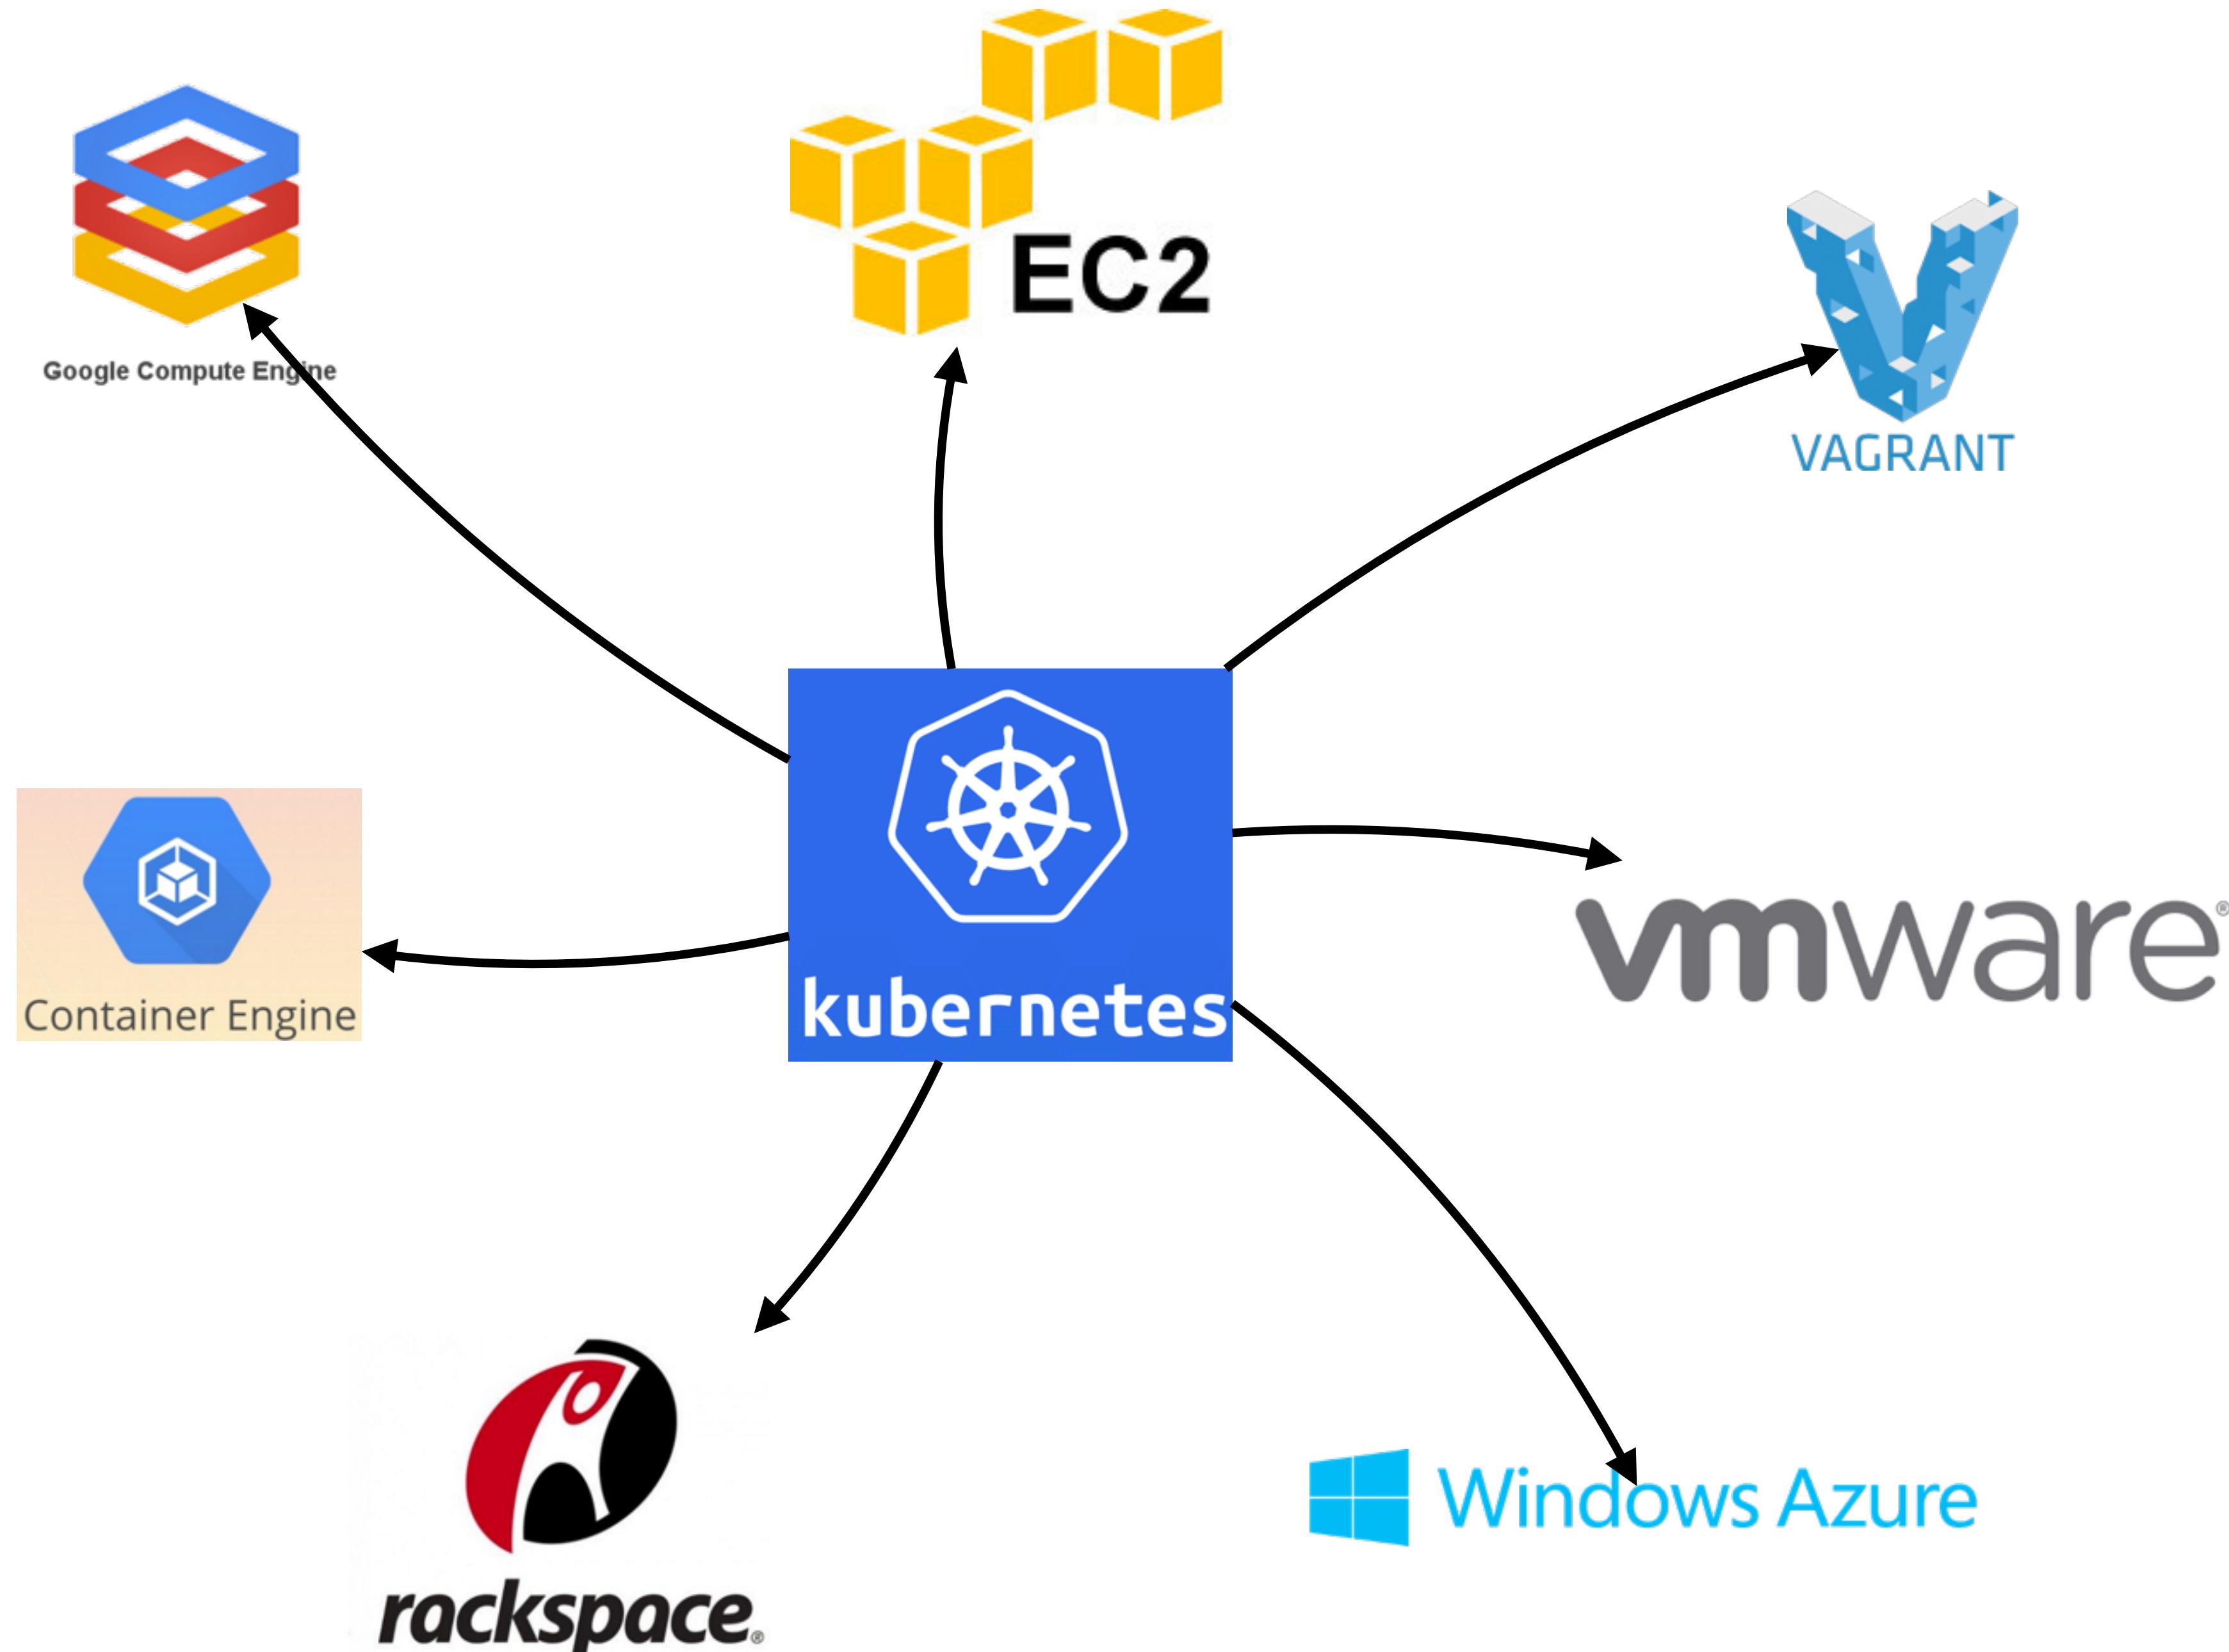
# Kubernetes Config

```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: wildfly-pod
5   labels:
6     name: wildfly
7 spec:
8   containers:
9     - image: jboss/wildfly
10    name: wildfly-pod
11    ports:
12      - containerPort: 8080
```

```
1 apiVersion: v1
2 kind: ReplicationController
3 metadata:
4   name: wildfly-rc
5   labels:
6     name: wildfly
7 spec:
8   replicas: 2
9   template:
10  metadata:
11  labels:
12  name: wildfly
13 spec:
14   containers:
15     - name: wildfly-rc-pod
16       image: jboss/wildfly
17     ports:
18       - containerPort: 8080
```

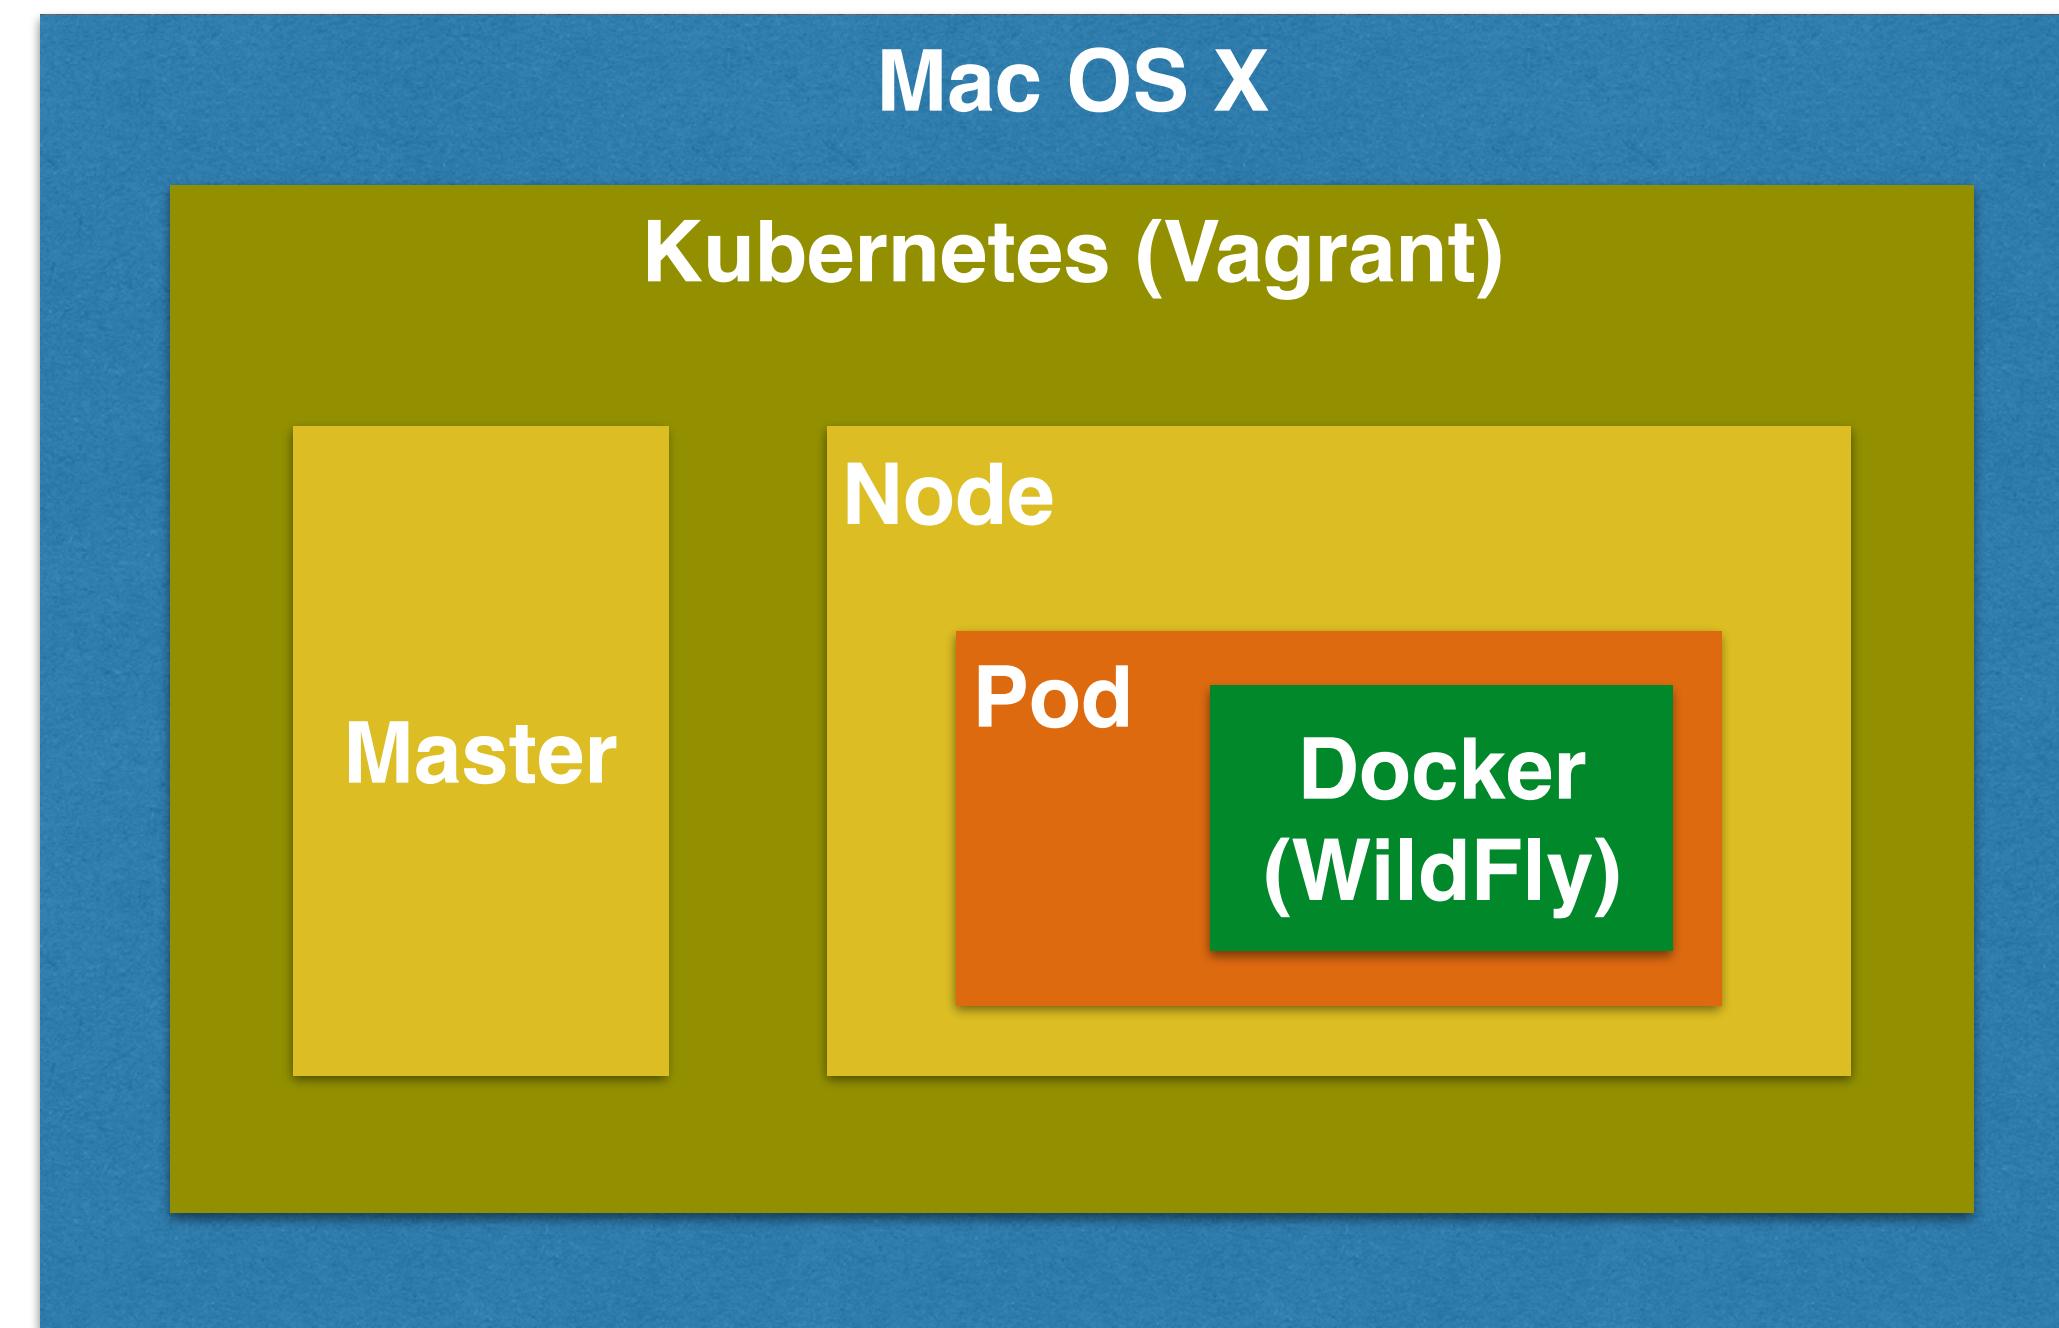
```
export KUBERNETES_PROVIDER=vagrant  
./cluster/kube-up.sh
```





# A Pod with One Container

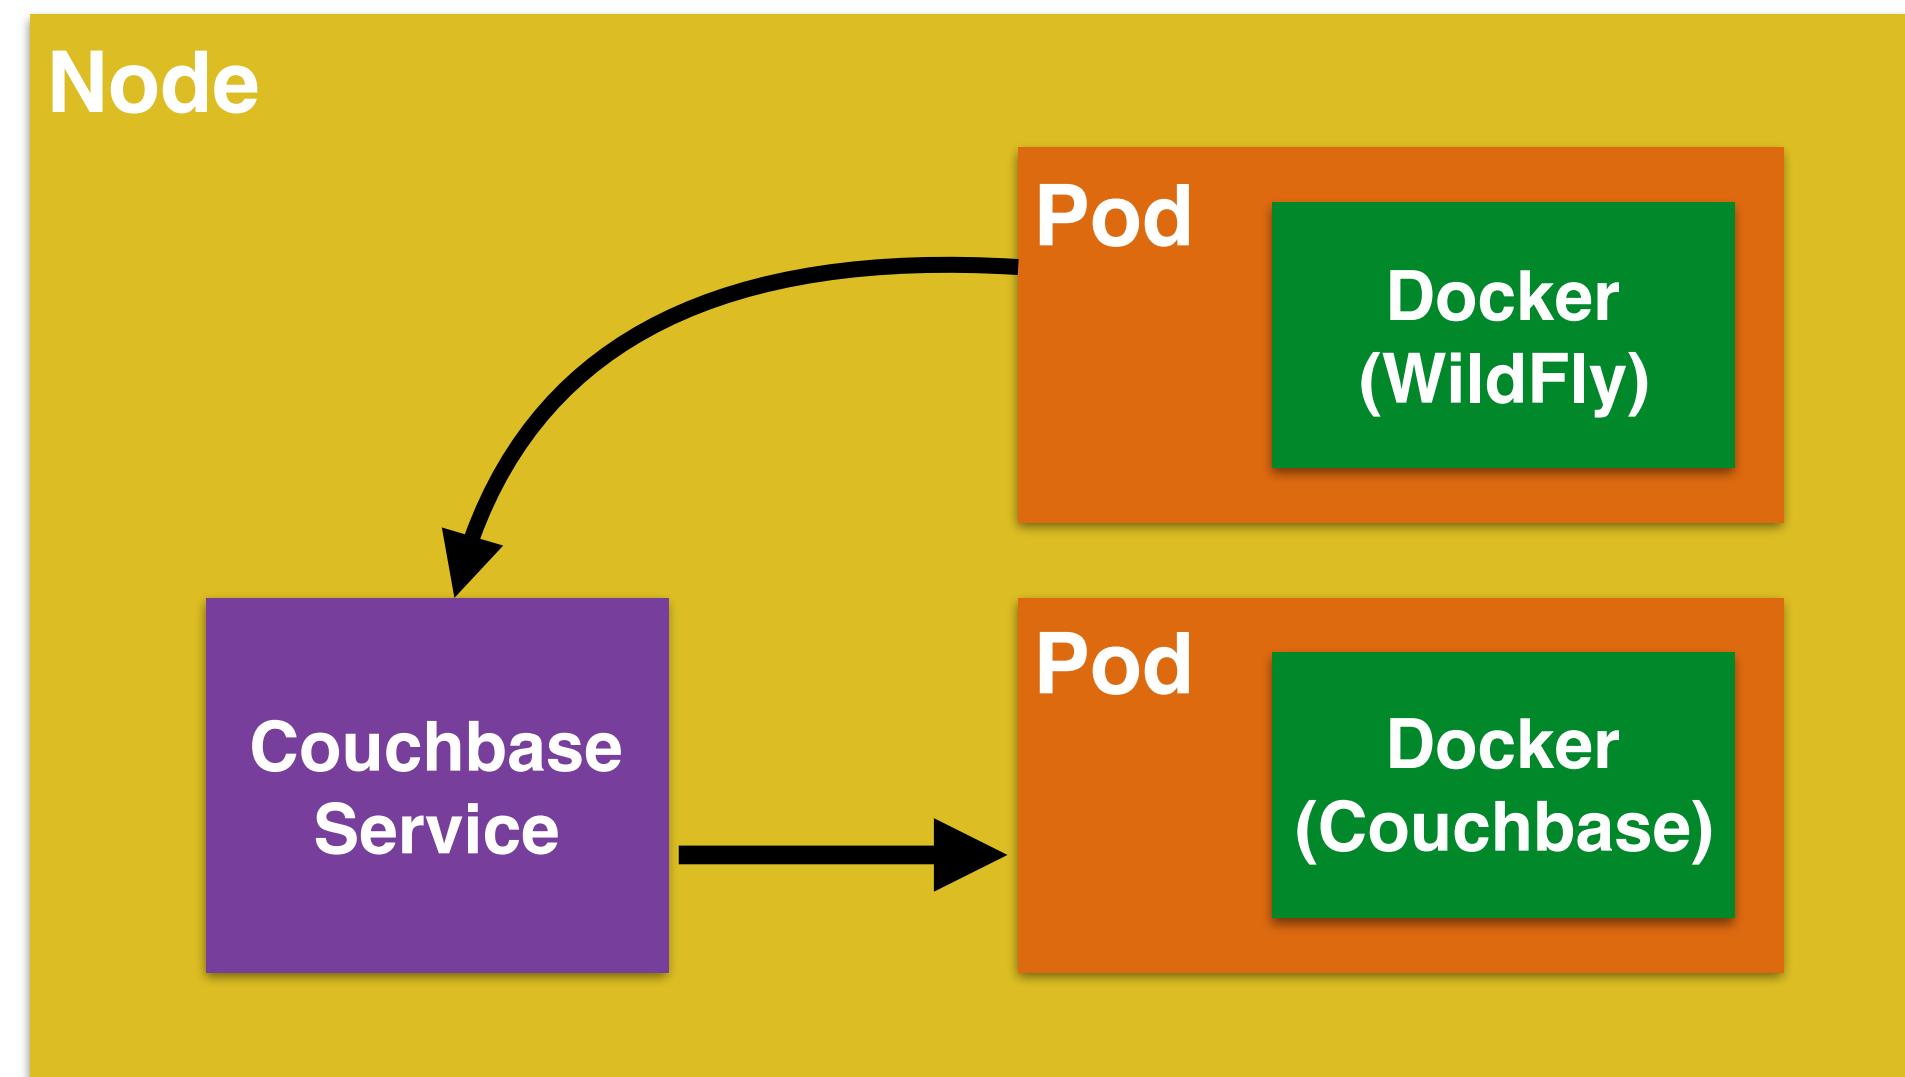
```
1 apiVersion: v1
2 kind: Pod
3 metadata:
4   name: wildfly-pod
5   labels:
6     name: wildfly
7 spec:
8   containers:
9     - image: jboss/wildfly
10    name: wildfly-pod
11    ports:
12      - containerPort: 8080
```



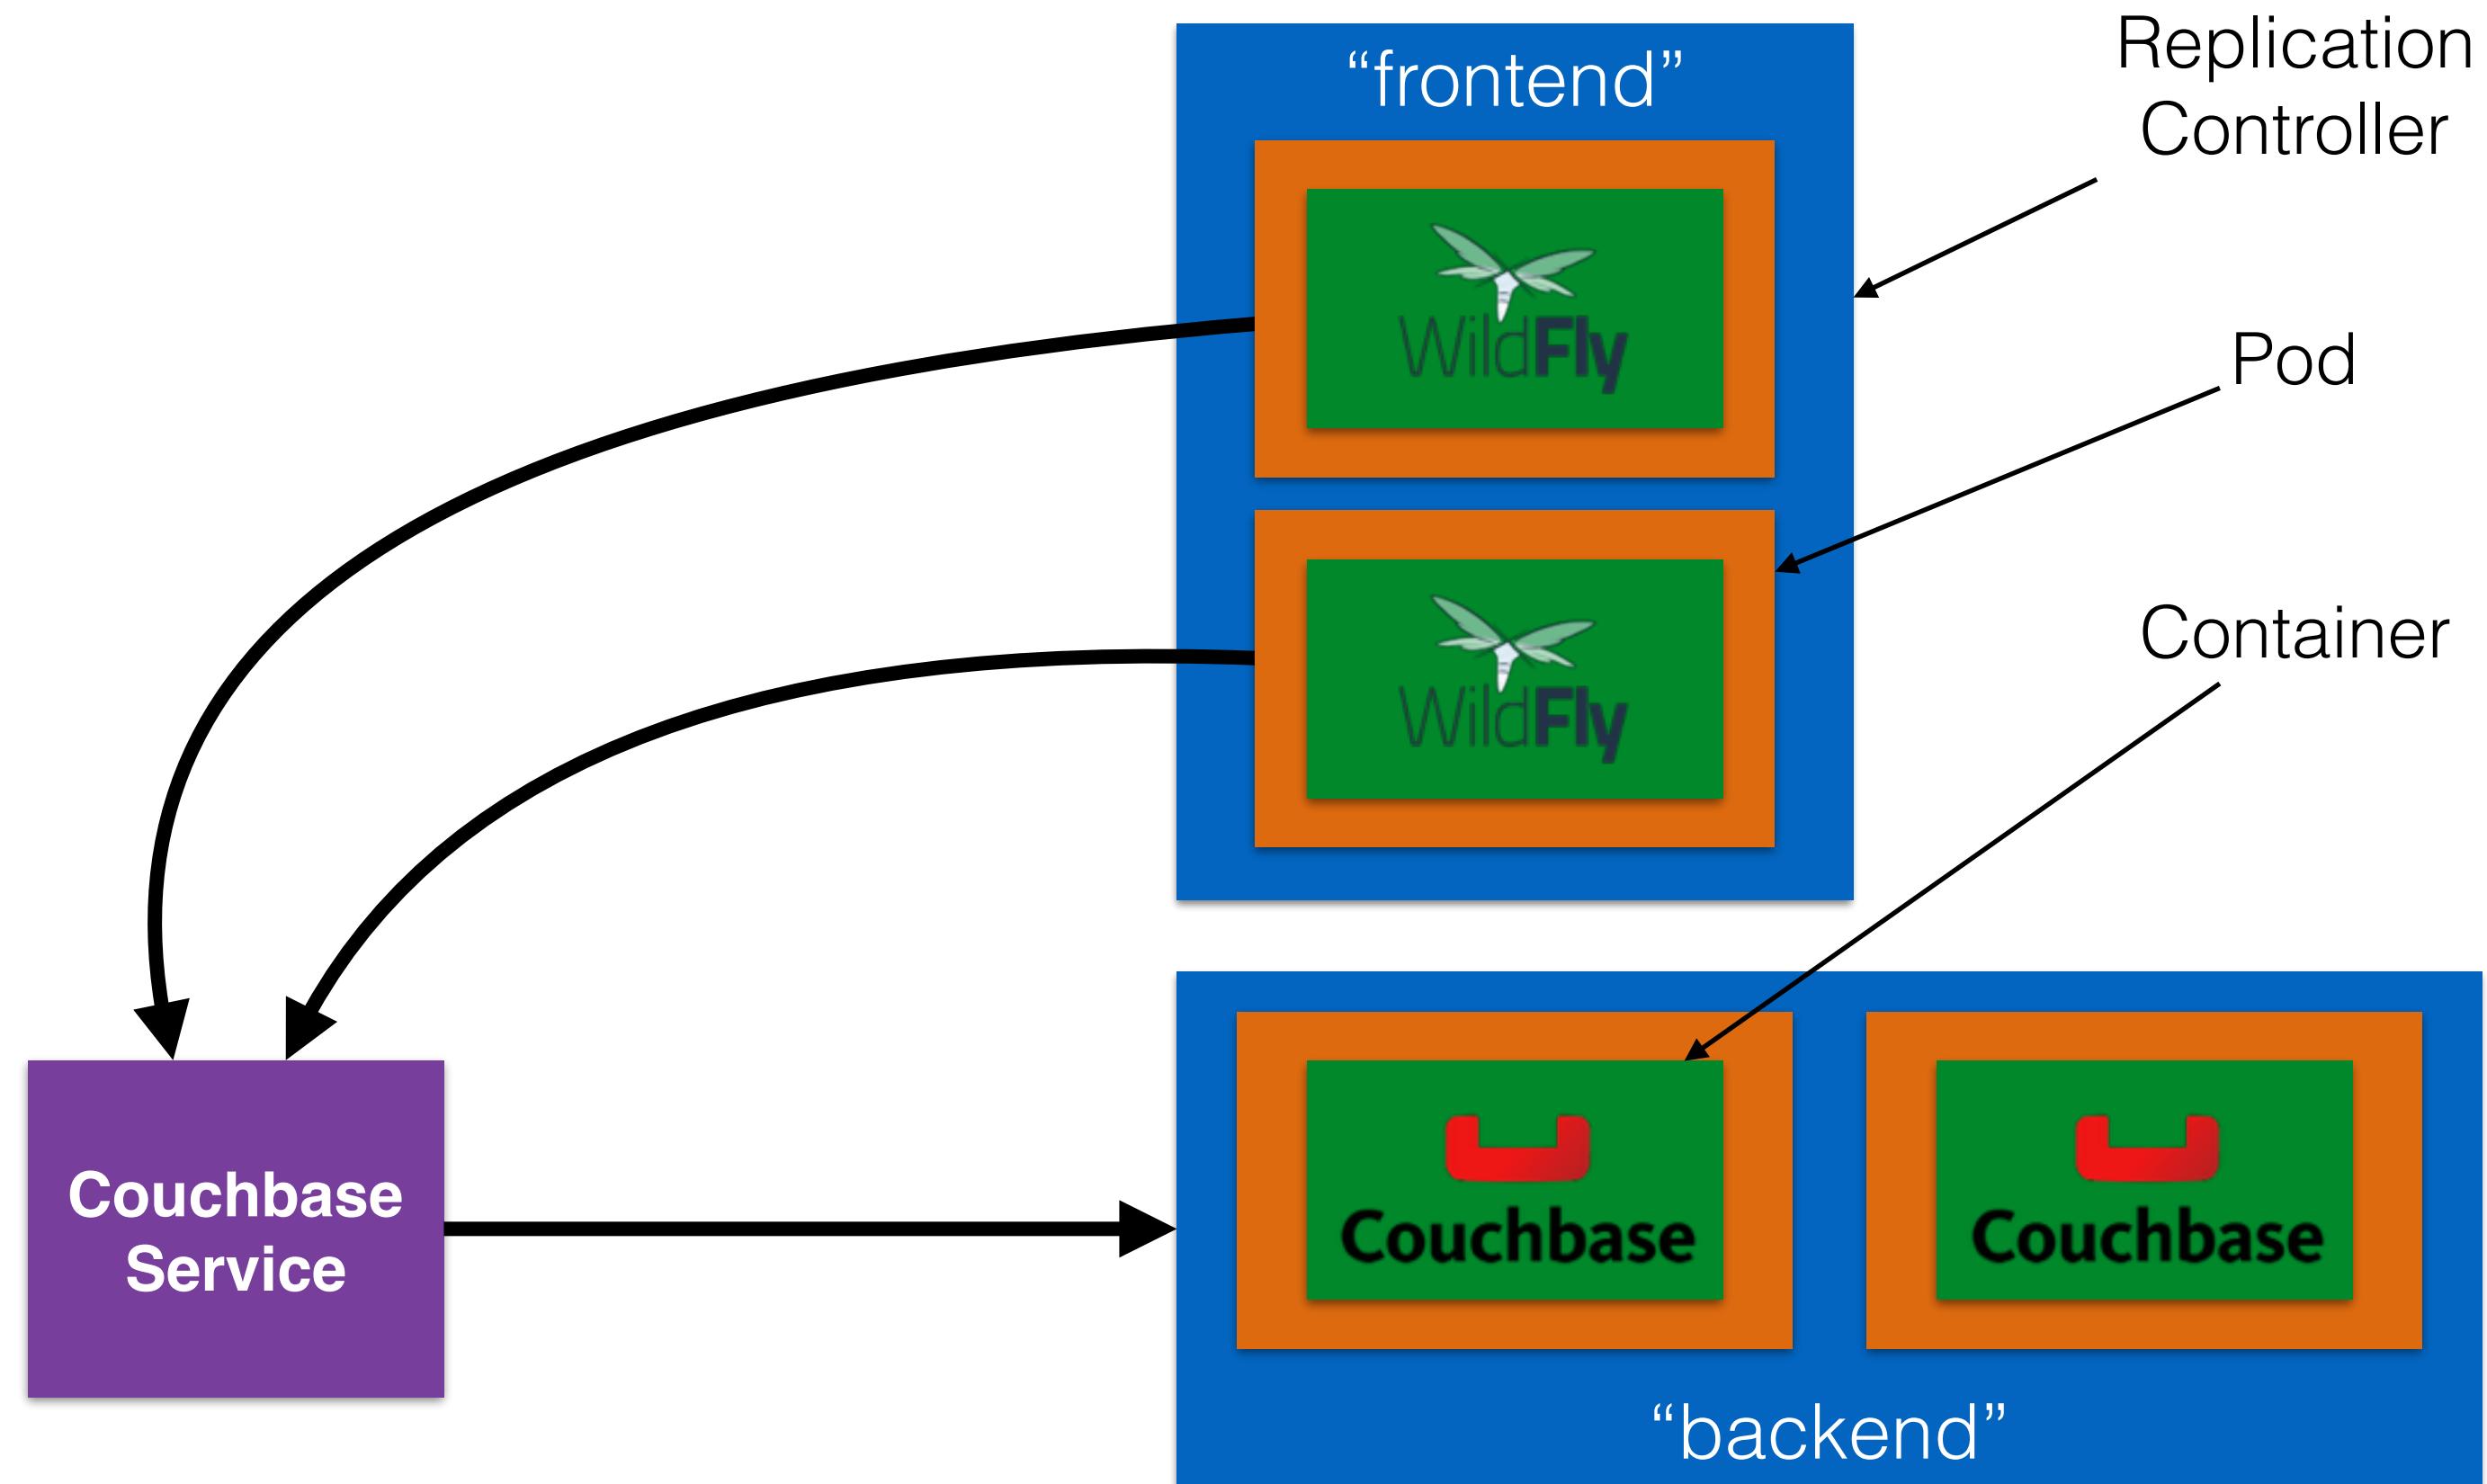
# Services

- Abstract a set of pods as a single IP and port
  - Simple TCP/UDP load balancing
- Creates environment variables in other pods
- Stable endpoint for pods to reference
  - Allows list of pods to change dynamically

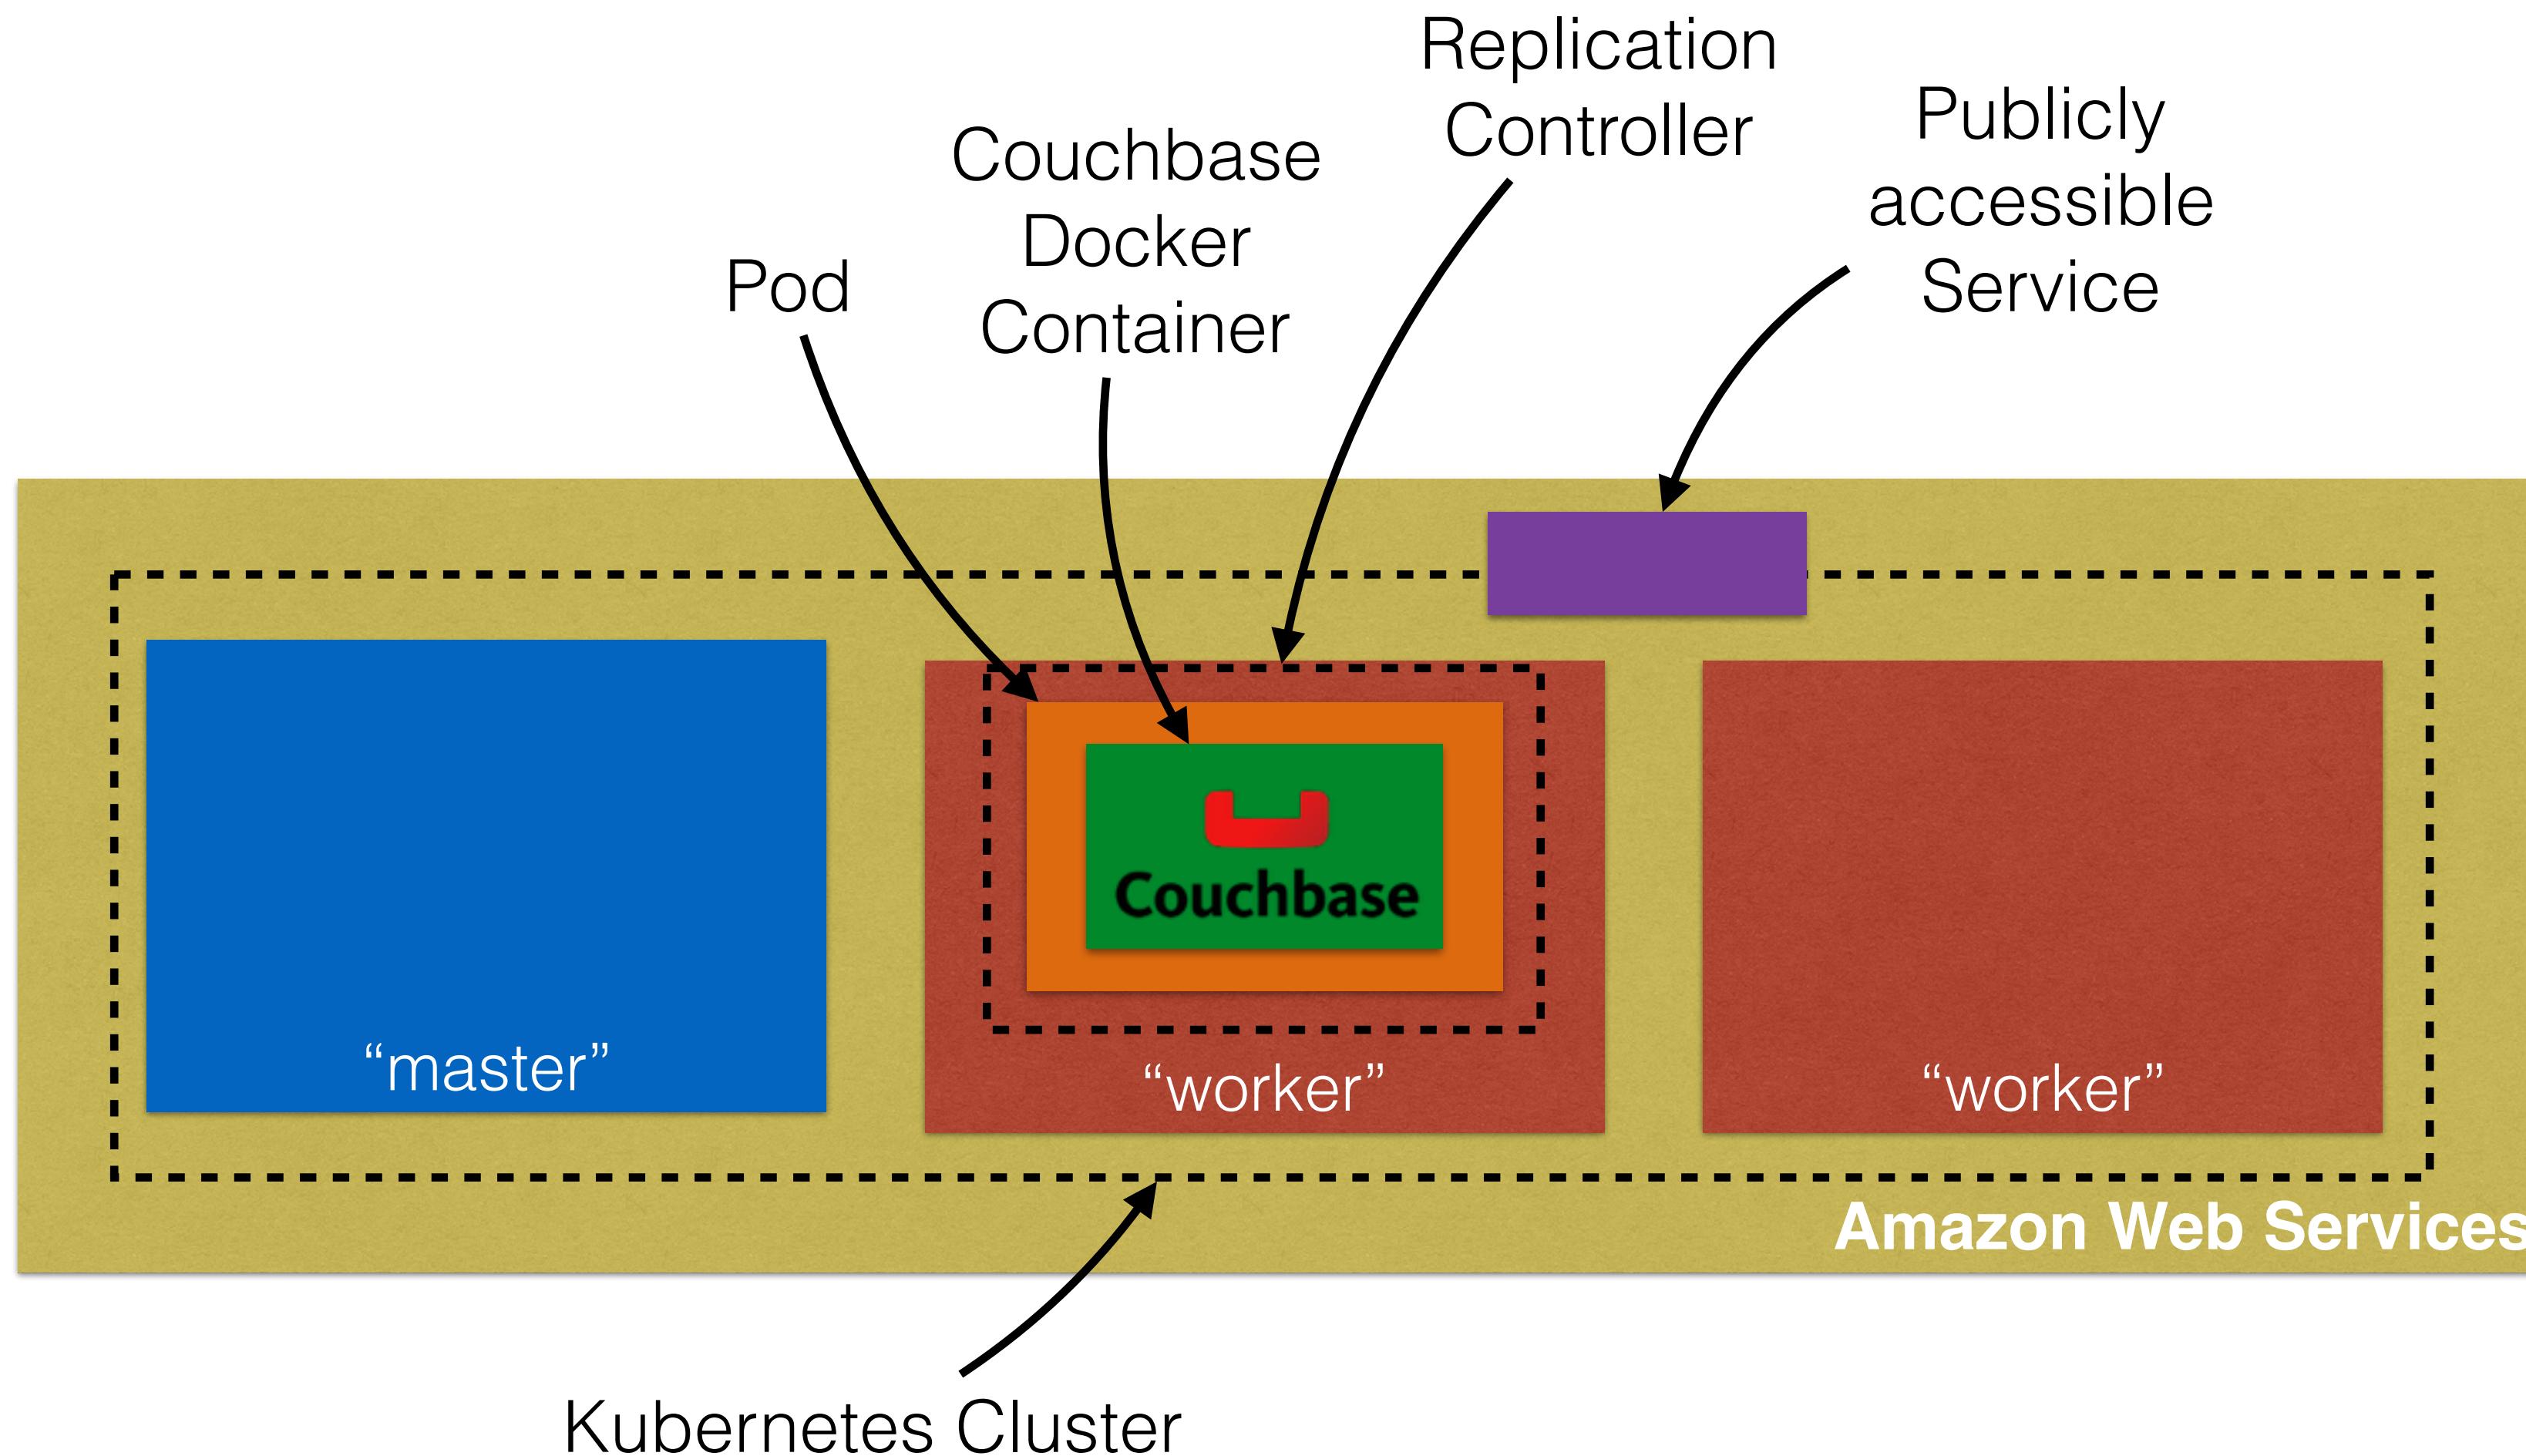
# Services



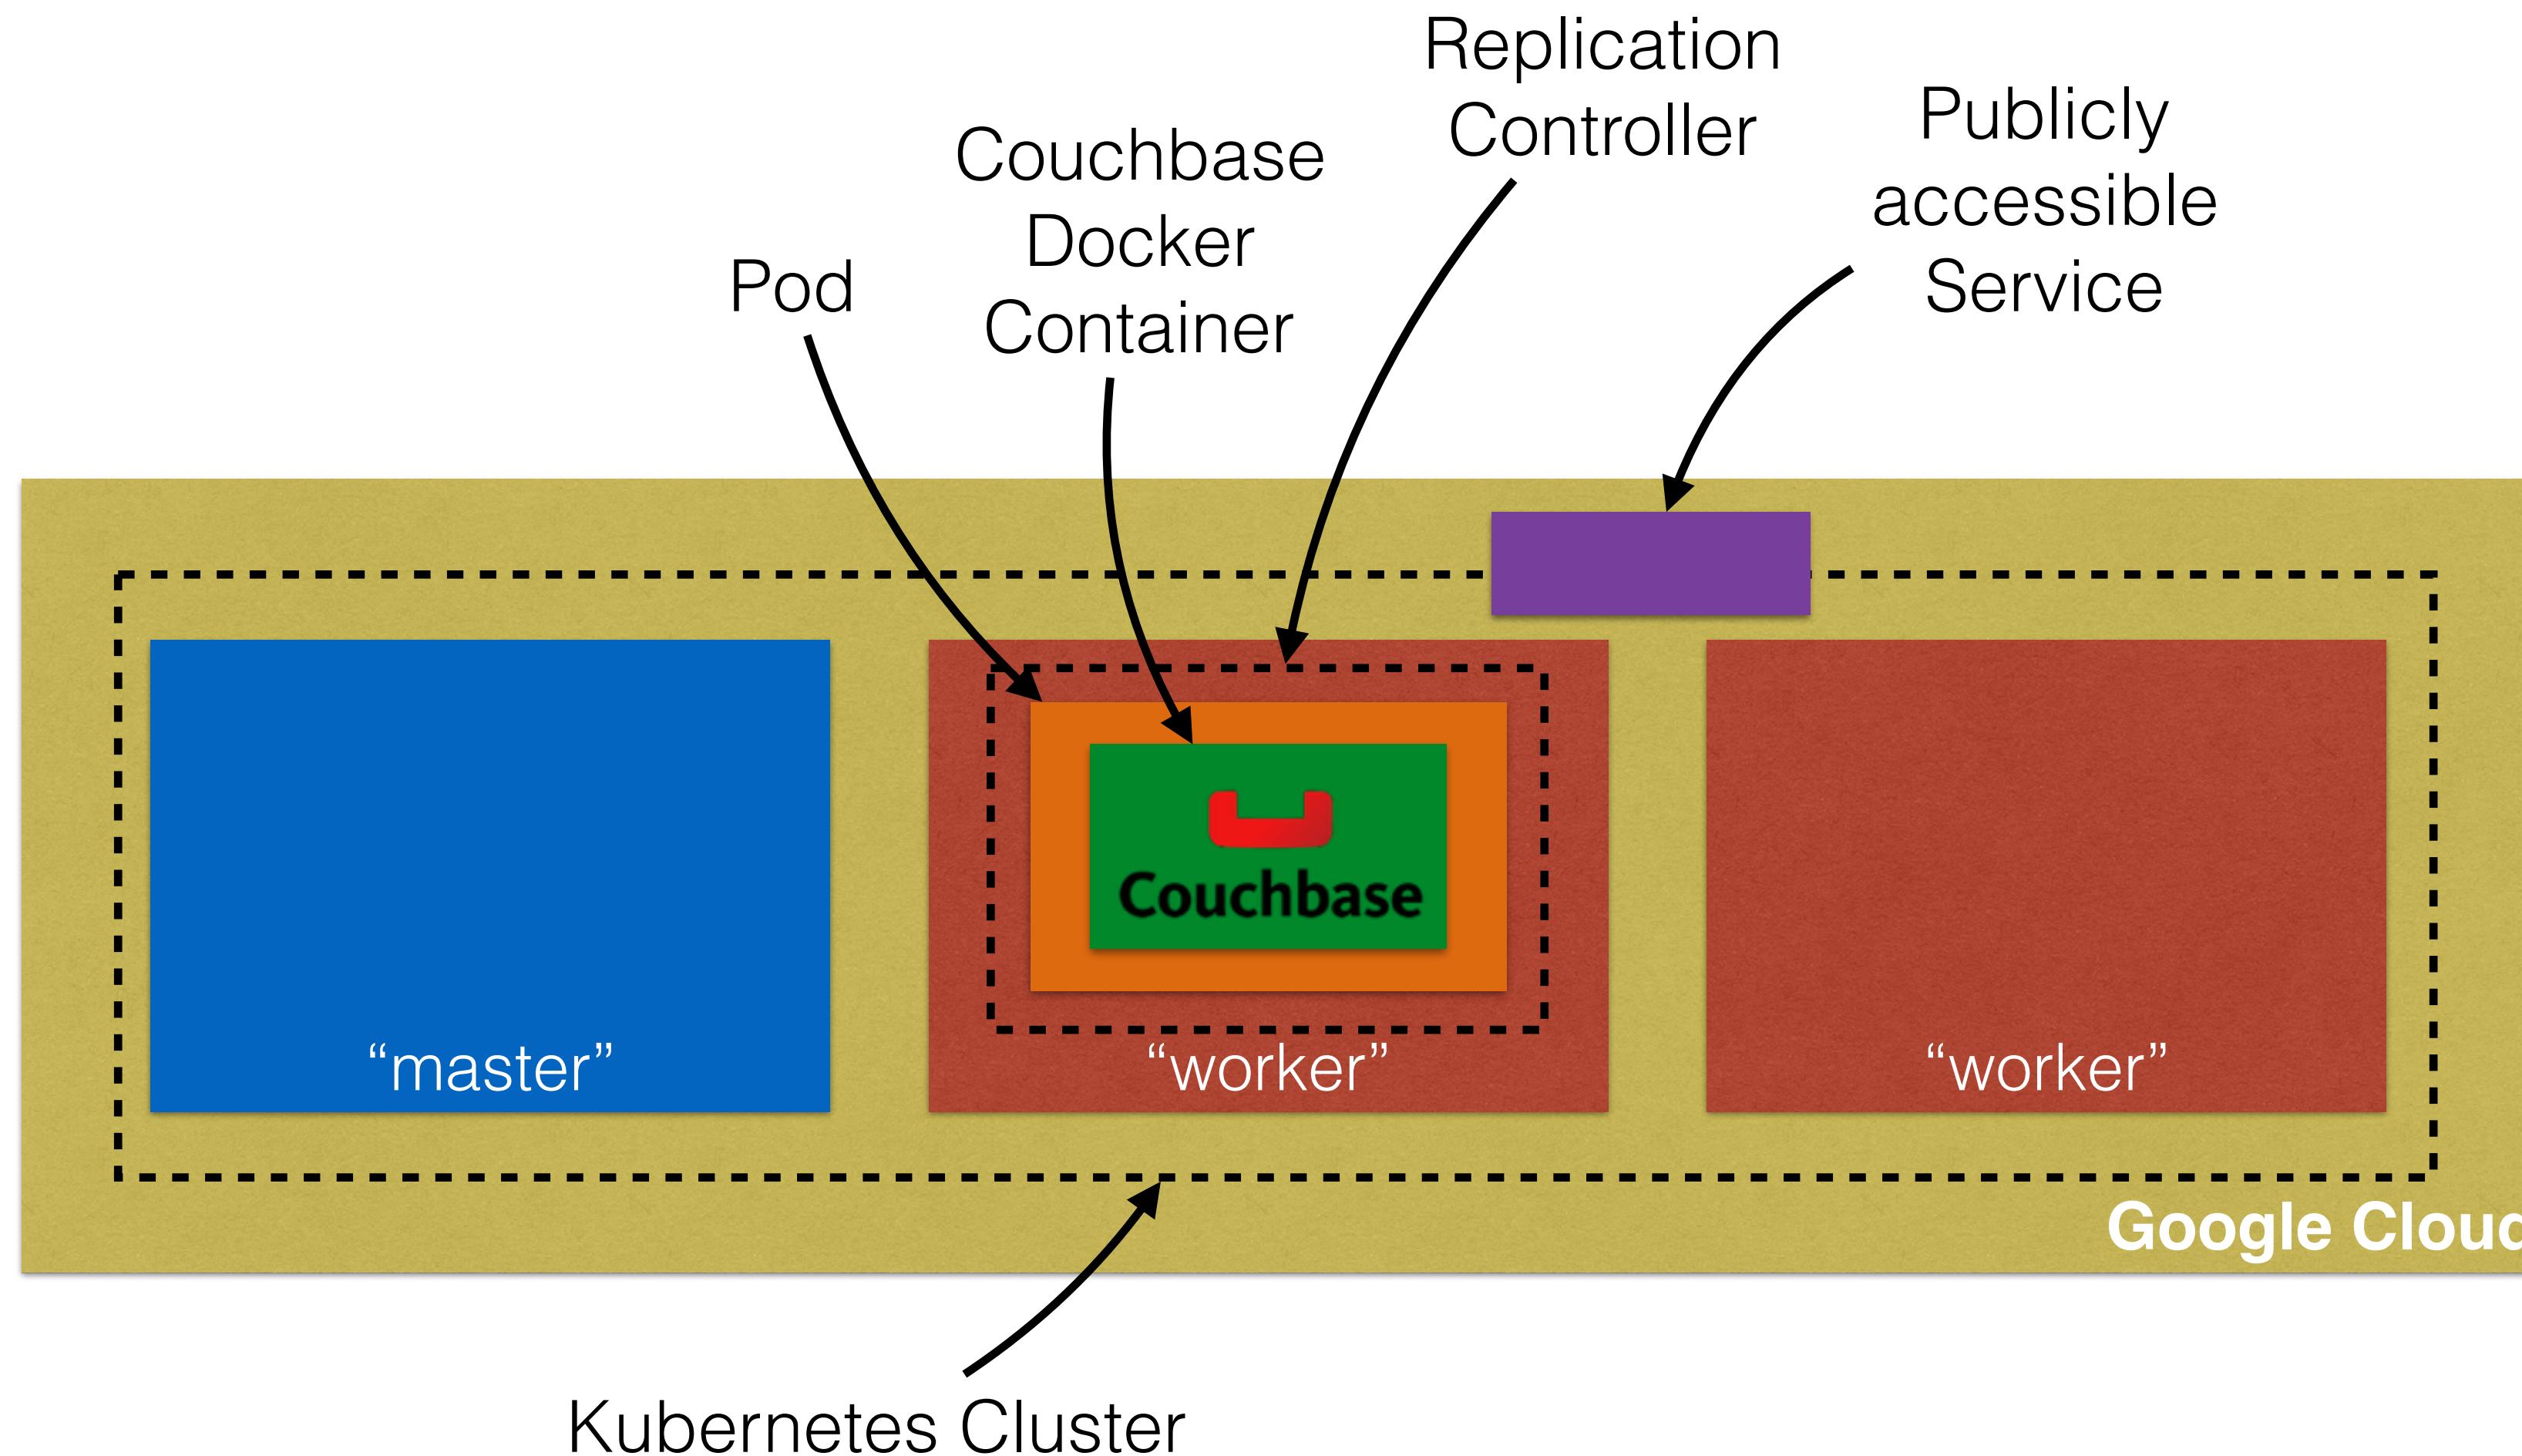
# Services



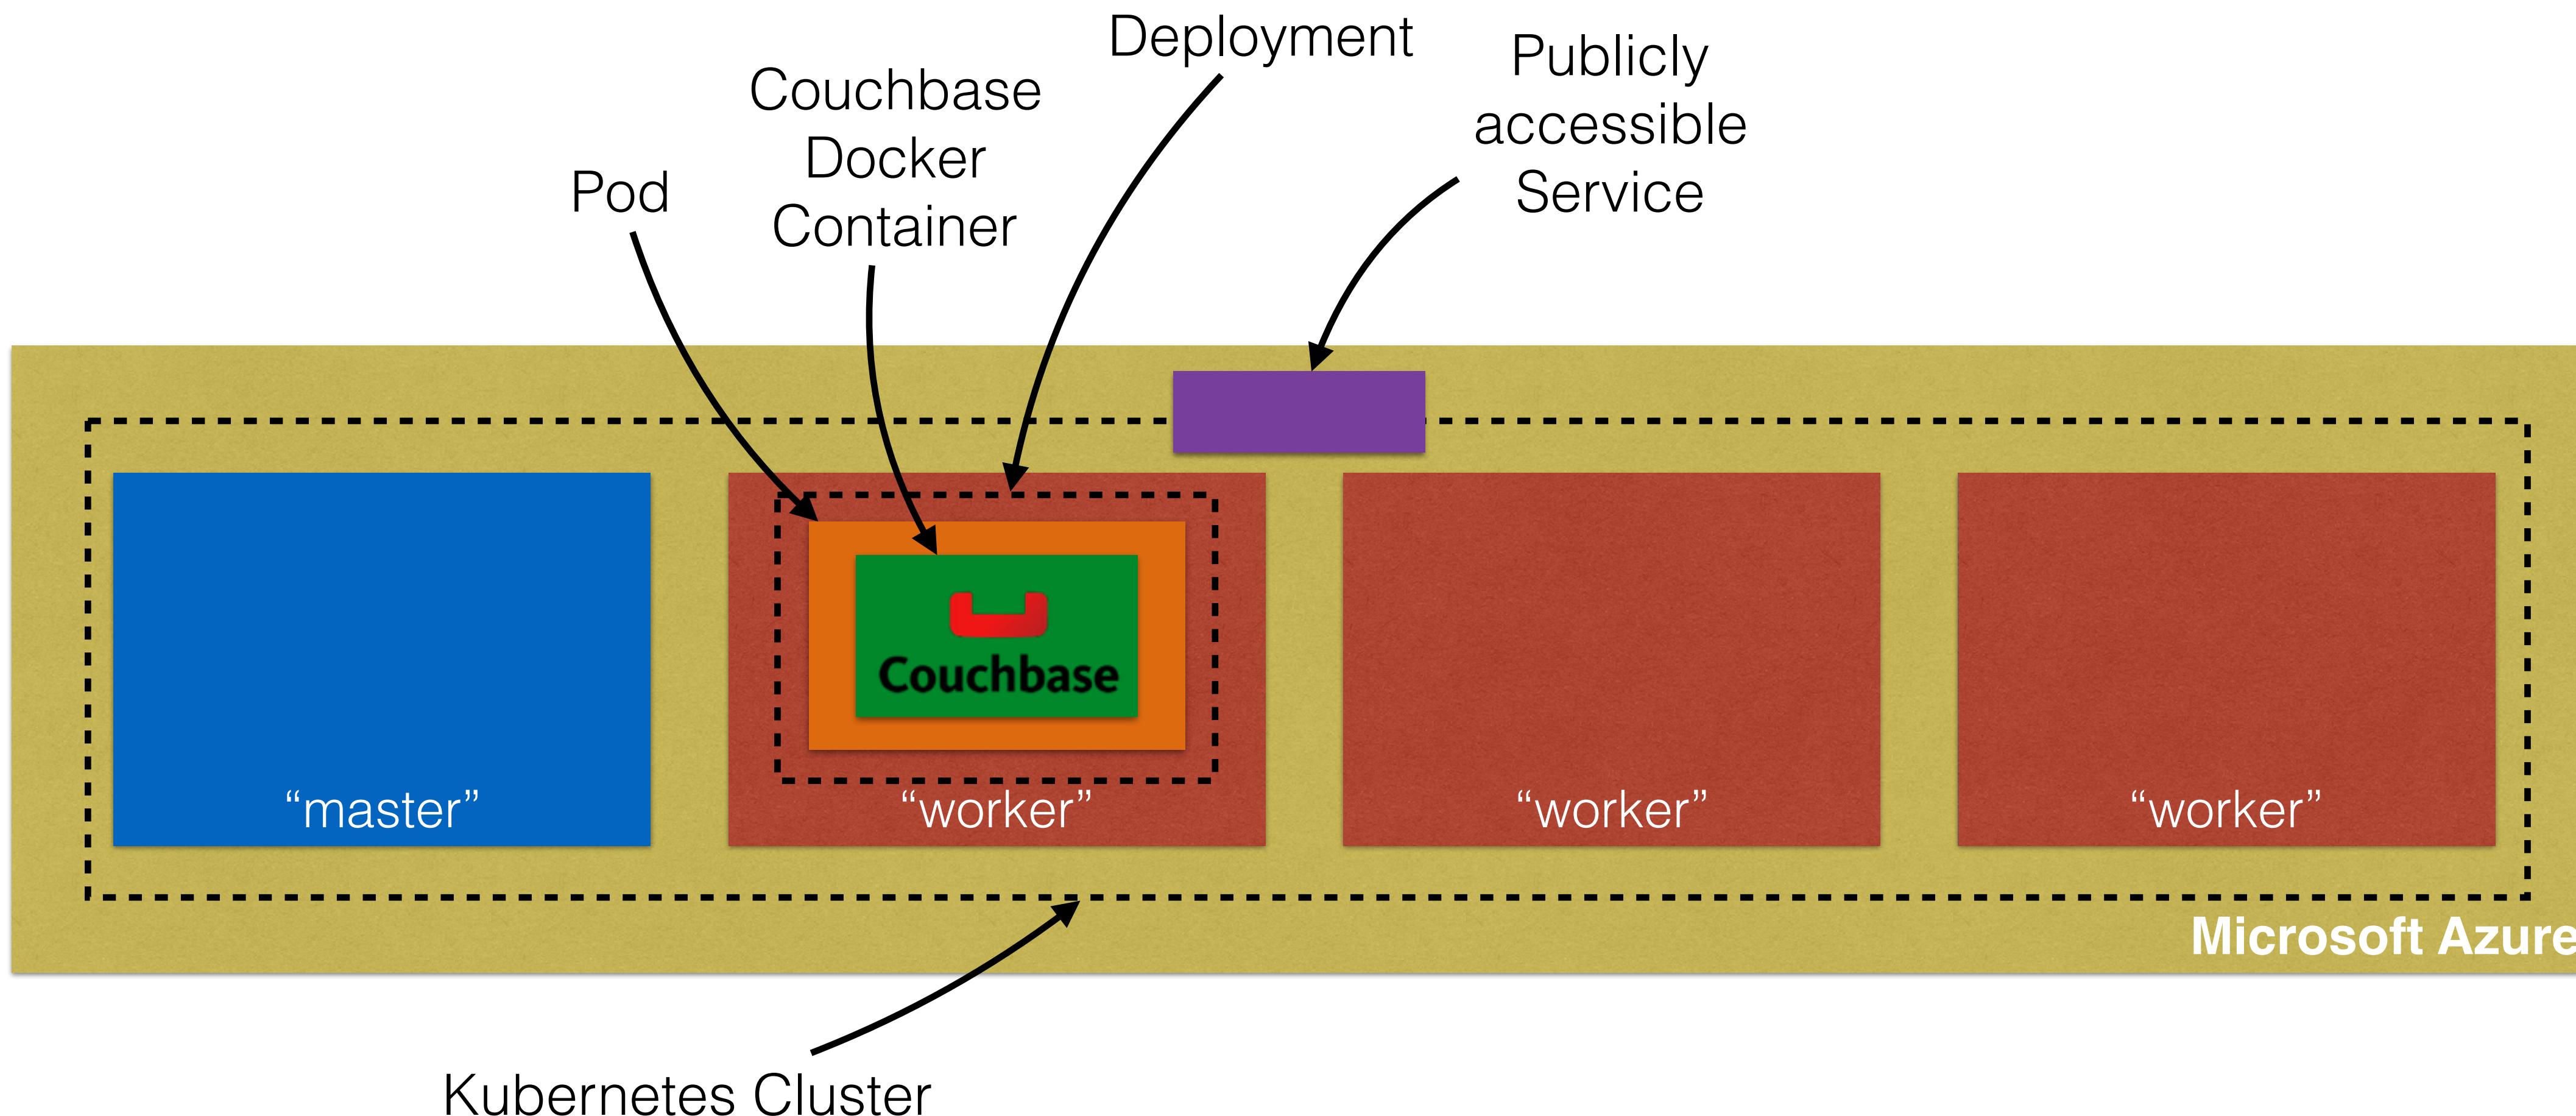
# Services - AWS



# Services - GCE



# Services - Azure

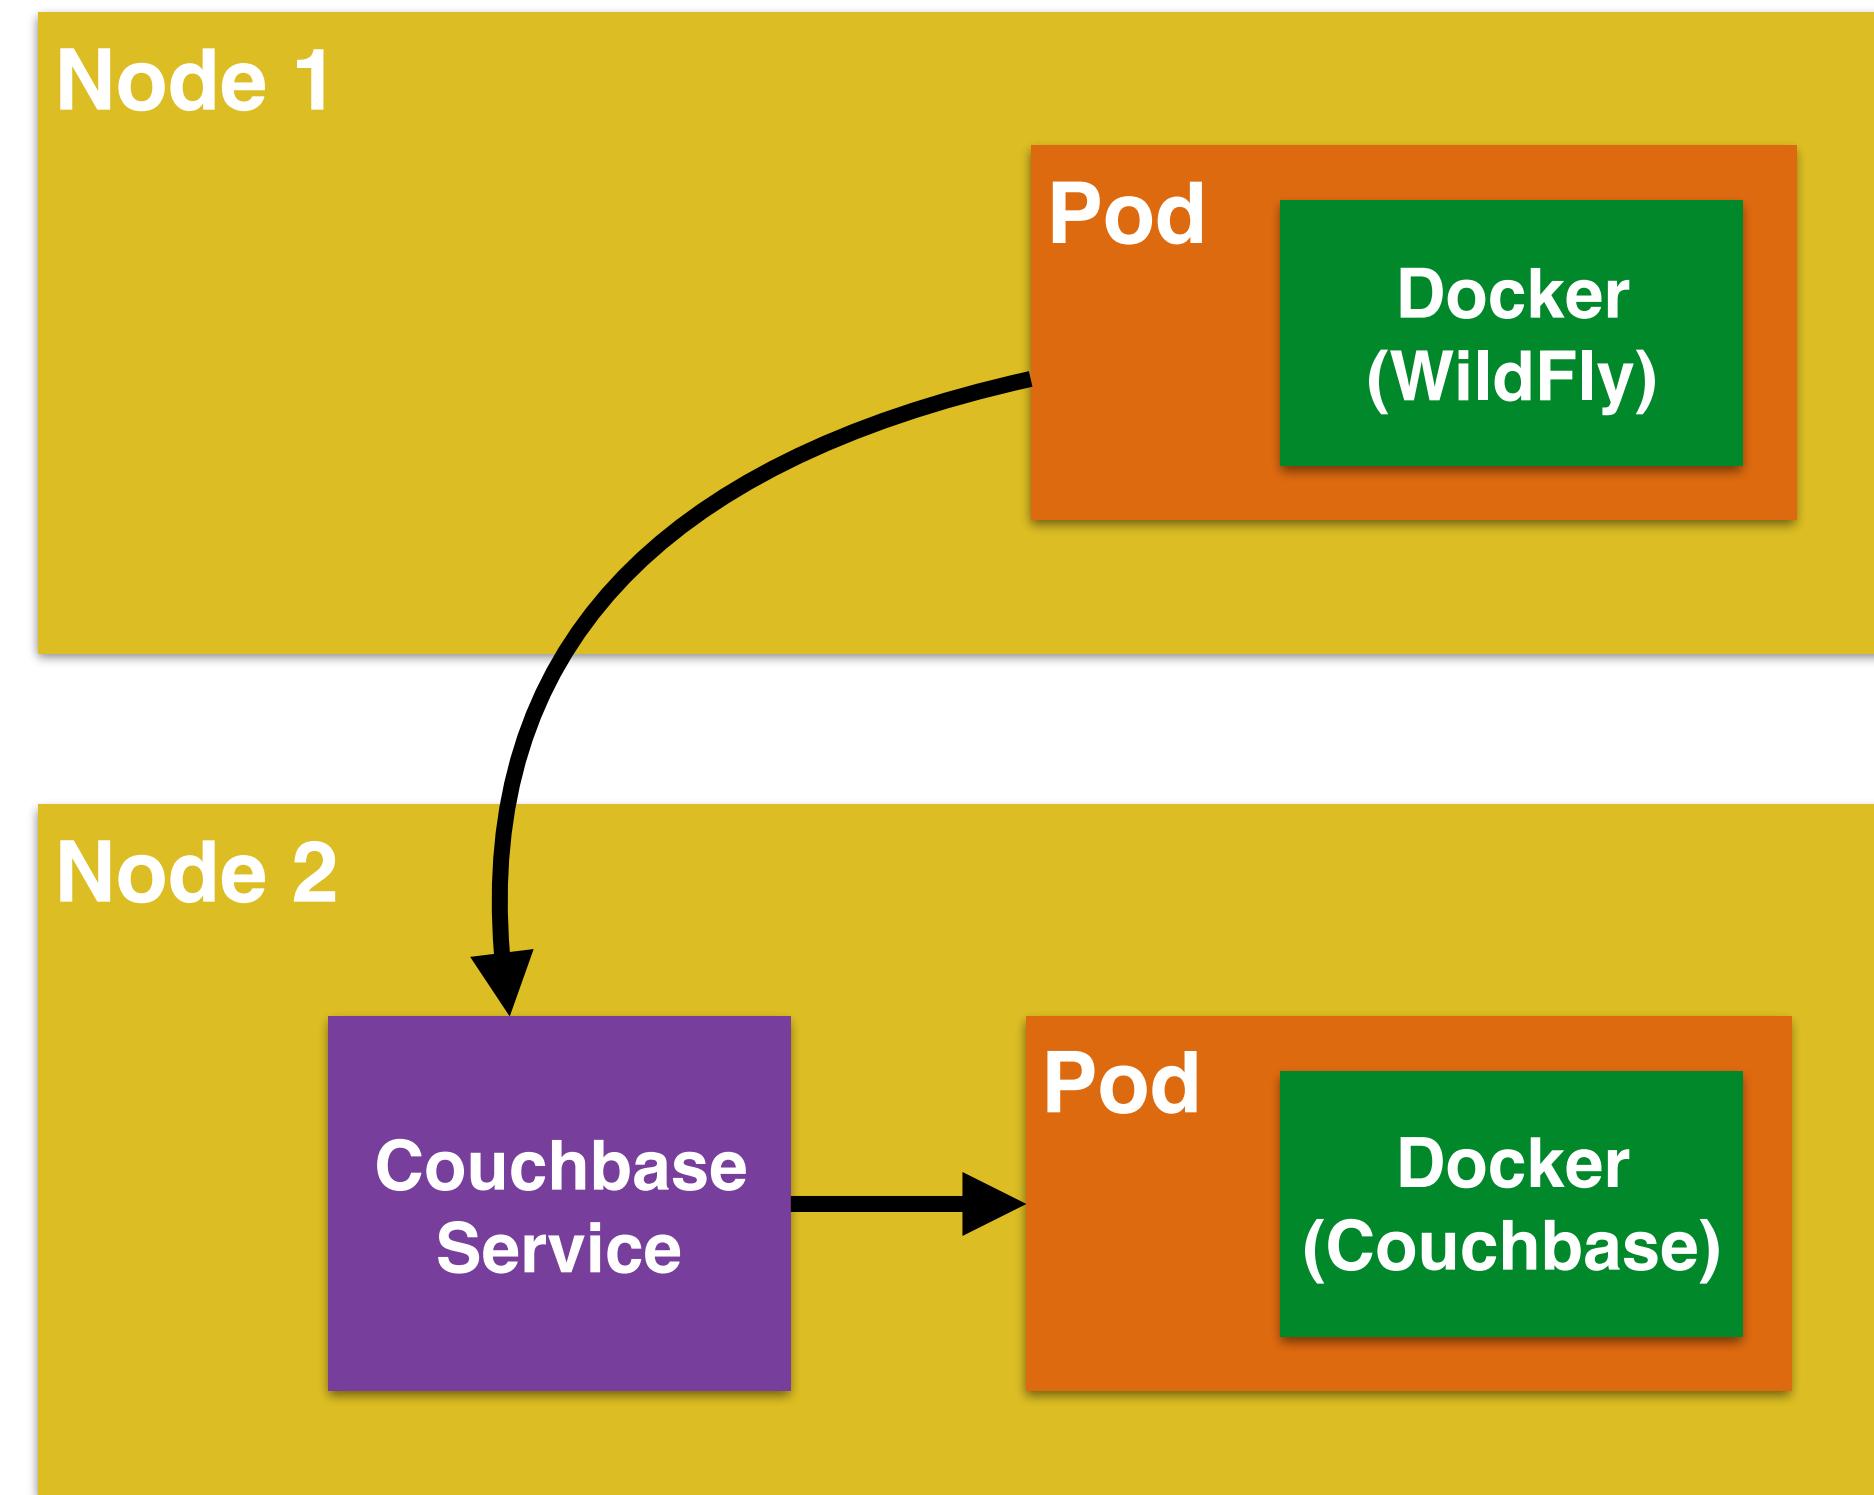


# Services

```
1 apiVersion: v1
2 kind: ReplicationController
3 metadata:
4   name: couchbase-rc
5   labels:
6     name: couchbase-rc
7     context: docker-k8s-lab
8 spec:
9   replicas: 1
10  template:
11    metadata:
12      name: couchbase-rc-pod
13      labels:
14        name: couchbase-rc-pod
15        context: docker-k8s-lab
16    spec:
17      containers:
18        - name: couchbase-rc-pod
19          image: couchbase/server
20          volumeMounts:
21            - mountPath: "/opt/couchbase/var"
22              name: mypd
23          ports:
24            - containerPort: 8091
25              hostPort: 8091
26            - containerPort: 8092
27            - containerPort: 8093
28            - containerPort: 11210
```

```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: couchbase-service
5   labels:
6     name: couchbase-service-pod
7     context: docker-k8s-lab
8 spec:
9   ports:
10    - port: 8091
11    - port: 8092
12    - port: 8093
13    - port: 11210
14  # label keys and values that must be present on pods
15  selection:
16    name: couchbase-rc-pod
17    context: docker-k8s-lab
```

# Service across Two Nodes



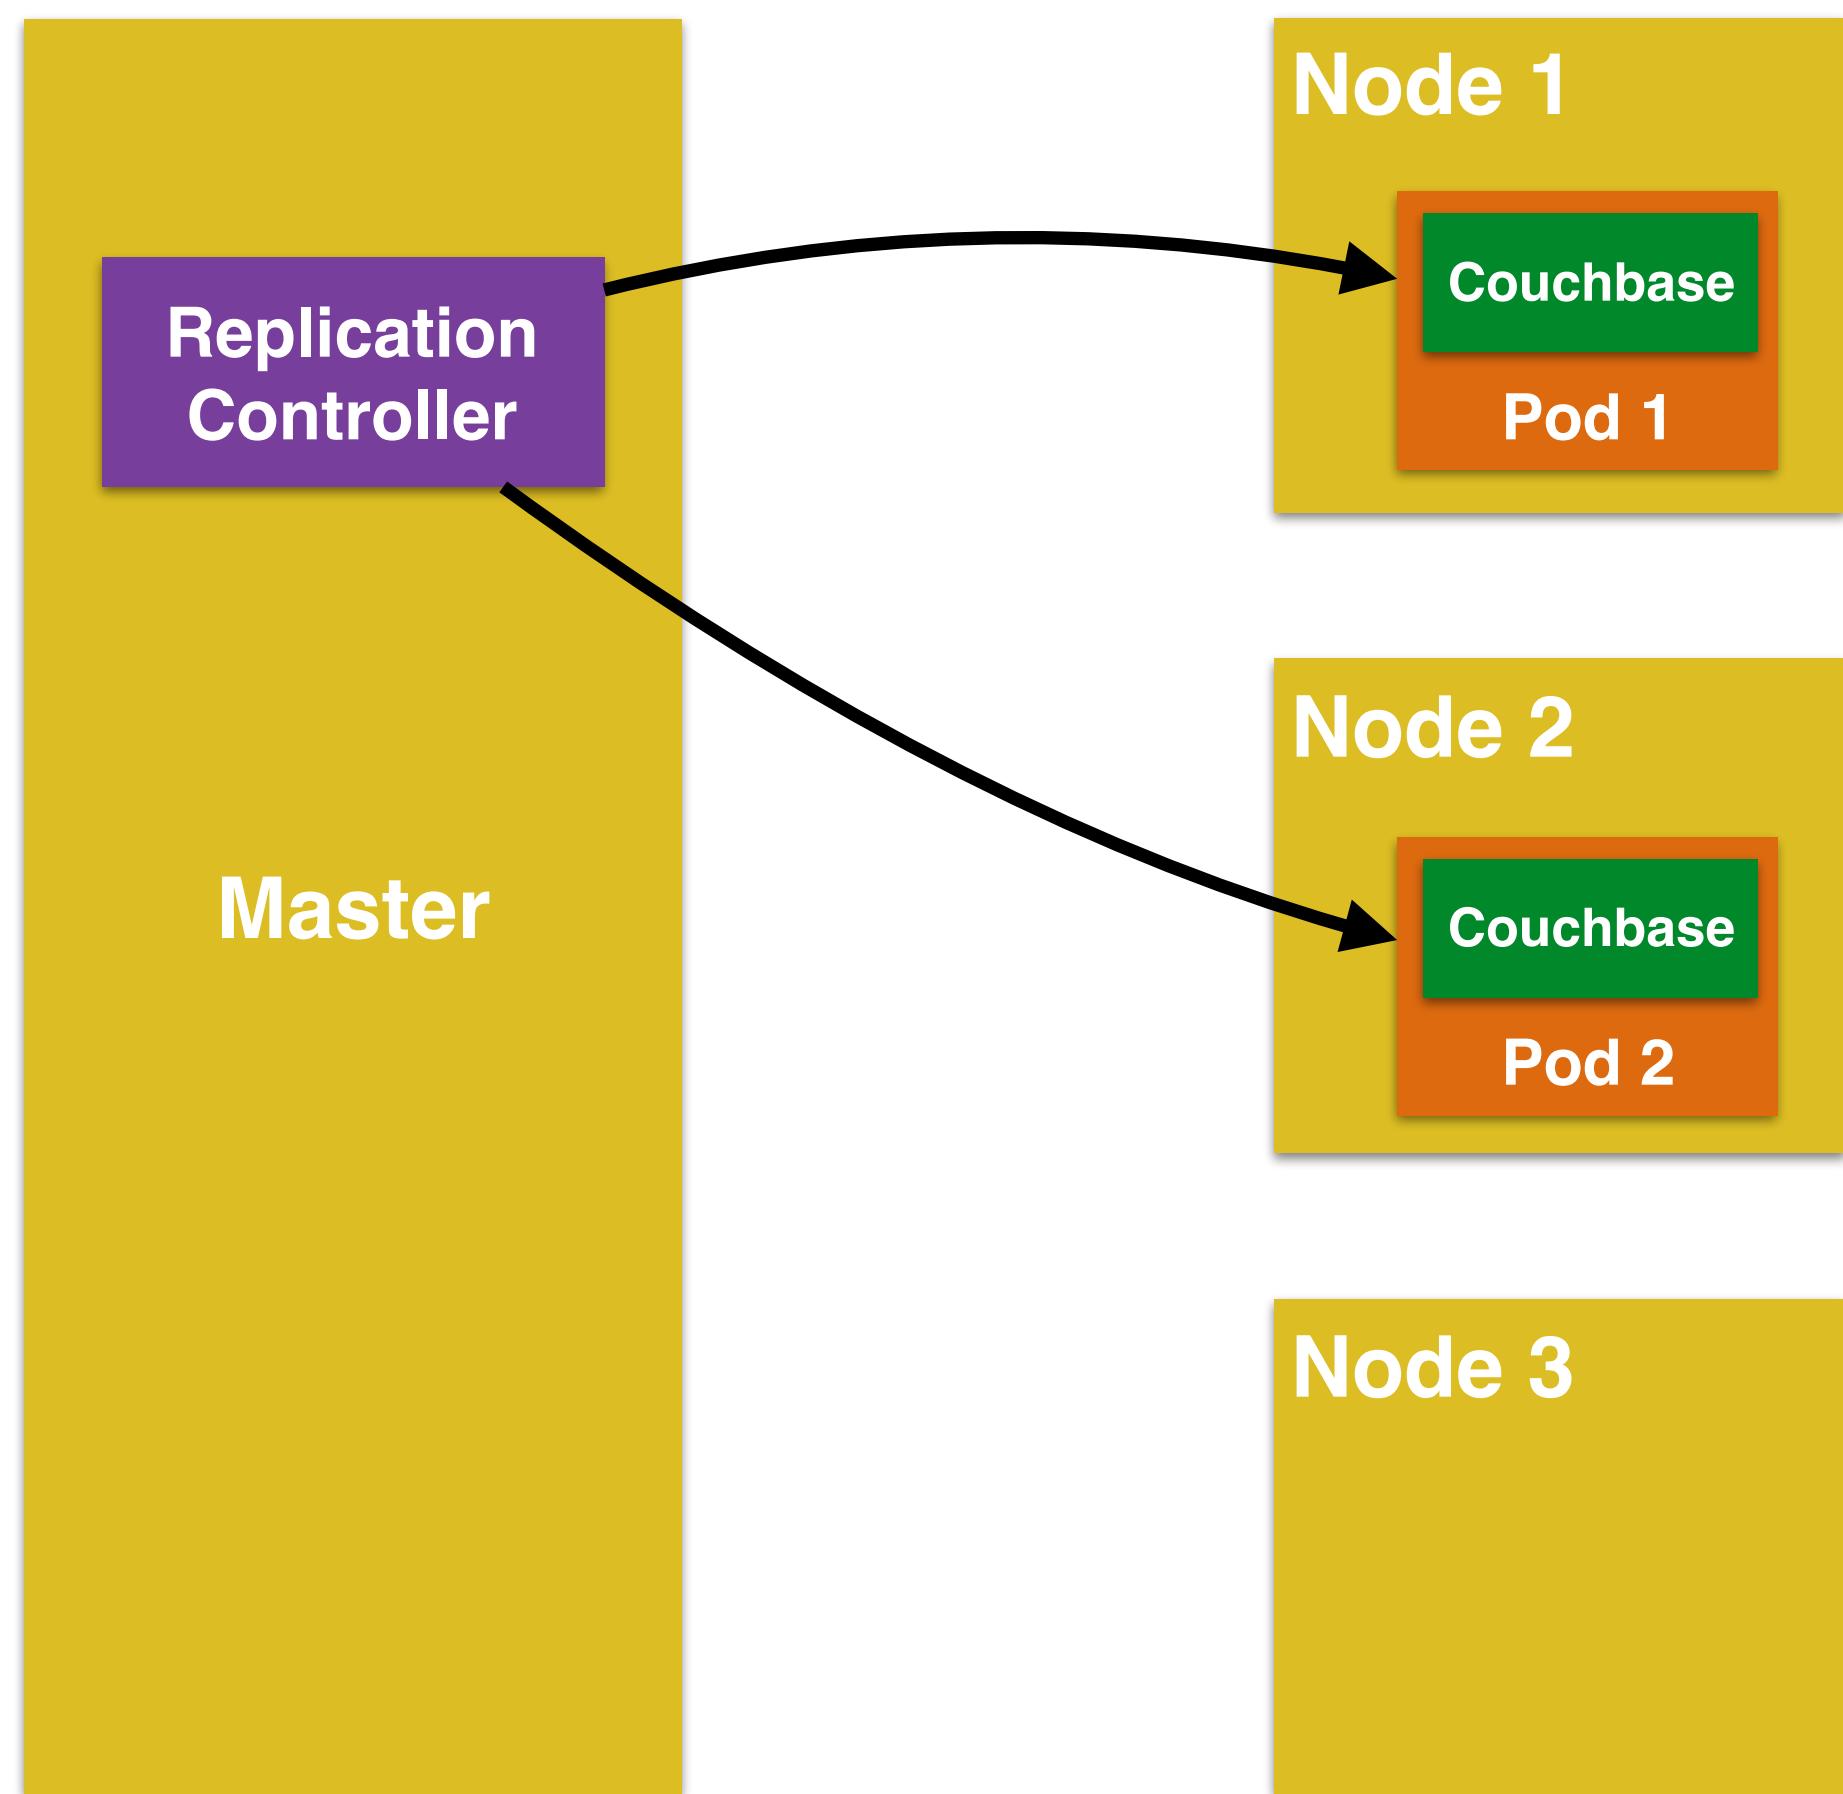
# Replication Controller

- Ensures that a specified number of pod "replicas" are running
  - Pod templates are cookie cutters
  - Rescheduling
  - Manual or auto-scale replicas
  - Rolling updates
- Recommended to wrap a Pod or Service in a RC
- Only appropriate for Pods with `Restart=Always` policy (default)

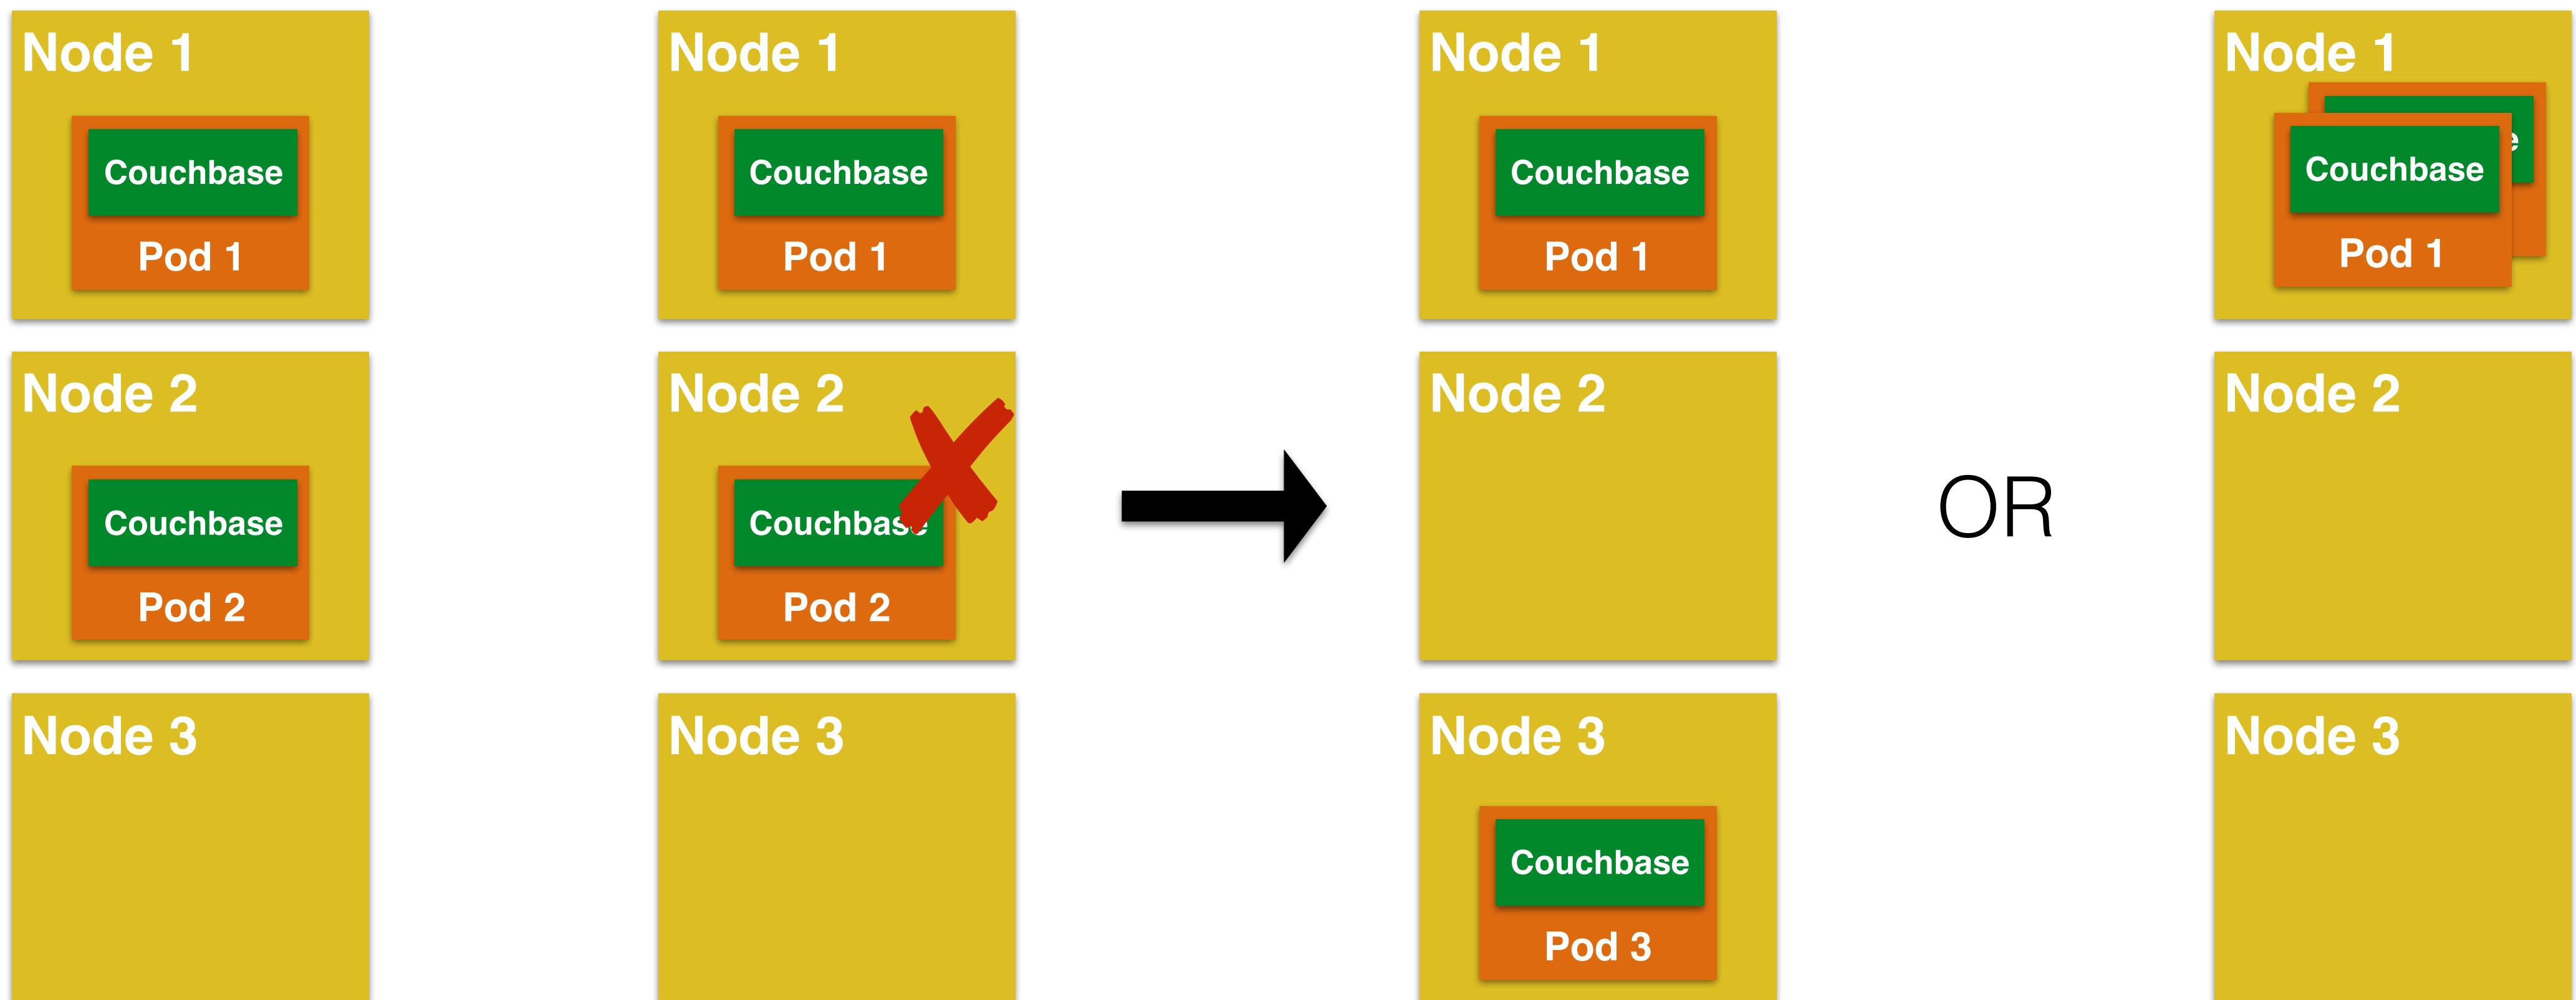
# Replication Controller Configuration

```
1  apiVersion: v1
2  kind: ReplicationController
3  metadata:
4    name: wildfly-rc
5    labels:
6      name: wildfly
7      context: docker-k8s-lab
8  spec:
9    replicas: 1
10   template:
11     metadata:
12       labels:
13         name: wildfly
14     spec:
15       containers:
16         - name: wildfly-rc-pod
17           image: arungupta/wildfly-mysql-javaee7:k8s
18           ports:
19             - containerPort: 8080
```

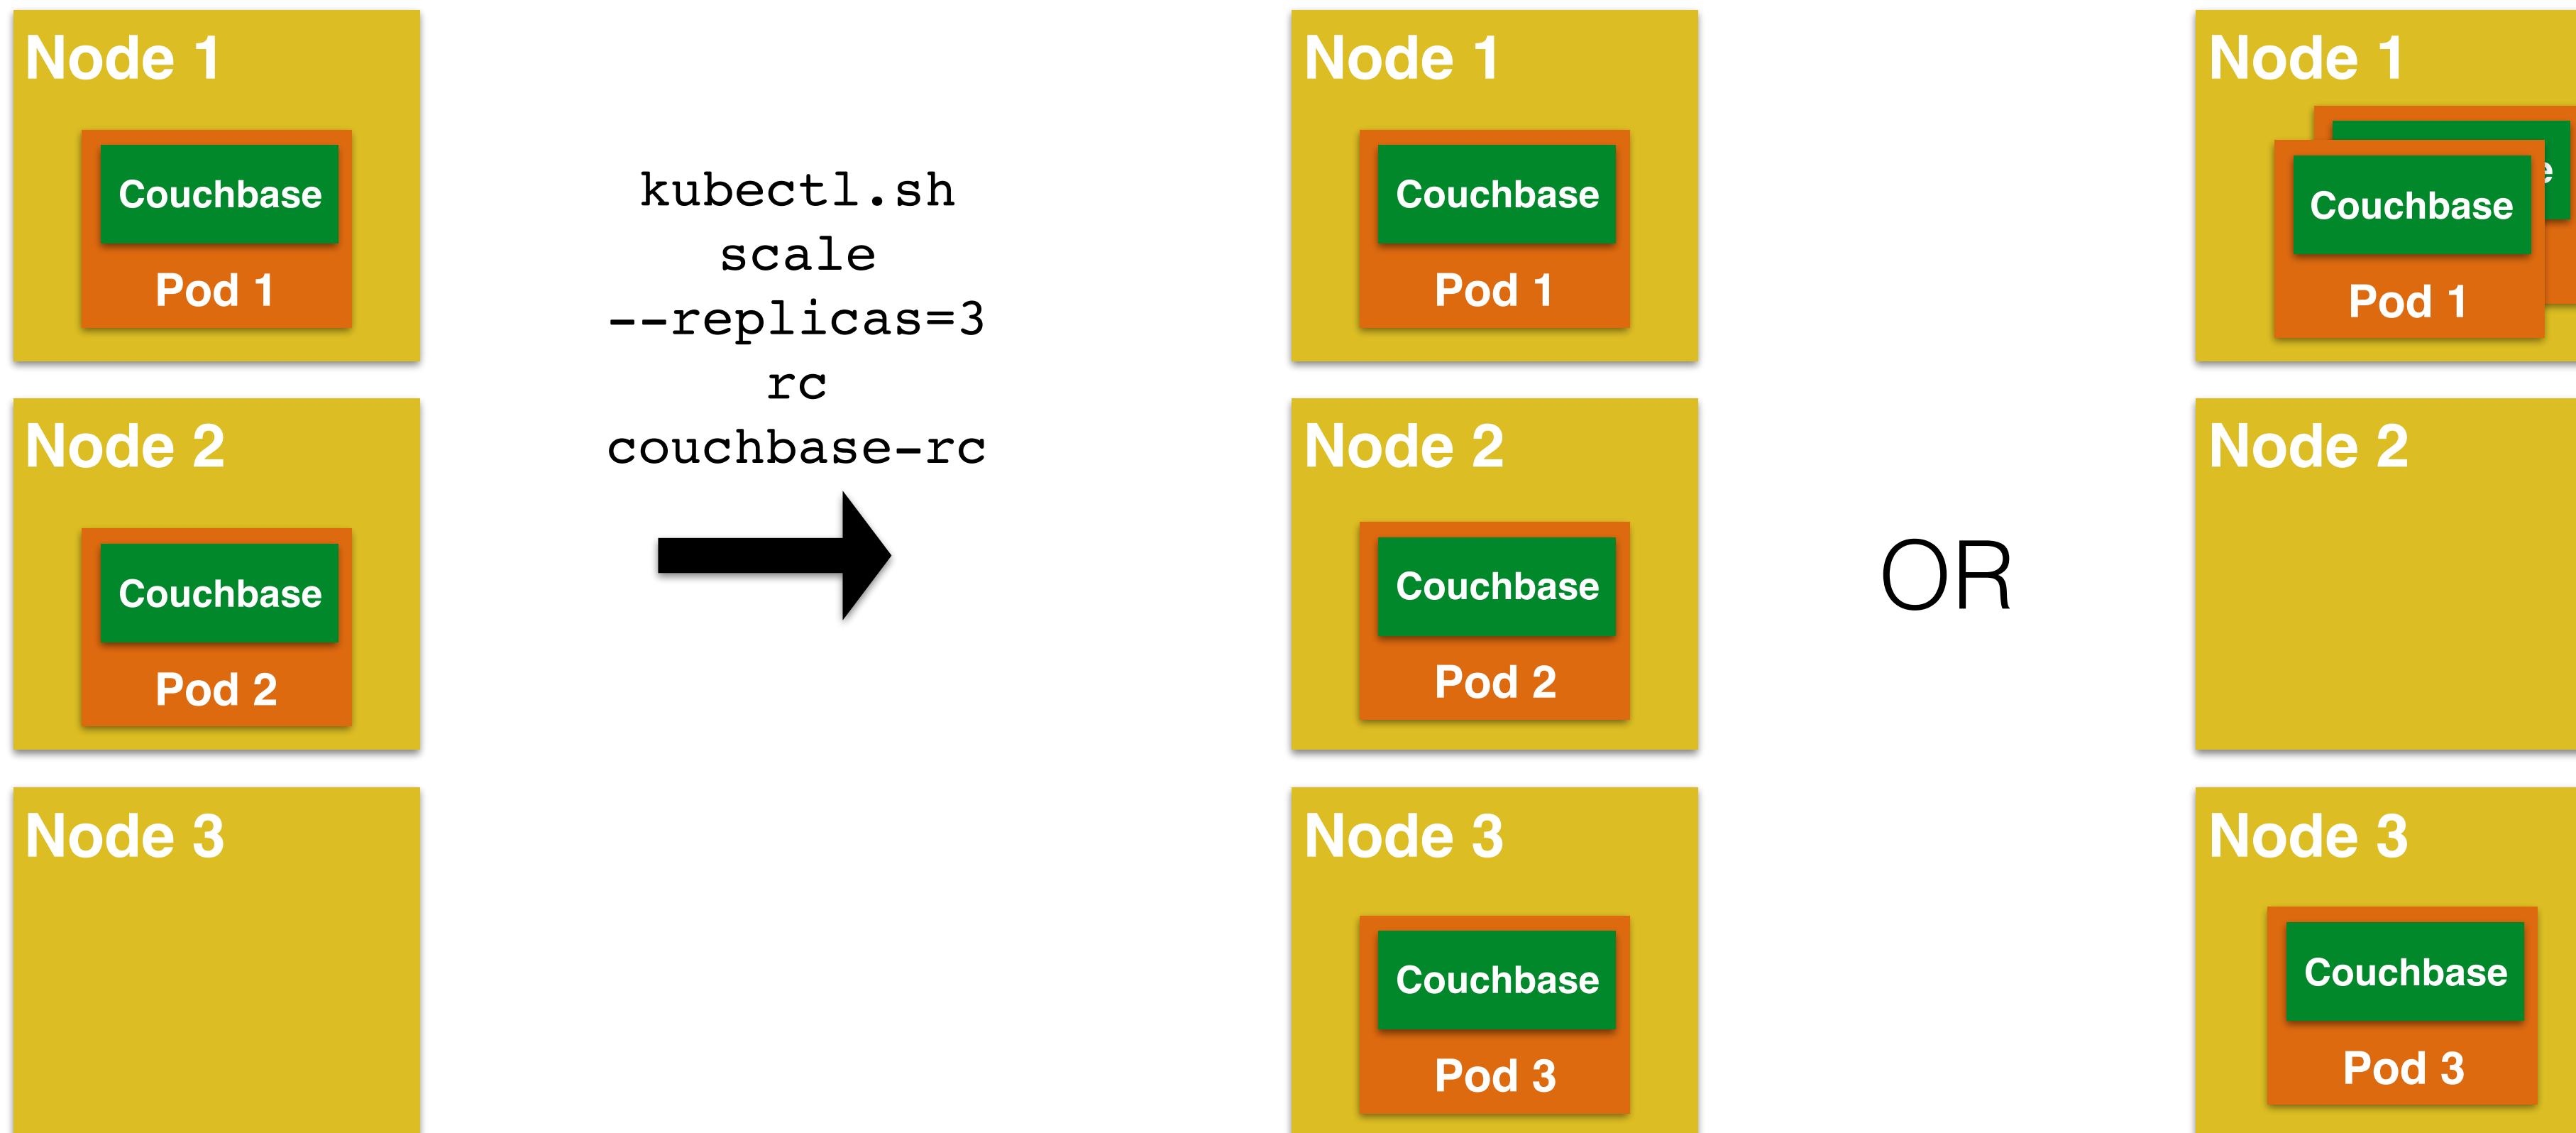
# Replication Controller



# Replication Controller: Automatic Rescheduling



# Replication Controller: Scaling



# Sample Production Deployment

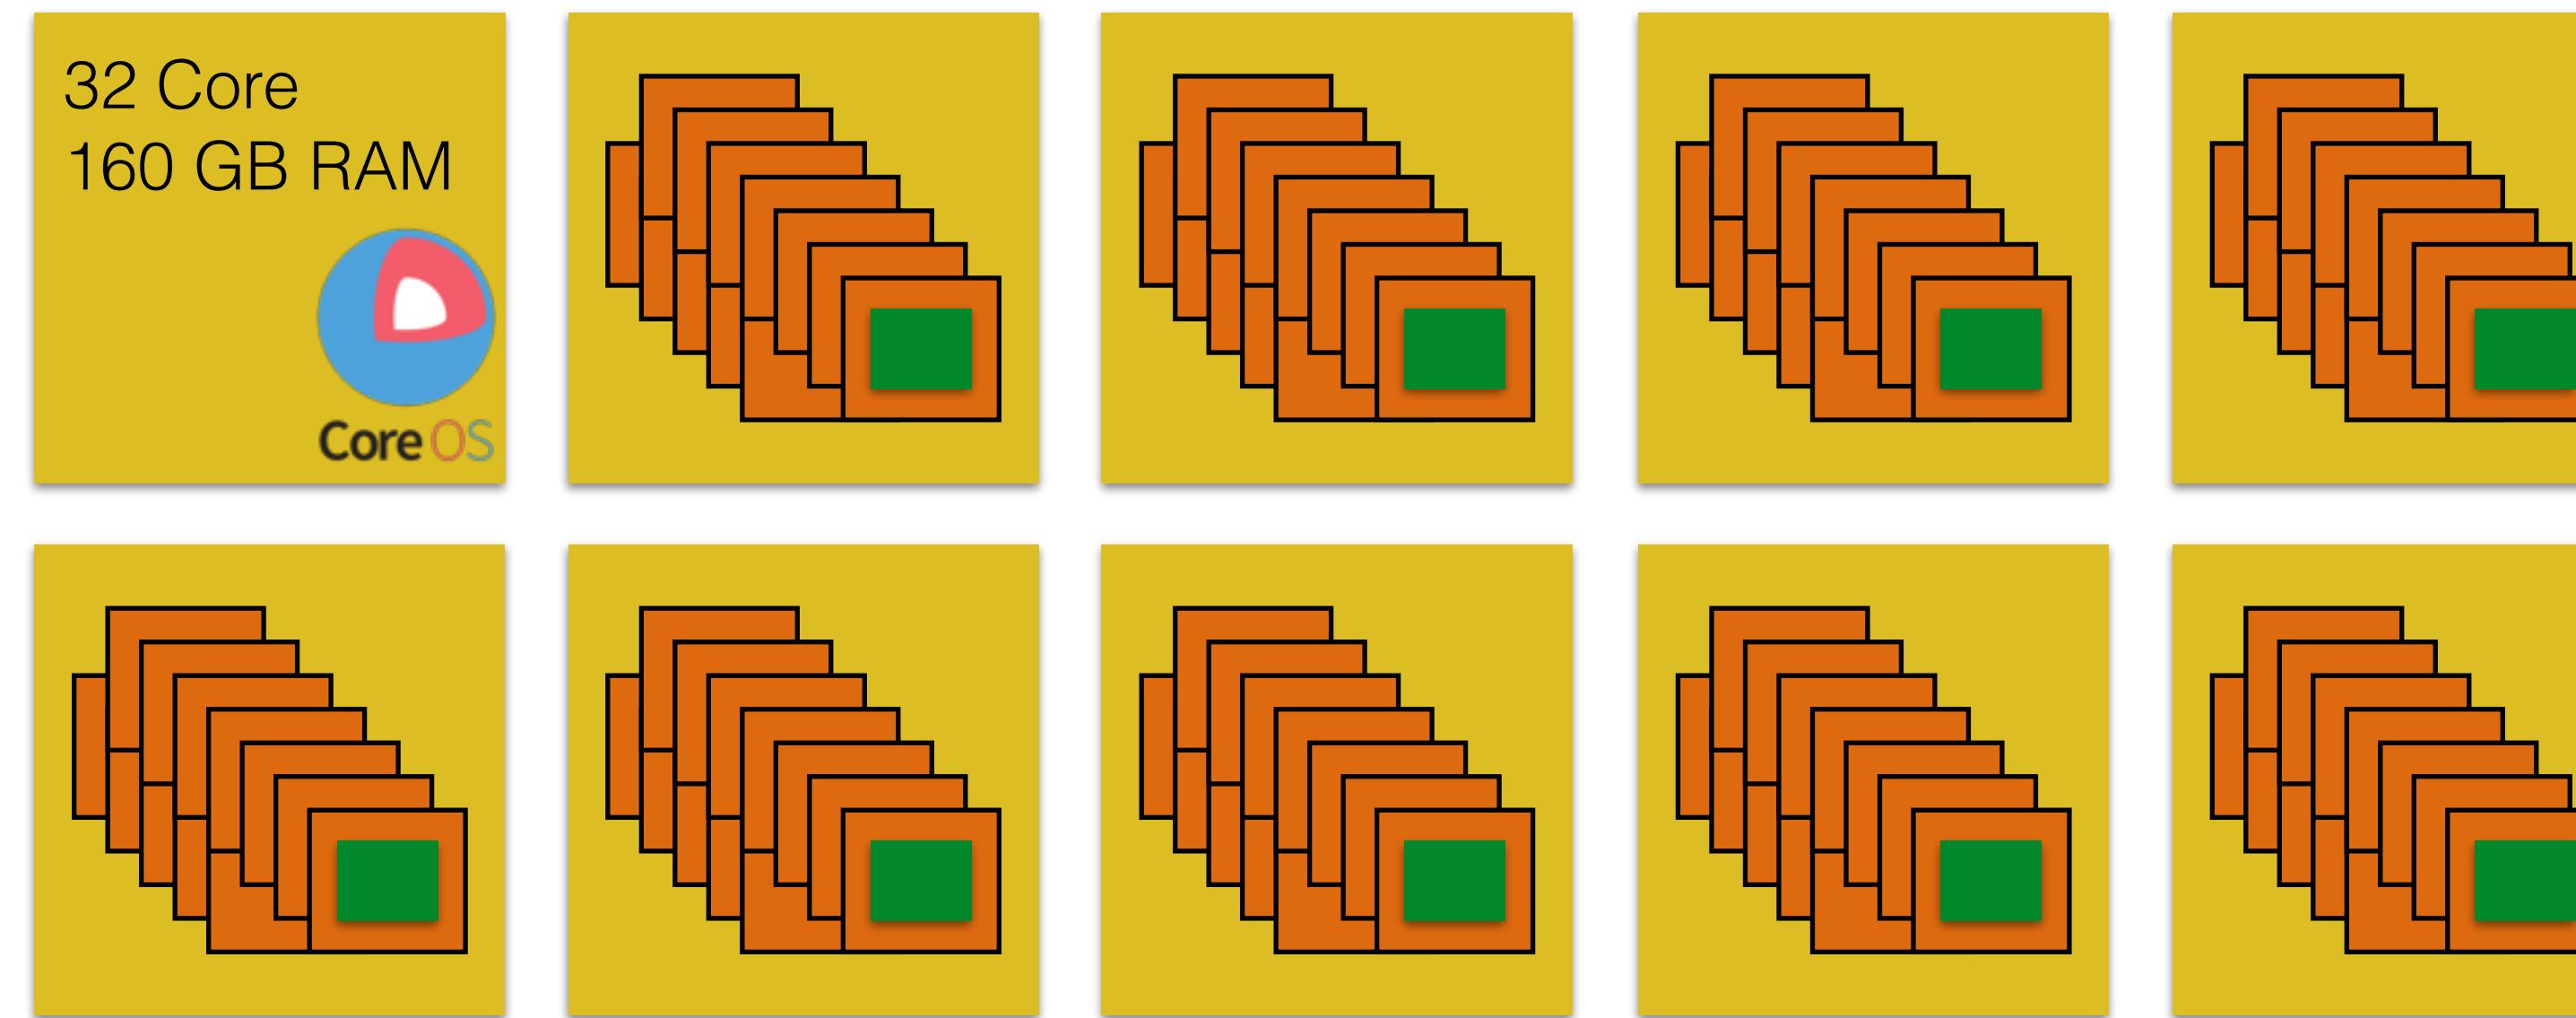
- [www.wombatsoftware.de](http://www.wombatsoftware.de)
- [shopadvisors.de](http://shopadvisors.de): E-commerce optimization and monitoring tools for increase of sales



**WOMBAT**  
SOFTWARE



# Sample Production Deployment



Load	Containers
Normal	400
Peak	600

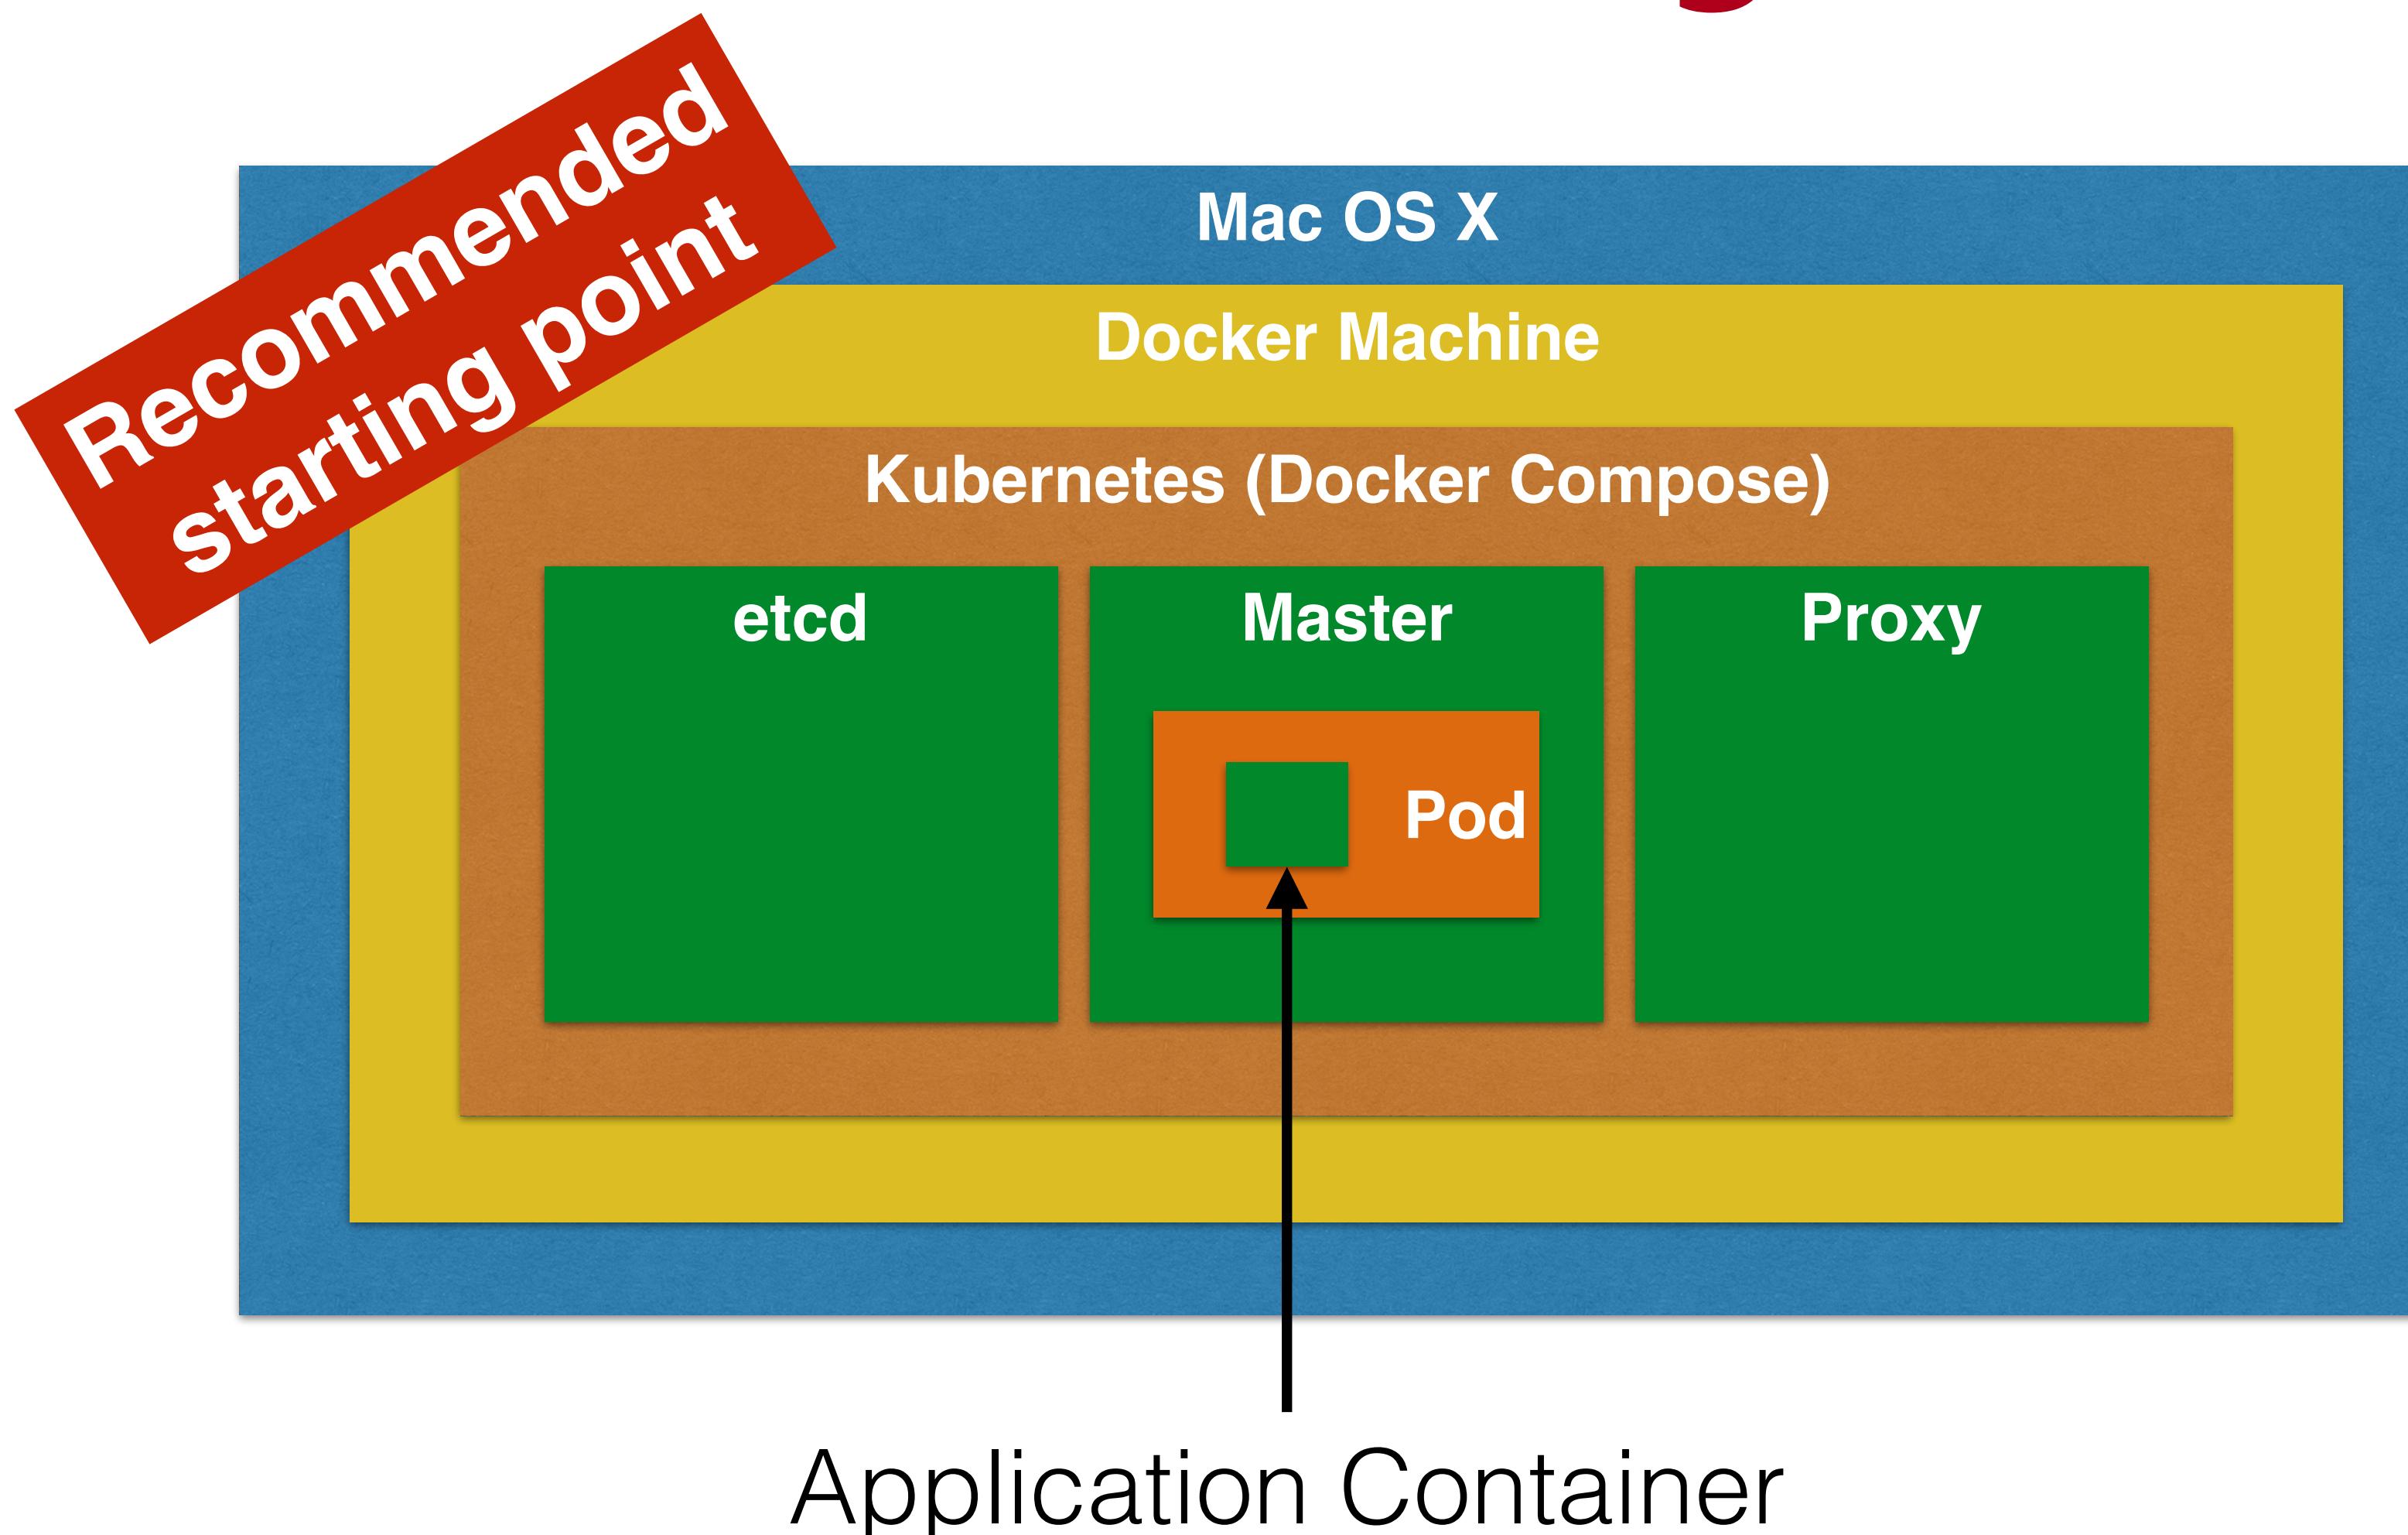
# Health Checks

- Restarts Pod, if wrapped in RC
- Application-level health checks
  - HTTP
  - Container Exec
  - TCP Socket
- Health checks performed by Kubelet

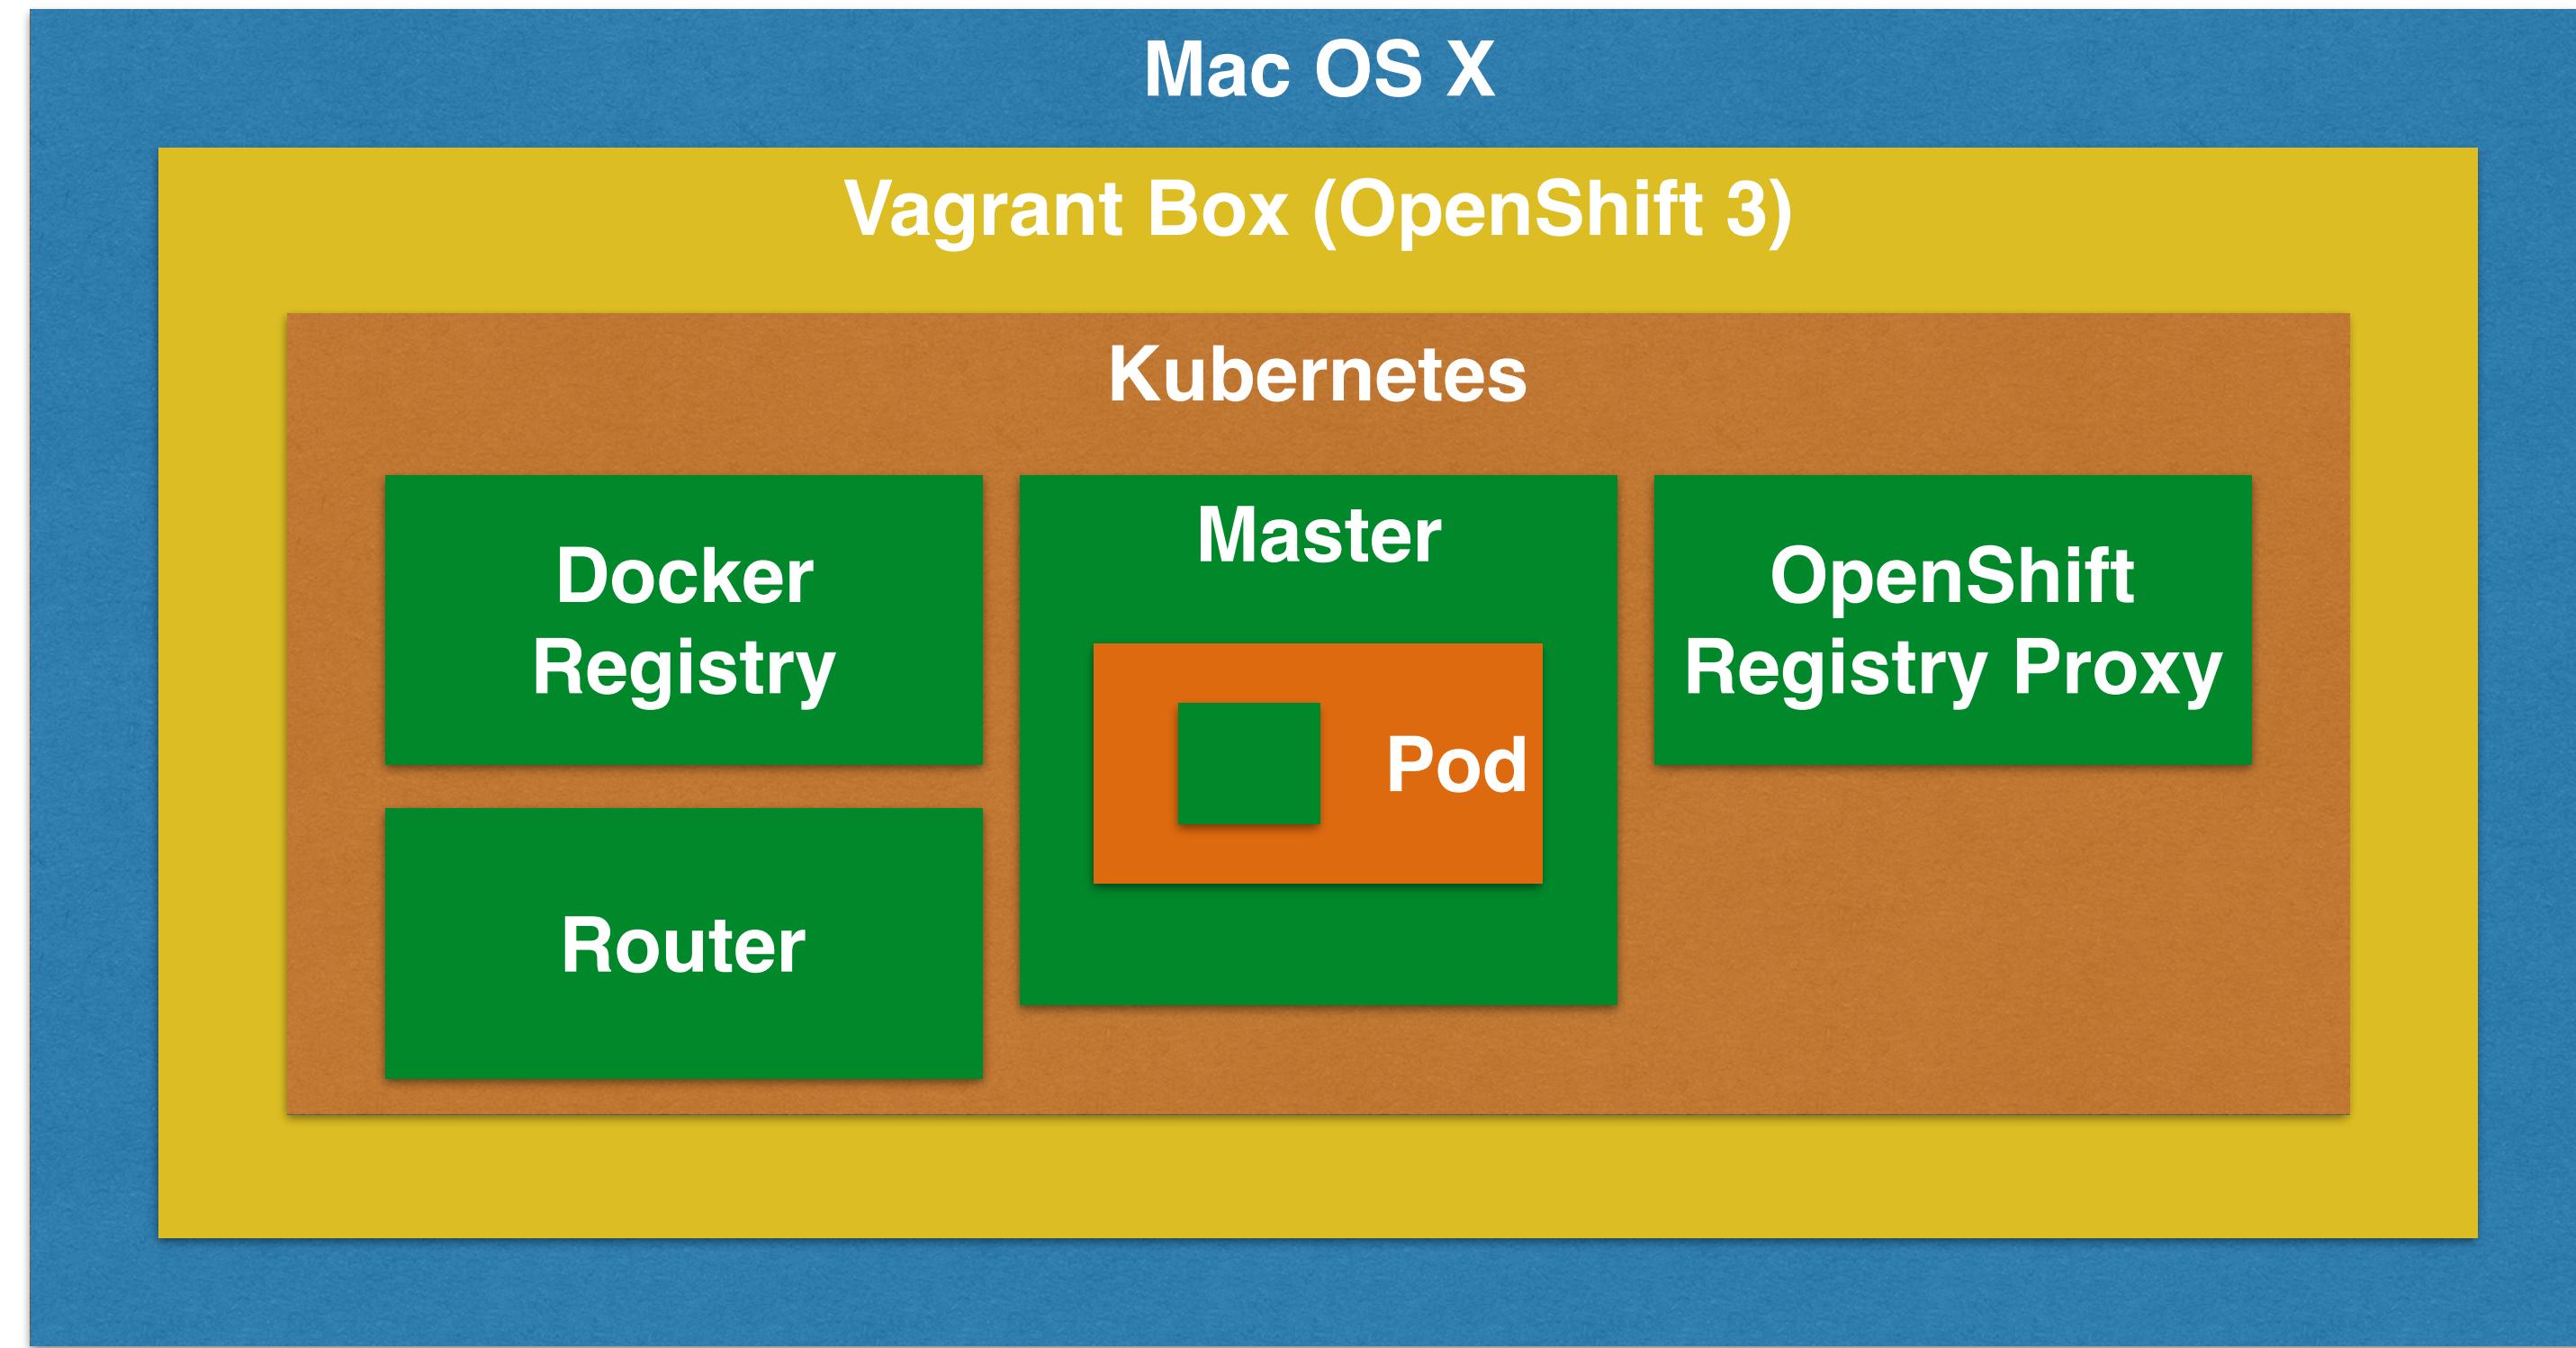
# Kubernetes using Docker

```
1 etcd:
2   image: gcr.io/google_containers/etcd:2.0.9
3   net: "host"
4   entrypoint: /usr/local/bin/etcd --addr=127.0.0.1:4001 --bind-addr=0.0.0.0:4001 --data-dir=/var/lib/etcd
5 master:
6   image: gcr.io/google_containers/hyperkube:v0.21.2
7   net: "host"
8   volumes:
9     - /var/run/docker.sock:/var/run/docker.sock
10  entrypoint: /hyperkube kubelet --api_servers=http://localhost:8080 --v=2 --address=0.0.0.0 --port=8080
11 proxy:
12   image: gcr.io/google_containers/hyperkube:v0.21.2
13   net: "host"
14   privileged: true
15   entrypoint: /hyperkube proxy --master=http://127.0.0.1:8080 --v=2
```

# Kubernetes using Docker



# OpenShift 3



# References

- [github.com/javaee-samples/docker-java](https://github.com/javaee-samples/docker-java)
- [kubernetes.io](https://kubernetes.io)
- Containers recipe: [couchbase.com/containers](https://couchbase.com/containers)