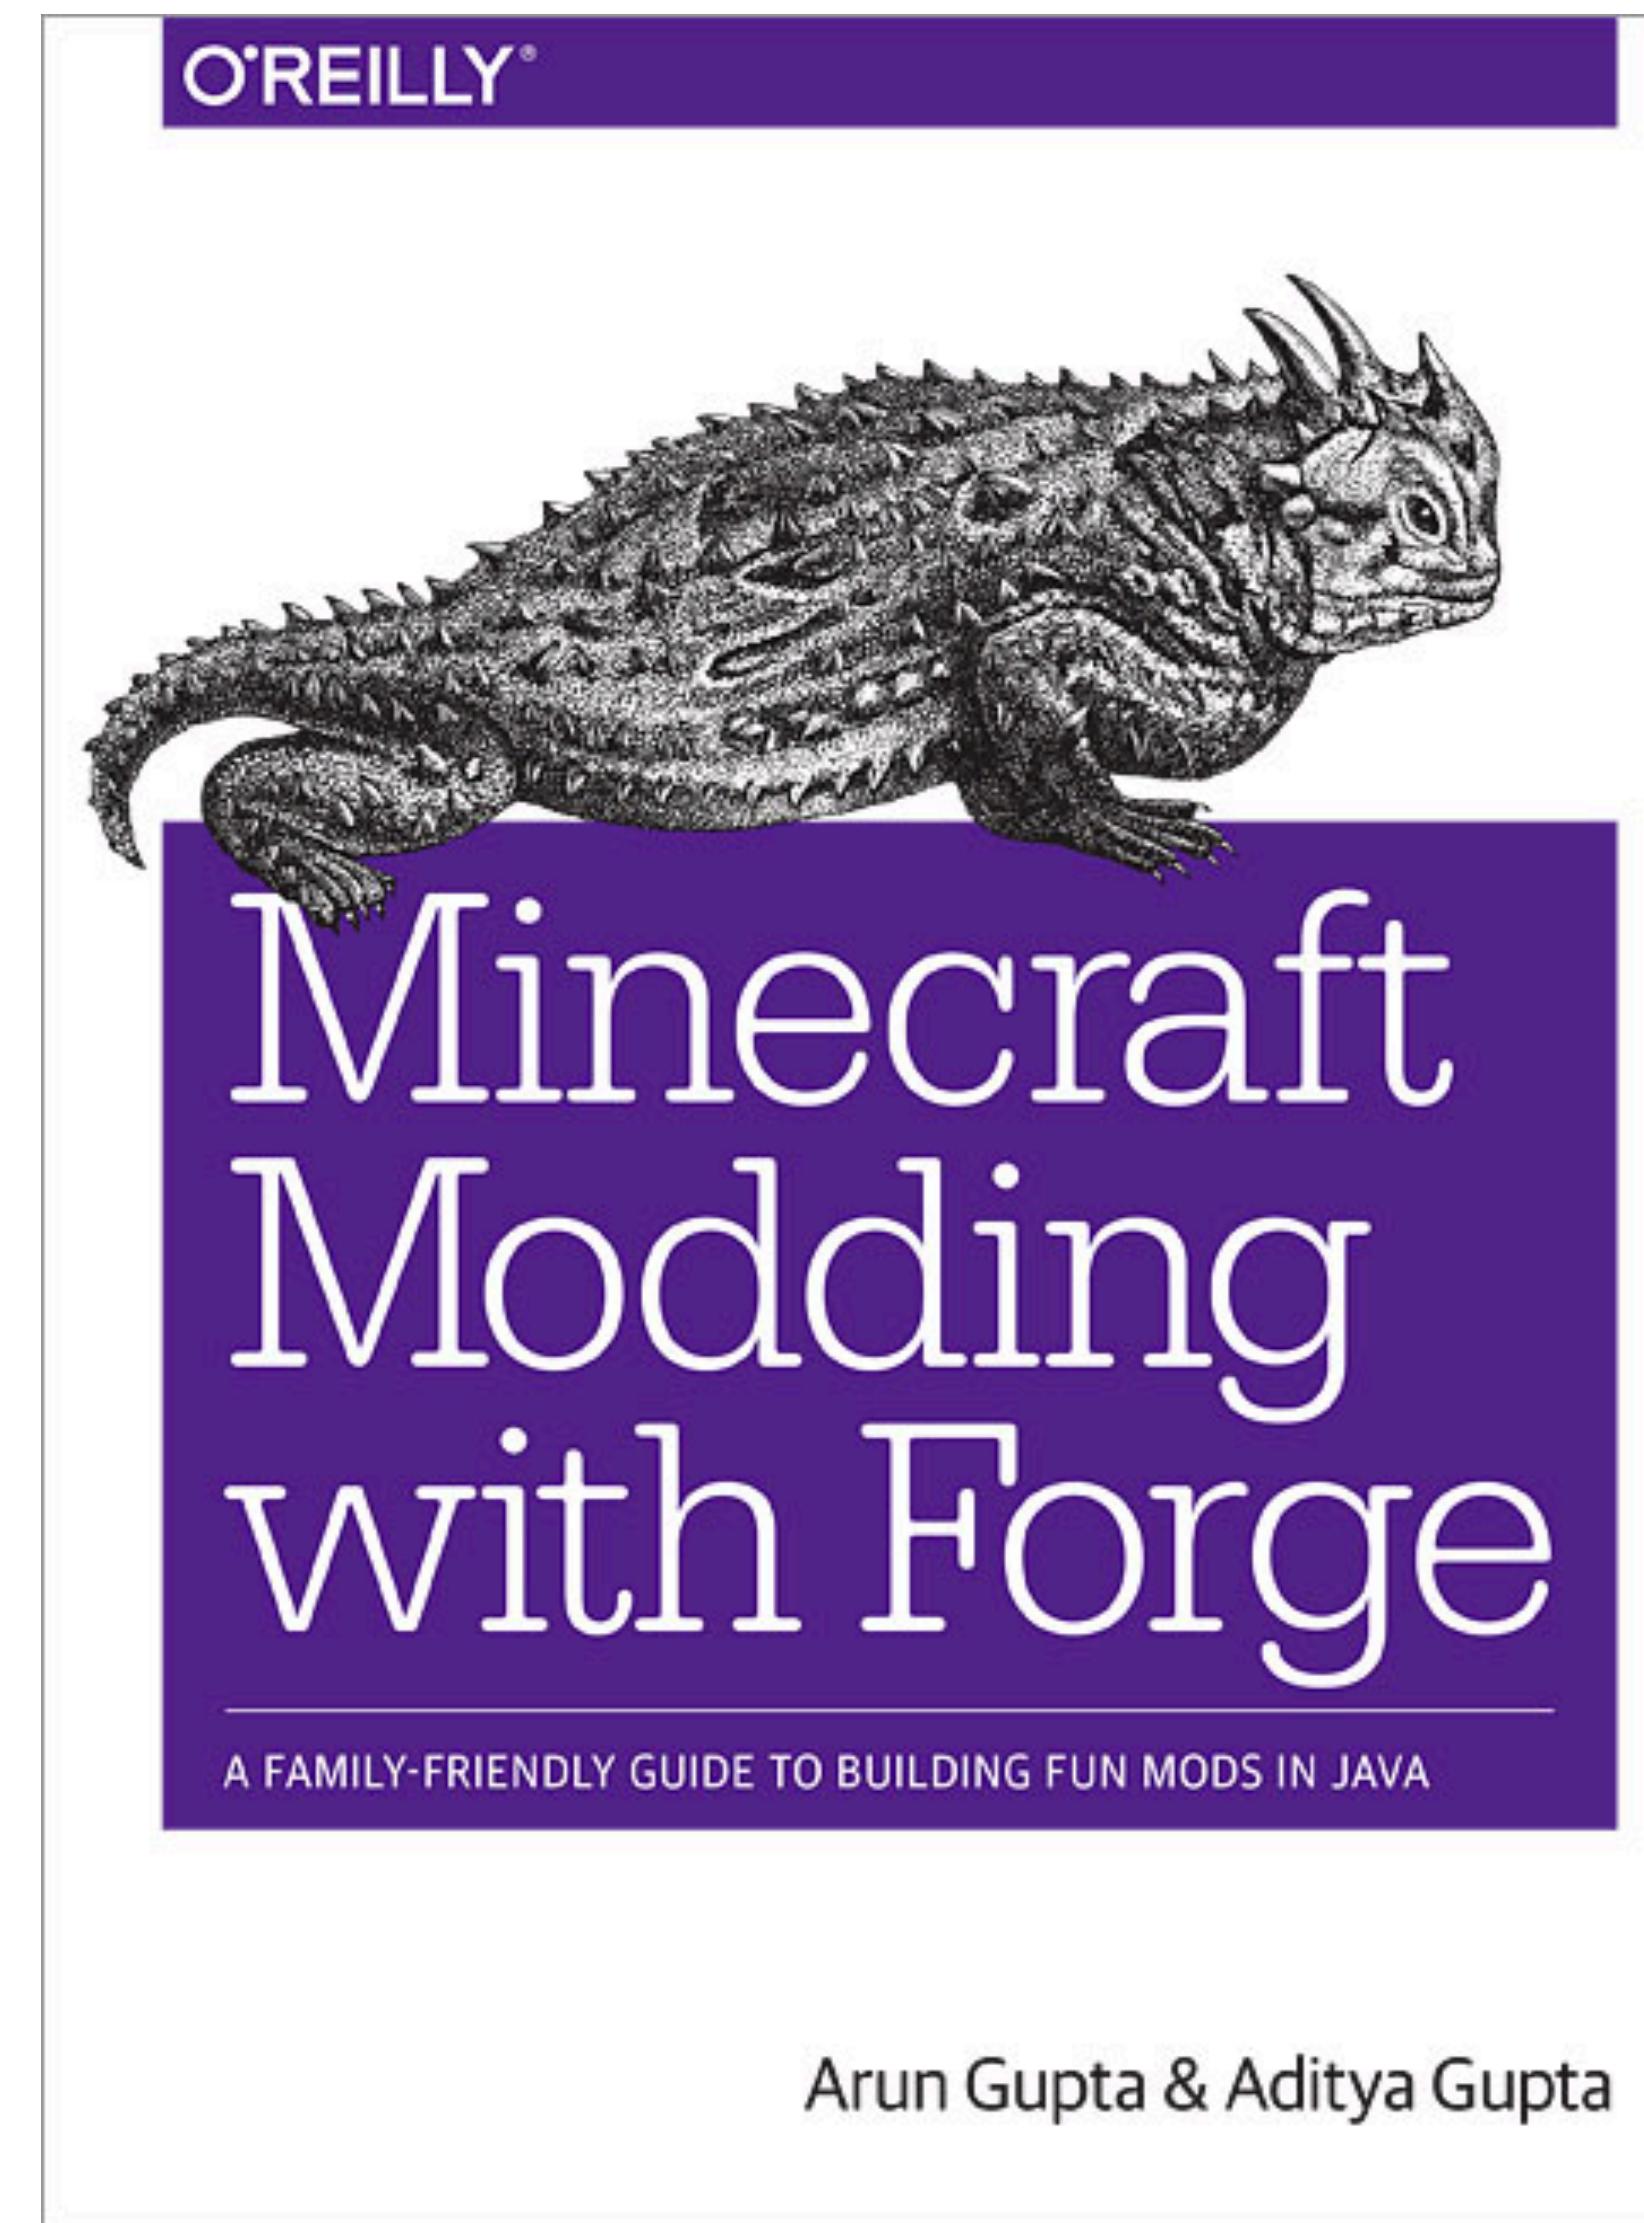




Getting Started with Kubernetes

Arun Gupta, @arungupta
VP Developer Advocacy, Couchbase



CONTENTS

- » What Is Kubernetes?
- » Kubernetes Architecture
- » Starting With Kubernetes
- » Run Your First Container
- » Scale Applications...and more!

Get More Refcardz! Visit DZone.com/Refcardz

WHAT IS KUBERNETES?

Kubernetes (kubernetes.io) is an open-source orchestration system for managing containerized applications across multiple hosts, providing basic mechanisms for the deployment, maintenance, and scaling of applications.

Kubernetes, or “k8s” or “kube” for short, allows the user to declaratively specify the desired state of a cluster using high-level primitives. For example, the user may specify that they want three instances of the Couchbase server container running. Kubernetes’ self-healing mechanisms, such as auto-restarting, re-scheduling, and replicating containers then converge the actual state towards the desired state.

Kubernetes supports Docker and Rocket containers. An abstraction around the containerization layer will allow for other container image formats and runtimes to be supported in the future.

KEY CONCEPTS OF KUBERNETES**POD**

A Pod is the smallest deployable unit that can be created, scheduled, and managed. It’s a logical collection of containers that belong to an application.

Each resource in Kubernetes is defined using a configuration file. For example, a Couchbase pod can be defined with the following .yaml file:

A Replication Controller creating two instances of a Couchbase pod can be defined as:

```
apiVersion: v1
kind: ReplicationController
metadata:
  name: couchbase-controller
spec:
  # Two replicas of the Pod to be created
  replicas: 2
  # Identifies the label key and value on the Pod that
  # this Replication Controller is responsible for
  # managing
  selector:
    app: couchbase-rc-pod
  # ‘cookie cutter’ used for creating new pods when
  # necessary
  template:
    metadata:
      labels:
        # label key and value on the pod.
        # These must match the selector above.
        app: couchbase-rc-pod
    spec:
      containers:
        - name: couchbase
          image: couchbase
          ports:
            - containerPort: 8091
```

SERVICE

Each Pod is assigned a unique IP address. If the Pod is inside a Replication Controller, then the pod is recreated but may be given

Kubernetes

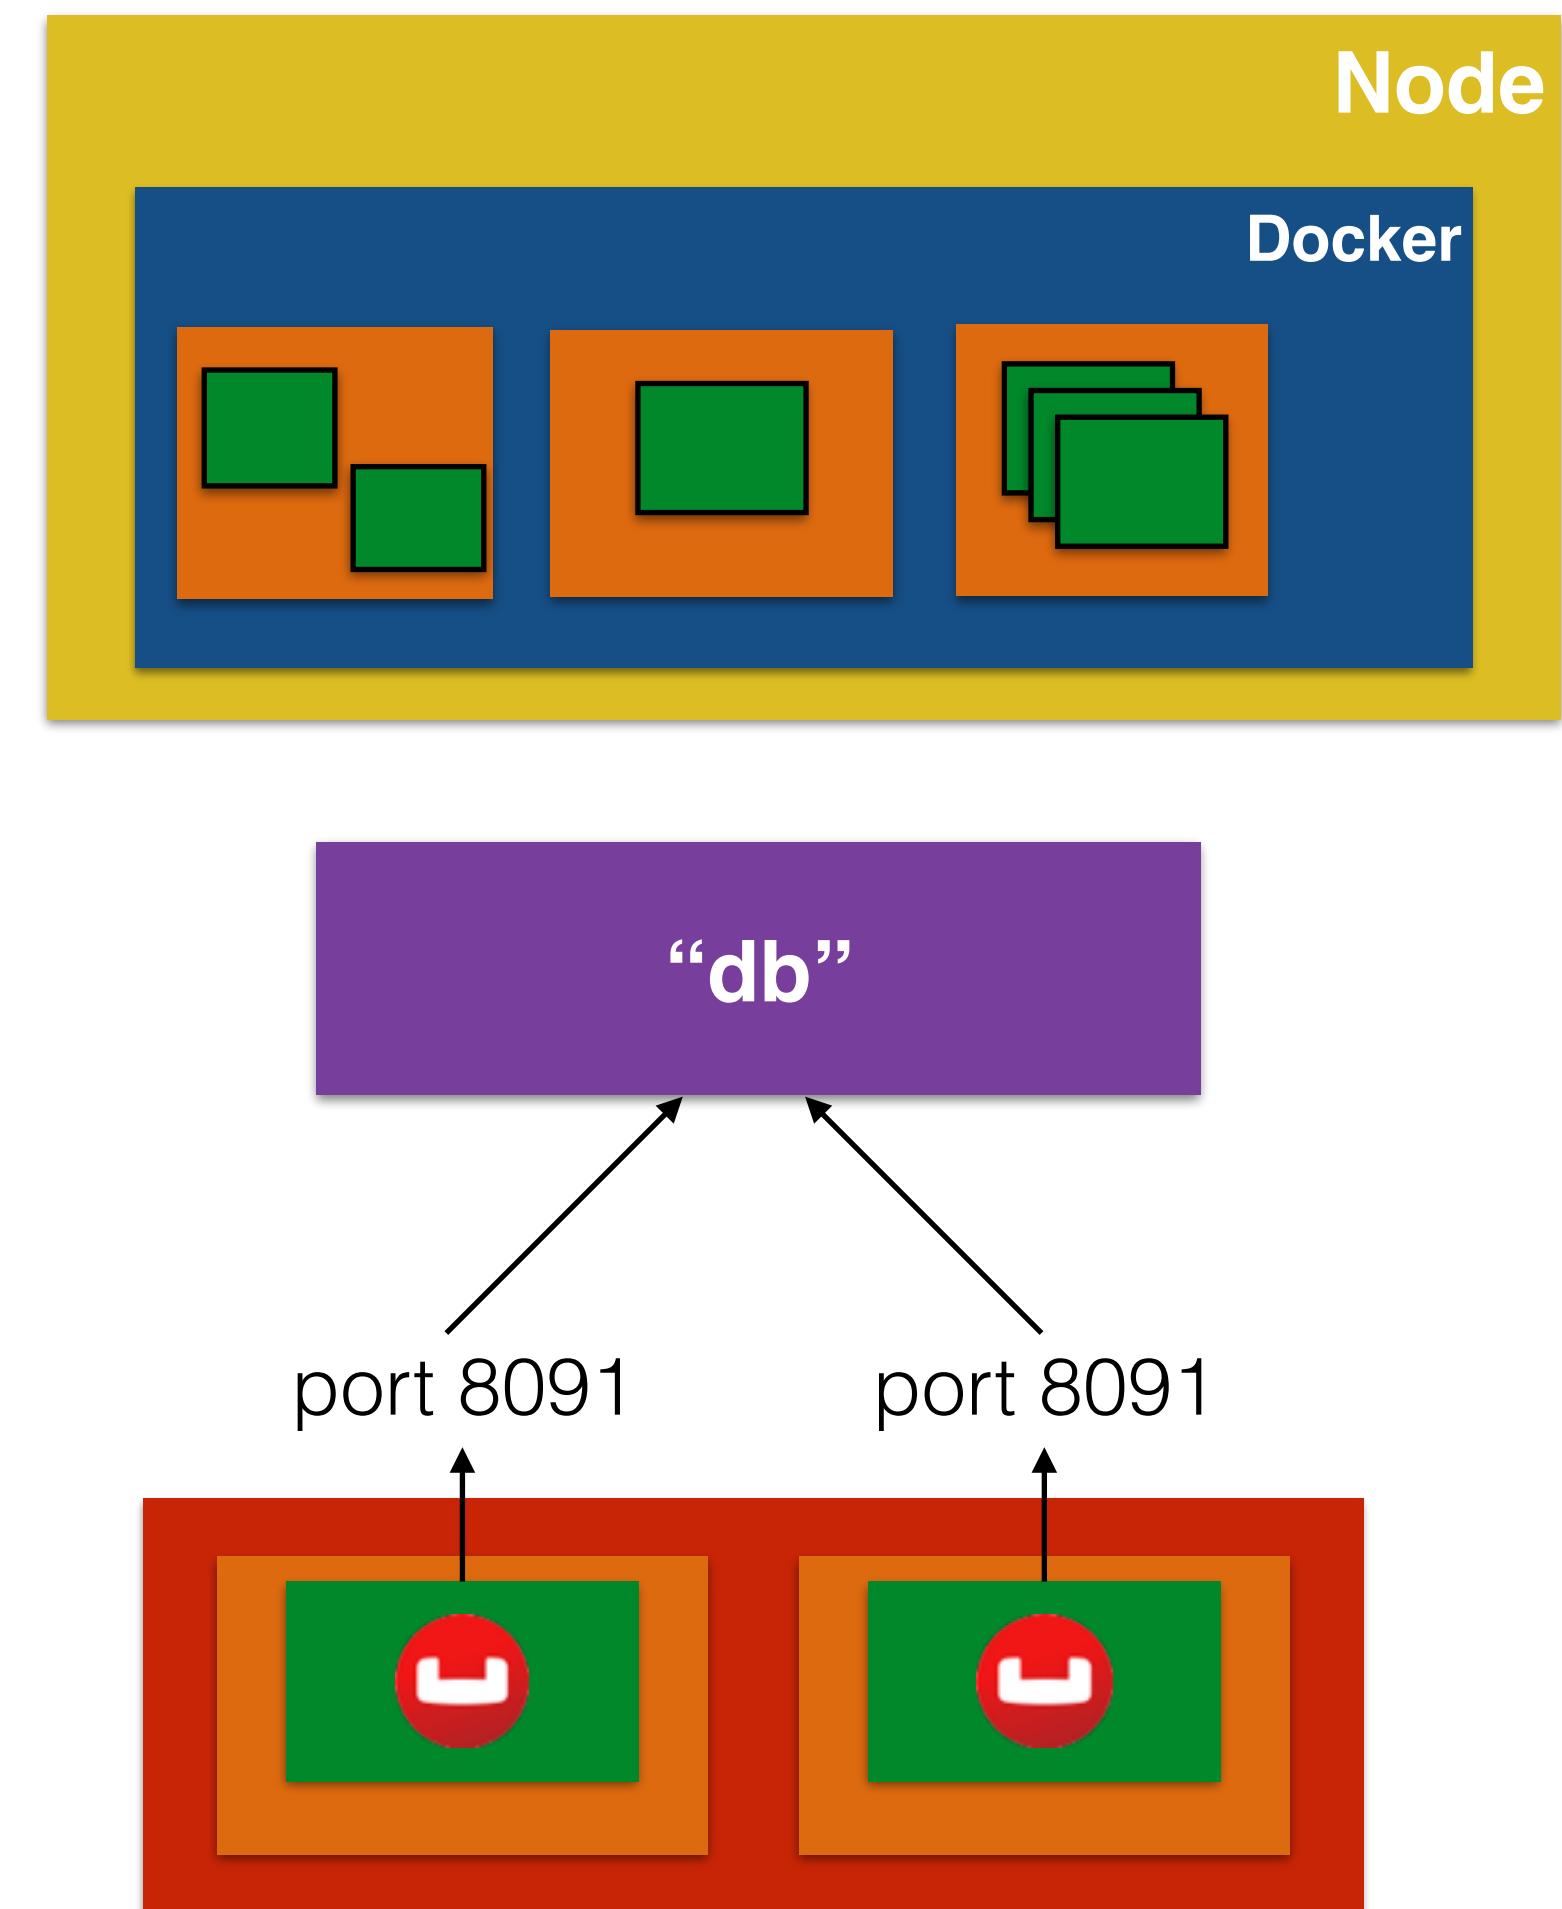
BY ARUN GUPTA

Kubernetes

- Open source orchestration system for Docker containers
- Provide declarative primitives for the “desired state”
 - Self-healing
 - Auto-restarting
 - Schedule across hosts
 - Replicating

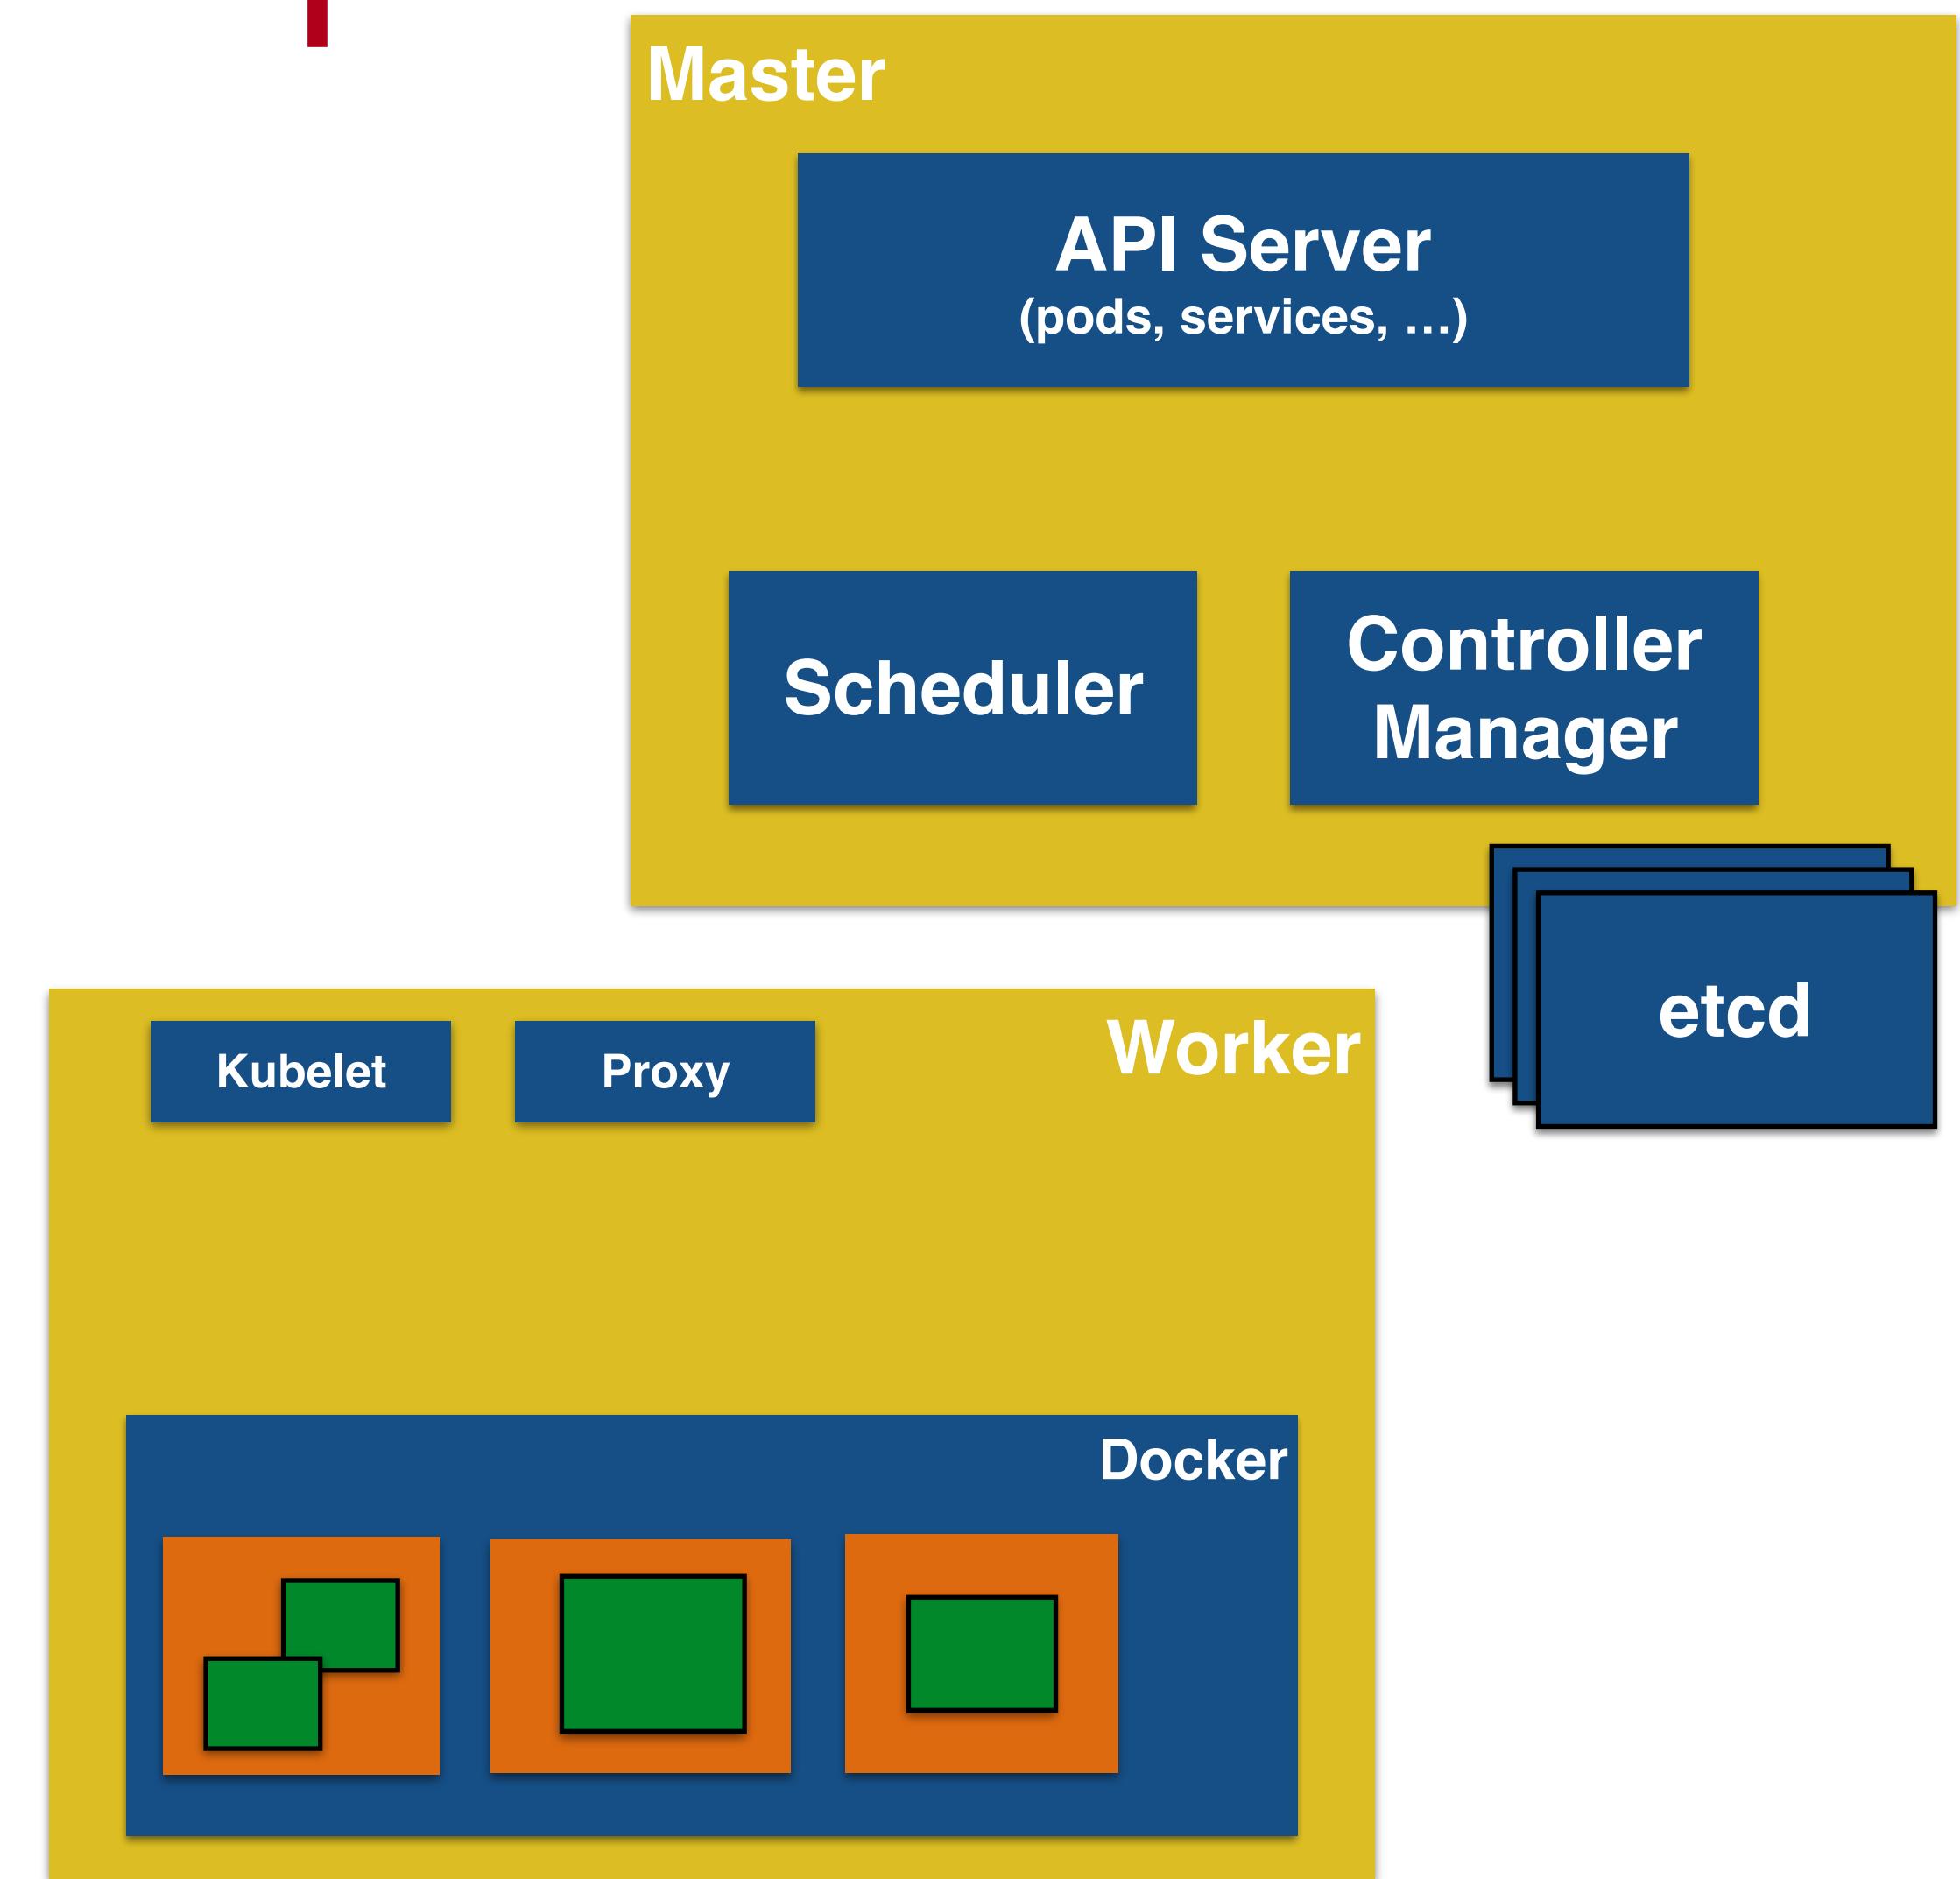
Kubernetes Concepts

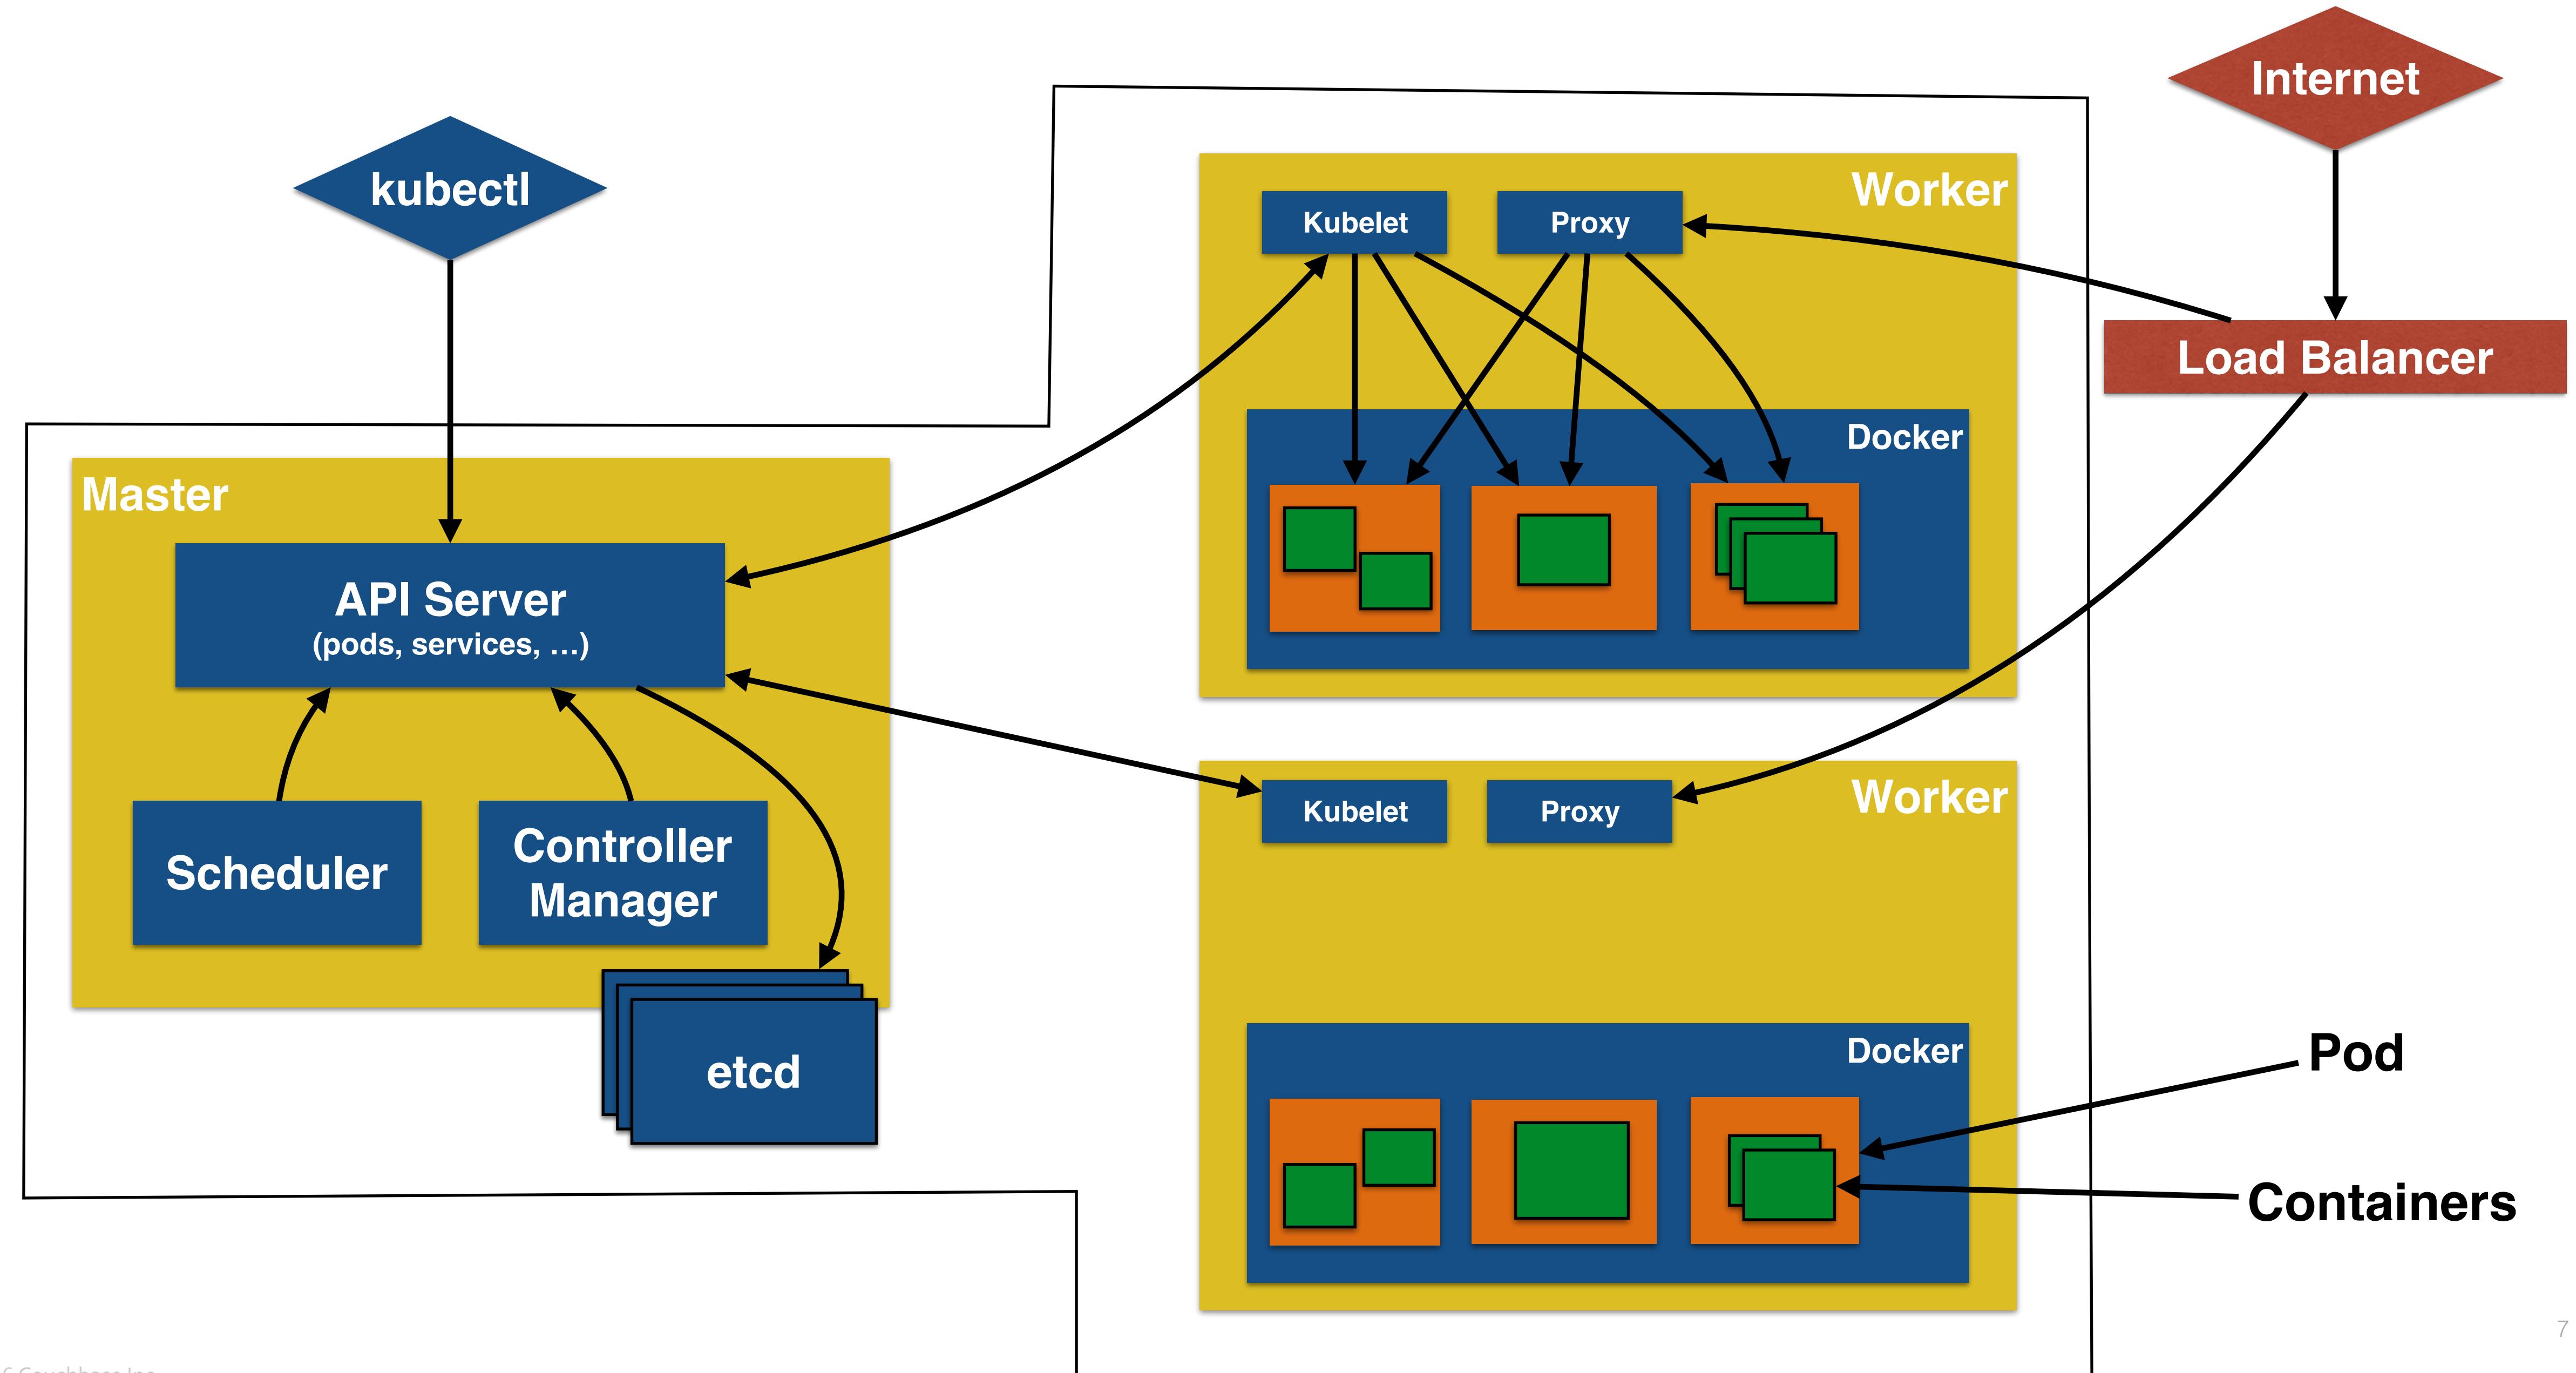
- **Pods**: colocated group of containers that share an IP, namespace, storage volume
- **Replica Set**: manages the lifecycle of pods and ensures specified number are running (next gen Replication Controller)
- **Service**: Single, stable name for a set of pods, also acts as LB
- **Label**: used to organize and select group of objects



Kubernetes Components

- **Node**: Machine or VM in the cluster
- **Master**: Central control plane, provides unified view of the cluster
 - **etcd**: distributed key-value store used to persist Kubernetes system state
- **Worker**: Docker host running *kubelet* (node agent) and *proxy* services
 - Runs pods and containers
 - Monitored by *systemd* (CentOS) or *monit* (Debian)







kubectl

- Controls the Kubernetes cluster manager
- `kubectl get pods or minions`
- `kubectl create -f <filename>`
- `kubectl update or delete`
- `kubectl scale --replicas=3 rc/<name>`

Kubernetes Pod Configuration

```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: wildfly-pod
5    labels:
6      name: wildfly-pod
7  spec:
8    containers:
9      - name: wildfly
10     image: jboss/wildfly
11     ports:
12       - containerPort: 8080
```

Replication Controller

- Ensures that a specified number of pod "replicas" are running
 - Pod templates are cookie cutters
 - Rescheduling
 - Manual or auto-scale replicas
 - Rolling updates
- Generally wrap a pod in a RC
- Only appropriate for pods with `Restart=Always` policy (default)

Kubernetes Replication Controller Configuration

```
1  apiVersion: v1
2  kind: ReplicationController
3  metadata:
4    name: wildfly-rc
5  spec:
6    replicas: 2
7    selector:
8      app: wildfly-rc-pod
9    template:
10      metadata:
11        labels:
12          app: wildfly-rc-pod
13      spec:
14        containers:
15          - name: wildfly
16            image: jboss/wildfly
17        ports:
18          - containerPort: 8080
```

Services

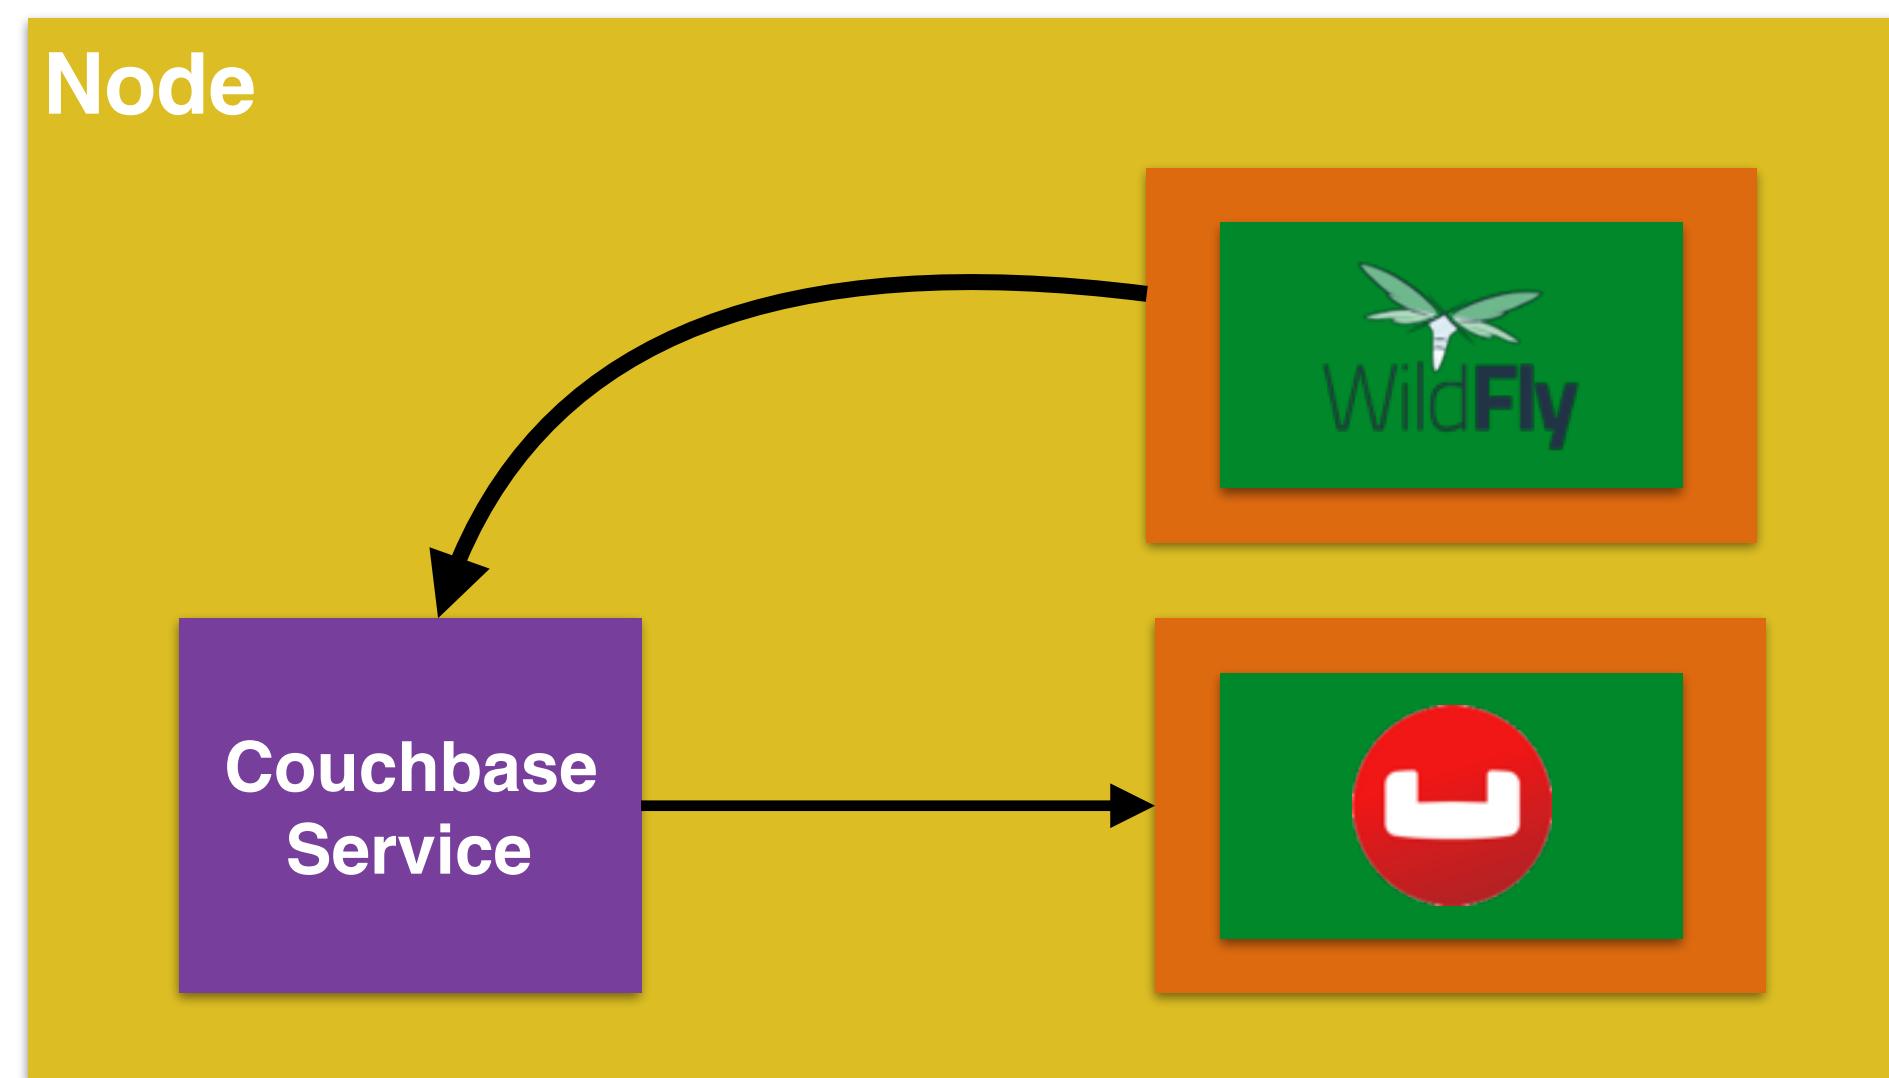
- Abstract a set of pods as a single IP and port
 - Simple TCP/UDP load balancing
- Creates environment variables in other pods or DNS resolution
- Stable endpoint for pods to reference
 - Allows list of pods to change dynamically

Kubernetes Service Configuration

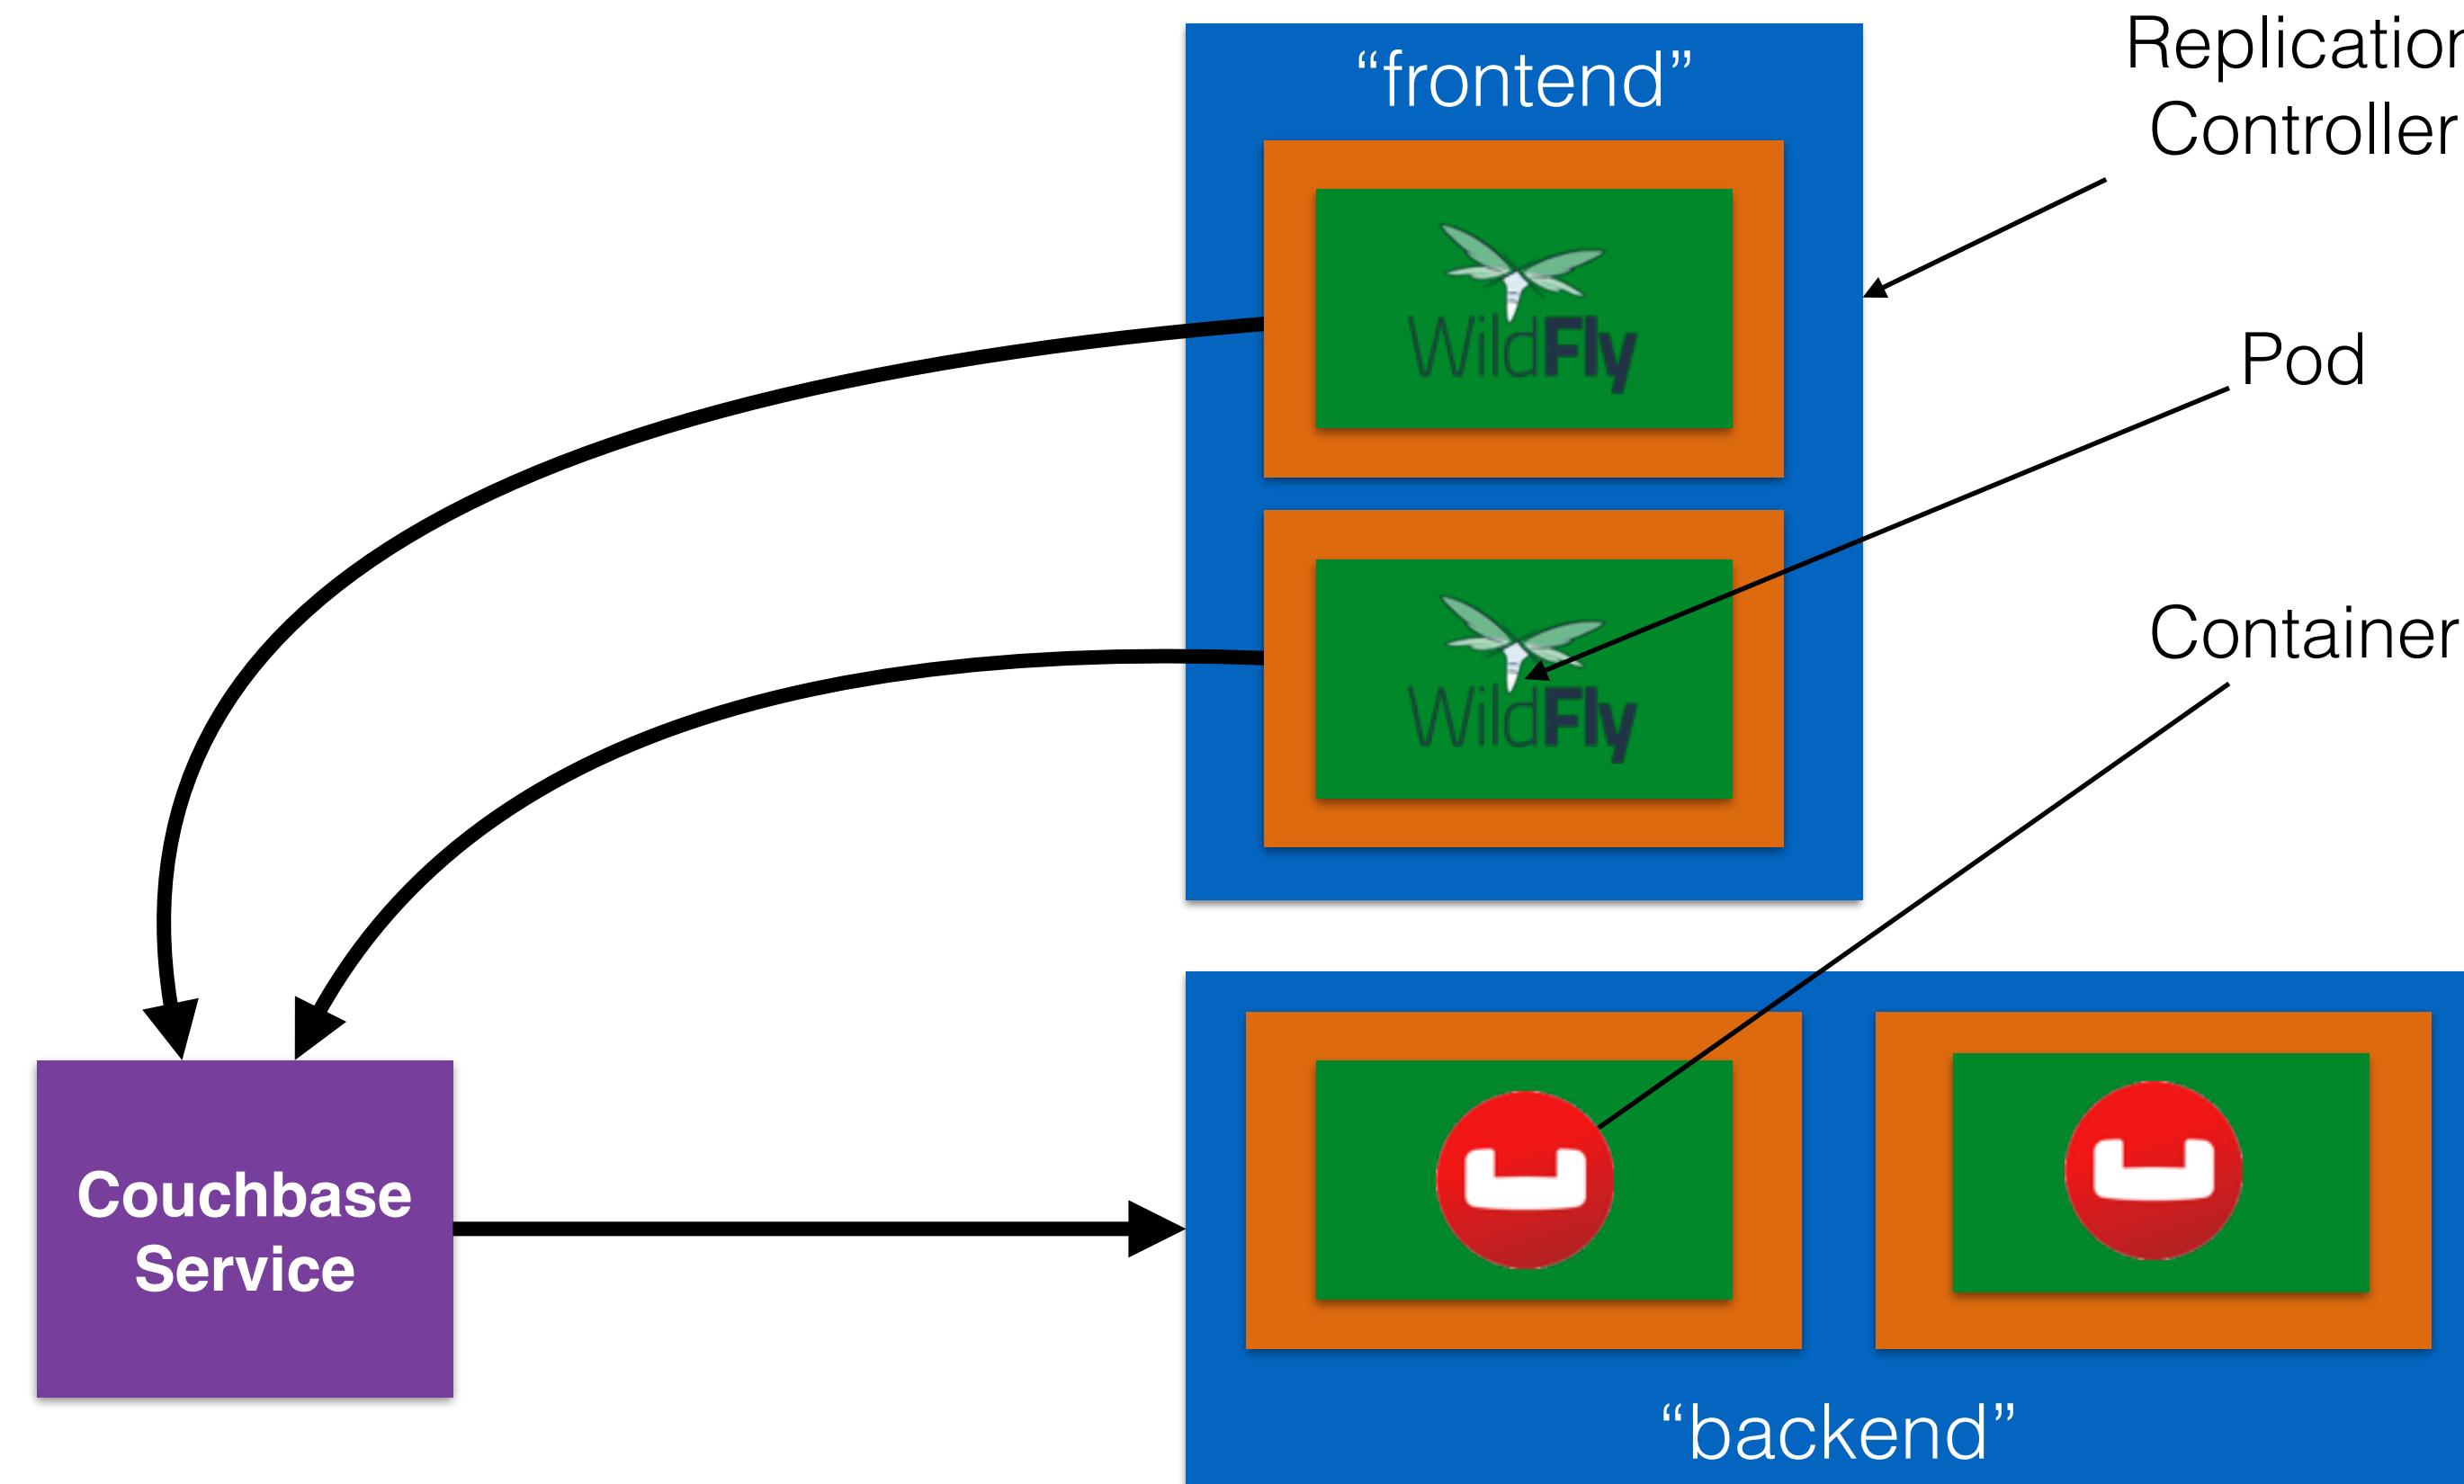
```
1 apiVersion: v1
2 kind: Service
3 metadata:
4   name: couchbase-service
5 spec:
6   selector:
7     app: couchbase-rc-pod
8   ports:
9     - name: admin
10    port: 8091
11    - name: query
12      port: 8093
13 ---
```

```
14 apiVersion: v1
15 kind: ReplicationController
16 metadata:
17   name: couchbase-rc
18 spec:
19   replicas: 2
20   selector:
21     app: couchbase-rc-pod
22   template:
23     metadata:
24       labels:
25         app: couchbase-rc-pod
26   spec:
27     containers:
28       - name: couchbase
29         image: couchbase
30     ports:
31       - containerPort: 8091
```

Couchbase Service



Service and Replication Controller

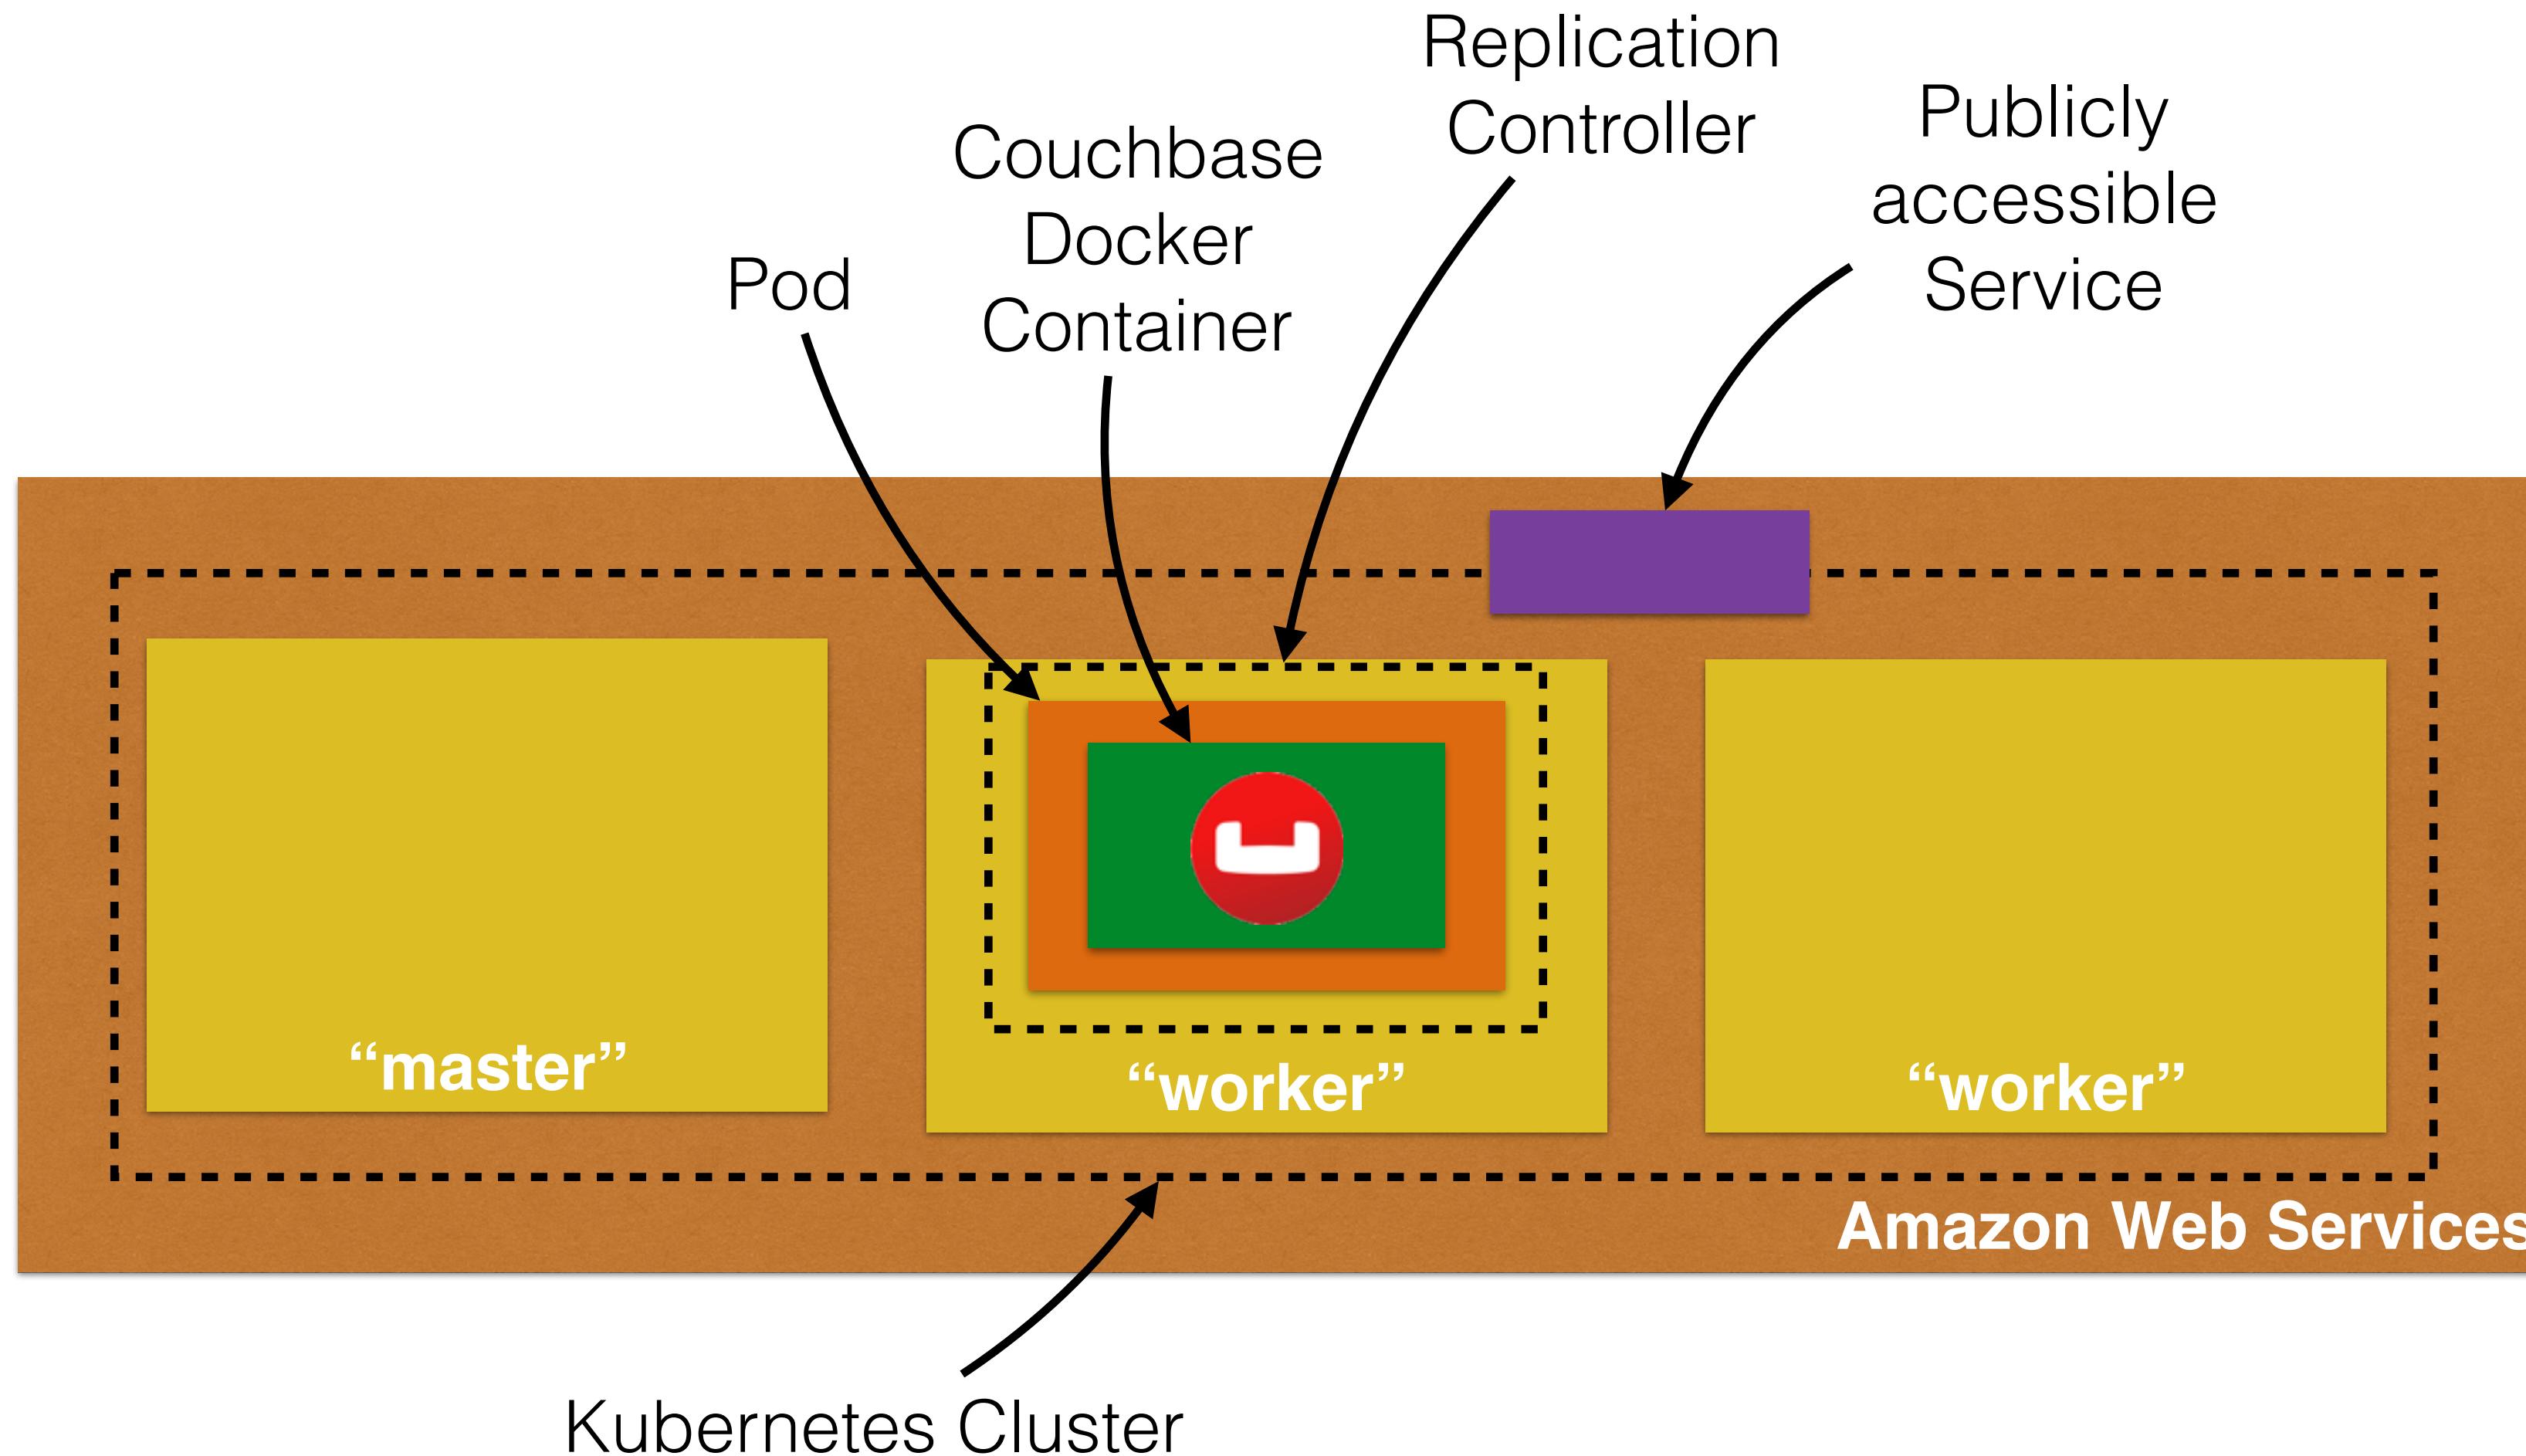


Exposing Service

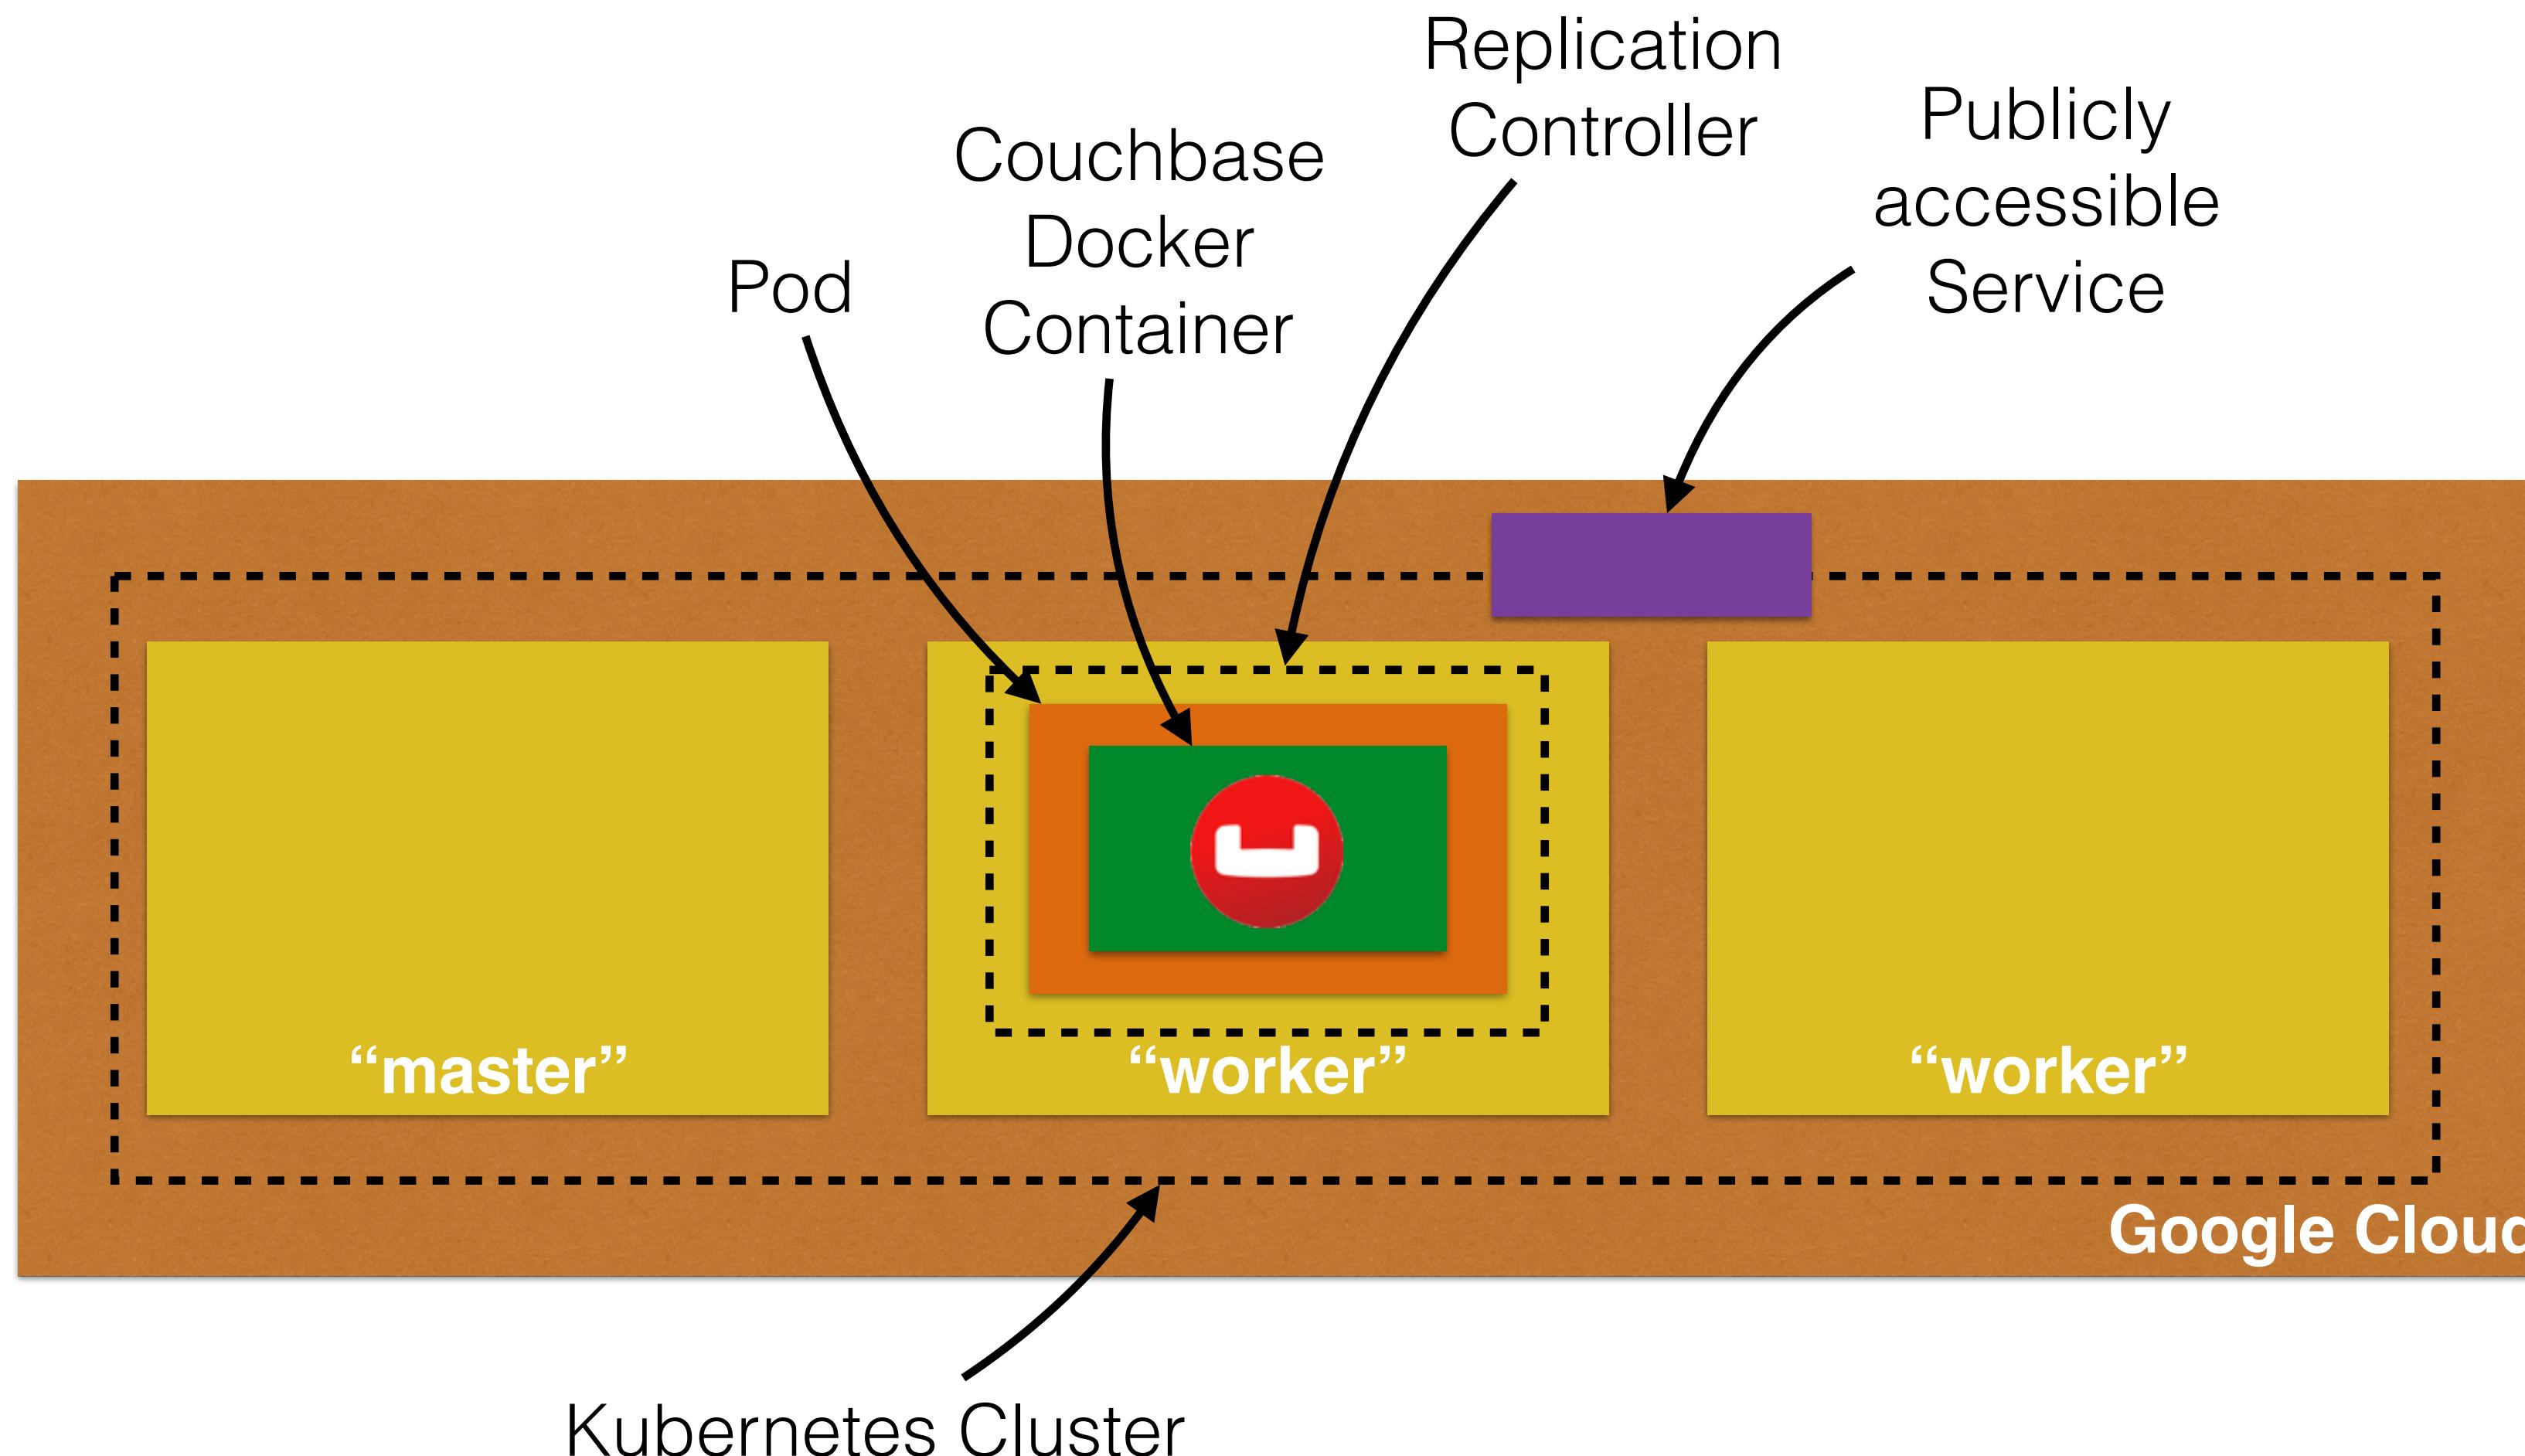
- Service may be exposed outside cluster or on Internet using **type**
 - **ClusterIP** (default)
 - **NodePort**: A port on each node

```
spec:  
  type: NodePort  
  ports:  
    - name: admin  
      port: 8091  
      nodePort: 30001  
    - name: query  
      port: 8093  
      nodePort: 30002
```
 - **LoadBalancer**: On cloud providers that support external LB

Exposing Service on Amazon



Exposing Service on GCE



Replica Set

- Next generation Replication Controller
- Set-based selector requirement
- Useful for Horizontal Pod Autoscaling

Replica Set Configuration

```
1  apiVersion: extensions/v1beta1
2  kind: ReplicaSet
3  metadata:
4    name: wildfly-rs
5  spec:
6    replicas: 2
7    selector:
8      matchLabels:
9        app: wildfly-rs-pod
10     matchExpressions:
11       - {key: tier, operator: In, values: ["backend"]}
12       - {key: environment, operator: NotIn, values: ["dev"]}
13   template:
14     metadata:
15       labels:
16         app: wildfly-rs-pod
17         tier: backend
18         environment: dev
19     spec:
20       containers:
21         - name: wildfly
22           image: jboss/wildfly
23         ports:
24           - containerPort: 8080
```

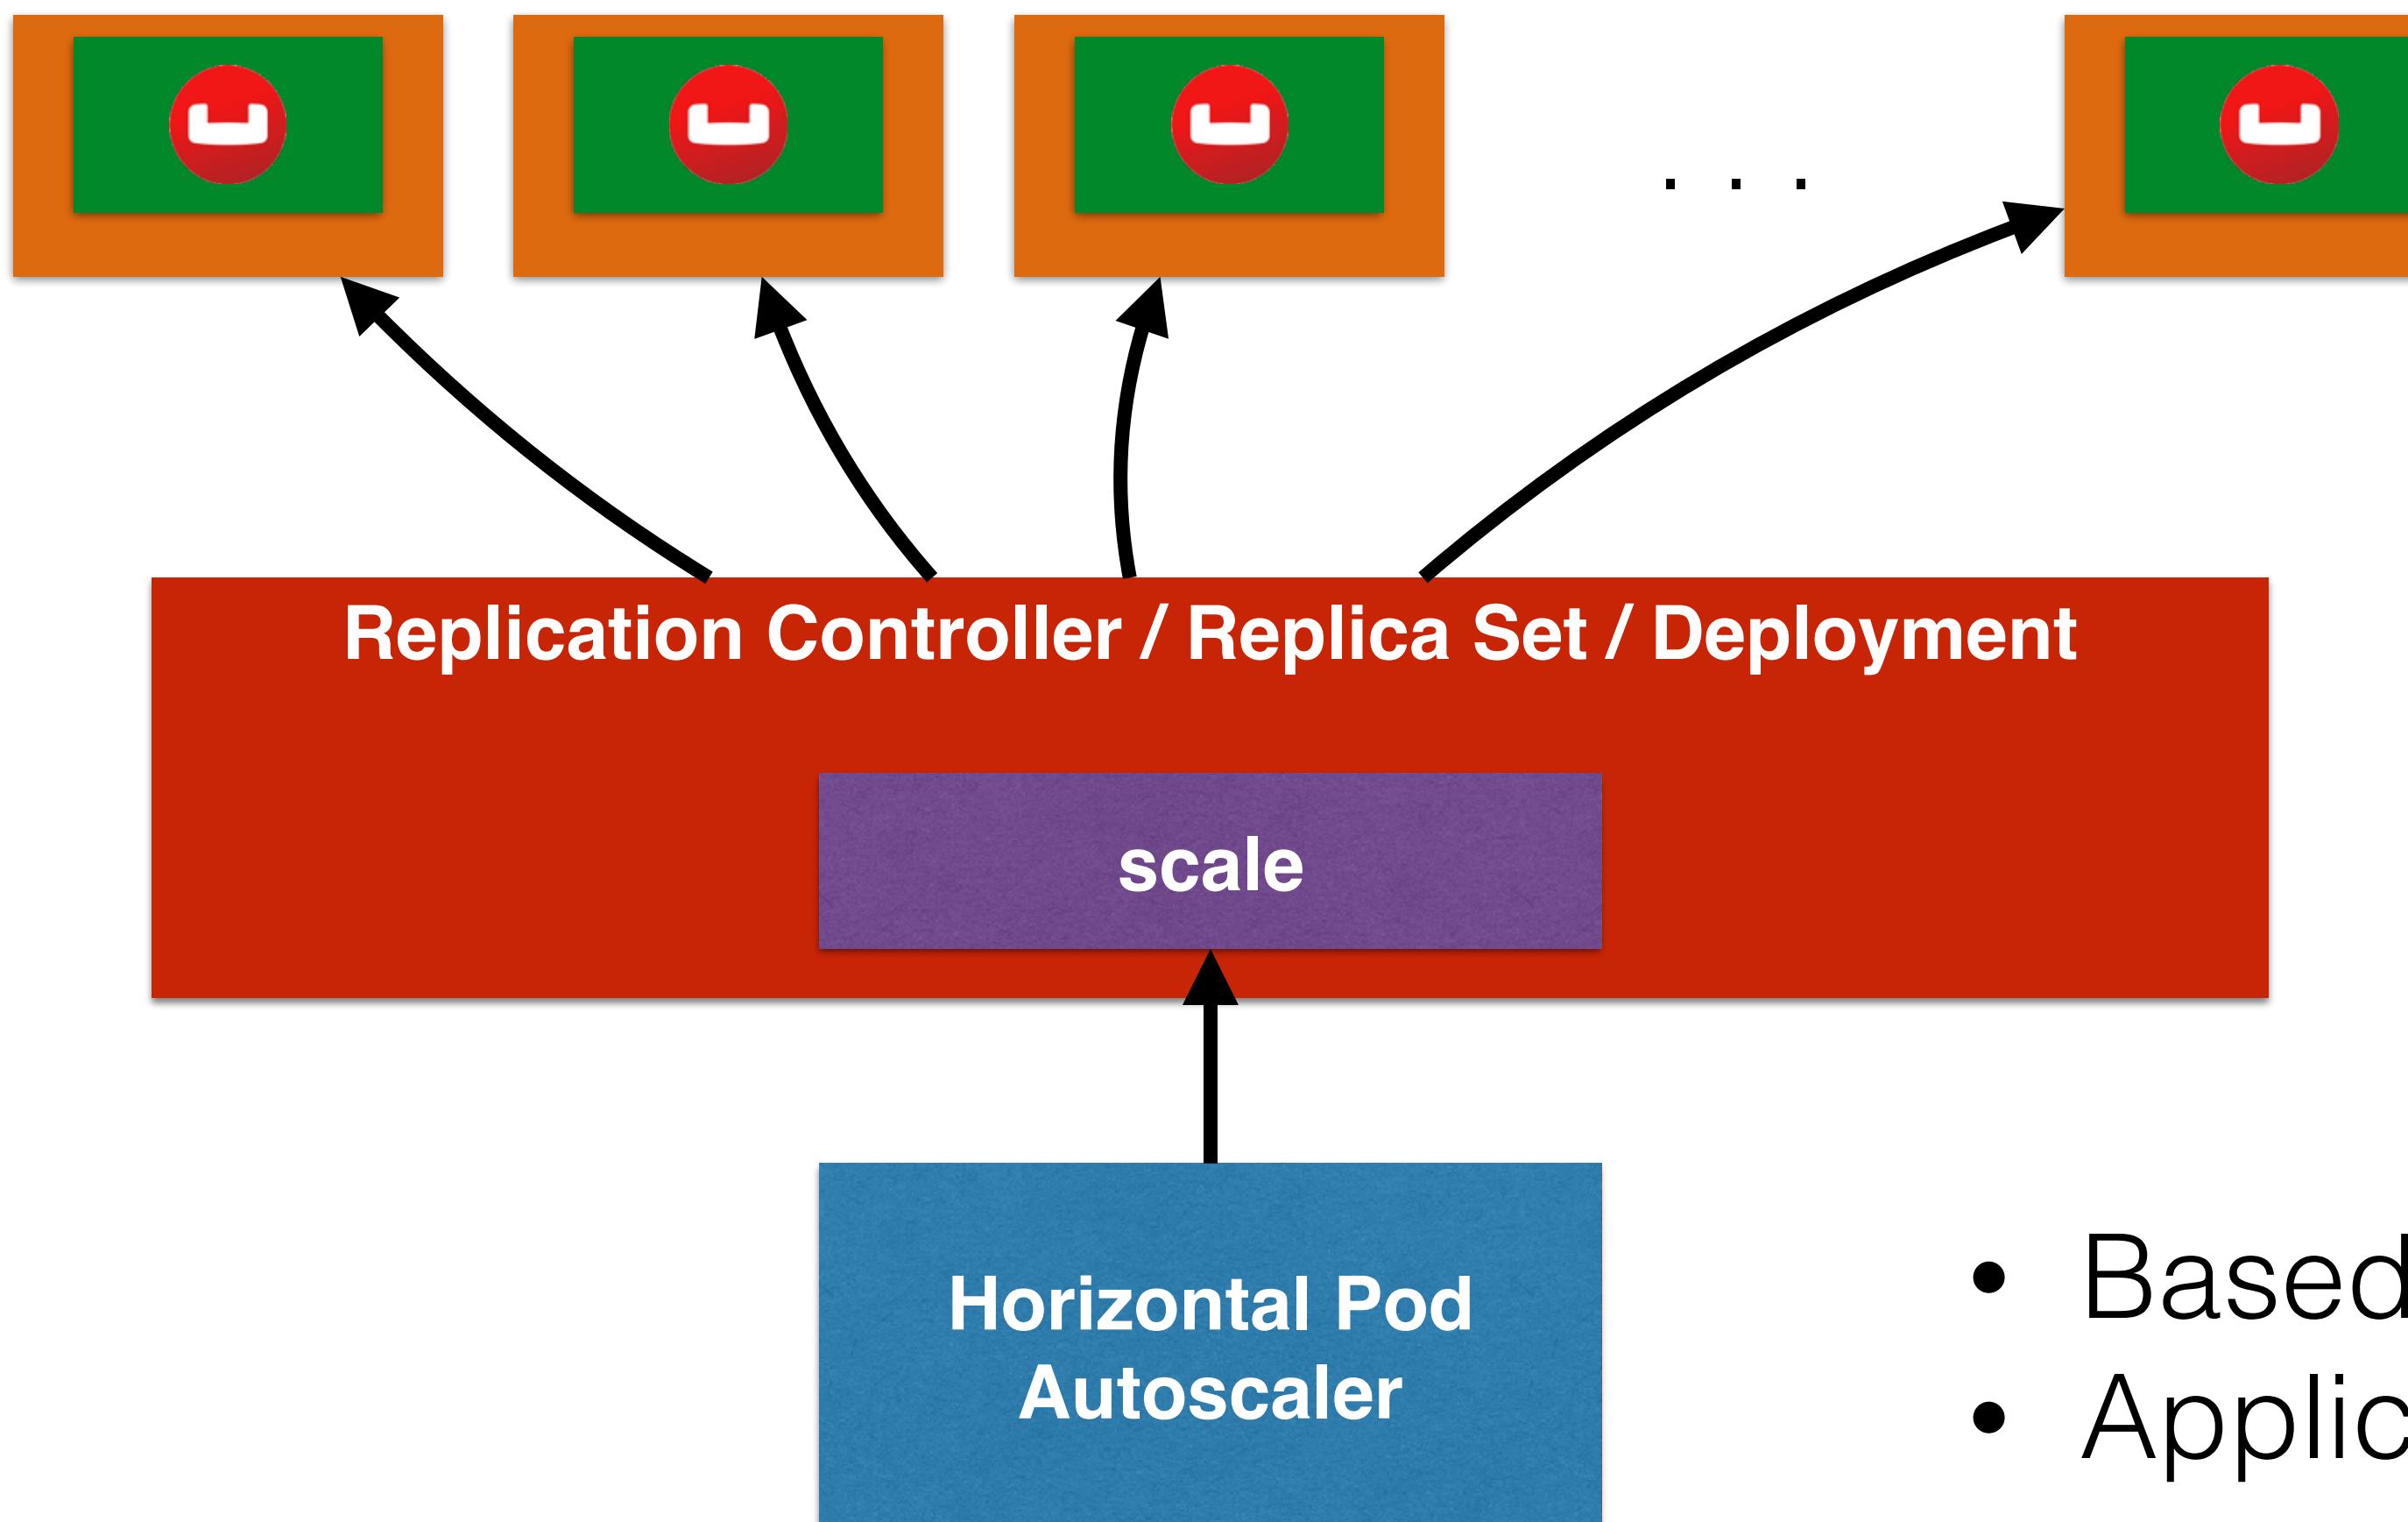
Deployment

- Declarative updates for pods and replica sets
 - For example: rolling updates
- Differences from `kubectl`
 - Declarative instead of imperative
 - Server-side and so is faster
 - More features, e.g. rollback to previous version

Deployment Configuration

```
1  apiVersion: extensions/v1beta1
2  kind: Deployment
3  metadata:
4    name: wildfly-deployment
5  spec:
6    replicas: 3
7    template:
8      metadata:
9        labels:
10          app: wildfly
11        spec:
12          containers:
13            - name: wildfly
14              image: jboss/wildfly
15            ports:
16              - containerPort: 8080
```

Horizontal Pod Autoscaling



- Based on observed CPU utilization
- Application provided metrics (health)

Horizontal Pod Autoscaling

- Typical usage
 - `kubectl autoscale deployment | rc | rs --min=<PODS> --max=<PODS> --cpu-percent=<CPU>`
- Autoscale a deployment with number of pods between 2 and 10
 - `kubectl autoscale deployment db --min=2 --max=10`
- Autoscale a Replication with maximum number of pods 5, target CPU utilization of 80%
 - `kubectl autoscale rc --max=5 --cpu-percent=80`

HPA Configuration

Volumes

- Directory accessible to the containers in a pod
- Volume outlives any containers in a pod
- Common types
 - hostPath
 - nfs
 - awsElasticBlockStore
 - gcePersistentDisk

Volume Configuration

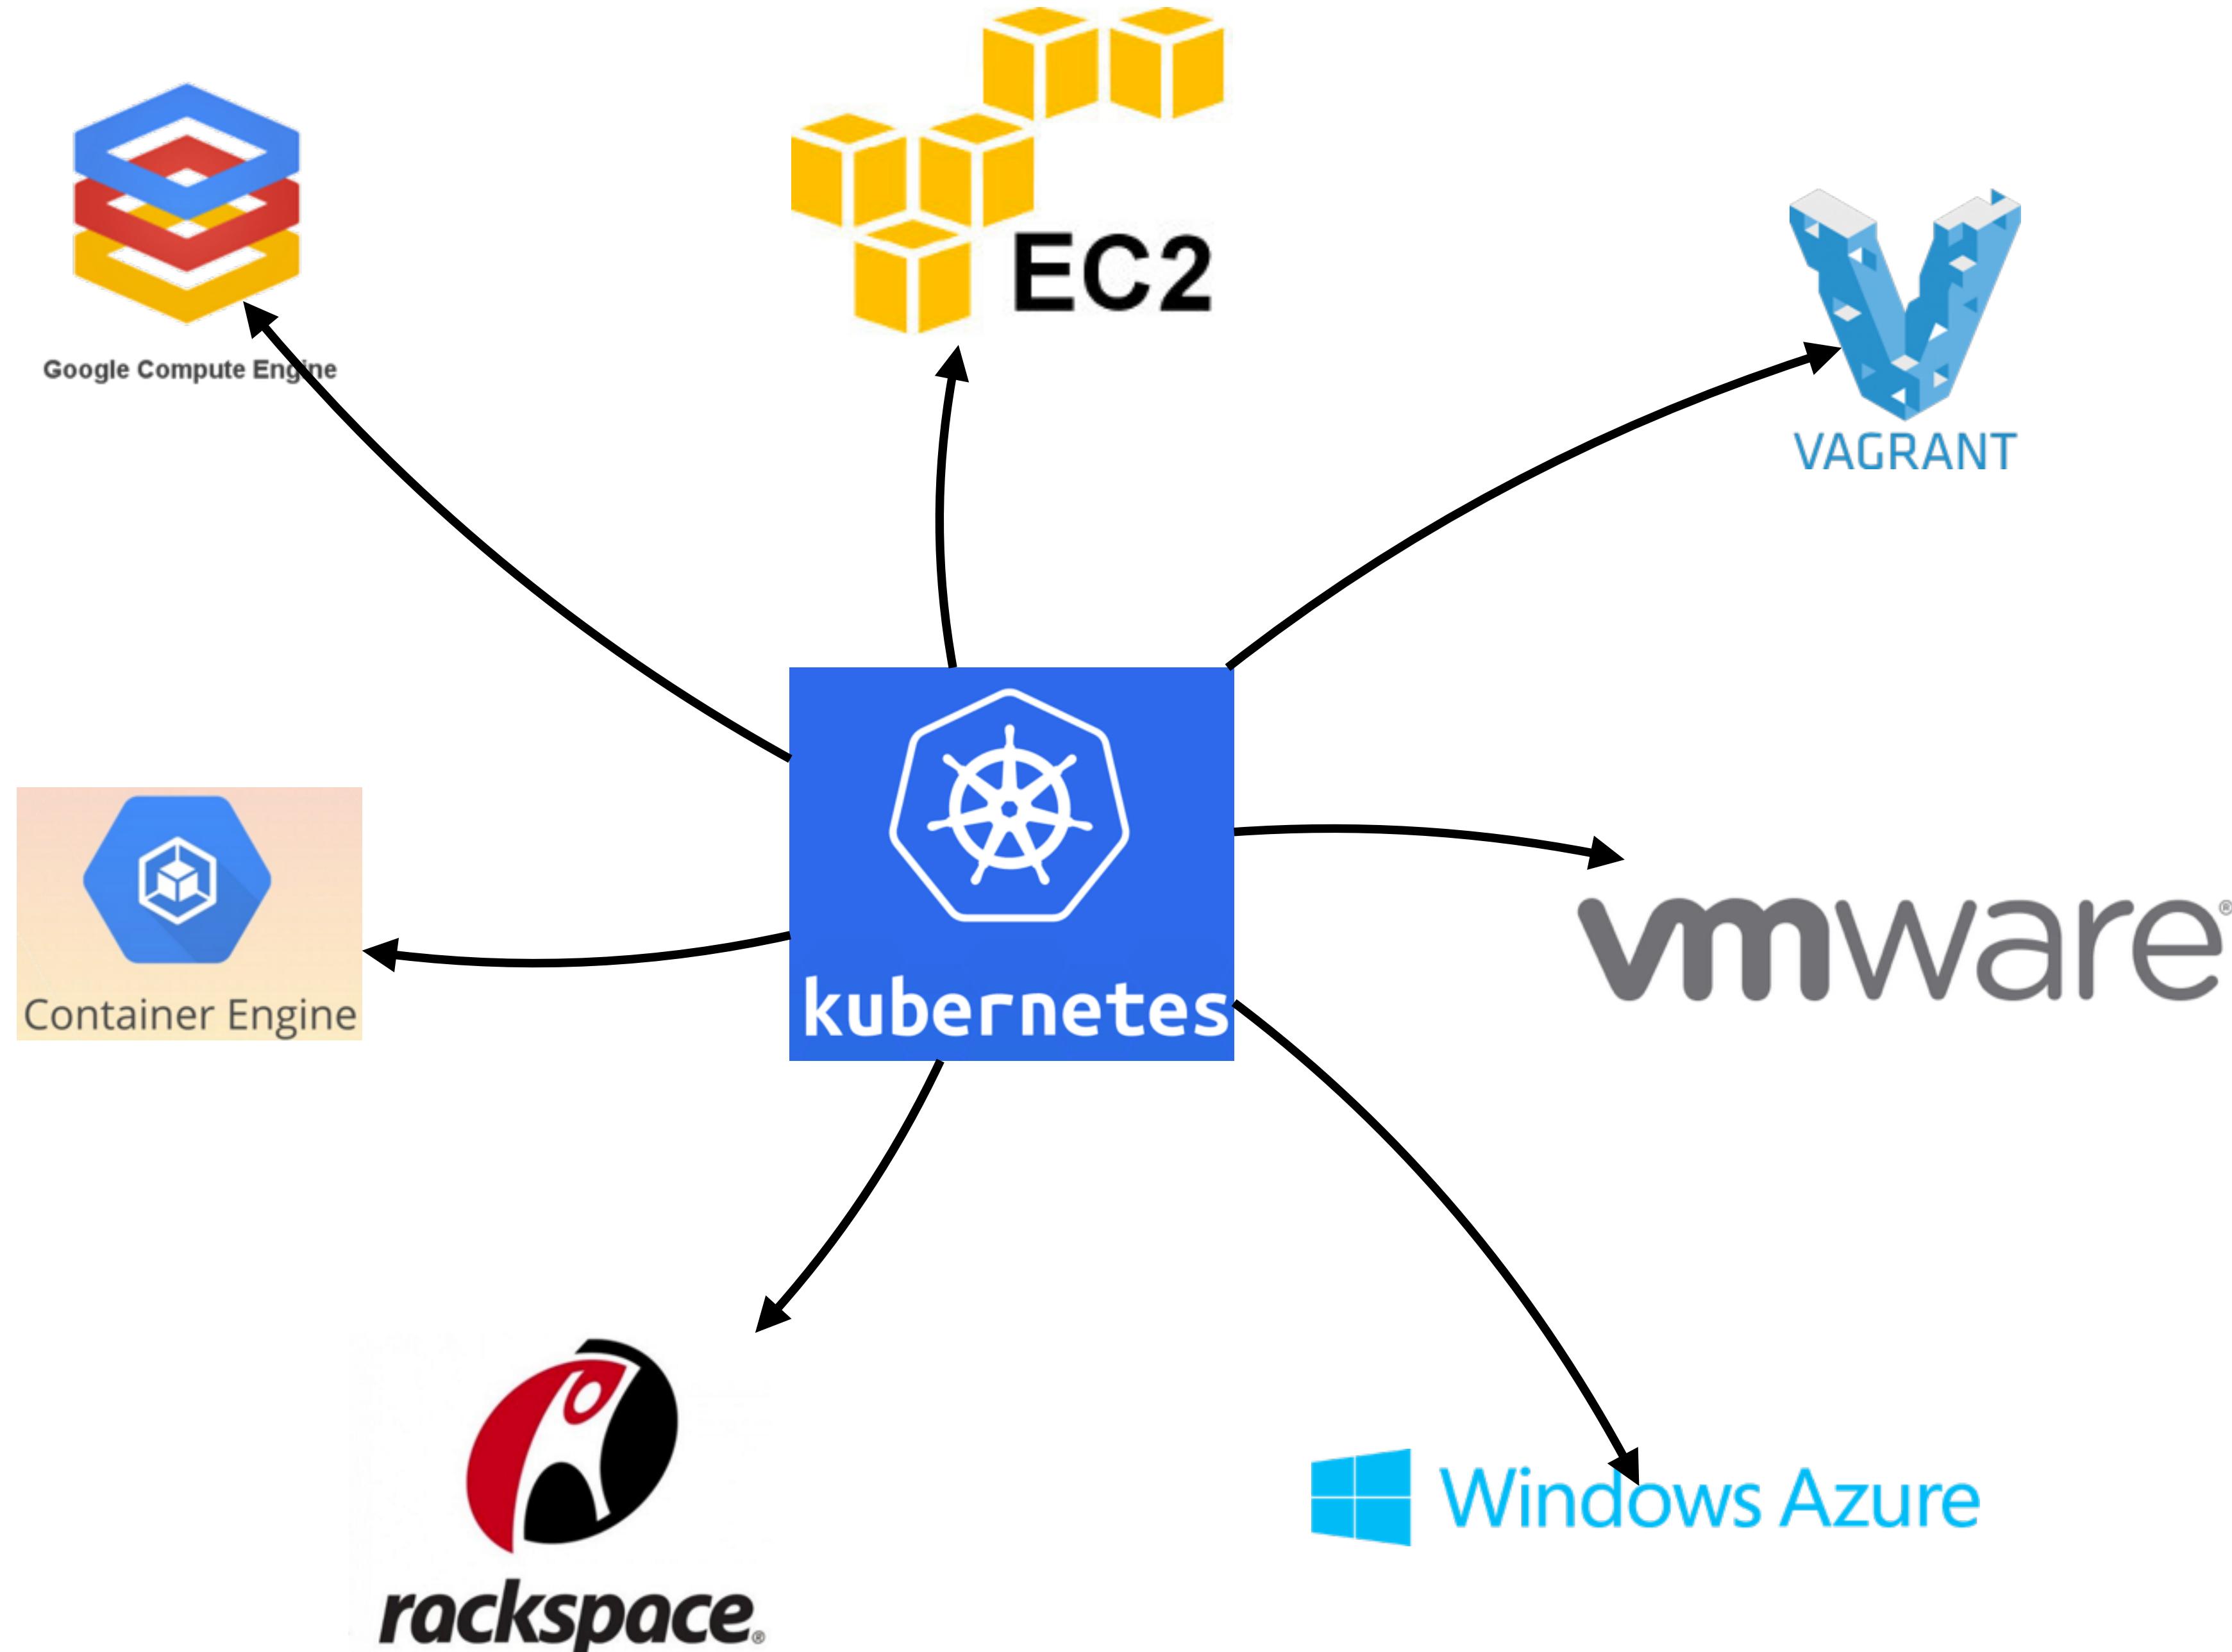
```
1  apiVersion: v1
2  kind: Pod
3  metadata:
4    name: couchbase-pod
5    labels:
6      name: couchbase-pod
7  spec:
8    containers:
9      - name: couchbase
10        image: arungupta/couchbase-oreilly:k8s
11    ports:
12      - containerPort: 8091
13    volumeMounts:
14      - mountPath: /var/couchbase/lib
15        name: couchbase-data
16    volumes:
17      - name: couchbase-data
18        hostPath:
19          path: /opt/data
```

Multitenancy - Namespace

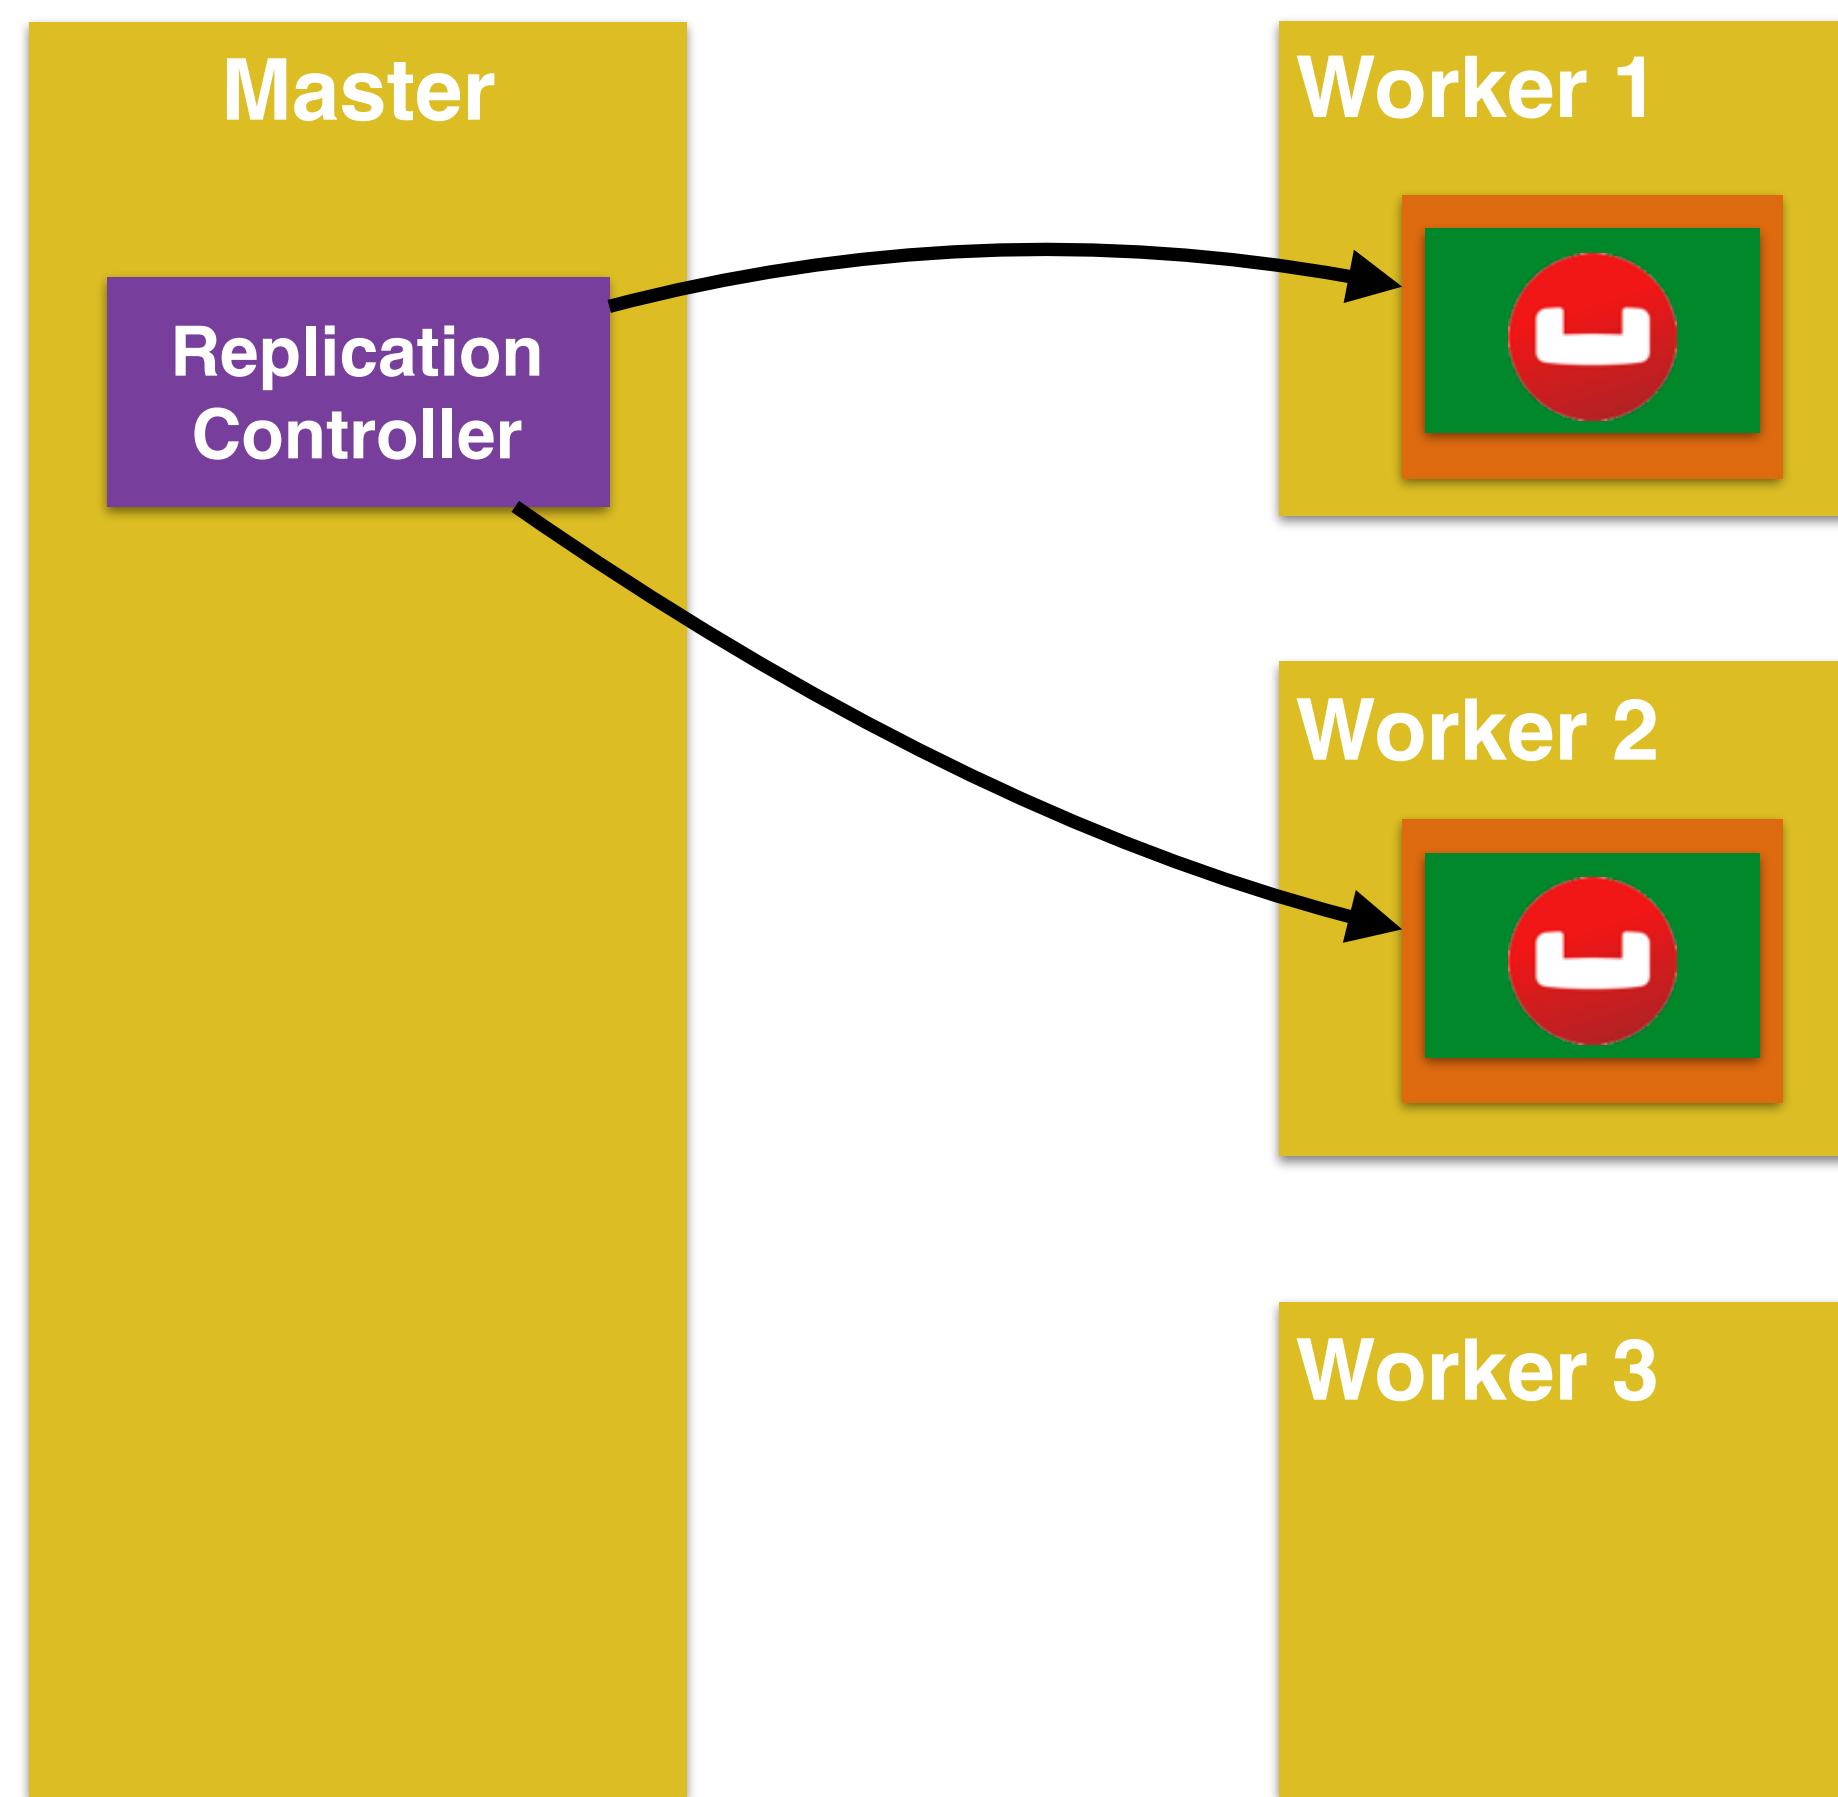
- Namespace allows to partition resources into a logical group
- Each namespace provides:
 - **scope** for resources to avoid collisions
 - **policies** to ensure appropriate authority to trusted users
 - **constraints** for resource consumption

Multitenancy - Resource Quota

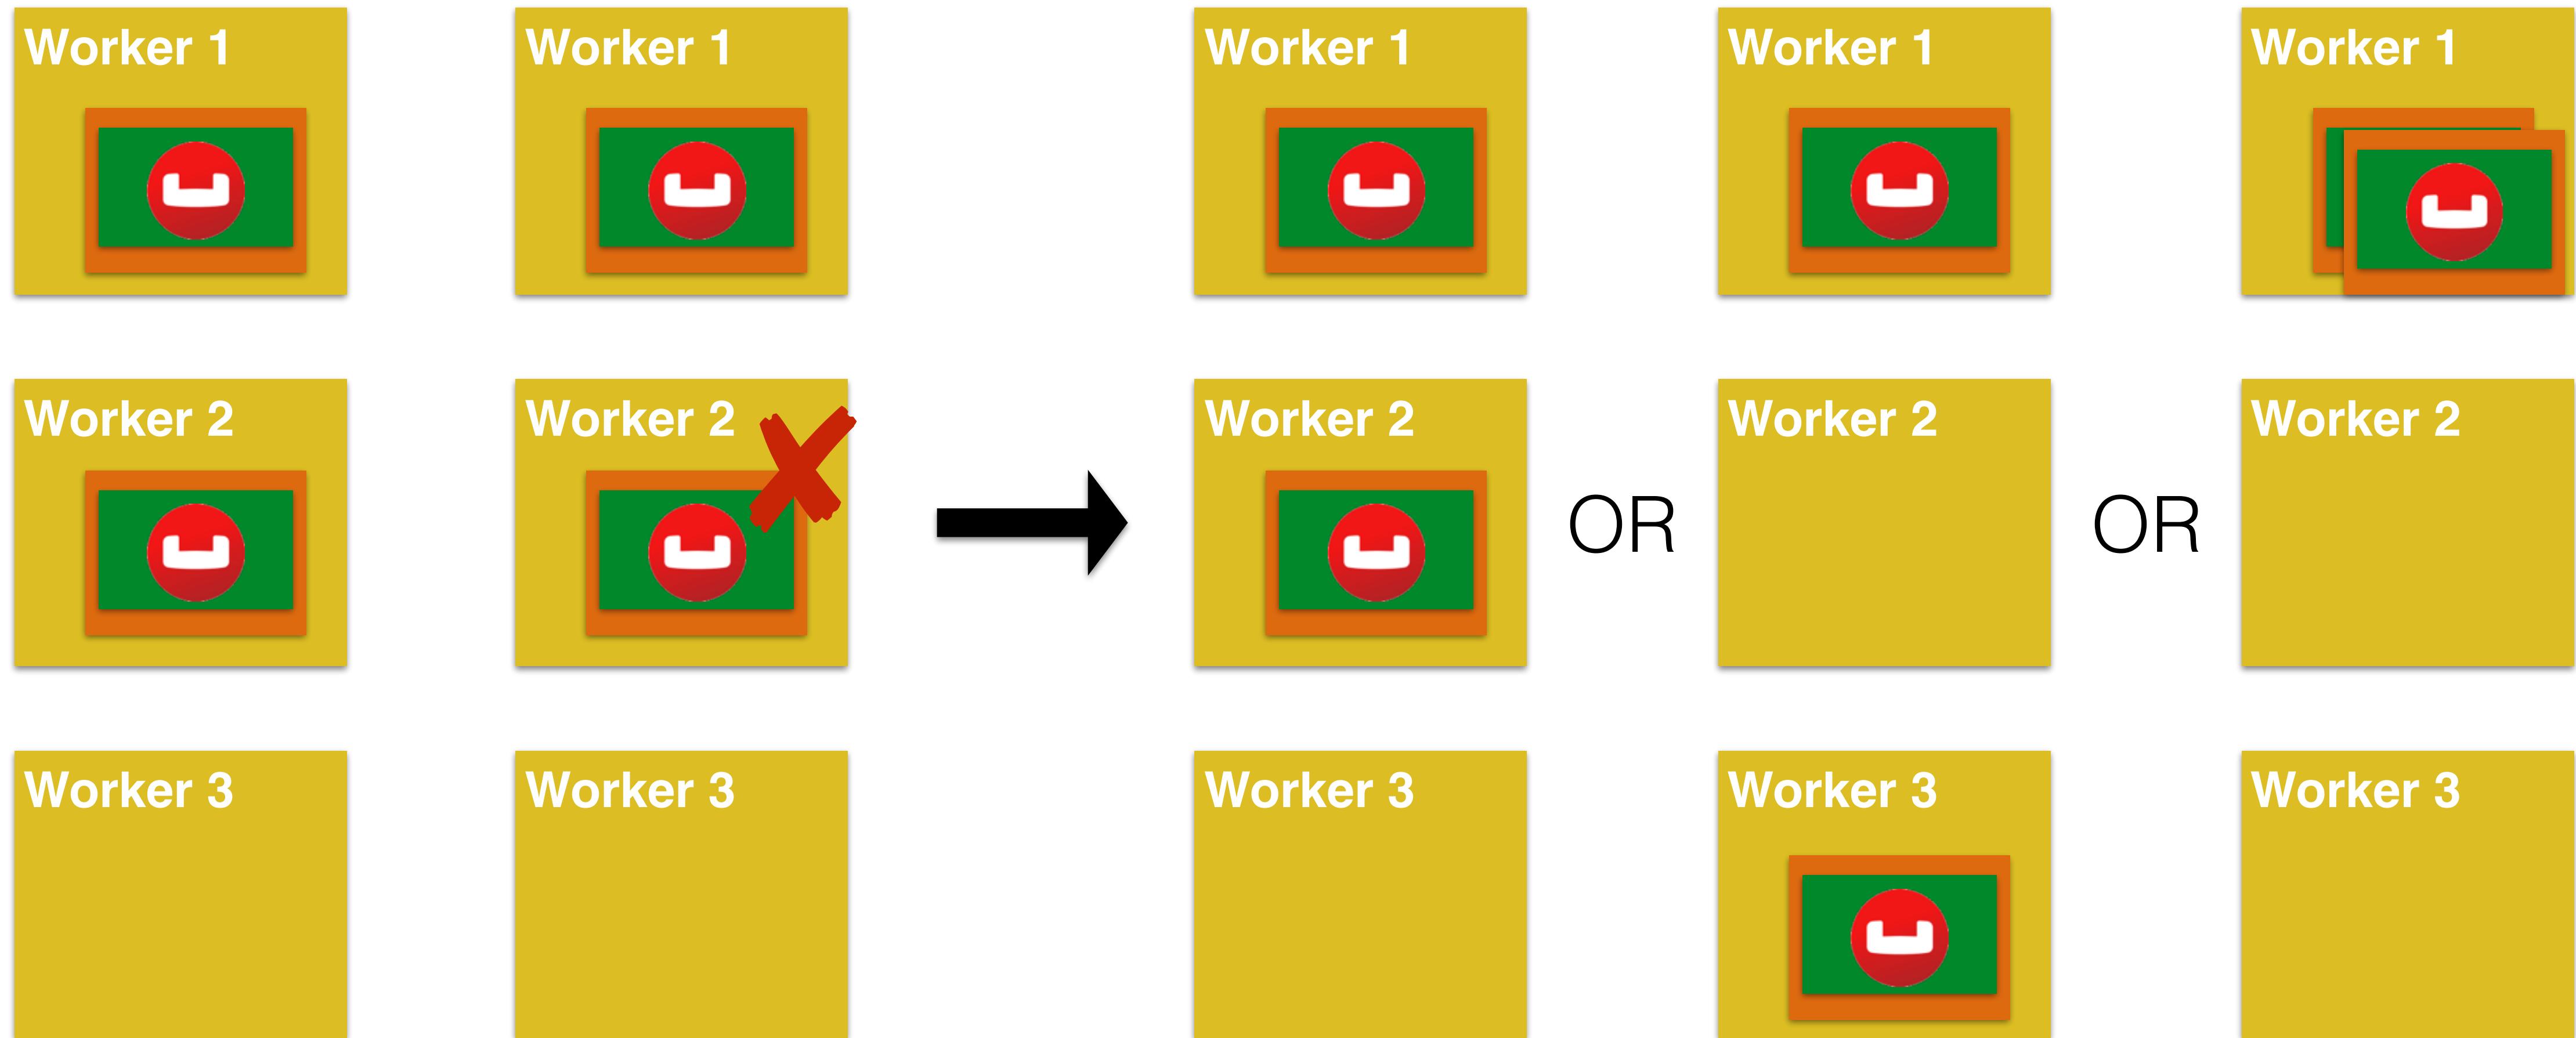
Multitenancy - Limit



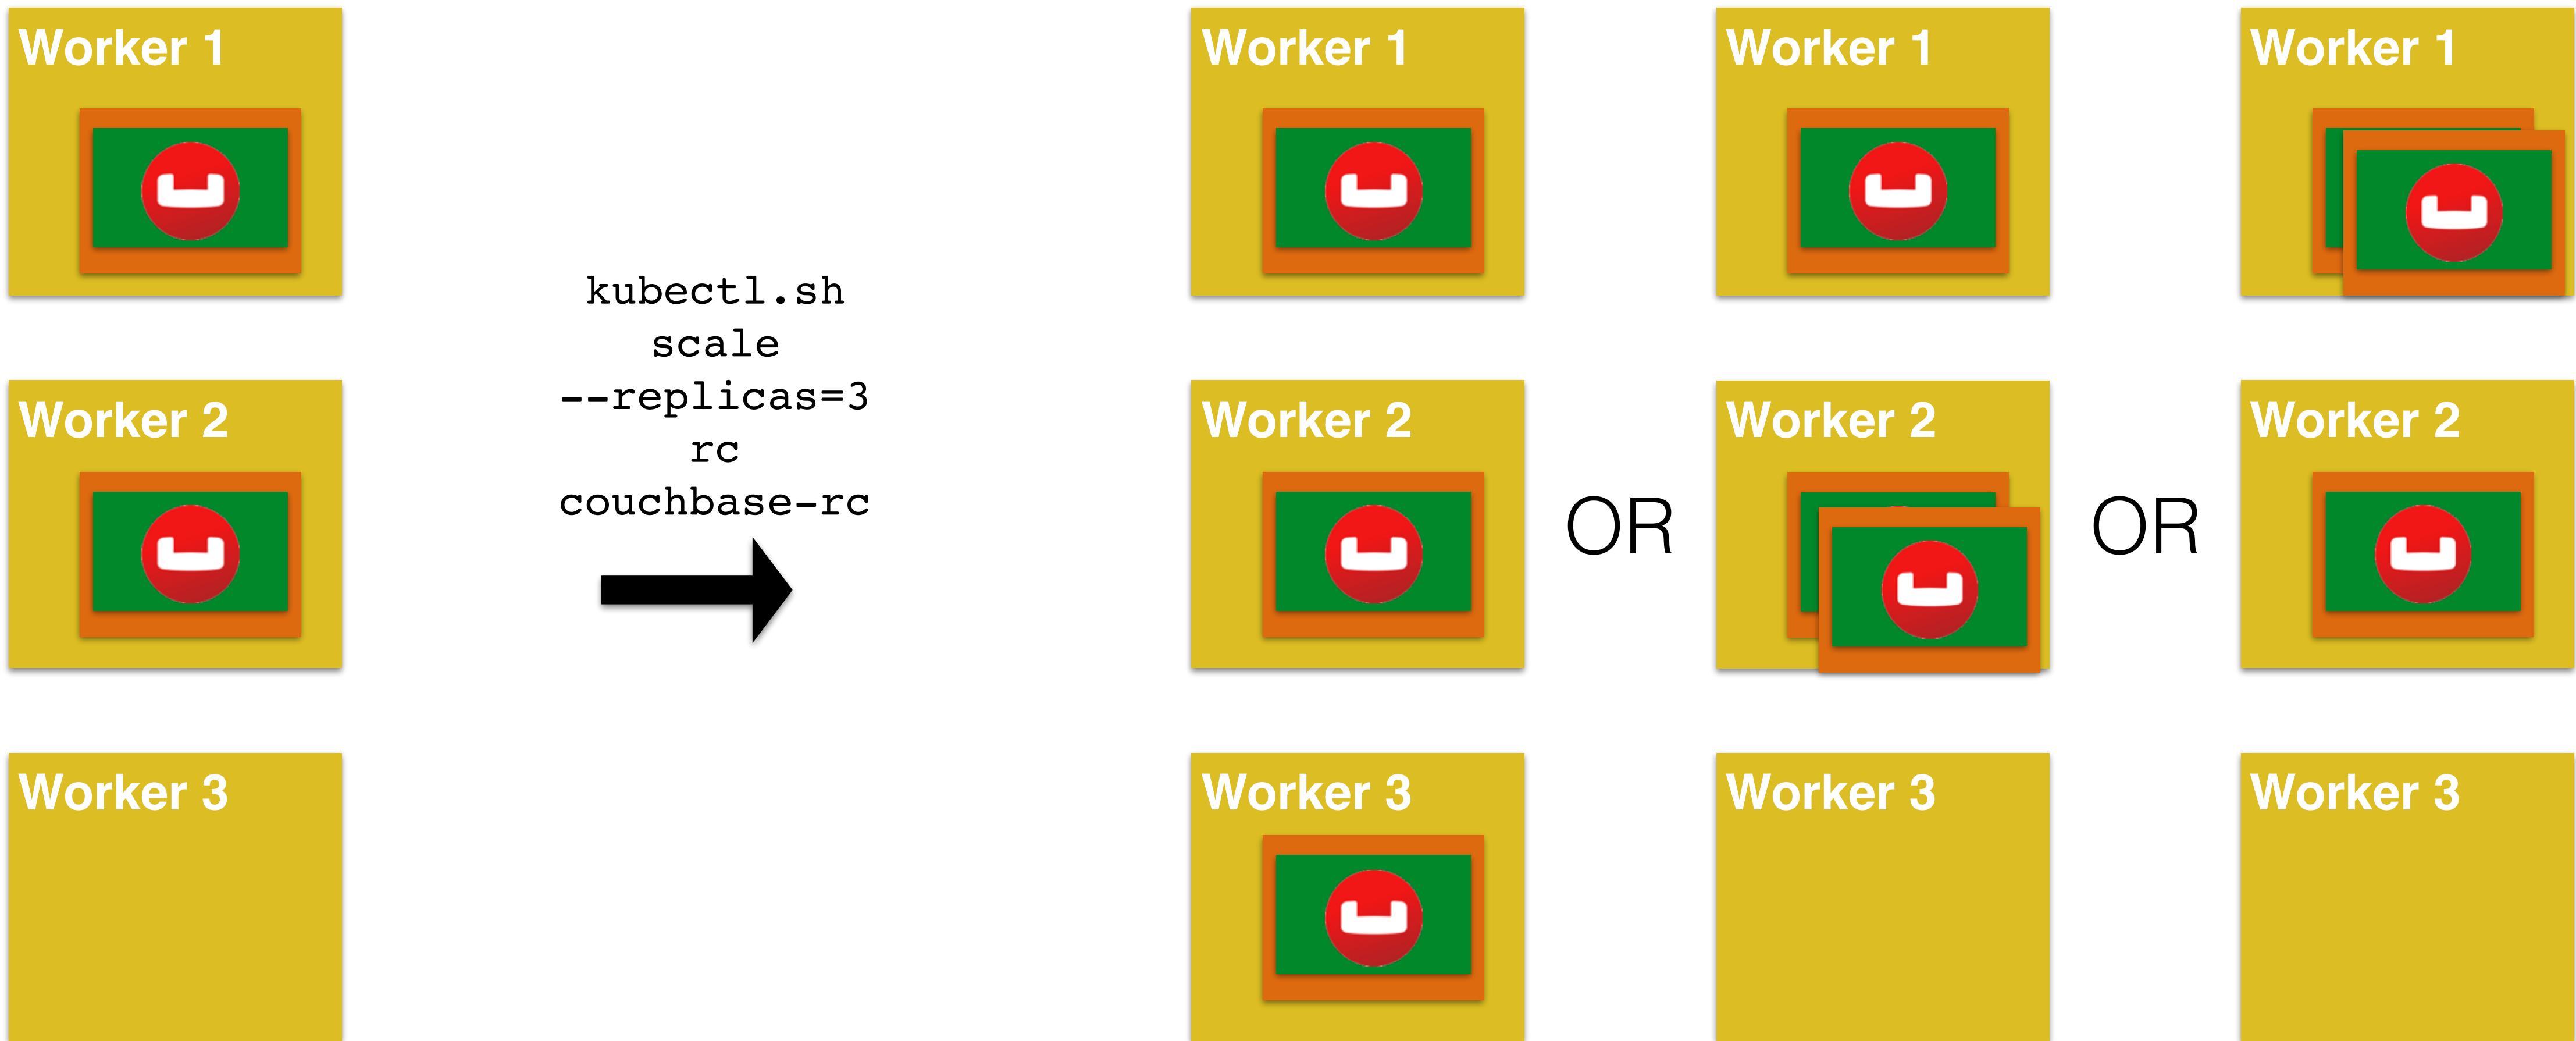
Replication Controller



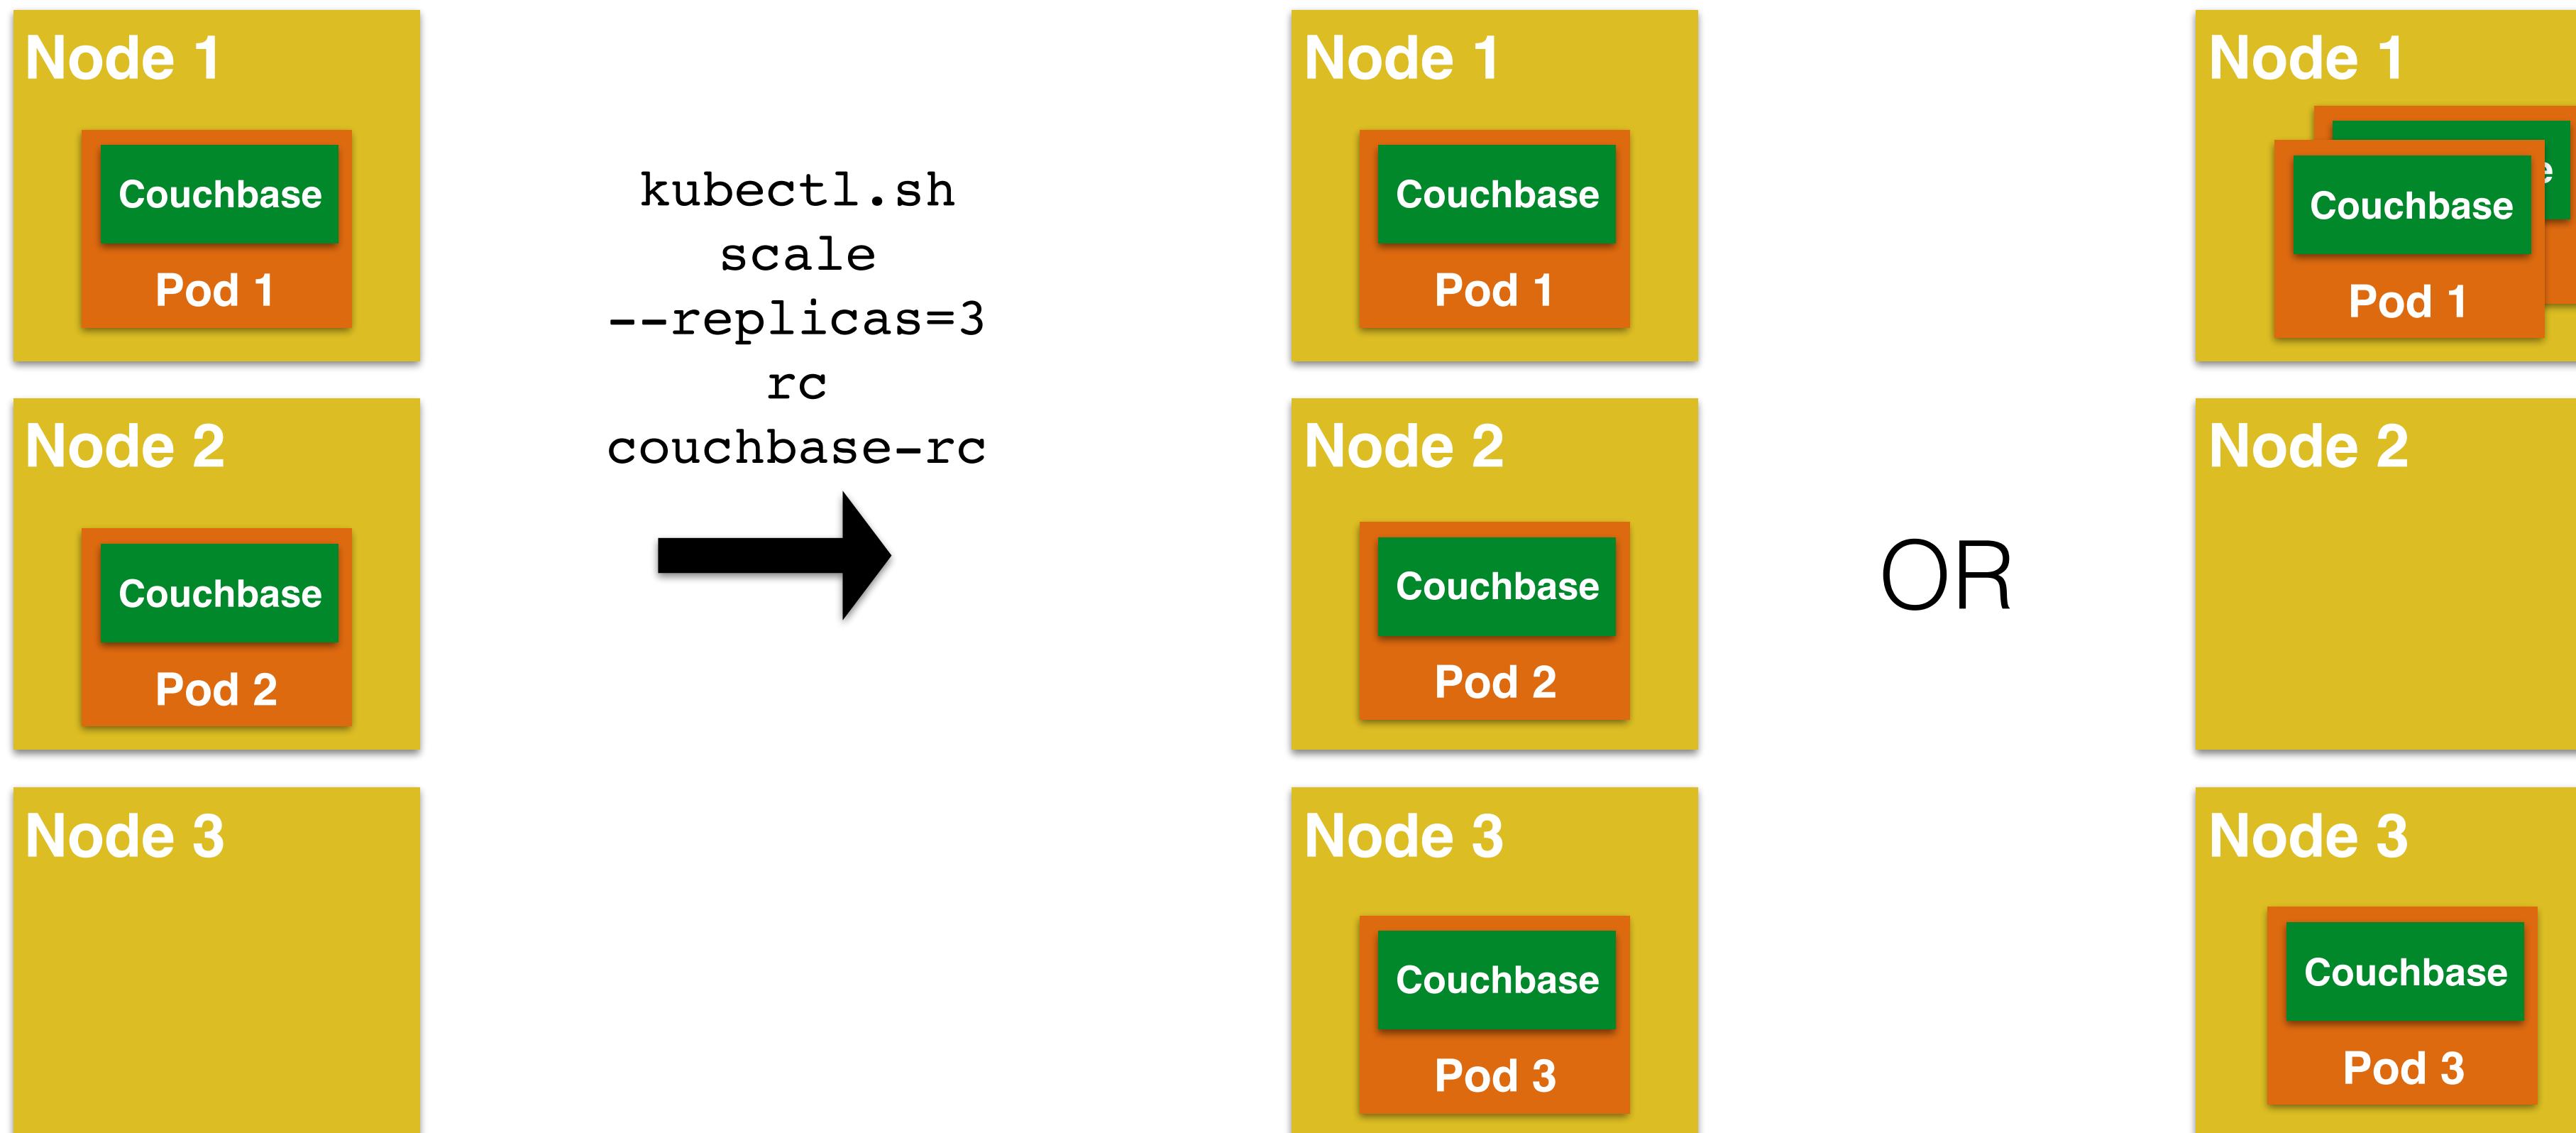
RC: “Actual” vs “Desired” State



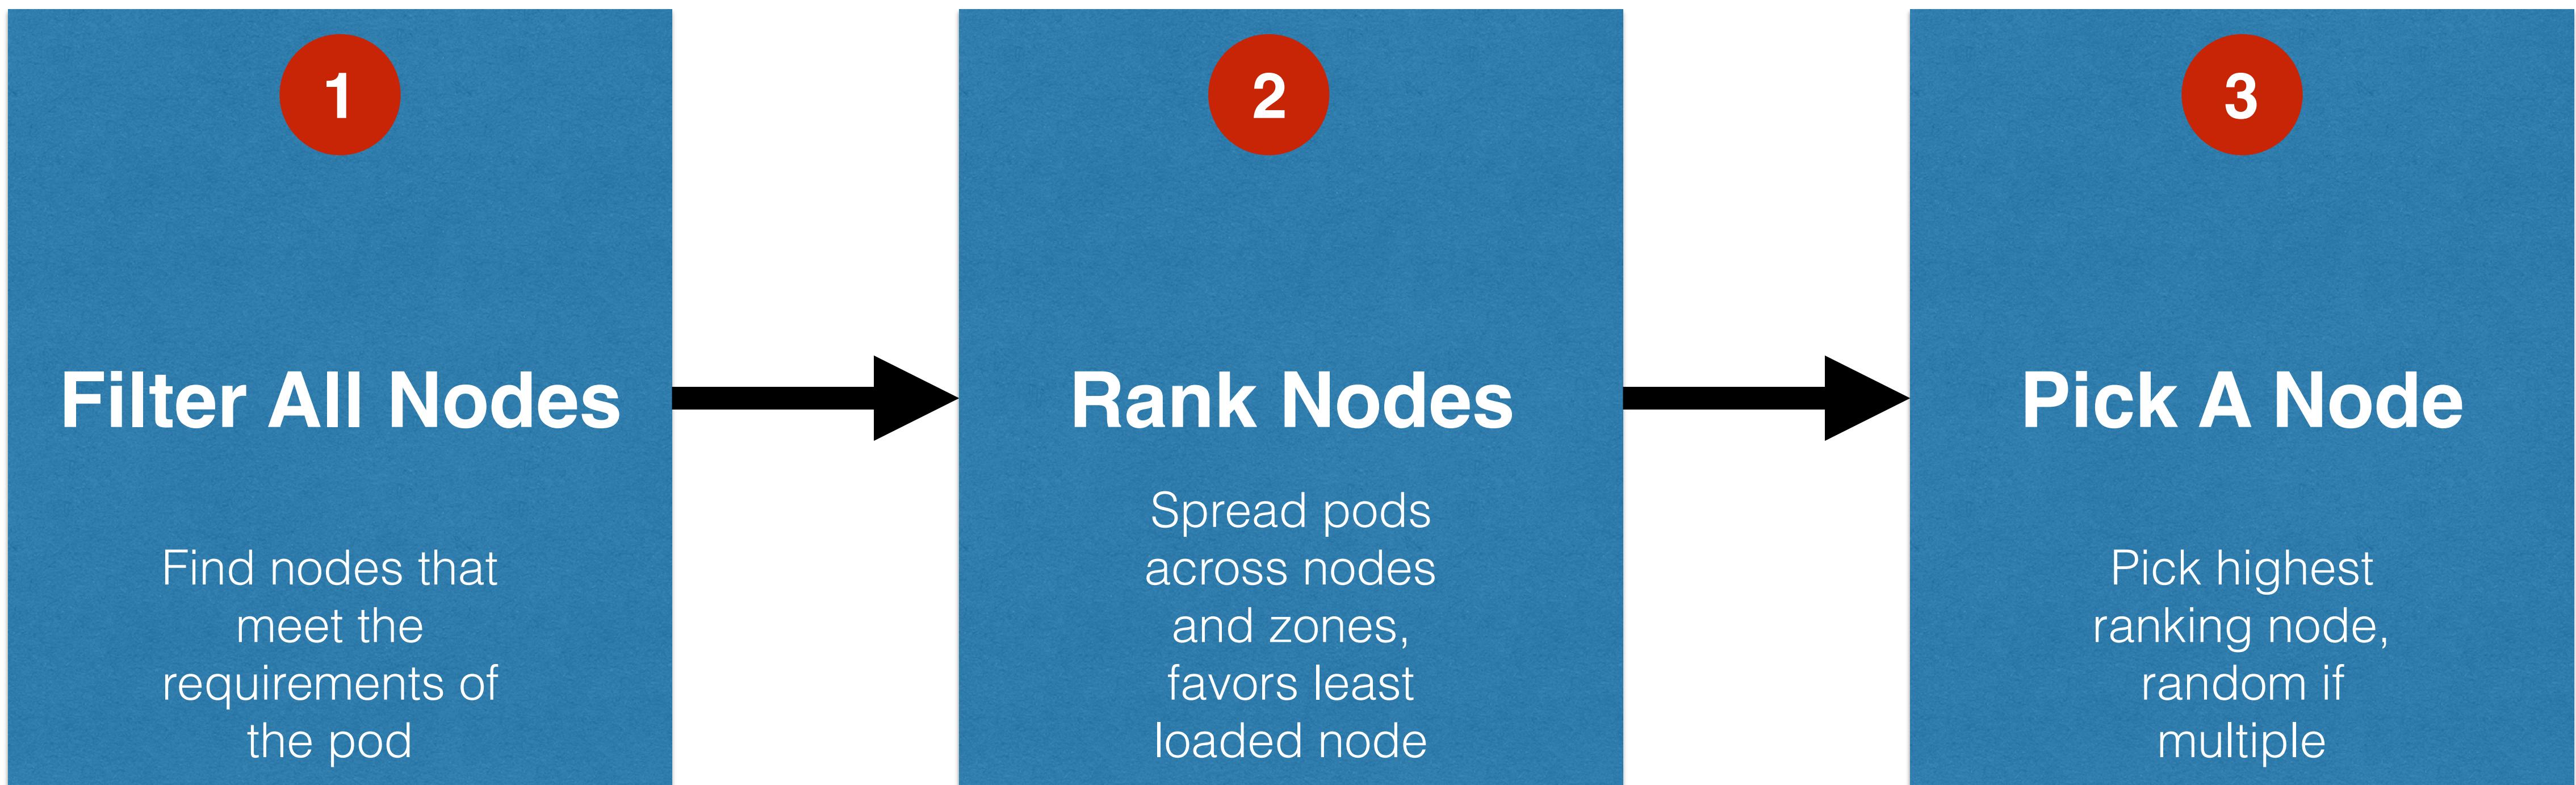
RC: Scale Pods



Replication Controller: Scaling



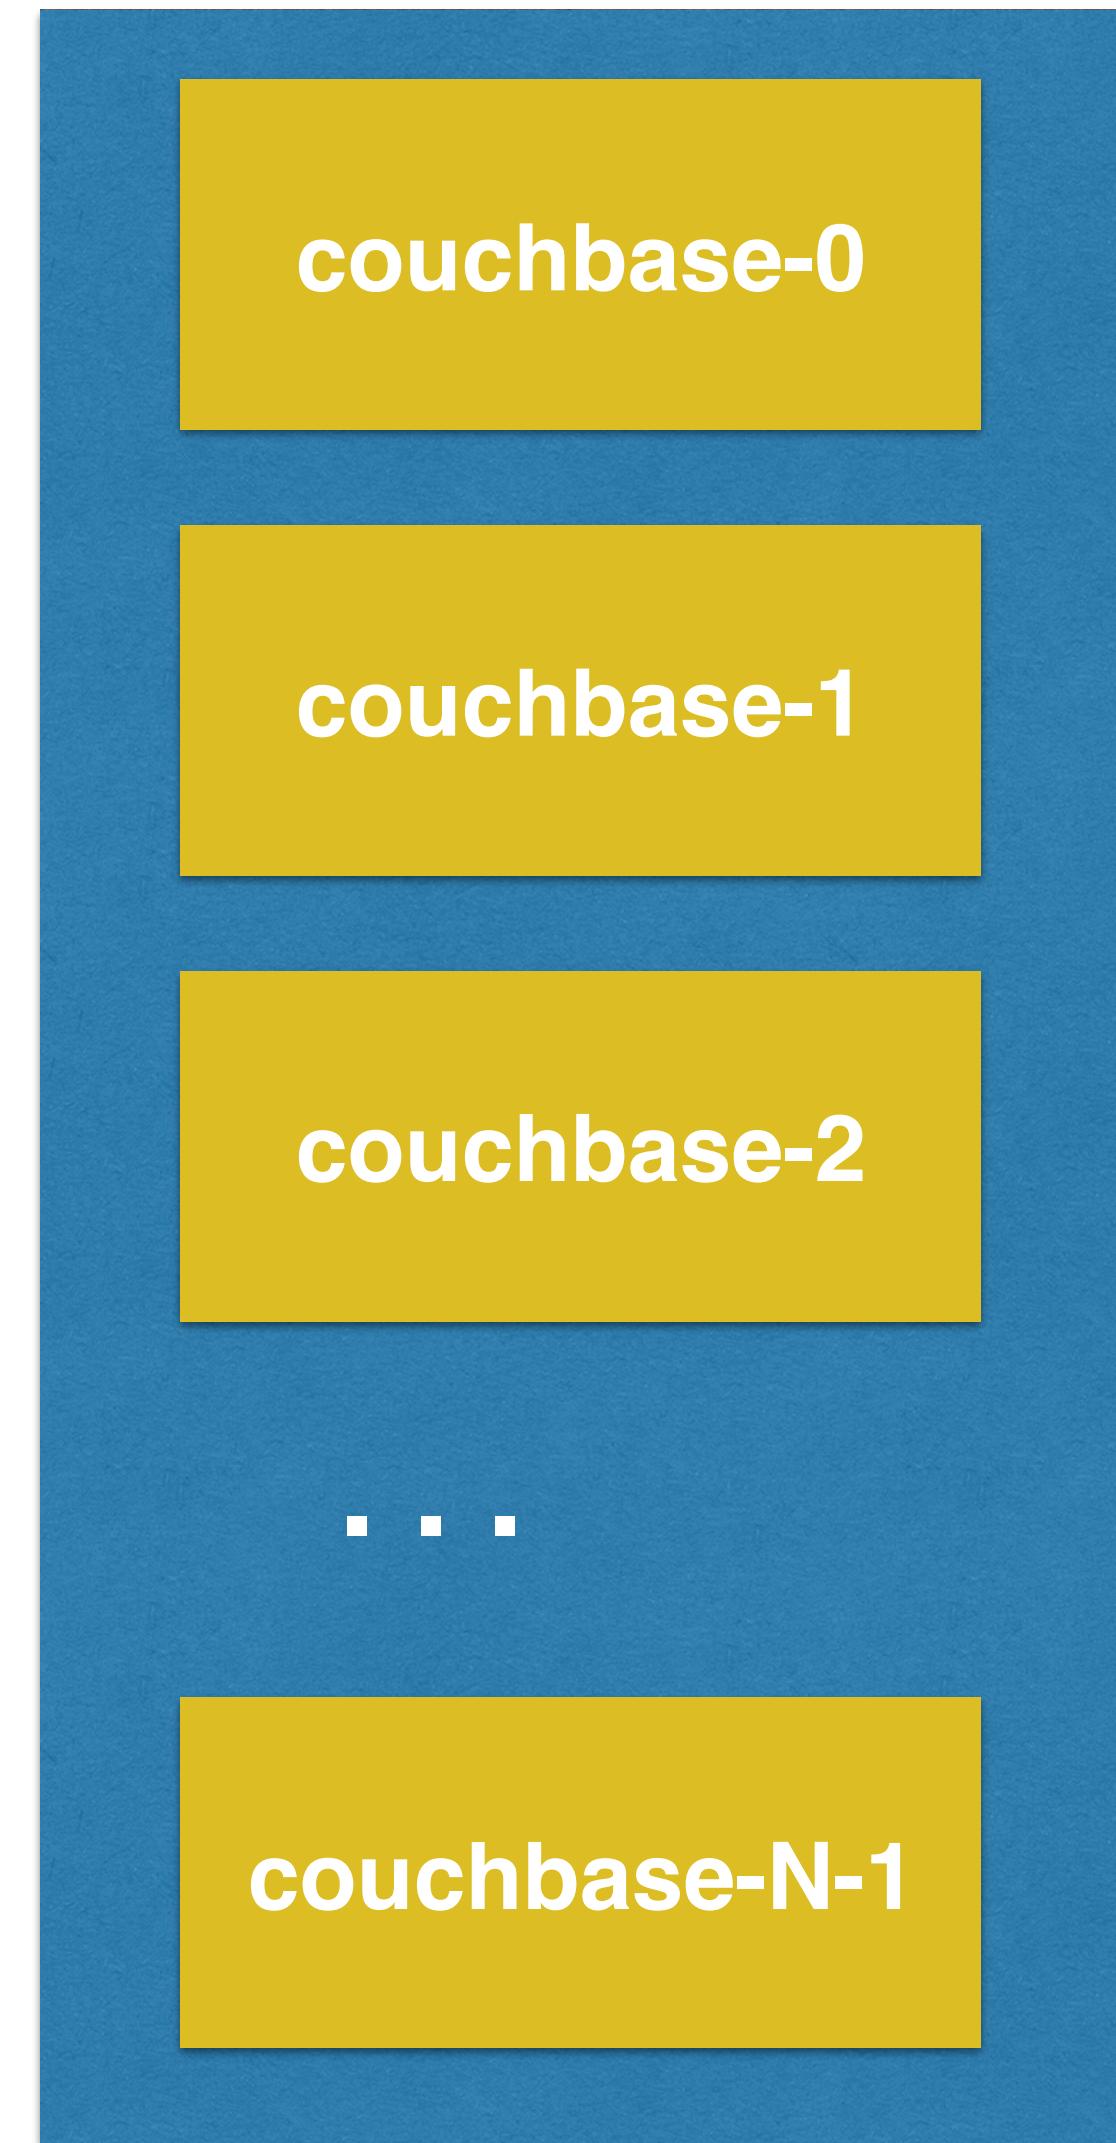
Kubernetes Scheduling Algorithm



Batch Jobs

Pet Set

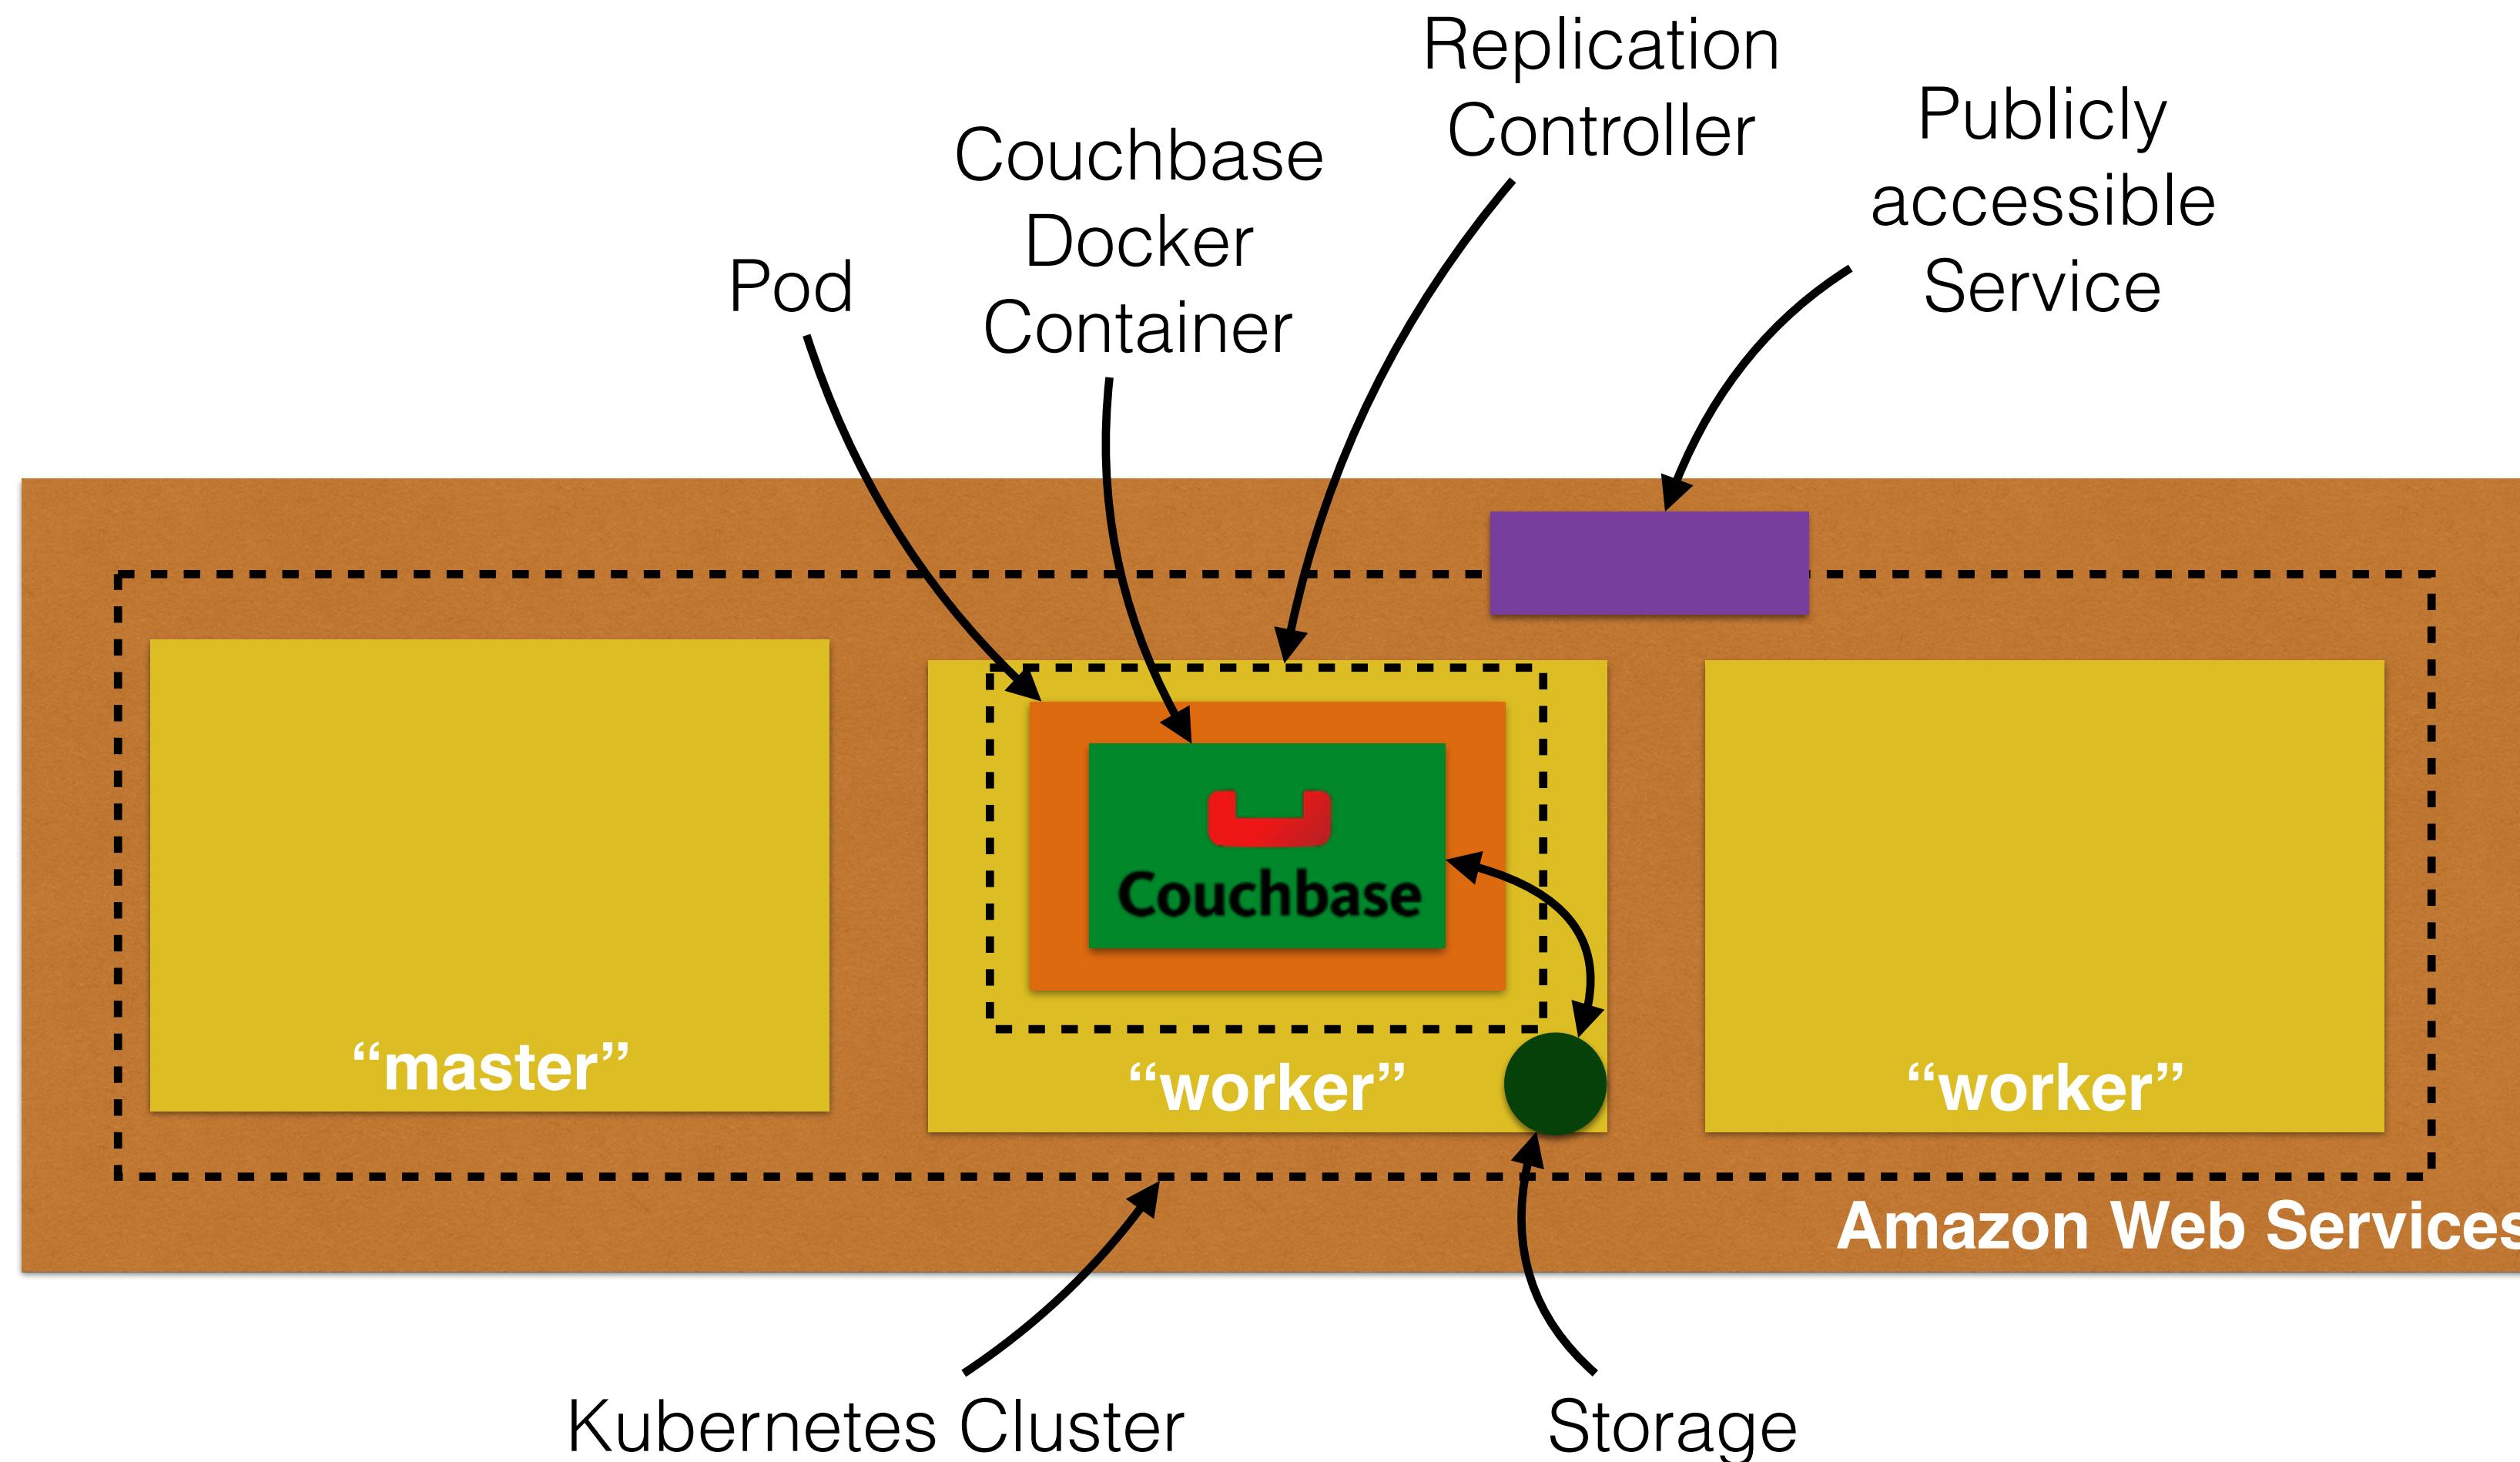
- Alpha resource introduced in 1.3
- Stateful pods
- PetSet has 0..N-1 Pets
- Each Pet has a deterministic name, and a unique identity
 - Identity = stable hostname, ordinal index, stable storage linked to ordinal & hostname
- Each Pet has at most one pod
- Each Pet Set has at most one Pet with a given identity



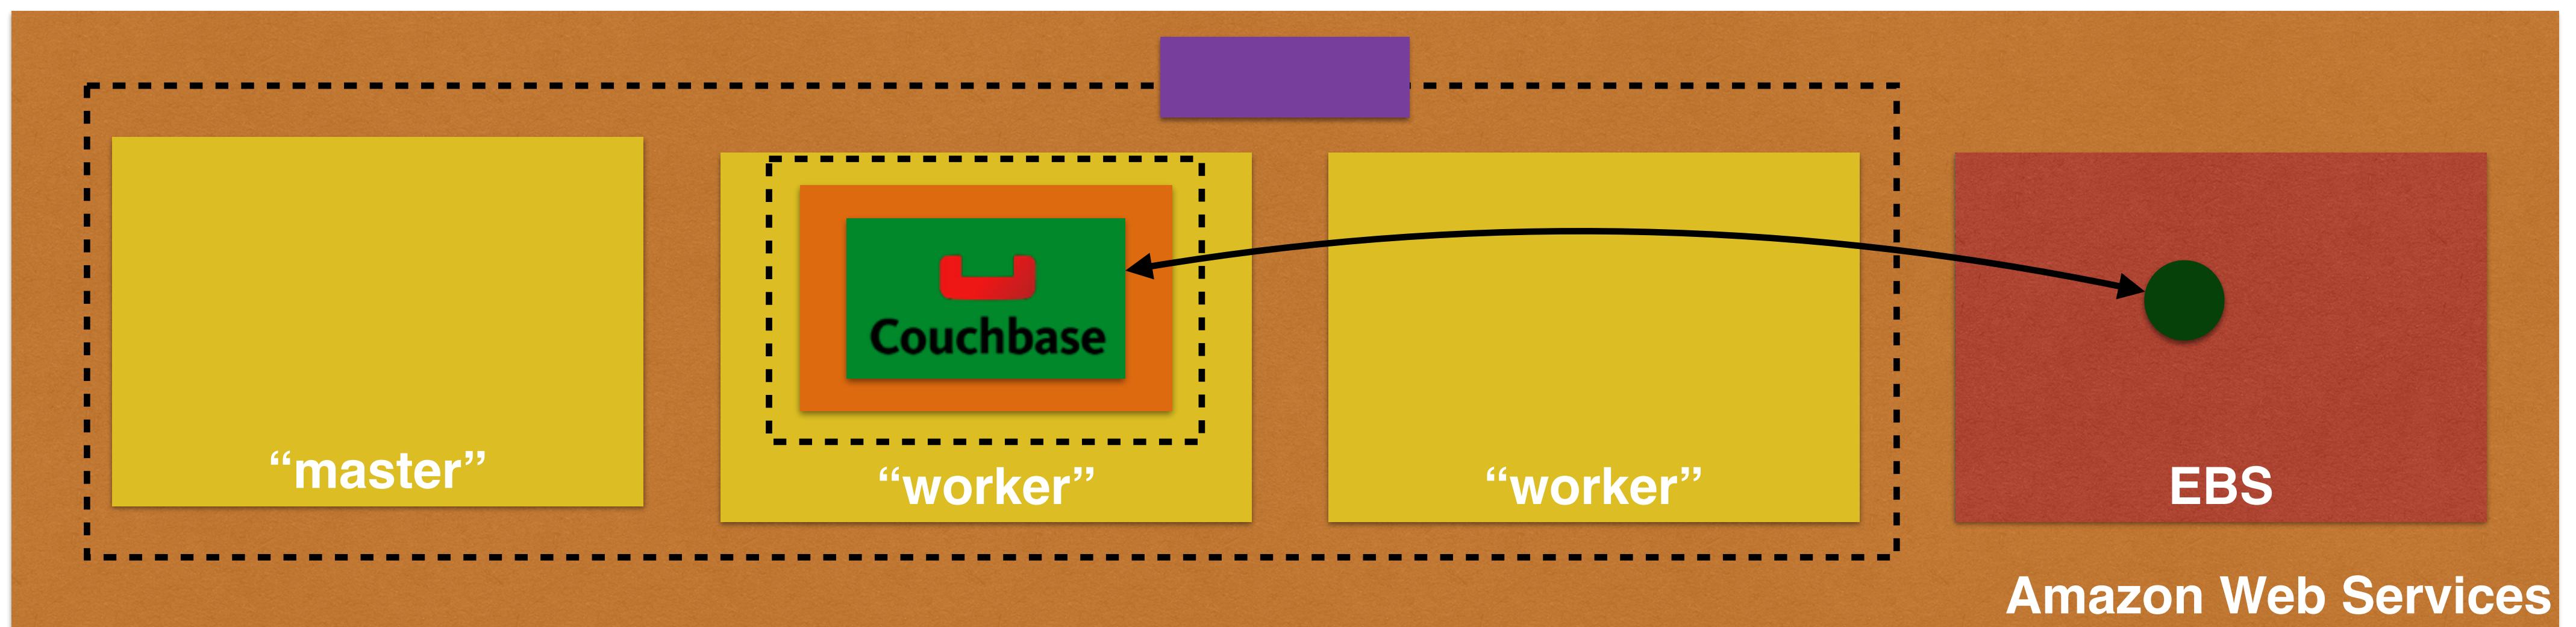
Health Checks

- Restarts Pod, if wrapped in RC
- Application-level health checks
 - HTTP
 - Container Exec
 - TCP Socket
- Health checks performed by Kubelet

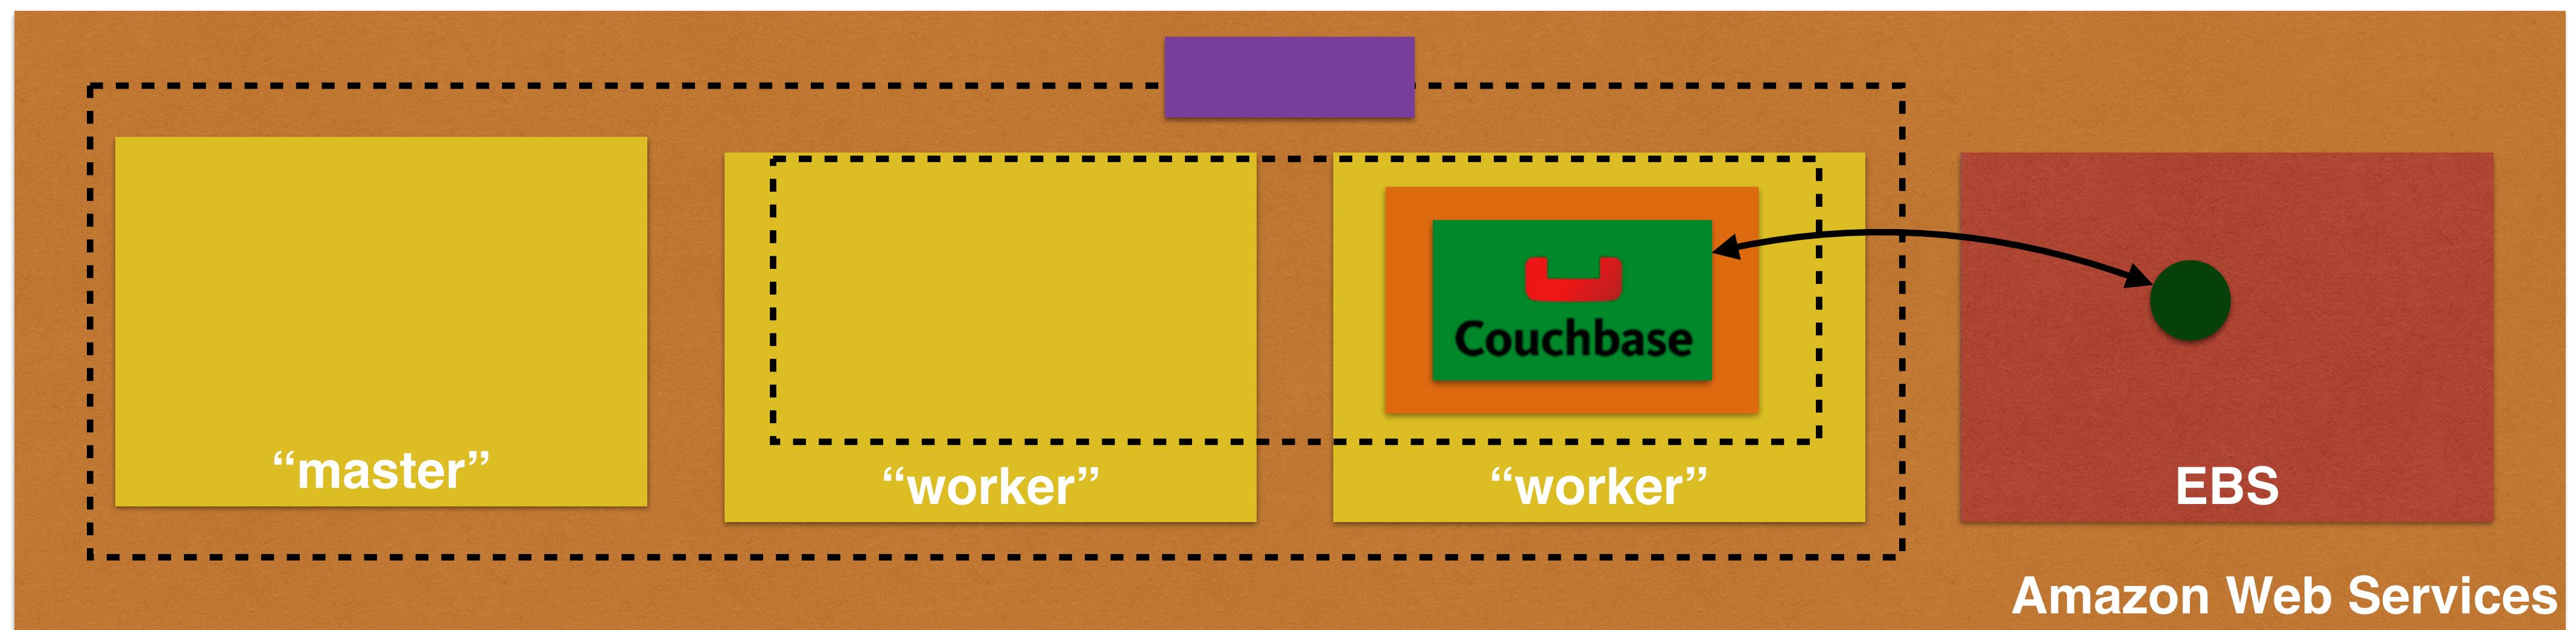
Stateful Containers



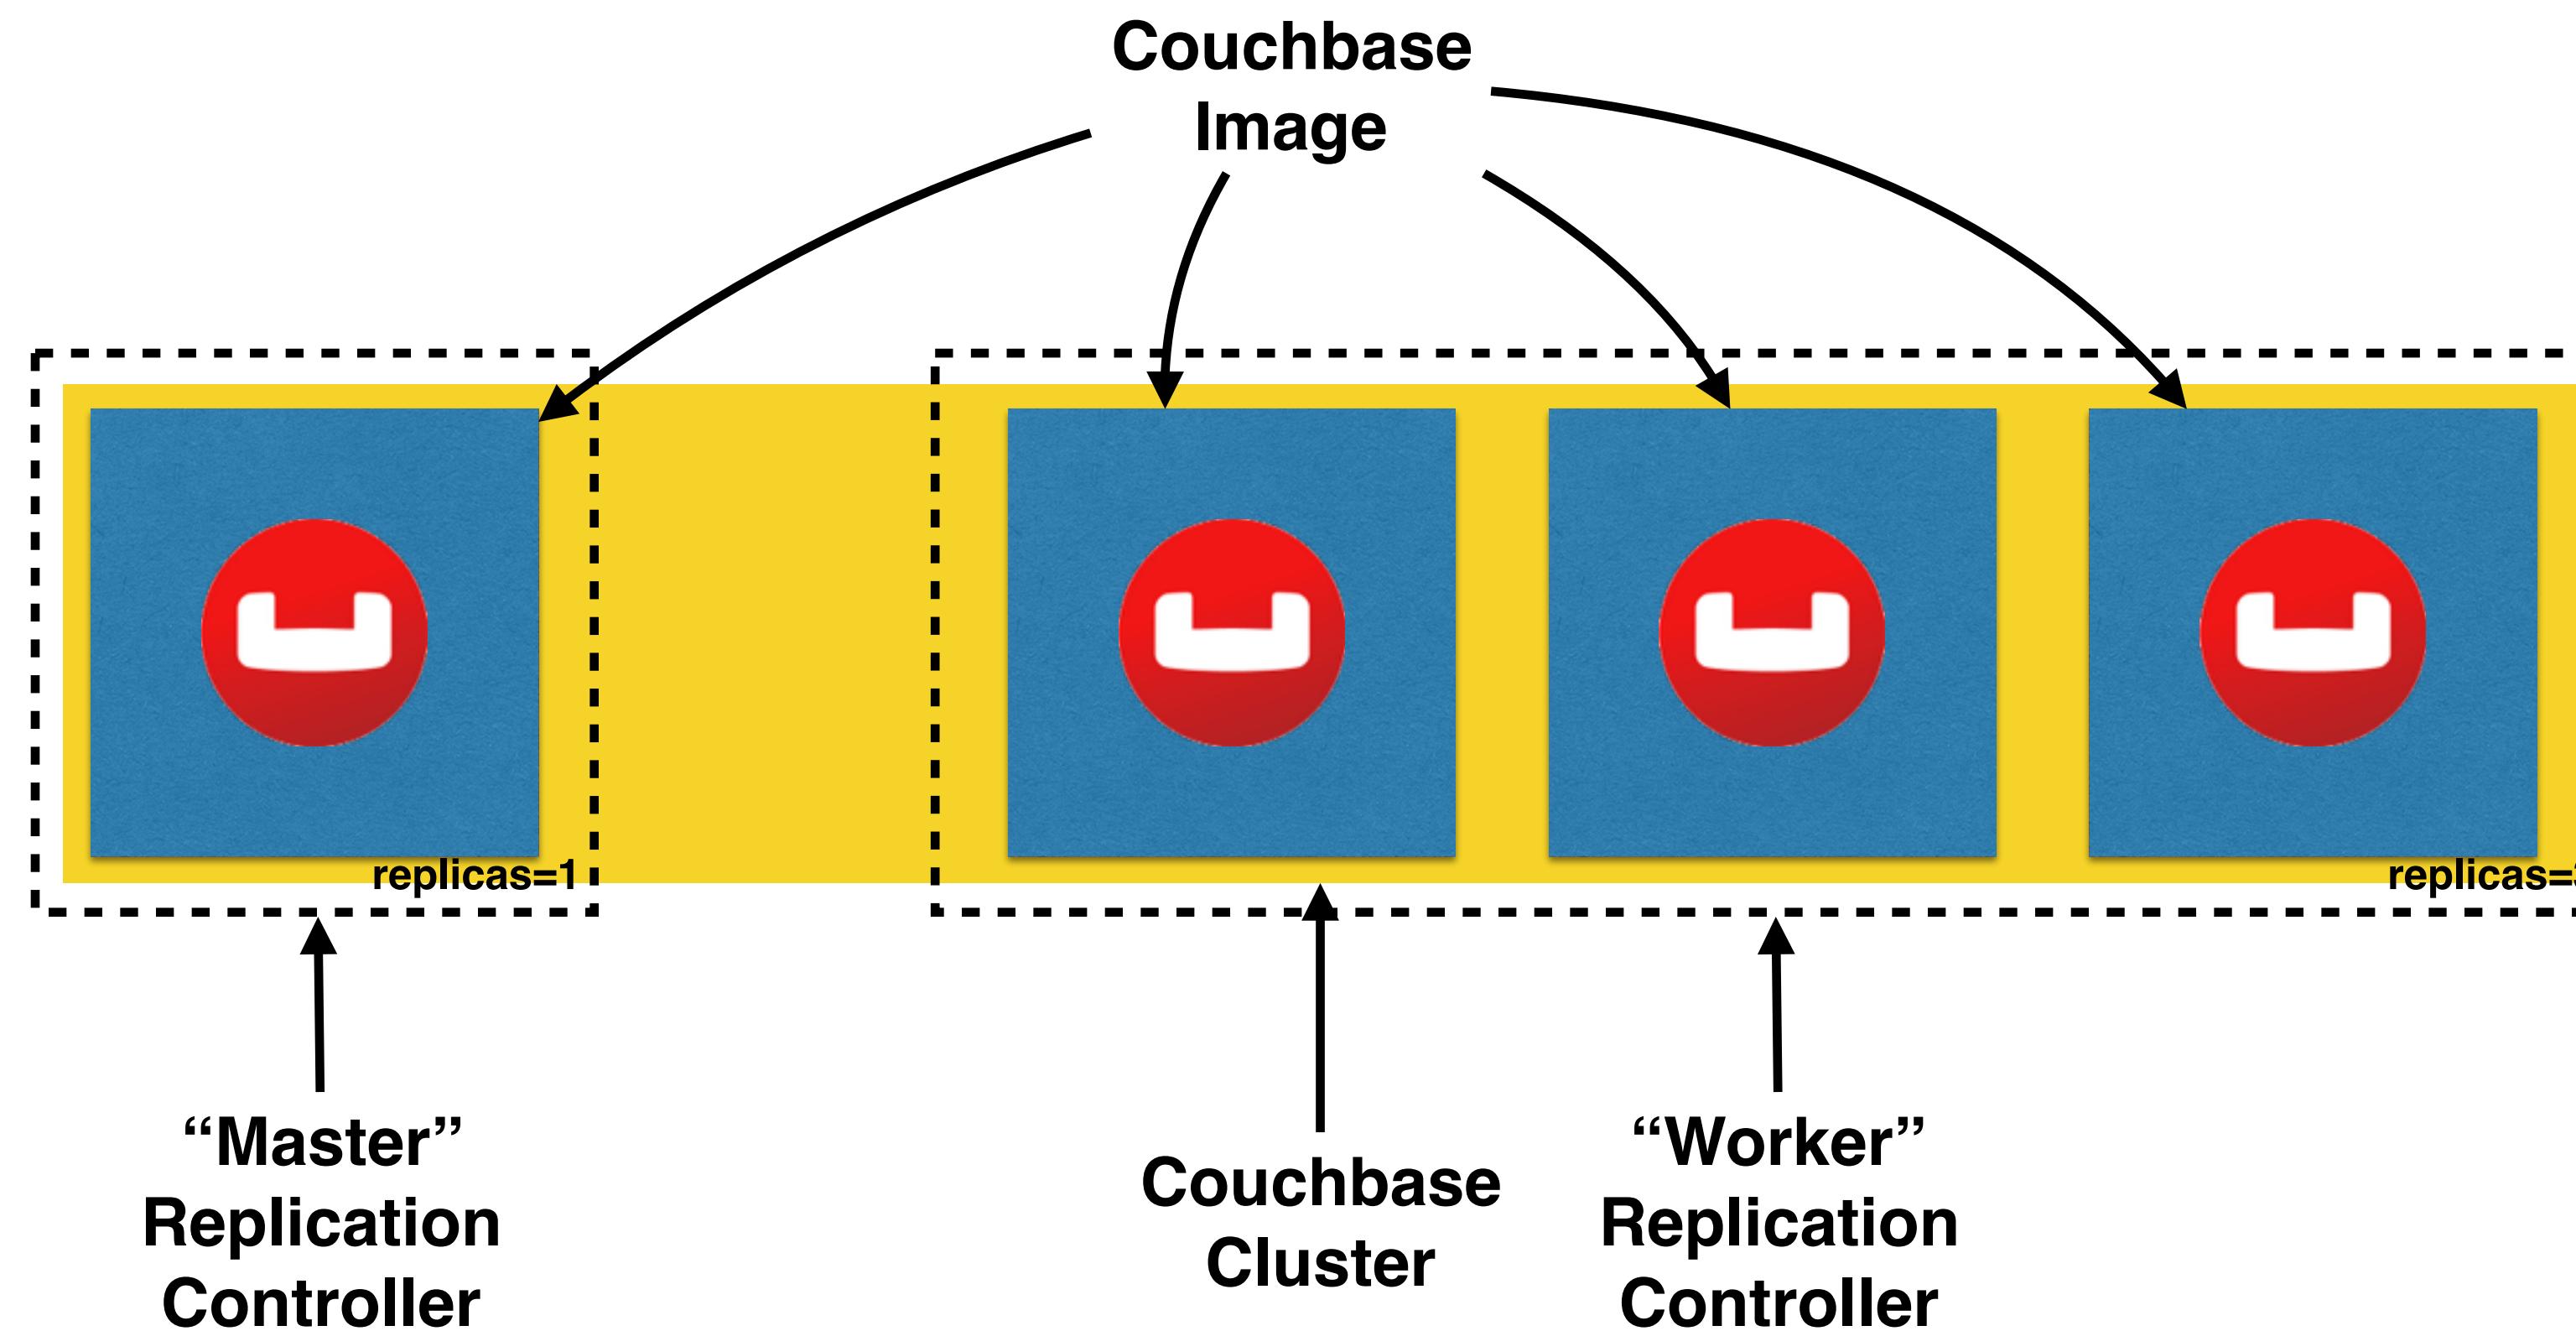
Stateful Containers



Stateful Containers



Couchbase Cluster on Kubernetes



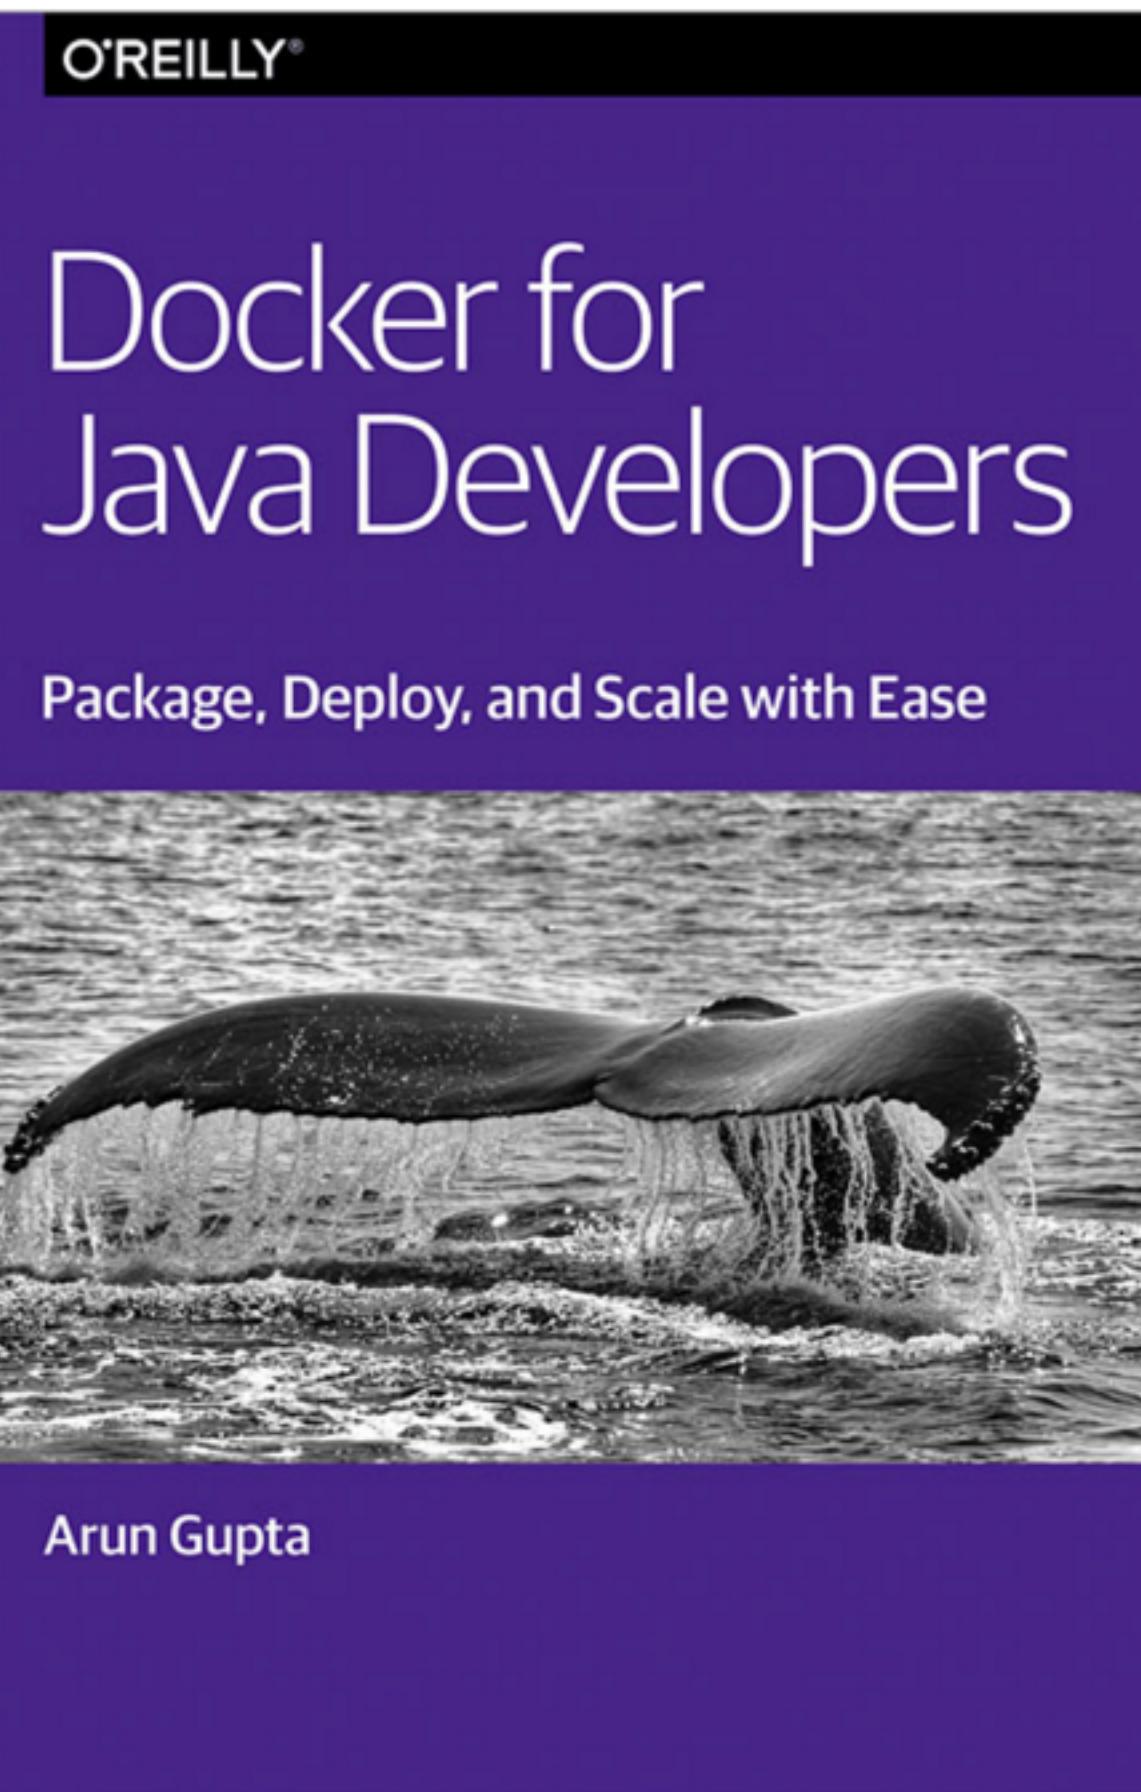
Rolling Update

- Replication Controller
- Deployment

Kubernetes and Maven

- fabric8-maven-plugin
- Brings your Java applications to Kubernetes and OpenShift
- Build Docker images
- Creates Kubernetes and OpenShift resource descriptors
- Opinionated defaults, XML configuration as well
- Reference: github.com/fabric8io/fabric8-maven-plugin

bit.ly/dockerjava



bit.ly/kubejava



References

- github.com/javaee-samples/docker-java
- kubernetes.io
- Containers recipe: couchbase.com/containers